Concept map: Getting, understanding and cleaning data

Probability distributions
Naive Bayes
Assumptions of linear regression
Standard normal and the z-table
Probability mass function (PMF)
Enough data
P-values
Probability density function (PDF)
Q-Q plot
Density estimation
Sampling noise and bias
Cumulative distribution function (CDF)
Power transformer
Population, sample, parameter and statistic
Kernel density estimation (KDE)
Kurtosis and moments
Normal distribution
Function transformer
Percentiles, quartiles and box plots
Frequency tables
Feature selection
Skewness
Bessel's correction
Binning and binarization
Percentile outlier method
Student's t-distribution
Z-score outlier method
Capping (winsorization)
Standardization
Missing indicator
Correlation significance test
Variance
K-nearest neighbours
Poor-quality data
Outliers
IQR outlier method
Descriptive statistics
Correlation
Covariance and covariance matrix
Inferential statistics
KNN imputer
Trimming outliers
PCA
Missing values
Univariate analysis
Measures of central tendency
One-way ANOVA
Ordinary least squares (closed form)
Simple linear regression
Iterative imputation (MICE)
Correlation and causation
Random sample imputation
Complete case analysis
Exploratory data analysis
Simple imputation (mean, median, mode, constant)
Bivariate and multivariate analysis
Expected value and variance of a random variable
XGBoost
Pandas Profiling
Data mining
Train-test split
Column transformer
Grid and random search
Setup: conda, Jupyter and Colab
CSV files
Precision, recall and F1
Cost-sensitive learning
Balanced random forest
ANN for classification
Imbalanced data
ROC curve and AUC
Random under- and oversampling
Confusion matrix
Web scraping
Date and time features
JSON and SQL data
Discrete and continuous data
Random variables
Labelled data
Supervised learning
APIs
Deployment

Links: needs | needs (draft) | is a kind of | fixes | compared with | used in | used in (draft)