Concept map: Getting, understanding and cleaning data

Probability distributions
Naive Bayes
Assumptions of linear regression
Standard normal and the z-table
Probability mass function (PMF)
P-values
Enough data
Labelled data
Supervised learning
Probability density function (PDF)
Density estimation
Sampling noise and bias
Q-Q plot
Cumulative distribution function (CDF)
Power transformer
Population, sample, parameter and statistic
Kernel density estimation (KDE)
Kurtosis and moments
Normal distribution
Function transformer
Percentiles, quartiles and box plots
Frequency tables
Feature selection
Bessel's correction
Skewness
Binning and binarization
Percentile outlier method
Student's t-distribution
Z-score outlier method
Capping (winsorization)
Standardization
Correlation significance test
Missing indicator
Variance
K-nearest neighbours
Poor-quality data
Outliers
IQR outlier method
Descriptive statistics
Correlation
Covariance and covariance matrix
Inferential statistics
KNN Imputer
Trimming outliers
Missing values
Univariate analysis
Measures of central tendency
One-way ANOVA
PCA
Ordinary least squares (closed form)
Simple linear regression
Iterative imputation (MICE)
Random sample imputation
Complete case analysis
Correlation and causation
Simple imputation (mean, median, mode, constant)
Exploratory data analysis
Bivariate and multivariate analysis
Expected value and variance of a random variable
XGBoost
Pandas Profiling
Data mining
Train-test split
Random under- and oversampling
Precision, recall and F1
Balanced random forest
Column transformer
CSV files
Imbalanced data
Grid and random search
Confusion matrix
ROC curve and AUC
Setup: conda, Jupyter and Colab
Cost-sensitive learning
Web scraping
Date and time features
Discrete and continuous data
JSON and SQL data
Random variables
APIs
Deployment

Links: needs · needs (draft) · is a kind of · fixes · compared with · used in · used in (draft)