

every observation gets weight  $1/n$

fit the weak learner  
with `sample_weight`

say:  $\alpha = \eta \times \ln \frac{1 - \text{error}}{\text{error}}$

mistakes  $\times e^\alpha$ , then normalise

**estimator**  
which weak learner  
(default: a stump)

**learning\_rate** =  $\eta$   
shrinks every say  
(default: 1.0)

**n\_estimators**  
at most this many stages  
(default: 50)

next  
stage

stops sooner if a stage has error 0,  
or error 0.5 or more