

1. Vanishing or exploding gradients

weight initialisation Note 1029

activation functions Note 1027

batch normalisation Note 1031

gradient clipping Note 1018

3. Slow training

optimizers Note 1032

learning rate scheduler

2. Not enough data

transfer learning section 3.1

unsupervised pre-training

4. Overfitting

L1, L2 penalties Note 1026

dropout Note 1024

early stopping Note 1022