



one context vector  
forgets long sentences

still one word at a  
time: slow to train

training from scratch  
needs huge data

still an LSTM, not  
yet a transformer

red: the problem each stage left, which the next stage fixed