



pruning



small weights (red, dotted)
removed

quantization



$0.73419 \times 128 \approx 94$: each weight a small
integer, 4 times less memory

knowledge distillation



the small model learns
to copy the large one