


Training: 1 decoder run



all 5 positions used at once

Inference: 5 decoder runs



grey: recomputed, not used
orange: the one position
that gives the next word