

input layer
one node per column

hidden layers

output layer
the prediction

x_1

x_2

x_3

$\hat{y}$

each circle is a **neuron** (perceptron)

each line carries a **weight**

