

gradient boosting
(any differentiable loss)

```
graph TD; A[gradient boosting<br/>(any differentiable loss)] --> B[XGBoost<br/>gradient boosting + optimisations]; B --> C[Flexibility<br/>- any operating system<br/>- many languages<br/>- works with other libraries<br/>- many problem types<br/>- custom loss and metrics]; B --> D[Speed<br/>- parallel split search<br/>- column blocks<br/>- cache awareness<br/>- out-of-core training<br/>- distributed training<br/>- GPU training]; B --> E[Performance<br/>- regularised objective<br/>- missing values handled<br/>- approximate split finding<br/> (weighted quantile sketch)<br/>- tree pruning with γ];
```

XGBoost

gradient boosting + optimisations

Flexibility

- any operating system
- many languages
- works with other libraries
- many problem types
- custom loss and metrics

Speed

- parallel split search
- column blocks
- cache awareness
- out-of-core training
- distributed training
- GPU training

Performance

- regularised objective
- missing values handled
- approximate split finding
(weighted quantile sketch)
- tree pruning with γ