

AI vs ML vs DL

Contents

1. Overview	2
2. Artificial Intelligence	2
2.1 Intelligence is many abilities	2
2.2 Narrow AI and general AI	3
3. Symbolic AI and expert systems	3
3.1 How an expert system works	4
3.2 Why expert systems fail	5
4. Machine Learning	5
4.1 Learning rules from data	5
4.2 Teaching a machine to recognise dogs	7
4.3 Why ML took off	7
5. Deep Learning	8
5.1 Neural networks	8
5.2 Features: chosen by us or learned	8
5.3 Layers build up understanding	9
5.4 More data helps DL more	10
6. Choosing between ML and DL	11
7. Summary	12
8. Key terms	12

1. Overview

Key point: Deep Learning is a part of Machine Learning, and Machine Learning is a part of Artificial Intelligence.




Figure 1: AI, ML and DL as nested fields

Figure 1 shows the three fields as nested circles. Every DL system is also an ML system, and every ML system is also an AI system.

2. Artificial Intelligence

Key point: AI is the field of building machines that show intelligence. Every AI system in use today handles one narrow task.

Artificial Intelligence (AI) is the field of building machines that show intelligence. To see what that involves, we first need to look at what intelligence is.

2.1 Intelligence is many abilities

Key point: Intelligence is not one ability but a combination of many, and some of them cannot even be defined precisely.




Figure 2: Components of intelligence

Figure 2 shows some of the abilities that make up human intelligence. We use different ones for different jobs: logic to write code, problem solving to crack a puzzle, imagination to create something new.

Some of these abilities, such as logic, are well defined. Others, such as creativity or emotional intelligence, have no precise definition. An ability we cannot define is very hard to build into a machine.

2.2 Narrow AI and general AI

Key point: All AI today is narrow AI. General AI, a machine with every human ability, does not exist yet.

- **Narrow AI:** a system that performs one specific task, such as recognising faces, translating text or playing chess.
- **Artificial general intelligence (AGI):** a single machine with all the abilities in Figure 2, like a human.

AGI is the long-term goal of the field. Every AI system we use today is narrow AI.

3. Symbolic AI and expert systems

Key point: Early AI worked by having humans write down rules. This works for problems with clear rules and fails for everything else.




Figure 3: Timeline of approaches to AI

Figure 3 shows how approaches to AI changed over time. Serious work on AI began in the 1950s.

The first approach was **symbolic AI**: humans write down the knowledge a machine needs, as explicit rules. The machine then follows those rules.

Extra: The start of AI is usually dated to Alan Turing’s 1950 paper *Computing Machinery and Intelligence*, which asked “Can machines think?”. The middle two eras in Figure 3 are approximate.

3.1 How an expert system works

Key point: An expert system stores an expert’s knowledge as rules and uses a program to apply them.

The best-known product of symbolic AI is the **expert system**. Figure 4 shows how one is built and used.




Figure 4: How an expert system works

1. We collect knowledge from a **human expert**, such as a doctor or a chess master.
2. We write that knowledge as rules in a **knowledge base**.

3. An **inference engine**, a program, applies those rules to answer questions.

Chess-playing computers are a classic example.

3.2 Why expert systems fail

Key point: Expert systems need clear, fixed rules. Most real-world problems do not have them.

Expert systems work well when a problem has clear, fixed rules, as in chess or logic puzzles. They fail when the rules are fuzzy.

Take the question: *does this photo contain a dog?* To answer it with rules, we would have to cover:

- hundreds of breeds,
- every size, colour, ear shape and tail,
- every camera angle and every kind of lighting.

Nobody can write rules for all of that. Speech recognition fails for the same reason. Machine Learning was developed to solve exactly this kind of problem.

4. Machine Learning

Key point: In Machine Learning, we do not write the rules. We give the machine examples with answers, and it works out the rules itself.

Machine Learning (ML) is a branch of computer science that uses statistical techniques to find patterns in data.

4.1 Learning rules from data

Key point: Traditional programming turns rules into answers. ML turns answers into rules.




Figure 5: Traditional programming vs Machine Learning

Figure 5 compares the two approaches:

- **Traditional programming:** we give the computer data and rules, and it produces answers.
- **Machine Learning:** we give the computer data and the correct answers, and it produces the rules.

Writing every rule by hand is called **explicit programming**. ML needs none of it.

Finding the rules from examples is called **learning**. Once a machine has learned, it can **predict**: give an answer for new data it has never seen.

Extra: What does “finding a pattern” look like? Suppose we record how many hours 12 students studied and the marks each one scored. Plotted together (Figure 6), the points rise from left to right, and the line through them says: each extra hour of study adds about 10 marks. That line *is* the pattern. A student who studies 5 hours can now be predicted to score about 65, even though we never saw that student. Finding the best line through data like this is one of the statistical techniques ML uses (covered later as *linear regression*).




Figure 6: A pattern in data: marks rise with hours studied (example data)

4.2 Teaching a machine to recognise dogs

Key point: We stop writing rules and start providing labelled examples.

Back to the dog problem from Section 3.2:

Approach	What we do	Result
Symbolic AI	Write rules for every breed, colour and shape	Impossible to finish
ML	Show thousands of photos labelled “dog” or “not dog”	The machine learns what a dog looks like and can classify new photos

Our job changes from writing rules to providing good data. This is how children learn too: from examples (“this is a dog, this is not”), not from a rulebook.

4.3 Why ML took off

Key point: ML is old, but it only became practical once we had enough data and fast enough computers.

ML methods have existed for decades. They became dominant in industry over the last 20 to 30 years because of two changes:

1. **Large amounts of data**, produced by the internet and phones.
2. **Fast, cheap hardware** that can learn from that data.

5. Deep Learning

Key point: DL is ML that uses neural networks. Its big advantage: it finds the features by itself, and it keeps improving as we give it more data.

Deep Learning (DL) is a subset of ML that took off after about 2010.

The process is the same as in ML. We give data to an algorithm and **train** it: the algorithm makes predictions, measures how wrong they are, and adjusts itself to be less wrong. Then we use it to predict on new data. What changes is the algorithm.

5.1 Neural networks

Key point: Neural networks are inspired by the brain, but they do not work like the brain.

DL uses **neural networks**, which are loosely inspired by the neurons in the brain. How the brain works is still not fully understood, so a neural network is not a copy of it. It is a mathematical model that borrows one idea: many simple units connected together.

The smallest unit of a neural network is the **perceptron**, an artificial neuron. It is covered in detail later in the course.

5.2 Features: chosen by us or learned


Key point: In ML, we choose the features. In DL, the network learns them.

A **feature** is one piece of information about an example that a model uses to make its decision. For example, to predict whether a student will get a job in campus placements, our data might look like this:

Student	CGPA	IQ	Certifications	Placed?
A	8.9	120	3	Yes
B	6.1	105	0	No
C	7.5	112	2	Yes

Here CGPA, IQ and Certifications are the features, and “Placed?” is the answer we want to predict.

Machine Learning



Deep Learning




Figure 7: Feature selection in ML vs DL (student placement example)

Figure 7 shows the difference:

- **ML:** we decide which features matter and supply them. Choosing well requires a good understanding of the data. A useful feature we leave out can never be used by the model.
- **DL:** we supply the raw data, and the network works out which information matters.

This makes DL valuable when nobody knows what the right features are. For example, nobody can list the features that make a photo a dog photo.

5.3 Layers build up understanding

Key point: Each layer of a neural network combines what the previous layer found into something bigger.

A neural network is organised in **layers** of neurons. Each layer works on the output of the layer before it. Adding layers lets the network detect more complex patterns, and a network with many layers is called *deep*.

Figure 8 shows a network recognising a handwritten digit:




Figure 8: Layers of a neural network recognising a handwritten 7

1. The first layer detects **edges**: short straight segments.
2. The next layer combines edges into **shapes**: a horizontal bar and a slanted line.
3. The next layer combines the shapes into the answer: 7.

Extra: Nobody tells the layers to look for “edges” or “shapes”. When researchers inspected trained networks, they found that early layers tend to respond to edges and later layers to larger parts. Figure 8 is a simplified version of this.

5.4 More data helps DL more

Key point: ML stops improving after a point. DL keeps improving as data grows.




Figure 9: Performance as the amount of data grows

Figure 9 shows how performance changes as we add data:

- **ML** improves at first, then levels off.
- **DL** keeps improving as data grows.

This is why DL now outperforms ML on tasks with very large datasets: image classification, object detection, and text and speech tasks.

6. Choosing between ML and DL

Key point: Lots of data, especially images, text or speech: use DL. Small or tabular data: use ML.

DL needs large amounts of data. With small datasets it performs worse than ML, as the shaded region of Figure 9 shows.

Many organisations, such as banks and insurance companies, do not have that much data. For them, ML remains the standard choice.

Our data	Use
Small or medium, especially tables of numbers	ML
Very large, especially images, text or speech	DL

7. Summary

	AI	ML	DL
What it is	Machines that show intelligence	Machines that learn rules from data	ML with many-layered neural networks
Who writes the rules	Humans (symbolic AI)	The machine	The machine
Who chooses the features	No learning involved	We do	The network
Data needed	None	Moderate	Very large
Strong at	Problems with clear rules	Tables of data	Images, text, speech

- $AI \supset ML \supset DL$.
- Expert systems fail on problems whose rules cannot be written down.
- ML: data + answers $\rightarrow$ rules. No explicit programming.
- DL learns its own features.
- With more data, DL keeps improving while ML levels off.
- With little data, use ML.

8. Key terms

Term	Meaning
Artificial Intelligence (AI)	The field of building machines that show intelligence
Narrow AI	AI that performs one specific task
Artificial general intelligence (AGI)	A machine with all human abilities; does not exist yet
Symbolic AI	AI built from rules written by humans
Expert system	Symbolic AI system: an expert's rules plus a program that applies them
Knowledge base	The set of rules inside an expert system
Inference engine	The program that applies an expert system's rules
Machine Learning (ML)	Using statistical techniques to let a machine find patterns in data
Explicit programming	Writing every rule a computer follows by hand
Learning	Finding rules (patterns) from examples
Predict	Give an answer for new, unseen data
Deep Learning (DL)	ML using neural networks with many layers

Term	Meaning
Train	Let a model learn by repeatedly reducing its errors
Neural network	A model of many connected simple units, loosely inspired by the brain
Perceptron	The smallest unit of a neural network; an artificial neuron
Feature	One piece of information about an example that a model uses
Layer	One stage of a neural network, building on the previous stage