


Types of Machine Learning

Contents

1. Overview	2
2. Supervised learning	3
2.1 Learning from inputs and outputs	3
2.2 Numerical and categorical data	3
2.3 Regression and classification	4
3. Unsupervised learning	5
3.1 Learning from inputs only	5
3.2 Clustering	5
3.3 Dimensionality reduction	6
3.4 Anomaly detection	7
3.5 Association rule learning	8
4. Semi-supervised learning	9
5. Reinforcement learning	10
5.1 Agent, environment, policy and reward	10
5.2 Example: fire or water	11
5.3 Where RL is used	12
6. Summary	13
7. Key terms	13

Where this fits: Pipeline map steps 0 (Foundations), 2 (Get data), 6 (Reduce dimensions), 8 (Model). See the [Course map](#).



- **Builds on:** Machine learning ([Note 2](#)).
- **Leads to:** K-nearest neighbours ([Note 6](#)); Applications of ML ([Note 8](#)); Feature selection (Video 9, coming); Feature extraction ([Note 23](#)); Simple linear regression ([Note 50](#)); Multiple linear regression ([Note 53](#)).

1. Overview

Key point: Based on how much supervision they need, ML algorithms fall into four types: supervised, unsupervised, semi-supervised and reinforcement learning.

ML algorithms can be grouped in three different ways:

- by **how much supervision** they need (this Note),
- by whether they learn all at once or keep learning over time (batch vs online),
- by how they generalise to new data (instance-based vs model-based).

Supervision means correct answers that guide the algorithm while it learns. Figure 1 shows the four types this produces and their sub-types.




Figure 1: ML types by amount of supervision

2. Supervised learning

Key point: In supervised learning, the data has both inputs and outputs. The algorithm learns how the output depends on the inputs, then predicts the output for new inputs.

2.1 Learning from inputs and outputs

Key point: Input columns + output column → learn the relationship → predict the output for new rows.

In **supervised learning**, every row of the data has both **input** values and the correct **output**. The algorithm learns the mathematical relationship between them.

Example: campus placements. We have data on 5,000 past students (Figure 2):

- **Inputs:** IQ and CGPA.
- **Output:** whether the student was placed.




Figure 2: Supervised learning: learn from labelled rows, predict for a new row

After learning, the algorithm can take a new student, say IQ 73 and CGPA 7.3, and predict *yes* or *no*. Most of the ML used in industry is supervised learning.

Extra: The output column is also called the **target** or the **label**. Data that includes it is called **labelled data**.

2.2 Numerical and categorical data

Key point: Data is either numerical (numbers) or categorical (categories).

Before splitting supervised learning further, we need the two basic types of data:

Type	What it holds	Examples
Numerical	Numbers	Age, weight, height, IQ
Categorical	Categories	Gender, nationality, phone brand

2.3 Regression and classification

Key point: Numerical output: regression. Categorical output: classification.

Supervised learning has two sub-types. Which one we have depends only on the **output column**:

- **Regression:** the output is numerical. *Example:* predicting a student's salary package (4.5 LPA, 3.0 LPA, ...) from IQ and CGPA.
- **Classification:** the output is categorical. *Example:* predicting placed / not placed from IQ and CGPA.




Figure 3: Regression vs classification

In Figure 3, regression draws a line that gives a number for any input. Classification draws a boundary that splits the inputs into categories.

Problem	Output	Type
Price of a house	A number	Regression
Is this email spam?	Yes / No	Classification
Will it rain today?	Yes / No	Classification
Is there a dog in this image?	Yes / No	Classification

3. Unsupervised learning

Key point: In unsupervised learning, the data has inputs only. There is nothing to predict, so the algorithm finds structure instead: groups, simpler columns, oddities or links.

3.1 Learning from inputs only

Key point: No output column means no prediction. The algorithm describes the data instead.

In **unsupervised learning**, the data has only input columns. For example, we have the IQ and CGPA of students, but no placement column.

Without an output, we cannot predict anything. Unsupervised learning performs four other jobs, covered in Sections 3.2 to 3.5.

3.2 Clustering

Key point: Clustering finds groups of similar rows, without being told what the groups are.

Clustering splits the data into groups (**clusters**) of similar rows. Figure 4 shows students plotted by IQ and CGPA; the algorithm found three groups on its own.




Figure 4: Clustering: groups found without labels

What the groups give us:

- **A category for each student:** for example, high IQ but low CGPA. A new student can be placed in one of the groups.

- **Labels for free:** we can number the groups (0, 1, 2) and then use them as an output column for supervised learning.
- **Customer segments:** an e-commerce site can group its customers by how they behave and treat each group differently.

With two columns we could spot the groups by eye. Clustering also works with hundreds of columns, where no human can see the groups.

3.3 Dimensionality reduction

Key point: Dimensionality reduction cuts down the number of input columns while keeping the information. It speeds up learning and lets us plot high-dimensional data.

Each input column is a **dimension**. Images and text can have thousands of input columns, which causes two problems:

1. Algorithms become **slow**.
2. After a point, extra columns **stop improving** the results.

Dimensionality reduction removes the extra columns.

Example: house prices. Number of rooms and number of washrooms carry related information. We can combine them into one column, area (Figure 5).




Figure 5: Combining two related columns into one

Creating a new column from existing ones like this is called **feature extraction**. The data has one column fewer and loses almost no information.

Visualisation. A graph can show at most 3 dimensions. To see data with hundreds of columns, we reduce it to 2 or 3 columns and plot that.

Example: handwritten digits. Each digit is an 8 x 8 grid of pixels, which makes 64 columns. Figure 6 reduces them to 3 columns: images of the same digit land close together.

1,797 handwritten digits: 64 columns reduced to 3




Figure 6: Handwritten digits: 64 columns reduced to 3

The technique used for Figure 6 is **PCA** (principal component analysis), covered in detail later in the course. The Notebook for this Note (`notebook.ipynb`) shows Figure 6 as a 3D plot that we can rotate.

Extra: The best-known digits dataset, MNIST, uses 28 x 28 pixel images, which gives 784 columns. The idea is the same.

3.4 Anomaly detection

Key point: Anomaly detection learns what normal data looks like and flags anything far from it.

Anomaly detection finds rows that do not fit the pattern of the rest. Typical uses:

- spotting defects in manufacturing,
- catching credit card fraud,
- flagging unusual loan applications.




Figure 7: Anomaly detection: one point far from the rest

In Figure 7, almost all transactions are small and close to home. The red one, large and 42 km away, is far from every normal point, so the system flags it.

3.5 Association rule learning

Key point: Association rule learning finds items that tend to occur together.

Association rule learning looks for “if this, then that” patterns in data. Its classic use is deciding which products a shop places next to each other.

Example: a supermarket. We scan all the bills from the last one or two years (Figure 8). Suppose milk appears in 8 of 100 bills, and eggs appear in 6 of those 8. Then people who buy milk usually buy eggs as well, so the shop places the two together.




Figure 8: Association rules: from past bills to shelf placement

A famous case: a large US retailer found that customers buying baby diapers often also bought beer. Placing the two side by side increased sales. Hidden patterns like this are hard to spot by hand and easy for an algorithm.

Extra: The diapers-and-beer story is usually told about Walmart. It most likely comes from a 1992 data analysis for the Osco Drug chain, and its details are disputed. It is still the standard example of association rules.

4. Semi-supervised learning

Key point: Semi-supervised learning labels a few rows by hand and lets the algorithm label the rest.

Labels are **expensive**. Collecting inputs is easy; for example, we can download thousands of images in minutes. But someone has to look at each image and write down what is in it, which costs time and money.

Semi-supervised learning works with data where only a small part is labelled. We label a few rows, and the algorithm labels the rest automatically.

Example: Google Photos (Figure 9).

1. The app groups photos by the faces in them, without knowing any names. This step is unsupervised.
2. We tag one photo with a name, such as “Dad”.
3. The name spreads to every photo in that group.




Figure 9: Semi-supervised learning: one label spreads to a whole group

One human label replaced hundreds.

5. Reinforcement learning

Key point: In reinforcement learning, there is no data. An agent learns by acting, getting rewards or punishments, and improving its rules.

5.1 Agent, environment, policy and reward

Key point: The agent acts, the environment rewards or punishes, and the agent updates its policy. Repeat.

Reinforcement learning (RL) starts with no data at all. The algorithm learns from its own experience, getting better over time.

The parts of an RL system:

- **Agent:** the learner, for example a self-driving car or a game-playing program.
- **Environment:** the world the agent acts in, for example the road or the game board.
- **Policy:** the agent's rule book, saying which action to take in each situation.
- **Reward / punishment:** feedback from the environment after each action, good or bad.

Figure 10 shows the loop. The agent observes the environment, takes an action according to its policy, receives a reward or punishment, and updates its policy. The goal is to collect as much reward and as little punishment as possible.




Figure 10: The reinforcement learning loop

5.2 Example: fire or water

Key point: A bad result changes the policy, so the next choice is better.

Figure 11 shows an agent that can walk to fire or to water.




Figure 11: An agent learning to choose water over fire

1. Its policy says *go to the fire*. It does, and gets a punishment.
2. It updates its policy to *go to the water*.
3. It goes to the water and gets a reward.

This is how people learn many things too: nobody hands us a dataset for living in a new city; we try, make mistakes and adjust. Training a pet with treats works the same way.

5.3 Where RL is used

Key point: RL beat the world champion at Go and is used for self-driving cars, but it is harder to build than the other types.

- **Self-driving cars:** the car learns to drive by acting on the road and receiving feedback.
- **Games:** Go is far more complex than chess, and beating top humans was thought to be years away. In 2016, DeepMind's agent **AlphaGo** beat world champion Lee Sedol in four out of five games.

RL is harder to set up than the other types, but its use is growing fast.

6. Summary

Type	Data	Goal	Example
Supervised: regression	Inputs + numerical output	Predict a number	Salary package from IQ and CGPA
Supervised: classification	Inputs + categorical output	Predict a category	Placed or not
Unsupervised: clustering	Inputs only	Find groups	Customer segments
Unsupervised: dimensionality reduction	Inputs only	Fewer columns	Rooms + washrooms → area
Unsupervised: anomaly detection	Inputs only	Flag unusual rows	Card fraud
Unsupervised: association rules	Inputs only	Find items that go together	Milk and eggs
Semi-supervised	A few labels, many unlabelled rows	Label the rest	Google Photos faces
Reinforcement	No data; rewards from an environment	Learn the best actions	AlphaGo

- To identify a supervised problem, look at the **output column**: number → regression, category → classification.
- No output column → unsupervised.
- Few labels → semi-supervised.
- No data, only feedback → reinforcement.

7. Key terms

Term	Meaning
Supervision	Correct answers that guide an algorithm while it learns
Supervised learning	Learning from data with inputs and outputs, to predict outputs
Input / output	The columns we know / the column we want to predict
Target, label	Other names for the output column
Labelled data	Data that includes the output column
Numerical data	Data made of numbers
Categorical data	Data made of categories
Regression	Supervised learning with a numerical output
Classification	Supervised learning with a categorical output
Unsupervised learning	Learning from inputs only, to find structure

Term	Meaning
Clustering	Splitting data into groups of similar rows
Cluster	One group found by clustering
Dimension	One input column
Dimensionality reduction	Reducing the number of input columns while keeping the information
Feature extraction	Creating a new column from existing ones
PCA	Principal component analysis, a dimensionality reduction technique
Anomaly detection	Finding rows that do not fit the pattern of the rest
Association rule learning	Finding items that tend to occur together
Semi-supervised learning	Learning from a few labelled rows and many unlabelled ones
Reinforcement learning	Learning by acting and receiving rewards or punishments
Agent	The learner in reinforcement learning
Environment	The world the agent acts in
Policy	The agent's rules for which action to take
Reward / punishment	Good / bad feedback after an action