


Outlier Detection with the IQR Method

Contents

1. Overview	2
2. The condition: a skewed feature	3
3. The fences	3
3.1 Quartiles and the IQR, by hand	3
3.2 From the IQR to the fences	4
4. Treating the outliers: trimming or capping	5
5. The placement data	6
6. Finding the fences and the outliers	7
7. Trimming in code	8
7.1 Why a new outlier appears after trimming	8
8. Capping in code	9
9. Learning the fences on the training set	11
10. Strengths and limits of the IQR method	11
11. Summary	13
12. Sources	13
13. Key terms	14

Where this fits: Pipeline map step 4 (Clean). See the [Course map](#).



- **Builds on:** Descriptive statistics ([Note 19](#)); Univariate analysis ([Note 20](#)); Outliers ([Note 41](#)); Percentile outlier method ([Note 41](#)).
- **Compare with:** Percentile outlier method ([Note 41](#)); Z-score outlier method ([Note 42](#)).

1. Overview

Key point: For a skewed feature, every value beyond the box-plot fences ($Q_1 - 1.5 \text{ IQR}$ and $Q_3 + 1.5 \text{ IQR}$) is an outlier; we then trim those rows or cap the values at the fences.

Note 41 listed three rules for detecting outliers, and Note 42 put the first one, the z-score method, to work on a normal feature. This Note covers the second rule: the **IQR method** (G-972), also called the **IQR proximity rule**. The IQR method is the rule to use when a **feature** (G-772) (an input variable, one column of the data table) is skewed. Each **observation** (G-1374) is one record (one row), here one student.

Figure 1 shows the whole method, in three steps:

1. check that the feature is skewed;
2. compute the two fences from the quartiles;
3. treat the values outside them by trimming or capping (both defined in Note 41, Section 7).




Figure 1: The IQR method: check the feature is skewed, compute the fences, then trim or cap

2. The condition: a skewed feature

Key point: The IQR method is meant for a feature that is skewed, not normal.

A feature is **skewed** (G-1817) when its values have a long tail on one side (Note 20). The z-score method of Note 42 does not fit such a feature, because it assumes a bell shape.

The IQR method makes no such assumption. The IQR method is built only on **percentiles** (G-1483), which do not care about the shape of the feature: the p -th percentile is the value below which p percent of the sorted data lies. Think of the middle half of a queue sorted by height: a single giant joining the end of the queue does not change who stands in the middle half. To use the IQR method we need two ideas from Note 20: the **box plot** (G-329), the picture of this method, and the IQR.

3. The fences

Key point: The lower fence is $Q_1 - 1.5 \times \text{IQR}$ and the upper fence is $Q_3 + 1.5 \times \text{IQR}$; values outside them are outliers.

3.1 Quartiles and the IQR, by hand

Key point: Sort the values; the median splits them into two halves; Q_1 and Q_3 are the middles of the halves, and the IQR is the distance between them.

Sort the values and cut the sorted list into four equal parts. The three cut points are the **quartiles** (G-1602): Q_1 (the 25th percentile), the **median** (G-1209, the 50th percentile) and Q_3 (the 75th percentile). They are taught in the [understanding your data Note](#) (section 7.2). The distance from Q_1 to Q_3 is the **interquartile range** (G-966), $\text{IQR} = Q_3 - Q_1$: the width of the middle half of the data ([univariate analysis Note](#), section 8).

Figure 2 finds them by hand for the exam marks of the first nine students of Section 5. Watch each step add one line:

1. **Sort:** 8, 11, 17, 23, 26, 38, 38, 39, 40.
2. **Median:** nine values, so the fifth one, 26, has four values on each side.
3. Q_1 : the middle of the lower half (8, 11, 17, 23). Four values have two in the middle, so we average them: $(11 + 17)/2 = 14$.
4. Q_3 : the middle of the upper half (38, 38, 39, 40): $(38 + 39)/2 = 38.5$.
5. **IQR:** $38.5 - 14 = 24.5$.

The last frame draws the box from Q_1 to Q_3 with the median inside it. Its two arms reach to the smallest and the largest mark, 8 and 40, because no mark lies beyond the fences of Section 3.2.




Figure 2: Quartiles by hand on the marks of nine students: sort, mark the median, take the middle of each half, and draw the box (idea after Khan Academy, “How to calculate interquartile range IQR”)

Extra: pandas computes quartiles a little differently. The steps above take the middle of each half. pandas’ `quantile` places the percentile between two neighbouring sorted values, in proportion to its position (**linear interpolation**, G-1092; Note 44). On these nine values pandas gives $Q_1 = 17$ and $Q_3 = 38$, not 14 and 38.5. On the 1,000 marks both ways give $Q_1 = 17$ and $Q_3 = 44$. Every number in the rest of this Note comes from pandas’ `quantile`.

3.2 From the IQR to the fences

Key point: Step 1.5 IQR out from each side of the box; a value beyond either point is an outlier.

The two box-plot **fences** (G-776) come from the [univariate analysis Note](#) (section 8.1). The IQR method uses exactly those fences as its lower and upper limits. Step by step:

1. **In words:** measure the width of the box, the IQR. Go one and a half box-widths below Q_1 for the lower fence, and one and a half above Q_3 for the upper fence.
2. **Formula:**

$$\text{lower} = Q_1 - 1.5 \times \text{IQR}, \quad \text{upper} = Q_3 + 1.5 \times \text{IQR}$$

3. **Example:** for the 1,000 placement exam marks of Section 5, $Q_1 = 17$ and $Q_3 = 44$, so $\text{IQR} = 44 - 17 = 27$ and $1.5 \times 27 = 40.5$:

$$\text{lower} = 17 - 40.5 = -23.5, \quad \text{upper} = 44 + 40.5 = 84.5$$

A mark below -23.5 or above 84.5 is an **outlier** (G-1420). Figure 3 builds the fences on the real marks; watch the orange arms grow 40.5 marks out of each side of the box, and only the far right tail turn red.




Figure 3: The IQR fences built step by step on the 1,000 exam marks: Q_1 and Q_3 , the IQR of 27, arms of 1.5 IQR , then the 15 outliers beyond 84.5

The plan for any skewed feature is therefore short:

1. compute Q_1 , Q_3 and the IQR;
2. compute the two fences;
3. trim or cap every value outside them.

Extra: The fences are robust: they hardly move when a few values are extreme. Q_1 and Q_3 depend only on the order of the middle values, so a few extreme values cannot drag them out, unlike the mean and standard deviation of Note 42. The factor 1.5 comes from John Tukey, who introduced the box plot (Tukey 1977). Some people also use 3: a point beyond $Q_3 + 3\text{IQR}$ (or below $Q_1 - 3\text{IQR}$) is called an “extreme” outlier, and one only beyond the 1.5 fences a “mild” one (NIST 7.1.6).

4. Treating the outliers: trimming or capping

Key point: Trimming deletes the outlier rows; capping replaces each outlier with the fence it crossed.

Both treatments, **trimming** (G-2019) and **capping** (G-345), come from the outliers Note (section seven, ways to treat outliers); here the fences are the limits.

5. The placement data

Key point: The data has 1,000 students; their placement exam marks are right-skewed (skewness 0.84), so the marks feature is the one for the IQR method.

The data is the placement data of Note 42: one observation per student of a college, 1,000 students, three columns.

cgpa	placement_exam_marks	placed
7.19	26	1
7.46	38	1
7.54	40	1
6.42	8	1

- **cgpa**: the student's CGPA (out of 10).
- **placement_exam_marks**: marks out of 100 in the aptitude test held before placement.
- **placed**: 1 if the student got a job offer, 0 if not; the **target** (G-1949), the output a model would predict.

Note 42 (Figure 5) showed the shapes: **cgpa** is a bell (skewness -0.01), while **placement_exam_marks** has a long tail to the right (skewness 0.84). So **placement_exam_marks** is the candidate for the IQR method.

The summary numbers of the marks feature tell the same story:

Count	Mean	Std	Min	25%	50%	75%	Max
1,000	32.23	19.13	0	17	28	44	100

A quarter of the students scored below 17, half below 28 and three quarters below 44. One student scored 0 and one scored the full 100.

Python: Loading the data and checking the marks feature.

```
import pandas as pd

df = pd.read_csv("data/placement.csv")
df.shape      # (1000, 3)

df["placement_exam_marks"].skew()      # 0.84
df["placement_exam_marks"].describe()  # the table above
```

6. Finding the fences and the outliers

Key point: $Q_1 = 17$ and $Q_3 = 44$ give an IQR of 27 and fences of -23.5 and 84.5 ; 15 students lie above the upper fence and none below the lower one.

Figure 4 shows the box plot of the marks above their histogram. The box runs from $Q_1 = 17$ to $Q_3 = 44$, so the IQR is 27, and the fences sit at -23.5 and 84.5 (Section 3).




Figure 4: The placement exam marks: box plot with Q_1 , median and Q_3 , the two fences, and the 15 outliers in red

- **Upper fence, 84.5:** 15 students scored above it, with marks from 86 to 100. In the box plot they are the red dots on the right.
- **Lower fence, -23.5 :** marks cannot be negative, so no student lies below it. With a right-skewed feature, all the outliers sit on the long-tail side.

The 15 outliers:

Mark	Students	Rows
86	3	40, 61, 771
87	4	283, 290, 311, 685
90	3	162, 324, 730
93, 94, 96, 97, 100	1 each	134, 9, 630, 846, 917

Python: The quartiles, the IQR and the fences.

```
marks = df["placement_exam_marks"]

percentile25 = marks.quantile(0.25)    # 17.0
```

```

percentile75 = marks.quantile(0.75)    # 44.0
iqr = percentile75 - percentile25      # 27.0

upper_limit = percentile75 + 1.5 * iqr  # 84.5
lower_limit = percentile25 - 1.5 * iqr  # -23.5

df[marks > upper_limit]    # 15 rows
df[marks < lower_limit]    # 0 rows
quantile(0.25) (G-127) returns the 25th percentile; quantile takes the fraction (0.25),
not the percent (25).

```

7. Trimming in code

Key point: Keeping only the rows inside the fences removes the 15 outliers and leaves 985 rows.

Trimming is a filter: keep the rows whose mark lies between the two fences. Since no mark is below the lower fence, keeping the marks below 84.5 is enough here.

Python: Trimming.

```

# keep the rows with lower <= mark <= upper
new_df = df[marks.between(lower_limit, upper_limit)]
new_df.shape    # (985, 3)

```

Figure 6 (middle row) shows the result. The histogram barely changes: only its thin right tail is gone. The box plot loses its row of red dots.

The mean drops from 32.23 to 31.34 and the skewness from 0.84 to 0.65.

7.1 Why a new outlier appears after trimming

Key point: The box plot of the trimmed data computes new fences from the 985 remaining rows, so a value that was inside before can now be outside.

The trimmed box plot in Figure 6 still shows one red dot, at a mark of 83. The dot is not a mistake. A box plot always computes its fences from the data it is given.

For the 985 trimmed rows, Q_3 drops from 44 to 43, so the new upper fence is $43 + 1.5 \times (43 - 17) = 82$. The mark 83, safely inside the old fence of 84.5, is now just outside the new one.

In this Note we detect once, with the fences of the original data, and treat once. (Trimming again here would remove the 83 and then stop: a third round finds nothing. The last cell of the Notebook shows this.)




Figure 5: Detecting again after each trim: the fence moves from 84.5 to 82, catches the 83, and then stays at 82 with nothing outside

In Figure 5, watch the dashed fence step left between rounds 1 and 2 and then stop moving.

8. Capping in code

Key point: Each mark above 84.5 becomes 84.5 (and each mark below -23.5 would become -23.5); all 1,000 rows stay.

Capping goes through the column value by value, exactly as in Note 42:

- above the upper fence: replace it with the upper fence;
- below the lower fence: replace it with the lower fence;
- otherwise: leave it as it is.

NumPy's **np.where** (G-124), called as `np.where(condition, value if true, value if false)`, does this for the whole column; two conditions need two calls, one inside the other.

Python: Capping with `np.where`.

```
import numpy as np
```

```
new_df_cap = df.copy()
col = new_df_cap["placement_exam_marks"]
new_df_cap["placement_exam_marks"] = np.where(
    col > upper_limit,
    upper_limit,
    np.where(col < lower_limit, lower_limit, col))
```

```
new_df_cap.shape    # (1000, 3): no row lost
df.copy() makes an independent copy, so the original df keeps its values for the
comparison below.
```

The result, compared with the original and the trimmed column:

	Original	Trimmed	Capped
Rows	1,000	985	1,000
Mean	32.23	31.34	32.14
Standard deviation	19.13	17.86	18.87
Maximum	100	83	84.5
Skewness	0.84	0.65	0.76

Figure 6 (bottom row) shows the capped column. The 15 outliers now all sit at 84.5, so the histogram rises in that one spot (the orange bar). The box plot has no dots: its whisker ends exactly at the fence.




Figure 6: Before and after: the original column, the trimmed column (985 rows) and the capped column (1,000 rows, outliers moved onto 84.5); histograms on the left, box plots on the right, the upper fence dashed

Extra: pandas has a one-line shortcut for capping, `clip` (G-70). `clip` gives exactly the same column as the two nested `np.where` calls.

```
marks.clip(lower_limit, upper_limit)
```

9. Learning the fences on the training set

Key point: The fences are learned from the data, so they should be learned on the training set only and then applied to the test set.

The fences follow the same train-only rule as the z-score limits of the [z-score Note](#) (section 11, learning the limits on the training set).

Extra: With an 80/20 split (`random_state=42`), the 800 training rows have the same $Q_1 = 17$ and $Q_3 = 44$, so the fences stay at -23.5 and 84.5 :

	Rows	Outliers (training fences)
Training set	800	14: marks 86 to 97
Test set	200	1: mark 100

Quartiles move very little between 800 and 1,000 rows, so here the result is the same. Trimming the training set leaves 786 rows; the Notebook shows both steps.

10. Strengths and limits of the IQR method

Key point: The method is simple, works on skewed features and is not pulled by the outliers it looks for; it is only a rule of thumb.

- **Simple:** two percentiles give both fences.
- **Shape-free:** it needs no bell shape, so it fits skewed features.
- **Robust:** the quartiles are not pulled by extreme values (Section 3.2, Extra).

To see the robustness, we replace the top mark, 100, by a typo that grows up to 1,000 and recompute both upper limits (Figure 7). Watch the two lines: the z-score limit of Note 42 chases the typo from 89.6 up to 141.2, while the IQR fence never leaves 84.5.




Figure 7: One typed-in extreme value: the z-score upper limit (mean + 3 std) climbs with it, the IQR upper fence stays at 84.5

Extra: Two things to keep in mind.

- **On a long tail, real values get flagged.** In a strongly skewed feature the far tail can be perfectly genuine (a few very high incomes, a few toppers; here, 15 real exam marks between 86 and 100). The fences flag it all the same. So we still decide, as in Note 41 (Section 4), whether those values are errors or real.
- **One side may never be used.** For a right-skewed feature with a natural lower bound, such as marks starting at 0, the lower fence can lie below every possible value, as it does here (-23.5). An unused fence is expected, not a bug.

11. Summary

Step	What we do	On the placement marks
Check the shape	plot the feature; the method is for skewed features	right-skewed, skewness 0.84
Quartiles	$Q_1 = 25\text{th}$, $Q_3 = 75\text{th}$ percentile, $\text{IQR} = Q_3 - Q_1$	$Q_1 = 17$, $Q_3 = 44$, $\text{IQR} = 27$
Detect	fences $Q_1 - 1.5 \text{ IQR}$ and $Q_3 + 1.5 \text{ IQR}$	-23.5 and 84.5 ; 15 outliers, all above
Trim	keep the rows inside the fences	985 rows left
Cap	move values beyond a fence onto it	1,000 rows; maximum 84.5
Better practice	learn the fences on the training set	same fences; 14 outliers in train, 1 in test

	Z-score method (Note 42)	IQR method (this Note)
Feature shape	roughly normal	skewed
Built on	mean and standard deviation	Q_1 and Q_3
Limits	$\mu \pm 3\sigma$	$Q_1 - 1.5 \text{ IQR}$, $Q_3 + 1.5 \text{ IQR}$
Pulled by outliers	yes	hardly

- The IQR method is for skewed features; the z-score method of Note 42 is for normal ones.
- The IQR is $Q_3 - Q_1$, the width of a box plot's box.
- The fences are $Q_1 - 1.5 \text{ IQR}$ and $Q_3 + 1.5 \text{ IQR}$: the same limits a box plot uses for its dots.
- Values outside the fences are outliers; we trim them (drop the rows) or cap them (set them to the fence).
- A box plot of trimmed data computes new fences, so a new dot can appear; we do not trim again.
- The fences should be learned on the training set only.

12. Sources

Built from

- CampusX, “Outlier Detection and Removal using the IQR Method | Handling Outliers Part 3”, YouTube, <https://www.youtube.com/watch?v=Ccv1-W5ilak>
- Khan Academy, “How to calculate interquartile range IQR”, YouTube, <https://www.youtube.com/watch?v=Ccv1-W5ilak>

Other references

- Tukey, J. W. (1977). *Exploratory Data Analysis*. Addison-Wesley.

- NIST/SEMATECH. *e-Handbook of Statistical Methods*, Section 7.1.6, “What are outliers in the data?”. <https://www.itl.nist.gov/div898/handbook/>

13. Key terms

Term	Meaning
Feature	An input variable: one column of the data table
Observation	One record: one row of the data table
Target	The output a model predicts
IQR method (IQR proximity rule)	Outlier detection that flags values beyond 1.5 IQR outside the box (Q_1 to Q_3); for skewed features
Percentile	The value below which a given share of the sorted data lies
Quartiles	Q_1 , the median and Q_3 : the 25th, 50th and 75th percentiles
Interquartile range (IQR)	$Q_3 - Q_1$: the width of the middle half of the data
Fences	$Q_1 - 1.5 \text{ IQR}$ and $Q_3 + 1.5 \text{ IQR}$; values beyond them are outliers
Trimming	Removing the rows that hold outliers
Capping	Replacing every value beyond a limit with the limit itself
Linear interpolation	Placing a percentile between two neighbouring sorted values, in proportion to its position
<code>quantile</code>	pandas method that returns a percentile, given as a fraction (0.25 for the 25th)