


Gradient Descent from Scratch

Contents

1. Overview	2
2. The idea	2
2.1 Which way to move	3
2.2 How far to move	4
2.3 When to stop	4
3. The slope of the loss with respect to b	4
4. A worked example: four points, b only	5
5. The learning rate	7
6. Both parameters: m and b together	8
7. Gradient descent as a class	9
8. The shape of the loss function matters	11
9. The effect of feature scaling	12
10. Summary	14
11. Sources	14
12. Key terms	14

Where this fits: Pipeline map steps 0 (Foundations), 8 (Model), 10 (Tune). See the [Course map](#).



- **Builds on:** Feature scaling ([Note 24](#)); Best-fit line and squared error ([Note 51](#)); Derivatives of one variable ([Note 51](#)).
- **Leads to:** Batch gradient descent ([Note 58](#)); Stochastic gradient descent ([Note 59](#)); Mini-batch gradient descent ([Note 60](#)); Ridge regression ([Note 63](#)); Logistic regression ([Note 70](#)); Gradient boosting ([Note 120](#)).
- **Compare with:** Ordinary least squares (closed form) ([Note 51](#)); Normal equation ([Note 55](#)).

1. Overview

Key point: Gradient descent finds the minimum of a function by starting anywhere and repeatedly taking a small step downhill, against the slope.

Gradient descent (G-862) is a first-order iterative optimisation algorithm for finding a local minimum of a differentiable function. In plainer words:

- **optimisation algorithm** (G-1398): it finds the values of parameters that make some function as small as possible;
- **iterative:** it does this in many small steps instead of with one formula;
- **first-order:** each step uses only the first derivative (the slope);
- **differentiable function:** the function must have a slope everywhere.

In ML, the function is the **loss function** (G-1130; for linear regression, the sum of squared errors) and the parameters are the model's coefficients. Nearly all of deep learning is trained with a version of gradient descent (Goodfellow Ch. 8).

Throughout, a **feature** (G-772) is an input variable (one column of the data table), an **observation** (G-1374) is one record (one row, here one data point), and the **target** (G-1949) is the output we predict. We learn gradient descent on linear regression because we can check the answer against OLS, whose normal equation becomes slow with many features (see [section seven of the multiple linear regression maths Note](#)); most models have no direct formula at all.

2. The idea

Key point: Standing on a curve with our eyes closed, the slope under our feet tells us which way is down. Step that way, and repeat.

2.1 Which way to move

Key point: If the slope is positive, the minimum is to the left, so decrease the parameter. If the slope is negative, increase it.

Start with a single parameter, the intercept b , and keep the slope of the line fixed for now. The loss $L(b)$, as b changes, is a U-shaped curve with its minimum at the best b .

Imagine standing somewhere on that curve, blindfolded. We cannot see where the bottom is. But we can feel the slope under our feet:

- **Slope positive** (the curve rises to the right): the bottom is to the left. Decrease b .
- **Slope negative** (the curve falls to the right): the bottom is to the right. Increase b .

In both cases, moving **against the sign of the slope** goes downhill. Moving against the slope is the whole idea.

Figure 1 shows the rule at four places on the loss curve of the example in Section 4. The two orange points sit where the curve rises to the right (slope positive), and their arrows point left. The two red points sit where it falls (slope negative), and their arrows point right. Every arrow points towards the bottom.




Figure 1: The loss $L(b)$ of the 4-point example. At four values of b , the black tangent shows the slope and the arrow shows the gradient descent step, minus 0.1 times the slope. Positive slopes (orange) step left, negative slopes (red) step right, and steeper slopes give longer steps.

2.2 How far to move

Key point: The step is the slope times a small number, the learning rate: $b_{\text{new}} = b_{\text{old}} - \eta \cdot \text{slope}$.

In words: subtract from the current value the slope multiplied by a small constant η (eta), the **learning rate** (G-1068).

$$b_{\text{new}} = b_{\text{old}} - \eta \frac{\partial L}{\partial b}$$

Two things happen automatically:

- The minus sign handles the direction (Section 2.1).
- Far from the bottom the slope is steep, so the steps are big; near the bottom the slope flattens, so the steps shrink. In Figure 1, the slope at $b = 100$ is 591 and the step is 59.1; at $b = 60$, closer to the bottom, the slope is 271 and the step only 27.1.

The learning rate keeps the steps from being too large. Without it, a slope of 590 would throw b 590 units away in one jump.

2.3 When to stop

Key point: Stop after a fixed number of steps (epochs), or when the steps become tiny.

Two common stopping rules:

1. **A fixed number of iterations**, called **epochs** (G-696): for example, 100 updates.
2. **The step size becomes negligible**: for example, when $b_{\text{new}} - b_{\text{old}}$ is below 0.0001, the algorithm has reached (almost) the bottom.

3. The slope of the loss with respect to b

Key point: For the sum of squared errors, $\frac{\partial L}{\partial b} = -2 \sum (y_i - mx_i - b)$.

Gradient descent needs the slope of the loss at the current b . For linear regression the loss is

$$L(m, b) = \sum_{i=1}^n (y_i - mx_i - b)^2$$

In words: differentiate each squared term with respect to b . The **chain rule** (G-371) brings down a factor 2, and the inner derivative of $-b$ is -1 .

$$\frac{\partial L}{\partial b} = -2 \sum_{i=1}^n (y_i - mx_i - b)$$

The OLS derivation used the same derivative. There we set it to zero and solved; here we only evaluate it at the current point and take a step.

The sum inside is just the residuals added up. Figure 2 evaluates it on the 4-point example of the next section, at $m = 78.35$ and $b = 100$. The line sits above all four points, so all four residuals are negative: -37.0 , -124.4 , -119.4 and -14.7 . They add up to -295.4 , and $-2 \times (-295.4) = 590.7$: a positive slope, so the next step lowers b .




Figure 2: The slope with respect to b on the 4-point example at $b = 100$. Left: the line and each point's residual. Right: the four residuals, their sum -295.4 , and -2 times the sum, 590.7 .

4. A worked example: four points, b only

Key point: Starting from $b = 100$ with learning rate 0.1 , three steps give 40.93 , 29.11 and 26.75 , closing in on the OLS value 26.16 .

The example data has 4 points. OLS gives slope $m = 78.35$ and intercept $b = 26.16$. To see gradient descent alone, we fix $m = 78.35$ and start b deliberately far away, at $b = 100$, with learning rate $\eta = 0.1$.




Figure 3: Gradient descent on b , step by step

Figure 3 shows the loss curve $L(b)$ and the steps.

1. **Step 1:** at $b = 100$ the slope is 590.7. The step is $0.1 \times 590.7 = 59.07$, so $b = 100 - 59.07 = 40.93$.
2. **Step 2:** at $b = 40.93$ the slope is 118.1. The step is 11.81, so $b = 29.11$.
3. **Step 3:** at $b = 29.11$ the slope is 23.6. The step is 2.36, so $b = 26.75$.
4. **Steps 4 and 5:** $b = 26.28$, then 26.18, almost exactly the OLS value 26.16.

Step	b before	Slope	Step = $0.1 \times \text{slope}$	b after
1	100.00	590.7	59.07	40.93
2	40.93	118.1	11.81	29.11
3	29.11	23.6	2.36	26.75
4	26.75	4.7	0.47	26.28
5	26.28	0.9	0.09	26.18

Figure 4 shows the same five steps from both sides at once. On the left, the line slides down towards the four points and its residuals (red) shrink. On the right, the point slides down the loss curve; the black tangent is the slope at the current b , and the green arrow is the step. Watch the arrow: 59.07, then 11.81, then 2.36. As the tangent flattens, the arrow shrinks.




Figure 4: Gradient descent on b , seen on the data (left) and on the loss curve (right). Each frame prints b , the slope and the step. Idea after StatQuest, “Gradient Descent, Step-by-Step”.

Each slope is about one fifth of the previous one, so each step shrinks too. Nothing in the rule tells the algorithm to slow down: the flattening curve does it. Here the distance to the bottom shrinks to exactly one fifth at every step (the Extra below shows why).

Extra: Why exactly one fifth? With m fixed, call b^* the best intercept, where the slope is zero: $\sum(y_i - mx_i - b^*) = 0$. Then

$$\frac{\partial L}{\partial b} = -2 \sum_{i=1}^n (y_i - mx_i - b) = 2n(b - b^*)$$

so the update becomes

$$b_{\text{new}} - b^* = (1 - 2n\eta)(b_{\text{old}} - b^*)$$

With $n = 4$ points and $\eta = 0.1$, the distance to b^* is multiplied by $1 - 0.8 = 0.2$ at every step: 73.84, then 14.77, then 2.95. With $\eta = 0.26$ (next section) the factor is $1 - 2.08 = -1.08$: the minus sign flips b to the other side, and the distance grows by 8% each step. Any η above 0.25 diverges on this data.

5. The learning rate

Key point: Too small and gradient descent crawls; too large and it overshoots and can run away. The right value is found by trying.

The learning rate is a **hyperparameter** (G-910): a setting we choose, not a value the algorithm learns. Figure 5 starts from the same $b = 100$ with three learning rates.




Figure 5: Three learning rates from the same start

- **0.01, too small:** after 10 steps b has only moved from 100 to 58. b would reach the bottom eventually, but slowly.
- **0.1, good:** the bottom is reached within about 5 steps.
- **0.26, too large:** each step jumps over the bottom to the other side, landing higher than before: $100 \rightarrow -53.6 \rightarrow 112.3 \rightarrow -66.9 \rightarrow \dots$. The loss grows with every step, and the algorithm never **converges** (G-469): it **diverges** (G-628). (The Extra in Section 4 shows why: above 0.25, every step overshoots by more than it gains.)

A common approach is to try a few values on a logarithmic scale, such as 0.1, 0.01 and 0.001 (Goodfellow §11.4.3), and keep the one where the loss falls fastest without jumping around.

6. Both parameters: m and b together

Key point: With two parameters, compute both slopes and update both at the same step. The point walks downhill across the loss bowl.

In practice both m and b are unknown. The loss $L(m, b)$ is then the bowl from the OLS Notes, and gradient descent needs one slope per parameter.

The slope with respect to b is as before. With respect to m , the inner derivative of $-mx_i$ is $-x_i$:

$$\frac{\partial L}{\partial m} = -2 \sum_{i=1}^n (y_i - mx_i - b) x_i$$

Each epoch updates both at once, using the current values of both:

$$m_{\text{new}} = m_{\text{old}} - \eta \frac{\partial L}{\partial m} \quad b_{\text{new}} = b_{\text{old}} - \eta \frac{\partial L}{\partial b}$$

The vector of these **partial derivatives** (G-1457) is called the **gradient** (G-865); it points uphill, in the direction where the loss rises fastest, so stepping against it goes downhill.

Stepping down against the gradient gives the method its name.

The gradient can be read one number at a time. At the starting point of Figure 6 it is $(-25,362, 28,641)$:

- **The sign** of each number gives the direction for that parameter: negative for m , so m must increase; positive for b , so b must decrease.
- **The size** says how much the loss reacts to that parameter right now. A number three times larger would mean a small change to that parameter moves the loss three times as much. Here the two sizes are similar, so the first step moves both by a similar amount: m by $+25.4$ and b by -28.6 .

Figure 6 runs this on 100 points, starting far away at $m = -127.8$, $b = 150$, with learning rate 0.001 for 30 epochs.




Figure 6: Gradient descent on m and b : the path and the loss per epoch

The left panel is a **contour plot** (G-468): a map of the bowl seen from above, with each ring joining points of equal loss and the darkest region at the bottom. The path heads straight down the bowl and arrives at $m = 27.2$, $b = -1.9$. The right panel shows the loss falling from about 4.2 million to 28,400, quickly at first and then levelling off.

7. Gradient descent as a class

Key point: A small class with `fit` (the update loop) and `predict` reproduces Linear-Regression's answer to within a few hundredths.

Python: Gradient descent for simple linear regression.

```
class GDRegressor:
    def __init__(self, learning_rate, epochs):
        self.m, self.b = 100.0, -120.0      # any start
```

```

self.lr, self.epochs = learning_rate, epochs

def fit(self, X, y):
    x = X.ravel()
    for _ in range(self.epochs):
        slope_b = -2 * np.sum(y - self.m * x - self.b)
        slope_m = -2 * np.sum((y - self.m * x - self.b) * x)
        self.b -= self.lr * slope_b
        self.m -= self.lr * slope_m
    return self

def predict(self, X):
    return self.m * X.ravel() + self.b

```

On 80 training points with learning rate 0.001 and 50 epochs:

	Slope m	Intercept b	Test R^2
OLS (LinearRegression)	28.126	-2.271	0.6345
Gradient descent	28.159	-2.300	0.6344

Gradient descent gets within 0.03 of the exact answer: an approximation, but a very close one. Figure 7 shows the run. The line starts steep ($m = 100$, $b = -120$) and swings onto the OLS line within about 20 epochs; the loss falls from about 1.5 million to about 24,000 and then levels off.




Figure 7: The GDRegressor class at work on the 80 training points: the line after each epoch (orange) closing in on the OLS line (dashed), and the loss per epoch on a log scale.

Extra: Both slopes are computed from the old values before either is updated. Updating b first and then using the new b for the slope of m is a slightly different algorithm. On this data it gives almost the same answer ($m = 28.157$, notebook).

8. The shape of the loss function matters

Key point: Linear regression's loss is convex: one bowl with no false dips, and gradient descent with a small enough learning rate finds its bottom. Other losses can trap it in a local minimum or slow it on a flat stretch.

Gradient descent only follows the local slope, so the shape of the loss function decides how well it works (Figure 8).




Figure 8: Convex and non-convex loss functions

- **Convex:** a function is **convex** (G-476) when a straight line between any two points on its curve never goes below the curve. Such a function has no false dips: any local minimum is also the global one (Boyd and Vandenberghe §4.2.2). The sum of squared errors of linear regression is convex, so gradient descent with a small enough learning rate reaches the best answer (Section 5 shows what a too-large one does).
- **Non-convex, with a local minimum:** a straight line between two points can cut through the curve. Then there can be several dips: a **local minimum** (G-1110; lowest only in its neighbourhood) and a **global minimum** (G-848; lowest of all). Started near the local one, gradient descent settles there and stops, unaware that a better answer exists.
- **A plateau** (G-1503): a nearly flat stretch has a tiny slope, so the steps become tiny and progress almost stops. With too few epochs, the algorithm halts before leaving the plateau. In more dimensions, a related trouble spot is a **saddle point** (G-1718), flat in one direction and curved in another.

Figure 9 runs gradient descent on the middle curve of Figure 8 from two starting points, with the same learning rate. Picture each run as a ball rolling downhill. Start A rolls into the global minimum (loss 0.23). Start B rolls into the local minimum (loss 1.44) and stays there:

the slope is zero, so the steps are zero. Where gradient descent ends depends on where it starts.




Figure 9: Two runs of gradient descent on the same non-convex curve (learning rate 0.07). Start A reaches the global minimum; start B stops in the local minimum. The ball-rolling picture follows 3Blue1Brown, “Gradient descent, how neural networks learn”.

Extra: Neural networks have non-convex losses (Goodfellow §8.2). One remedy is to start from several different points and keep the best result. Others, such as **momentum** (Goodfellow §8.3.2) and adaptive learning rates (Goodfellow §8.5), come with neural networks.

9. The effect of feature scaling

Key point: When features are on very different scales, the loss bowl becomes a long, narrow valley and gradient descent zigzags slowly. Scaling the features makes the bowl round and the descent fast.

With two features, the shape of the loss depends on their scales. Picture walking down a steep, narrow gorge: every step toward the bottom of the gorge throws you against the other wall, and you make little progress along it. Figure 10 compares the same made-up data, with feature 2 eight times larger in scale than feature 1, before and after standardising.




Figure 10: The effect of scaling the inputs on gradient descent

- **Unscaled:** the contours are long, narrow ellipses. The learning rate must be small enough to avoid overshooting in the steep direction, and that is usually too small to make real progress in the flat direction (Goodfellow §4.3.1). After 40 steps the weights are still 2.8 away from the best values.
- **Standardised:** the contours are nearly circles, the slope points straight at the minimum, and 40 steps reach it.

Features should therefore be scaled (**standardisation** (G-1874) or normalisation, from the feature scaling Notes) before using gradient descent.

10. Summary

Idea	In one line
Update rule	parameter $\leftarrow$ parameter $-\eta \times$ slope
Slope for b	$-2 \sum (y_i - mx_i - b)$
Slope for m	$-2 \sum (y_i - mx_i - b) x_i$
Learning rate η	too small: slow; too large: overshoots and diverges
Stopping	fixed number of epochs, or tiny steps
Convex loss	every local minimum is the global one (linear regression)
Non-convex loss	local minima, plateaus, saddle points can trap it
Feature scaling	round contours, faster convergence

- Gradient descent works for any differentiable loss, which is why it powers most of ML.
- On linear regression it reproduces OLS closely: slope 28.16 against 28.13 here.
- The steps shrink by themselves near the minimum, because the slope shrinks.

11. Sources

Built from

- CampusX, “Gradient Descent From Scratch | End to End Gradient Descent | Gradient Descent Animation”, YouTube, <https://www.youtube.com/watch?v=ORyFPJypKuU>
- StatQuest with Josh Starmer, “Gradient Descent, Step-by-Step”, YouTube, <https://www.youtube.com/watch?v=sDv4f4s2SB8> (the data and the loss curve side by side, with each step’s slope and step size, Figure 4)
- Sanderson, G. (3Blue1Brown), “Gradient descent, how neural networks learn | Deep Learning Chapter 2”, YouTube, <https://www.youtube.com/watch?v=IHZwWFHWa-w> (the ball rolling into different valleys, Figure 9; reading the sizes of the gradient’s components, section 6)

Other references

- **Goodfellow:** I. Goodfellow, Y. Bengio, A. Courville, *Deep Learning*, MIT Press, 2016 (deeplearningbook.org). §4.3.1 (poor conditioning), Ch. 8 (training deep models; §8.2 non-convexity, §8.3.2 momentum, §8.5 adaptive learning rates), §11.4.3 (grid search on a log scale).
- **Boyd and Vandenberghe:** S. Boyd and L. Vandenberghe, *Convex Optimization*, Cambridge University Press, 2004. §4.2.2 (local optima of convex problems are global).

12. Key terms

Term	Meaning
Feature	An input variable: one column of the data table
Observation	One record: one row of the data table
Target	The output we predict
Gradient descent	An iterative algorithm that finds a minimum by repeatedly stepping against the slope
Optimisation algorithm	A method for finding the parameter values that make a function as small (or large) as possible
Learning rate (η)	The number the slope is multiplied by to get the step size
Epoch	One full update of the parameters using the whole training set
Gradient	The vector of partial derivatives of the loss; it points in the direction of steepest increase
Converge	To settle at a minimum, with steps becoming negligible
Diverge	To move further away with each step, the loss growing instead of shrinking
Contour plot	A map of a surface seen from above, with lines joining points of equal height
Convex function	A function where a straight line between any two points of its curve never goes below the curve; every local minimum is the global one
Local minimum	A point lower than everything around it, but not the lowest overall
Global minimum	The lowest point of the whole function
Plateau	A nearly flat region of the loss, where steps become very small
Saddle point	A flat point that curves up in one direction and down in another