

Ridge Regression: Five Key Points

Contents

1. Overview	2
2. Point 1: coefficients shrink but never reach 0	2
3. Point 2: the largest coefficients shrink most	3
4. Point 3: bias rises, variance falls	5
5. Point 4: the loss curve rises and its lowest point moves to 0	6
5.1 With two coefficients: why “Ridge”?	7
6. Point 5: use Ridge with 2 or more inputs	8
7. Summary	9
8. Key terms	9

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Standardization ([Note 24](#)); Multiple linear regression ([Note 55](#)); Normal equation ([Note 55](#)); Gradient descent ([Note 57](#)); Bias-variance trade-off ([Note 62](#)); Regularisation ([Note 63](#)).
- **Compare with:** Lasso regression (Video 67, coming).

1. Overview

Key point: Five facts sum up Ridge. Coefficients shrink but never reach 0. The largest shrink most. Bias rises and variance falls with λ . The loss curve's lowest point slides towards 0. Use Ridge when there are 2 or more inputs.

The last three Notes built Ridge regression: the idea, the formulas, and gradient descent. This Note collects the five points that are easiest to mix up afterwards. They are also common interview questions.

#	Question	Short answer
1	What happens to the coefficients as λ grows?	they shrink towards 0, but never reach it
2	Which coefficients are affected most?	the largest ones
3	What happens to bias and variance?	bias rises, variance falls
4	What happens to the loss function?	its lowest point moves towards 0
5	When should Ridge be used?	with 2 or more input columns

All diabetes examples below use test size 0.2 and random state 2.

2. Point 1: coefficients shrink but never reach 0

Key point: As λ grows from 0 towards infinity, every coefficient gets closer to 0. None becomes exactly 0.

λ can be any number from 0 upwards. At $\lambda = 0$ there is no penalty, so Ridge is plain linear regression. Figure 1 trains Ridge on the diabetes data for four values of alpha (scikit-learn's name for λ).




Figure 1: Diabetes coefficients for alpha 0, 10, 100 and 1000; each panel has its own scale

alpha	Largest coefficient (size)	Test R^2
0	895.6	0.44
10	63.7	0.15
100	7.5	0.01
1000	0.8	-0.01
10,000	0.08	-0.01

From alpha 10 on, each tenfold increase makes the coefficients about ten times smaller. But even at alpha 10,000, every coefficient is still a small non-zero number.

The reason is the slope formula from the Ridge maths Note: $m = \frac{\text{top}}{\text{bottom} + \lambda}$. Adding λ to the bottom makes the fraction smaller, but a fraction with a non-zero top never becomes 0.

Extra: This is the key difference from Lasso (the next Note), which can set coefficients exactly to 0 and so removes inputs from the model.

3. Point 2: the largest coefficients shrink most

Key point: A large coefficient falls quickly as λ grows; a small one hardly changes.

Figure 2 follows each diabetes coefficient as alpha goes from 0 to 2.




Figure 2: Coefficient paths for alpha from 0 to 2

Input	alpha 0	alpha 0.01	alpha 0.1	alpha 1
s1	-895.6	-383.7	-72.9	14.4
s5	861.1	659.9	484.4	260.1
s2	561.2	152.7	-80.6	-22.6
age	-9.2	-6.4	6.6	42.2

- **s1, s5 and s2**, the largest, drop steeply even for alpha 0.01.
- **age**, the smallest, barely moves at first.

The penalty is $\lambda \sum \beta_j^2$, and the square grows fast: a coefficient of 900 adds $810,000\lambda$ to the loss, while one of 9 adds only 81λ . So reducing the big coefficients lowers the loss far more, and that is where Ridge cuts first.

Extra: A small coefficient can even grow for a while, as age does here (from -9 to 42 at alpha 1). The inputs s1, s2 and s5 are strongly correlated, so as their huge opposing coefficients are reduced, part of their effect moves to other inputs. Once alpha is large enough, all coefficients head to 0 together.

4. Point 3: bias rises, variance falls

Key point: A small λ gives low bias and high variance (overfitting); a large λ gives high bias and low variance (underfitting). The best λ is in between.

Recall the bias-variance Note: **bias** is error from a model too simple to follow the pattern, and **variance** is how much the model changes from one training sample to another.

Figure 3 fits a degree-15 polynomial with Ridge to 80 points from a curve, $y = 0.7x^2 - 2x + 3$ plus noise. For each alpha, it trains the model on 200 resampled versions of the training data and measures bias and variance on 20 test points.




Figure 3: Bias², variance and expected test error against alpha

alpha	Bias ² (+ noise)	Variance	Expected test error
0.01	1.90	69.39	71.29
1	0.83	3.10	3.93
10	0.97	0.19	1.16
100	2.49	0.18	2.67
1000	4.60	0.30	4.90

- **Small alpha:** the flexible curve bends with each sample, so the variance is huge: overfitting.
- **Large alpha:** the coefficients are squeezed so hard that the curve is too stiff, and bias takes over: underfitting.
- **In between** (here about alpha 10): both are low, and the total error is smallest.

So, as in the bias-variance Note, λ moves a model along the trade-off. Choose λ where variance has fallen a lot but bias has not yet risen much.

Extra: The curve here is standardised after `PolynomialFeatures`, as the Ridge intuition Note advises. On the left of Figure 3, bias^2 first falls slightly as α grows. The wild curves at tiny α also move the average curve away from the truth, the same edge effect noted in the bias-variance Note.

5. Point 4: the loss curve rises and its lowest point moves to 0

Key point: Adding λm^2 lifts the loss most where m is large, so the bottom of the curve slides towards $m = 0$.

To see this in one picture, take the one-input example and hold the intercept at -2.29 , its linear regression value. Then the loss depends only on the slope m :

$$L(m) = \sum_{i=1}^n (y_i - mx_i + 2.29)^2 + \lambda m^2$$

Figure 4 draws this curve while λ grows from 0 to 100.




Figure 4: The loss against the slope for $\lambda = 0, 10, 50$ and 100

λ	Slope at the lowest point
0	27.8
10	25.0
50	17.7

λ	Slope at the lowest point
100	13.0

The added term λm^2 is 0 at $m = 0$ and grows fast away from it. So the curve rises and narrows, more on the far side from 0, and its lowest point moves towards $m = 0$. Whatever method finds that lowest point (the formula or gradient descent) therefore returns a smaller slope.

5.1 With two coefficients: why “Ridge”?

Key point: With two coefficients, the answer lies where the error contours first touch a circle around the origin.

With two coefficients β_1 and β_2 , the loss has two parts: the squared error and the penalty $\lambda(\beta_1^2 + \beta_2^2)$. The squared error alone is lowest at the OLS answer; its contours are ellipses around it (Figure 5).




Figure 5: The squared-error contours and the penalty circle

The penalty is the same at every point of a circle around the origin. Ridge looks for the point with the lowest squared error for a given penalty: the point where the ellipses first touch the circle. That point is always closer to the origin than the OLS answer, so both coefficients are smaller.

A larger λ means a smaller circle. Seeing Ridge as a hard limit on the coefficients is the **constrained** view of the problem; the details come in a later Note.

Extra: The answer to “why is it called Ridge?” is often given with this picture: the solution always lies on the edge of the circle. Historically, the name comes from the formula of the maths Note: adding λ to the diagonal of $X^T X$ puts a “ridge” along that diagonal.

6. Point 5: use Ridge with 2 or more inputs

Key point: Ridge helps when there are several inputs, especially correlated ones. With a single input there is little to gain.

Regularisation fights overfitting caused by many, large coefficients. With one input there is only one slope, so there is little room to overfit. With 2 or more inputs, and especially with many or correlated inputs, the coefficients can become extreme, and Ridge keeps them in check.

7. Summary

#	Point
1	As λ grows, all coefficients shrink towards 0 but never reach it
2	The largest coefficients shrink fastest; small ones barely change at first
3	Larger λ : higher bias, lower variance. Pick λ in between
4	The loss curve rises and its lowest point slides towards 0
5	Use Ridge with 2 or more inputs, especially correlated ones

8. Key terms

Term	Meaning
Coefficient path	How each coefficient changes as the regularisation strength grows
Bootstrap sample	A sample of the same size drawn from the data with replacement
Constrained form	Writing regularisation as a hard limit on the size of the coefficients