

Naive Bayes: The Mathematics

Contents

1. Overview	2
2. Notation	2
3. Step 1: Bayes' theorem without the evidence	3
4. Step 2: the chain rule	3
5. Step 3: the naive assumption	4
6. The formula and the MAP rule	4
7. When the assumption fails	5
8. Summary	6
9. Sources	6
10. Key terms	6

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Classification problems ([Note 3](#)); Probability density function (PDF) ([Note 20](#)); Normal distribution ([Note 42](#)); Independent and mutually exclusive events ([Note 84](#)); Bayes' theorem ([Note 86](#)).

1. Overview

Key point: Starting from Bayes' theorem, the chain rule writes the joint probability exactly; the naive assumption (conditional independence) then simplifies it to $P(C_k) \times \prod P(x_i | C_k)$. The predicted class is the one that maximises this.

The previous Note used the Naive Bayes recipe on a small example. This Note derives it in general, step by step, and makes the “naive” assumption precise.

$$\begin{aligned}
 P(C_k | x) &\propto P(x_1, x_2, \dots, x_n, C_k) && \text{Bayes, evidence dropped} \\
 &= P(x_1 | x_2, \dots, x_n, C_k) P(x_2 | x_3, \dots, x_n, C_k) \cdots P(x_n | C_k) P(C_k) && \text{chain rule (exact)} \\
 &\approx P(x_1 | C_k) P(x_2 | C_k) \cdots P(x_n | C_k) P(C_k) && \text{naive assumption: drop the red parts} \\
 &= P(C_k) \prod_{i=1}^n P(x_i | C_k) && \text{the Naive Bayes formula} \\
 \hat{y} &= \arg \max_k P(C_k) \prod_{i=1}^n P(x_i | C_k) && \text{maximum a posteriori (MAP) rule}
 \end{aligned}$$

Figure 1: The derivation in five lines

2. Notation

Key point: Inputs $x = (x_1, \dots, x_n)$; K classes, C_1 to C_K for k up to K .

- A row has n input values, $x = (x_1, x_2, \dots, x_n)$. In the cricket example, $n = 3$: toss, venue and outlook.
- The output is one of K classes, $C_1, \dots, C_K$. In the cricket example, $K = 2$: win and loss.

The goal: for each class C_k , compute $P(C_k | x)$, then pick the largest.

3. Step 1: Bayes' theorem without the evidence

Key point: $P(C_k | x)$ is proportional to $P(x | C_k) P(C_k) = P(x_1, \dots, x_n, C_k)$.

By Bayes' theorem,

$$P(C_k | x) = \frac{P(x | C_k) P(C_k)}{P(x)}$$

The evidence $P(x)$ is the same for every class, so it does not affect which class is largest (the intuition Note). Dropping it, we write $\propto$, “is proportional to”:

$$P(C_k | x) \propto P(x | C_k) P(C_k)$$

From the Bayes' theorem proof, $P(x | C_k)P(C_k) = P(x \cap C_k)$, the probability of the inputs and the class together. Writing the inputs out, with commas meaning “and”:

$$P(C_k | x) \propto P(x_1, x_2, \dots, x_n, C_k)$$

4. Step 2: the chain rule

Key point: Peel off one variable at a time with $P(A, B) = P(A | B) P(B)$. This is exact.

Apply $P(A \cap B) = P(A | B) P(B)$ with $A = x_1$ and $B = (x_2, \dots, x_n, C_k)$:

$$P(x_1, \dots, x_n, C_k) = P(x_1 | x_2, \dots, x_n, C_k) P(x_2, \dots, x_n, C_k)$$

The last factor has the same form with one variable fewer, so apply the rule again, with $A = x_2$:

$$P(x_2, \dots, x_n, C_k) = P(x_2 | x_3, \dots, x_n, C_k) P(x_3, \dots, x_n, C_k)$$

Repeating until only x_n and C_k are left, and finally $P(x_n, C_k) = P(x_n | C_k) P(C_k)$:

$$P(x_1, \dots, x_n, C_k) = P(x_1 | x_2, \dots, x_n, C_k) P(x_2 | x_3, \dots, x_n, C_k) \cdots P(x_n | C_k) P(C_k)$$

This repeated splitting is the **chain rule** of probability. Nothing has been assumed yet: it is exact.

5. Step 3: the naive assumption

Key point: Assume each input depends only on the class, not on the other inputs: $P(x_i \mid \text{other inputs}, C_k) = P(x_i \mid C_k)$.

The factors of the chain rule are hard to estimate. For the cricket match, $P(x_1 \mid x_2, x_3, C_k)$ is “the probability the toss was lost, given the venue was Mumbai, the weather was sunny and the match was won”. Few or no training rows match all those conditions, so the estimate is 0 or unreliable (the intuition Note).

Naive Bayes assumes **conditional independence**: once the class is known, each input is independent of the others. In symbols, for any inputs,

$$P(x_i \mid x_{i+1}, \dots, x_n, C_k) = P(x_i \mid C_k)$$

This is the independence of the independent events Note, $P(A \mid B) = P(A)$, but holding **within each class**. Every chain-rule factor loses its other inputs (the red parts in Figure 1):

$$P(x_1, \dots, x_n, C_k) \approx P(x_1 \mid C_k) P(x_2 \mid C_k) \cdots P(x_n \mid C_k) P(C_k)$$

6. The formula and the MAP rule

Key point: Score each class by $P(C_k) \prod P(x_i \mid C_k)$ and predict the largest: the maximum a posteriori rule.

Written with a product sign ($\prod$ is to multiplication what $\sum$ is to addition):

$$P(C_k \mid x) \propto P(C_k) \prod_{i=1}^n P(x_i \mid C_k)$$

To get actual probabilities, divide by the sum over all classes: $P(C_k \mid x) = \frac{1}{Z} P(C_k) \prod_i P(x_i \mid C_k)$, where Z is that sum, which equals the evidence $P(x)$.

The prediction is the class with the largest posterior:

$$\hat{y} = \arg \max_{k \in \{1, \dots, K\}} P(C_k) \prod_{i=1}^n P(x_i \mid C_k)$$

(arg max means “the k that gives the maximum”.) This is the **maximum a posteriori (MAP) rule**. On the cricket example, it gives 0.040 for win and 0.056 for loss, so $\hat{y} = \text{loss}$.

7. When the assumption fails

Key point: The chain rule is exact; the naive product is only exact when the inputs really are independent within each class. Strongly related inputs get counted twice.

A simulation with two binary inputs and 200,000 rows shows the difference:

Inputs within a class	Exact $P(x_1, x_2 \mid C)$	Chain rule	Naive product
independent	0.476	0.476	0.477
x_2 is a copy of x_1	0.800	0.800	0.639

When the inputs are independent, the naive product matches. When x_2 simply repeats x_1 , the naive product multiplies the same evidence twice and is wrong.

The same information counted again and again makes Naive Bayes over-confident




Figure 2: Copying a column makes Naive Bayes over-confident

Figure 2 shows the effect on the cricket prediction: each extra copy of the “toss” column pushes $P(\text{loss})$ further, from 58% with one copy to 99.4% with five, although no new information was added.

In practice, Naive Bayes often still picks the right class even when its independence assumption is wrong, because only the **order** of the scores matters for the prediction. Its probabilities, however, tend to be too extreme, so they should not be trusted as exact (Domingos and Pazzani 1997; scikit-learn user guide §1.9).

Extra: Multiplying many probabilities, each below 1, produces very small numbers (the log loss Note). Implementations therefore add logarithms instead: $\log P(C_k) + \sum_i \log P(x_i \mid C_k)$. The largest sum identifies the same class.

8. Summary

Step	Result
Bayes, drop the evidence	$P(C_k x) \propto P(x_1, \dots, x_n, C_k)$
Chain rule (exact)	$\prod_i P(x_i x_{i+1}, \dots, x_n, C_k) \times P(C_k)$
Naive assumption	$P(x_i \dots, C_k) = P(x_i C_k)$
Formula	$P(C_k) \prod_i P(x_i C_k)$
MAP rule	predict $\arg \max_k$ of the formula

9. Sources

- **Domingos and Pazzani 1997:** Domingos, P. and Pazzani, M. “On the Optimality of the Simple Bayesian Classifier under Zero-One Loss.” *Machine Learning* 29, 103–130, 1997.
- **scikit-learn user guide:** Section 1.9, “Naive Bayes”, scikit-learn 1.9.

10. Key terms

Term	Meaning
Chain rule of probability	Writing a joint probability as a product of conditional probabilities, one variable at a time
Conditional independence	Independence that holds once a third variable (here the class) is known
Proportional to ($\propto$)	Equal up to a constant factor that is the same for every class
$\arg \max$	The value of the variable that makes an expression largest
MAP rule	Maximum a posteriori: predict the class with the largest posterior probability