


Visualising Decision Trees with dtreeviz

Contents

1. Overview	2
2. The problem with the default tree drawing	2
3. Installing and calling dtreeviz	3
4. Reading a dtreeviz tree	4
4.1 A fully grown tree	6
5. Regression trees in dtreeviz	6
6. Useful view options	7
6.1 Left-to-right trees	7
6.2 The prediction path of one observation	8
6.3 Node numbers	8
6.4 Simple boxes: fancy=False	8
6.5 The path in plain English	9
6.6 Feature importance	9
7. The whole tree on one plot	10
7.1 One input: all cuts on one axis	10
7.2 Two inputs: a step surface in 3D	10
8. Summary	12
9. Sources	12
10. Key terms	12

Where this fits: Pipeline map steps 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** Model-based learning ([Note 6](#)); Decision surface and boundary ([Note 91](#)); Entropy, information gain and Gini ([Note 97](#)).
- **Leads to:** Ensemble learning ([Note 101](#)); Bagging ([Note 105](#)); Random forest ([Note 108](#)); Permutation importance ([Note 114](#)); AdaBoost ([Note 115](#)).
- **Compare with:** Feature scaling ([Note 24](#)).

1. Overview

Key point: dtreeviz draws a trained decision tree with the training data shown inside every node, so we can see why each split was chosen and where a prediction goes.

dtreeviz is a Python library for visualising decision trees trained with scikit-learn (and with XGBoost, LightGBM, Spark MLlib and TensorFlow Decision Forests; dtreeviz README). Compared with scikit-learn's own tree drawing, dtreeviz adds the following. (A **feature** is an input variable, one column of the data table; an **observation** is one record, one row of the table.)

- real feature and class names at every split;
- the training data of every node, as histograms (classification) or scatter plots (regression);
- the path a single new observation follows to its prediction, and a plain-English summary of it;
- views of the whole tree as cuts on a scatter plot, in 2D or 3D.

dtreeviz itself draws with Graphviz and matplotlib. The figures below instead rebuild each dtreeviz view from the fitted tree's own numbers, with Plotly and Graphviz. The Notebook ([notebook.ipynb](#)) lists the dtreeviz calls and computes everything the figures show.

2. The problem with the default tree drawing

Key point: scikit-learn's drawing shows feature indexes like `x[3]` and raw lists of numbers; it is hard to read on a real dataset.

scikit-learn can draw a fitted tree with `plot_tree` (built on matplotlib) or export it with `export_graphviz`. Both produce boxes like those in Figure 1, for a tree of depth 2 trained on the iris data (the [softmax regression Note](#)).




Figure 1: The default drawing of a depth-2 iris tree (scikit-learn’s export)

The drawing has several weaknesses:

- **Features appear as indexes.** The split “ $x_3 \leq 0.8$ ” does not say which feature x_3 is. With many features we cannot remember every index position. (x_3 is the petal width.)
- **Some information is noise for most readers**, such as the Gini value of every node (the [decision tree intuition Note](#), section 8).
- **Some information is hard to read:** value = [0, 49, 5] means 0 setosa, 49 versicolor and 5 virginica, but we must remember the class order.
- **The data is invisible.** We see the threshold 0.8 but not why the split falls there.

Extra: Passing `feature_names=` and `class_names=` to `plot_tree` or `export_graphviz` fixes the first problem, and `filled=True` colours the boxes by class (sklearn `plot_tree` reference). The data inside each node, however, only `dtreeviz` shows.

3. Installing and calling dtreeviz

Key point: Wrap the trained tree with `dtreeviz.model(...)`, passing the training data and the feature and class names, then call `.view()`.

`dtreeviz` is installed with `pip install dtreeviz`. `dtreeviz` also needs the **Graphviz** program `dot`, which lays out the boxes of the tree. In Jupyter the result appears inline; in a script, `.show()` opens it in a window and `.save("tree.svg")` writes a file.

Python: A first `dtreeviz` tree (`dtreeviz 2.x`).

```
import dtreeviz
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
```

```

iris = load_iris()
X, y = iris.data, iris.target
clf = DecisionTreeClassifier(max_depth=2).fit(X, y)

viz = dtreeviz.model(clf, X_train=X, y_train=y,
                    feature_names=iris.feature_names,
                    target_name="flower",
                    class_names=list(iris.target_names))
viz.view()           # draws the tree
viz.view(scale=2)    # the same, twice as big

```

`dtreeviz.model` needs four things: the trained tree, the data it was trained on, the names of the input features, and (for classification) the names of the classes.

Extra: Older tutorials use the `dtreeviz 1.x` interface, a single function: `from dtreeviz.trees import dtreeviz`, then `dtreeviz(clf, X, y, target_name=..., feature_names=..., class_names=...)`, with the options of section 6 passed to the same call. `dtreeviz 2.0` replaced it with `dtreeviz.model(...)` and methods such as `.view()` (`dtreeviz 2.0.0` release notes).

4. Reading a dtreeviz tree

Key point: Each decision node shows a stacked histogram of its split feature, with a triangle at the threshold; each leaf shows a pie of the classes that reached it.

Figure 2 shows the `dtreeviz` view of the same depth-2 iris tree, rebuilt with Plotly from the fitted tree.




Figure 2: The depth-2 iris tree drawn the dtreeviz way: data at every node

The root splits on **petal width** at **0.80**, shown by the black triangle. The histogram of all 150 flowers makes the reason obvious: every setosa flower (blue) has a petal width below 0.8, and every other flower is above it.

- Left of the threshold, **all 50 setosa** flowers go to a leaf. The leaf's pie is entirely blue: a pure leaf.
- Right of it, the **other 100 flowers** go to the next node.

The second node splits the remaining 100 flowers on petal width again, at **1.75**. Here the two colours overlap between about 1.4 and 1.9, so no threshold separates them perfectly. Because `max_depth=2`, the tree stops here and makes two leaves:

- left: **54 flowers**, mostly versicolor (49 versicolor, 5 virginica);
- right: **46 flowers**, mostly virginica (1 versicolor, 45 virginica).

The counts add up: $54 + 46 = 100$, and $100 + 50 = 150$.

4.1 A fully grown tree

Key point: Without `max_depth` the tree keeps splitting until every leaf is pure, even if a leaf holds one flower.

Without `max_depth`, the same data gives a tree of depth 5 with 9 leaves (Figure 4, in section 6). Every leaf is **pure**: it holds a single class. Some leaves hold 1, 2 or 3 flowers only.

Leaves that small are the overfitting of the [hyperparameters](#) **Note**: the tree has carved out a region for one flower. A big tree is also harder to view on screen, which is another reason to limit its growth.

5. Regression trees in dtreeviz

Key point: For a regression tree, each node shows a scatter plot of its split feature against the output, with the threshold and the mean of each side.

Everything works the same for a regression tree: we pass the `DecisionTreeRegressor` and give `target_name` (the name of the **target**, the output we predict) instead of class names. The nodes then show **scatter plots** instead of histograms.

Figure 3 shows a regression tree of depth 1 trained on all 506 districts of the Boston housing data (the [regression trees](#) **Note**, section 7).




Figure 3: A one-split regression tree on the Boston data, drawn as a scatter plot

The single split is on **RM** (average rooms per home) at **6.94**:

- **430 districts** have $RM \leq 6.94$; the tree predicts their mean value, **19.93** thousand dollars;
- **76 districts** have $RM > 6.94$; the tree predicts **37.24** thousand dollars.

So a new district with 7.5 rooms per home is predicted at 37.24 thousand dollars, and one with 6 rooms at 19.93.

6. Useful view options

Key point: `orientation`, `x`, `show_just_path`, `show_node_labels` and fancy change what `.view()` draws.

Figure 4 combines several of these options on the fully grown iris tree: drawn left to right, with node numbers, and with the path of one flower highlighted in red.




Figure 4: The fully grown iris tree, left to right, with node numbers and one flower's prediction path; leaves show setosa/versicolor/virginica counts

6.1 Left-to-right trees

Key point: `orientation="LR"` draws the tree from left to right, which fits wide trees better.

By default trees are drawn top-down (`orientation="TD"`). With `viz.view(orientation="LR")` the root is on the left and the leaves on the right, as in Figure 4. For a large tree this often fits a screen or page better.

6.2 The prediction path of one observation

Key point: `viz.view(x=row)` highlights the path that observation follows from the root to its leaf.

Passing a new observation as `x` highlights its **prediction path**: the nodes it passes through on the way to its leaf. In Figure 4 the flower has sepal length 5.9, sepal width 3.2, petal length 4.8 and petal width 1.8 cm.

Following the red path:

1. **#0**, petal width ≤ 0.80 ? No (1.8), so it goes to **#2**.
2. **#2**, petal width ≤ 1.75 ? No (1.8), so it goes to **#12**.
3. **#12**, petal length ≤ 4.85 ? Yes (4.8), so it goes to **#13**.
4. **#13**, sepal width ≤ 3.10 ? No (3.2), so it reaches leaf **#15**: versicolor.

The flower really is a versicolor, and the prediction is right. Note, though, that leaf **#15** holds a single training flower: the tree trusts one example here.

With `show_just_path=True`, `dtreeviz` draws only the nodes on the path and leaves out the rest of the tree, which helps when the tree is large.

6.3 Node numbers

Key point: `show_node_labels=True` prints each node's number; nodes are numbered depth-first, left branch before right.

With `show_node_labels=True`, every node shows its number, as in Figure 4. `scikit-learn` numbers the nodes **depth-first**: it starts at the root (0), goes down the left branch to its end (1), then returns up and continues with the next unvisited branch (2, 3, 4, 5, ...).

In Figure 4: **#0** is the root, **#1** its left leaf (setosa), and **#2** the right child. From **#2** the numbering first finishes the whole left branch, **#3** to **#11**, before returning to the right branch, **#12** to **#16**. These are the same numbers the arrays of the fitted tree use (`clf.tree_`), so they help when inspecting a node in code.

6.4 Simple boxes: `fancy=False`

Key point: `fancy=False` drops the histograms and scatter plots, keeping small boxes and the leaf pies.

When the histograms make a big tree too heavy, `viz.view(fancy=False)` draws each decision node as a plain box with its question, and keeps only the leaf pies (or leaf values for regression).

6.5 The path in plain English

Key point: `viz.explain_prediction_path(x)` lists, for each feature used, the range of values that leads to the prediction.

For a deep tree, a short text can be clearer than a picture. `viz.explain_prediction_path(x)` turns the path into one condition per feature. For the flower of Figure 4:

- $3.10 < \text{sepal width}$
- $\text{petal length} \leq 4.85$
- $1.75 < \text{petal width}$

Petal width was asked twice on the path (> 0.80 , then > 1.75), and the two conditions are merged into the tighter one. Any flower that meets these three conditions lands in the same leaf.

6.6 Feature importance

Key point: The tree's `feature_importances_` shows which features it relied on; for iris, petal width does almost all the work.

The **feature importance** of a feature (the [regression trees Note](#), section 7.5) is its share of all the impurity reduction in the tree. Figure 5 shows it for the fully grown iris tree.




Figure 5: Feature importance of the fully grown iris tree

Petal width carries 0.923 of the total, petal length 0.064 and sepal width 0.013. Sepal length is never used. The exact values depend on the training data. In the Notebook, over 20 different training splits of iris, petal width and petal length always come first and second, but the two sepal features swap places half of the time.

Extra: In `dtreeviz 2.x`, `viz.instance_feature_importance(x)` shows the importance computed only along one observation's prediction path (`dtreeviz` source, `trees.py`). The importance of the whole tree is the scikit-learn attribute `clf.feature_importances_`, which Figure 5 plots.

7. The whole tree on one plot

Key point: With one or two input features, dtreeviz draws all the cuts of a regression tree directly over the data: a staircase in 2D, a step surface in 3D.

7.1 One input: all cuts on one axis

Key point: Every split uses the same feature, so all the cuts are vertical lines on one scatter plot.

The `cars.csv` sample data from dtreeviz lists 392 cars with their fuel use (**MPG**, miles per gallon), weight (**WGT**, pounds), engine size (**ENG**, cubic inches) and cylinders (**CYL**).

With **WGT** as the only input and MPG as the output, every split of the tree is on WGT. A regression tree of depth 3 makes 7 cuts and 8 leaves, and each leaf predicts its mean MPG (Figure 6). In dtreeviz 2.x this view is `viz.rtree_feature_space(features=["WGT"])`.




Figure 6: A depth-3 regression tree on one input: every cut is on the weight axis

Heavier cars use more fuel: the staircase steps down from about 33 MPG for the lightest cars to about 13 MPG for the heaviest.

7.2 Two inputs: a step surface in 3D

Key point: With two inputs, each leaf is a rectangle of the input plane, and the prediction is a flat step over it.

With **WGT** and **ENG** as inputs, the cuts are lines on the WGT-ENG plane, and each leaf is a rectangle with a flat height, its mean MPG (Figure 7). In dtreeviz 2.x this view

is `viz.rtree_feature_space3D(features=["WGT", "ENG"])`. The root split here is on engine size, at 190.5 cubic inches.




Figure 7: A depth-3 regression tree on two inputs: a surface of flat steps over the data

This picture is exactly the idea of the [regression trees Note](#) (section 5): planes parallel to the axes cut the data into boxes, and each box predicts one mean.

8. Summary

dtreeviz call	What it shows	In this Note
<code>dtreeviz.model(clf, X_train, y_train, feature_names, target_name, class_names)</code>	wraps a trained tree with its data and names	(setup)
<code>viz.view()</code>	tree with histograms (classification) or scatter plots (regression)	Figures 2 and 3
<code>viz.view(scale=2)</code>	a bigger drawing	
<code>viz.view(orientation="LR")</code>	left-to-right tree	Figure 4
<code>viz.view(x=row)</code>	highlights one observation's prediction path	Figure 4
<code>viz.view(x=row, show_just_path=True)</code>	only the path's nodes	
<code>viz.view(show_node_labels=True)</code>	node numbers (depth-first)	Figure 4
<code>viz.view(fancy=False)</code>	plain boxes, no histograms	
<code>viz.explain_prediction_path(row)</code>	the path as one range per feature	section 6.5
<code>clf.feature_importances_</code>	each feature's share of the impurity reduction	Figure 5
<code>viz.rtree_feature_space(...)</code> , <code>viz.rtree_feature_space3D(...)</code>	all cuts of a regression tree over the data	Figures 6 and 7

- scikit-learn's default drawing names features by index and hides the data; dtreeviz shows names, data and thresholds together.
- The depth-2 iris tree splits on petal width at 0.80 (50 setosa) and at 1.75 (54 mostly versicolor, 46 mostly virginica).
- A regression split on the Boston data, $RM \leq 6.94$, predicts 19.93 (430 districts) or 37.24 (76 districts).
- A prediction path shows exactly which questions decided one observation's prediction.

9. Sources

- **dtreeviz README:** T. Parr et al., dtreeviz, README at github.com/parrr/dtreeviz (supported libraries).
- **dtreeviz 2.0.0 release notes:** github.com/parrr/dtreeviz/releases/tag/2.0.0 (the API re-organized around `dtreeviz.model`).
- **dtreeviz source, trees.py:** class `DTreeVizAPI` in github.com/parrr/dtreeviz, file `dtreeviz/trees.py` (methods `view`, `explain_prediction_path`, `instance_feature_importance`, `rtree_feature_space`, `rtree_feature_space3D`).
- **sklearn plot_tree reference:** scikit-learn API reference for `sklearn.tree.plot_tree` and `sklearn.tree.export_graphviz`. scikit-learn.org/stable/api/sklearn.tree.html

10. Key terms

Term	Meaning
dtreeviz	A Python library that draws decision trees with the training data shown at every node
Graphviz	The graph-drawing program (<code>dot</code>) that lays out tree diagrams for dtreeviz and <code>export_graphviz</code>
<code>export_graphviz</code>	scikit-learn function that writes a tree as Graphviz DOT text
Prediction path	The nodes an observation passes through, from the root to the leaf that predicts it
Pure leaf	A leaf whose training observations all belong to one class
Node number	A node's index in the fitted tree, assigned depth-first starting from 0 at the root
<code>cars.csv</code>	dtreeviz's sample data: 392 cars with MPG, weight, engine size and cylinders