


Voting Regressor

Contents

1. Overview	2
2. The core idea	2
3. Seeing it on a curve	2
4. VotingRegressor on the Boston housing data	4
4.1 The data and the base models	4
4.2 The voting regressor	5
4.3 Weights	5
4.4 One algorithm, different settings	5
5. Summary	7
6. Key terms	7

Where this fits: Pipeline map steps 8 (Model), 9 (Evaluate). See the [Course map](#).



- **Builds on:** Independent and mutually exclusive events ([Note 84](#)); Ensemble learning ([Note 101](#)).
- **Leads to:** Grid and random search ([Note 106](#)).
- **Compare with:** Train-test split ([Note 13](#)); Bagging ([Note 105](#)); OOB score ([Note 105](#)); Stacking and blending (Video 127, coming).

1. Overview

Key point: A voting regressor trains several regressors on the same data and predicts the mean of their outputs, or a weighted mean if we give weights.

The voting regressor is the regression version of the voting classifier (the [voting classifier Note](#)). Only the aggregation step changes: instead of a vote, we take the **mean**. This Note covers:

- the core idea and a demo with curves;
- scikit-learn's `VotingRegressor` on the Boston housing data, with **weights**;
- voting over one algorithm with different settings.

The Notebook (`notebook.ipynb`) runs every experiment. The Dash app `app.py` lets us tick base regressors and redraws their curves next to the voting regressor's.

2. The core idea

Key point: Train every regressor on the same data; to predict, average their numerical outputs.

We have base models M_1 , M_2 and M_3 , all regression algorithms, trained on the same dataset D . For a new query point x_q , each returns a number, and the voting regressor returns their mean (the [introduction to ensemble learning Note](#), section 3.3, works an example). With weights, it returns the weighted mean, exactly as in soft voting (the [voting classifier Note](#), section 4.4).

3. Seeing it on a curve

Key point: The voting regressor's curve runs between its base models' curves, smoothing out the extremes of each.

The demo data is a sine wave of 80 points between 0 and 5, with every fifth point pushed up or down at random: a non-linear pattern with a few outliers. We train on 72 points and keep 8 for testing. Figure 1 shows three base regressors and the voting regressor.




Figure 1: Three base regressors (dash-dot) and the voting regressor (solid blue) on noisy sine data; legend gives R^2 and MAE on the 8 test points

Adding models one at a time (R^2 on the 8 test points):

Base models	Each model's R^2	Voting regressor's R^2
linear regression	0.18	0.18
linear regression, SVR	0.18, 0.47	0.51
linear regression, SVR, decision tree	0.18, 0.47, 0.38	0.52

- With **linear regression alone**, the voting regressor is that same line: one model averaged with nothing. A straight line cannot follow a wave, so R^2 (the [regression metrics Note](#)) is only 0.18.
- **Adding SVR** (support vector regression, the [SVM intuition Note](#)): SVR follows the wave well. The voting curve runs exactly halfway between the line and the SVR curve, and its R^2 (0.51) beats both.
- **Adding a decision tree** (depth 5): the tree jumps to reach many single points, outliers included. The voting curve follows its jumps only a third of the way, so it is much smoother than the tree. R^2 rises a little, to 0.52.

The point is the picture: averaging pulls each model's extremes toward the others.

Extra: The legend also shows MAE (mean absolute error). By MAE, SVR and the tree (both 0.23) beat the voting regressor (0.31): with only 8 test points, and one of them an outlier, different metrics can rank models differently. A score from 8 points is a rough guide only; cross-validation, used below, is more reliable.

4. VotingRegressor on the Boston housing data

Key point: Three weak models (R^2 0.20, -0.07 and -0.41) combine into a voting regressor with R^2 0.43; weights lift it to 0.45.

4.1 The data and the base models

Key point: 506 districts, 13 inputs, median house price; linear regression, a decision tree and SVR, each scored by 10-fold cross-validation.

The Boston housing data (the [regression trees Note](#), section 7.1, describes it and why `load_boston` was removed from scikit-learn) has 506 rows and 13 input columns. The output, MEDV, is the median house price of each district. The Notebook reads it from `data/boston.csv`.

Python: Each base model alone.

```
boston = pd.read_csv("data/boston.csv")
X, y = boston.drop(columns="MEDV"), boston["MEDV"]

estimators = [("lr", LinearRegression()),
              ("dt", DecisionTreeRegressor(random_state=0)),
              ("svr", SVR())]
for name, model in estimators:
    scores = cross_val_score(model, X, y,
                             scoring="r2", cv=10)
    print(name, scores.mean())
scoring="r2" makes cross_val_score report  $R^2$  instead of accuracy.
```

Model	Linear regression	Decision tree	SVR
Mean R^2 (10 folds)	0.20	-0.07	-0.41

All three are poor: a negative R^2 means worse than always predicting the mean price. We could tune each one, but here we only want to compare them with voting.

4.2 The voting regressor

Key point: Passing the same three models to `VotingRegressor` doubles the best R^2 , from 0.20 to 0.43.

Python: The voting regressor.

```
from sklearn.ensemble import VotingRegressor
```

```
vr = VotingRegressor(estimators)
```

```
cross_val_score(vr, X, y, scoring="r2", cv=10).mean()    # 0.43
```

`VotingRegressor` takes the same list of (name, model) tuples as `VotingClassifier`, plus optional `weights` and `n_jobs`. There is no `voting` setting: regression always averages.

`n_jobs=-1` trains the base models in parallel on all CPU cores. It saves time on big data and makes little difference on small data like this.

4.3 Weights

Key point: Trying weights 1 to 3 for each model, the best is (2, 1, 1), giving R^2 0.45.

As for classification, we try each weight from 1 to 3 for each model, 27 combinations:

Weights (lr, dt, svr)	(1, 1, 1)	(3, 1, 1)	(3, 2, 2)	(2, 1, 1)
Mean R^2	0.429	0.433	0.445	0.447

The best weights give linear regression, the best single model, twice the say of the others. The gain over equal weights is small here.

4.4 One algorithm, different settings

Key point: Five trees of depth 1, 3, 5, 7 and unlimited score between -0.85 and 0.04; their vote scores 0.19.

We can also vote over one algorithm with different hyperparameter values. Five decision trees of different depths:

max_depth	1	3	5	7	None	Vote of all five
Mean R^2	-0.85	-0.23	0.04	-0.07	-0.07	0.19

Every tree alone is poor, yet their average is clearly better than each of them.

Extra: Why are all these scores so low? The Boston rows are stored grouped by town, and `cv=10` cuts the rows into 10 folds **in order**, without shuffling. Each fold

then tests the model on towns it never saw during training. Shuffling first, with `cv=KFold(n_splits=10, shuffle=True, random_state=42)`, gives far higher scores: linear regression 0.72, decision tree 0.76, SVR 0.19, voting 0.75, and the vote of the five trees **0.79**. Note that with shuffled folds the plain vote (0.75) no longer beats the decision tree (0.76): voting is worth trying, not guaranteed to win.

5. Summary

	Voting classifier	Voting regressor
Base models	classifiers	regressors
Aggregation	majority vote (hard) or average probability (soft)	mean of the predictions
Weights	<code>weights=...</code>	<code>weights=...</code> (weighted mean)
scikit-learn class	<code>VotingClassifier</code>	<code>VotingRegressor</code>

- A voting regressor predicts the mean (or weighted mean) of its base regressors' predictions.
- On the noisy sine data, the voting curve smooths out the tree's jumps: R^2 0.52 against 0.18, 0.47 and 0.38.
- On Boston, three poor models give a vote of 0.43 (0.45 with weights 2, 1, 1); five trees of different depths give 0.19.
- Shuffle the rows before cross-validation when the data is stored in a meaningful order.

6. Key terms

Term	Meaning
Voting regressor	A regressor that predicts the mean (or weighted mean) of several trained regressors' predictions
<code>n_jobs</code>	scikit-learn setting for how many CPU cores to use in parallel; -1 means all