

Bagging Classifier

Contents

1. Overview	2
2. Seeing bagging on decision surfaces	2
2.1 Bagging against a single tree	3
2.2 Other base models	4
2.3 More estimators, more observations	5
2.4 Pasting, random subspaces and random patches	5
3. BaggingClassifier in code	5
3.1 The data and a single tree	5
3.2 Bagging	6
3.3 Which observations and features each tree got	6
3.4 Other base models, pasting, subspaces and patches	7
4. The out-of-bag score	8
5. What works in practice	10
6. Summary	12
7. Sources	12
8. Key terms	13

Where this fits: Pipeline map steps 8 (Model), 9 (Evaluate), 10 (Tune). See the [Course map](#).



- **Builds on:** Simple imputation (mean, median, mode, constant) ([Note 37](#)); ML pipelines ([Note 38](#)); Bias-variance trade-off ([Note 62](#)); Overfitting ([Note 91](#)); Hyperparameter tuning ([Note 98](#)); Decision trees ([Note 100](#)).
- **Leads to:** Random forest ([Note 108](#)).
- **Compare with:** Boosting ([Note 101](#)); Voting ensembles ([Note 104](#)); Cross-validation ([Note 104](#)); Random forest ([Note 108](#)); Optuna ([Note 134](#)); Bayesian optimisation ([Note 134](#)).

1. Overview

Key point: scikit-learn's `BaggingClassifier` trains many copies of one classifier on random samples of the observations and/or features and takes their majority vote. Its settings choose the type: bagging, pasting, random subspaces or random patches.

The idea was explained in the [bagging Note](#): bootstrapping, then aggregation, in four variants. This Note applies it to classification:

- a demo of decision surfaces, comparing one model with its bagged version, for each of the four types (section 2);
- `BaggingClassifier` in code on a dataset of 10,000 observations, with its hyperparameters (section 3);
- the out-of-bag score (section 4);
- what works in practice, and tuning with `GridSearchCV` (section 5).

The Notebook (`notebook.ipynb`) runs every experiment. The Dash app `app.py` has a control for every bagging setting and redraws the decision surfaces.

2. Seeing bagging on decision surfaces

Key point: On the moons data, one fully grown tree overfits (test accuracy 0.856); bagging 100 trees smooths the decision boundary and reaches 0.912.

The demo data is the two-moons toy dataset: 500 **observations** (points; each is one row of the data table), 375 for training, with two **features** (input variables, the columns x_1 and x_2) and a class label as the **target** (the output we predict). It is the same data as in the [decision tree hyperparameters Note](#), section 4. The app asks for six settings, the main hyperparameters of `BaggingClassifier`:

- **base model** (G-260; **estimator**): decision tree (the default), KNN or SVM;
- **n_estimators**: how many base models;

- **max_samples:** how many observations each model gets (here out of 375);
- **bootstrap:** draw those observations with replacement (**True**) or without (**False**);
- **max_features:** how many features each model gets (here out of 2);
- **bootstrap_features:** draw those features with replacement or without.




Figure 1: Decision surfaces on the moons data: one fully grown tree, then 100-tree ensembles of each bagging type, and bagged KNN. Titles give the test accuracy

2.1 Bagging against a single tree

Key point: The single tree’s decision boundary wraps small boxes around single points; the bagged trees’ decision boundary is one smooth curve.

With a decision tree, 100 estimators, 50 observations each, drawn with replacement, and both features:

- **One tree** (Figure 1, top left): 0.856. Its **decision surface** (G-560; the feature plane coloured by the predicted class) has small boxes of one colour inside the other. A tree splits on one feature at a time, so each region it carves out is a rectangle, and a fully grown tree keeps splitting until it gives single training points their own rectangle. That is **overfitting** (G-1429). The tree is right on the training data, wrong on new data.

- **Bagging** (G-251; top middle): 0.912. The **decision boundary** (G-555), the line where the predicted class changes, is much smoother: each tree's stray boxes sit in different places, so the majority vote outvotes them. The **bias** (G-287) stays low and the **variance** (G-2073) falls, as the [bagging Note](#) predicted.

Figure 2 builds the same bagging classifier one tree at a time. Watch the left panel: every tree draws its own 50 points (large markers; bigger means drawn more than once) and cuts its own boxes. On the right, the vote of all trees so far turns from one tree's hard boxes into a smooth band of shared votes, and test accuracy climbs from 0.816 with one tree to 0.912 after 10 trees; from there on it only wobbles between 0.896 and 0.928 and ends at 0.912 with 100 trees.




Figure 2: Bagging builds up on the moons data. Left: the newest tree alone, trained on its 50 drawn points. Right: the share of trees so far that vote blue, with the ensemble's test accuracy

2.2 Other base models

Key point: Bagging helps unstable models such as trees; here KNN and the SVM gain nothing, or lose a little.

Any classifier can be bagged, but the gain depends on how **unstable** the model is (an **unstable model**, G-2156): how much its fit changes when the training data changes a little. Breiman (1996, sections 1 and 6.3) found that bagging helps unstable models such as trees, and can slightly degrade stable ones such as nearest-neighbour methods. Think of asking one calm friend the same question ten times: averaging the answers adds nothing new. The moons data shows the same:

- **KNN:** a single KNN ($k = 5$) already gives a smooth decision boundary and scores

0.912. Bagged, its decision boundary is a little smoother still (Figure 1, bottom right), with the same 0.912.

- **SVM:** a single SVM scores 0.896, but 100 bagged SVMs score only 0.872. The SVM behaves like a stable model here: bagging does not help it.

Instability is the reason decision trees are by far the usual base model for bagging.

2.3 More estimators, more observations

Key point: Beyond a point, more models or more observations per model change little.

With 500 trees instead of 100 (50 observations each), accuracy moves only from 0.912 to 0.920. With 500 trees of 150 observations each it is 0.904. More estimators usually help up to some number, then stop making a difference. Both settings are worth tuning.

2.4 Pasting, random subspaces and random patches

Key point: Pasting behaves like bagging here; sampling one of only two features hurts badly.

- **Pasting** (G-1463; 50 observations, `bootstrap=False`): 0.904, close to bagging's 0.912.
- **Random subspaces** (G-1618; all 375 observations, `bootstrap=False`, `max_features=1`, `bootstrap_features=True`): **0.648**, much worse than one tree. Each tree sees only one feature, so it can only cut along one axis. The surface becomes vertical stripes (Figure 1, bottom left).
- **Random patches** (G-1614; 50 observations with replacement, 1 feature): 0.880. Observation sampling brings back some variety, but one feature is still too little.

Feature sampling makes sense only when there are many features: 10, 20, 50, 100 or more. With two features, sample observations (bagging or pasting) instead.

3. BaggingClassifier in code

Key point: On 10,000 observations with 10 features, one tree scores 0.927; 500 bagged trees score 0.945.

3.1 The data and a single tree

Key point: `make_classification` builds 10,000 observations; 8,000 train, 2,000 test.

We make a dataset of 10,000 observations and 10 features with `make_classification` (the [perceptron code Note](#)), of which 3 features carry the information. 8,000 observations go to training and 2,000 to testing. A single fully grown decision tree scores **0.9265**: the number to beat.

3.2 Bagging

Key point: 500 trees, each trained on 2,000 observations (25%) drawn with replacement.

Python: A bagging classifier.

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
bag = BaggingClassifier(
    estimator=DecisionTreeClassifier(),
    n_estimators=500,
    max_samples=0.25,  # 25% of 8,000 rows
    bootstrap=True,    # with replacement
    random_state=42,
    n_jobs=-1)         # all CPU cores
```

```
bag.fit(X_train, y_train)
```

```
accuracy_score(y_test, bag.predict(X_test))  # 0.945
```

`max_samples` is a share of the observations when it is a decimal (0.25 means 25%) and a count when it is a whole number (50 means 50 observations). `max_features` works the same way for features.

Bagging scores **0.945**, against 0.9265 for one tree. Training takes a while: 500 trees are trained instead of one. `n_jobs=-1` spreads them over every CPU core.

Extra: Older code passes the base model as `base_estimator=`. The name `base_estimator` was deprecated in scikit-learn 1.2 and removed in 1.4; the current name is `estimator` (scikit-learn 1.2 release notes; `BaggingClassifier` API docs, versions 1.3 and 1.4).

3.3 Which observations and features each tree got

Key point: Two attributes record what each tree saw: `estimators_samples_` (its observations) and `estimators_features_` (its features).

After training, the classifier records what each base model saw:

- `bag.estimators_samples_[0]` holds the row numbers (in `X_train`) the first tree was trained on: 2,000 of them (25% of 8,000), starting 2523, 3113, 7114, ... Rows can repeat, because of replacement.
- `bag.estimators_features_[0]` holds its column numbers: all 10, `[0 1 2 ... 9]`, because we did not sample features.

With `max_samples=0.5`, each tree gets 4,000 observations instead, and the accuracy is 0.950.

3.4 Other base models, pasting, subspaces and patches

Key point: Bagged SVMs 0.913, pasting 0.946, random subspaces 0.942, random patches 0.938: all within a point or so of bagging, except the SVMs.

Every variant is the same class with different settings:

Variant	Settings (all with 500 estimators)	Test accuracy
Bagging	<code>max_samples=0.25,</code> <code>bootstrap=True</code>	0.945
Bagged SVMs	<code>estimator=SVC()</code> , otherwise as bagging	0.9125
Pasting	<code>max_samples=0.25,</code> <code>bootstrap=False</code>	0.946
Random subspaces	<code>max_samples=1.0,</code> <code>bootstrap=False,</code> <code>max_features=0.5,</code> <code>bootstrap_features=True</code>	0.9415
Random patches	<code>max_samples=0.25,</code> <code>bootstrap=True,</code> <code>max_features=0.5,</code> <code>bootstrap_features=True</code>	0.938




Figure 3: Test accuracy of one tree (grey, dashed line) and of each BaggingClassifier variant on the 10,000-observation data

In Figure 3, every tree-based variant clears the dashed line of one tree; only the bagged SVMs fall below it.

- The **SVMs** do worse than trees and take much longer: as on the moons data (section 2.2), bagging does not help the SVM.
- For **pasting**, we also set `verbose=1`, which prints training progress, and `n_jobs=-1`.
- For **random subspaces**, `estimators_samples_[0]` now has all 8,000 observations, and `estimators_features_[0]` has 5 feature numbers, for example `[9 2 9 7 7]`. Features 9 and 7 appear twice because `bootstrap_features=True` draws features with replacement.

4. The out-of-bag score

Key point: Every tree misses some observations: about 37% when it draws as many as the training set, about 78% with `max_samples=0.25`. Each observation is predicted by only the trees that never saw it; the share of observations they get right is a free estimate of test accuracy.

The observations a tree never drew are its **out-of-bag (OOB) rows** (G-1412). The 37% is counted per tree: every tree misses a different set of observations. Across many trees, almost every observation is drawn by some trees and missed by others.

A tree has never seen its out-of-bag observations, so for that tree they work like test data. **OOB evaluation** (G-1411) uses this, one observation at a time (Figure 4):

1. **Take one training observation.** Find the trees whose sample does not contain it.
2. **Let only those trees vote.** Their majority is the observation's **OOB prediction** (G-1392).
3. **Compare** the OOB prediction with the observation's true class.
4. **Repeat for every training observation.** The share predicted correctly is the **OOB score**.




Figure 4: The out-of-bag score on the moons training data, with 25 trees. For each starred observation, the trees that drew it are grey and may not vote; the others vote. Observation 1: 9 trees vote, 6 blue against 3 orange, right. Observation 3: 6 of 7 vote blue, but the true class is orange, wrong. Over all 375 observations: 337 right, OOB score 0.899, against 0.888 on the test set. Idea after StatQuest, “Random Forests Part 1”; the data are ours.

In Figure 4, watch the grey squares change from one observation to the next: a different third or so of the trees is free to vote each time (on average 37% of the trees miss a given observation here). The final OOB score, 0.899, is close to the accuracy on the 125 test observations, 0.888, although no test set was used to compute it.

When a tree draws as many observations as the training set, about 37% are never drawn (the [bagging Note](#), section 2.3); the trees of section 3 draw only 25%, so each misses about 78% (Figure 5a). With `oob_score=True` (it needs `bootstrap=True`), scikit-learn computes the OOB score during training, so no separate test set is needed:

```
Python: The out-of-bag score.
bag = BaggingClassifier(DecisionTreeClassifier(),
    n_estimators=500, max_samples=0.25,
    bootstrap=True, oob_score=True,
    random_state=42, n_jobs=-1)
bag.fit(X_train, y_train)
bag.oob_score_           # 0.9429
```




Figure 5: (a) Share of the 8,000 training observations a tree never draws, against `max_samples`: the theory curve $(1 - 1/n)$ raised to the number of draws, and the mean over 500 fitted trees. (b) The out-of-bag score against the test accuracy

In Figure 5, watch the curve fall as each tree draws more: the fewer rows a tree draws, the more rows are left over to score it on.

The OOB score, **0.943**, is close to the real test accuracy, **0.945**. The [OOB score Note](#) covers the OOB score in full, including when it can be trusted.

5. What works in practice

Key point: Try both bagging and pasting; start `max_samples` at 0.25 to 0.5; sample features only when there are many; let `GridSearchCV` decide.

Four rules of thumb, each with what our data says:

1. **Bagging and pasting trade bias for variance.** Drawing with replacement makes the samples more varied, so the trees are less alike and the ensemble's variance is lower, at the cost of a little more bias. Which effect wins depends on the data: here the two score almost the same (0.945 against 0.946, and 0.912 against 0.904 on the moons). Since it is only `bootstrap=True` or `False`, try both.

The variance part follows from theory: the variance of an average falls as the models become less correlated (ESL §15.2), and the Notebook's last cell confirms it on the sine data of the [bagging Note](#), section 3.2: changing only `bootstrap`, bagging's trees are less correlated (0.82 against 0.84), its variance lower (0.041 against 0.055) and its squared bias a little higher (0.0019 against 0.0013).




Figure 6: Bagging against pasting on 100 noisy sine datasets, changing only `bootstrap`: tree correlation, ensemble variance and squared bias

In Figure 6, read the trade from left to right: slightly less correlated trees, clearly lower variance, slightly higher bias.

2. **`max_samples` between 0.25 and 0.5** is a good place to start. Here 0.5 beat 0.25 (0.950 against 0.945), and the grid search below prefers 0.7: the best share depends on the data.
3. **Feature sampling** (random subspaces, random patches) is for **high-dimensional** data, with many features. With few features it hurts, as Figure 1 showed.
4. **Tune with `GridSearchCV` or `RandomizedSearchCV`** instead of guessing (the [KNN Note](#), section 4.2, and the [regression trees Note](#), section 7.2).

Python: Grid search over the bagging settings.

```
parameters = {"n_estimators": [50, 100, 500],
              "max_samples": [0.1, 0.4, 0.7, 1.0],
              "bootstrap": [True, False],
              "max_features": [0.1, 0.4, 0.7, 1.0]}
search = GridSearchCV(BaggingClassifier(random_state=42),
                      parameters, cv=5, n_jobs=-1)
search.fit(X_train, y_train)
search.best_params_
```

The search tries $3 \times 4 \times 2 \times 4 = 96$ combinations with 5-fold cross-validation: 480 fits, which takes several minutes (about 12 minutes on our 12-core machine). The best: **500 trees, 70% of the observations without replacement (pasting), 70% of the features**. Its cross-validation accuracy is 0.955, and its test accuracy 0.952, the best of all our settings.

Here pasting wins by a hair: the rules are a starting point; the search decides.

6. Summary

Hyperparameter	Default	Meaning
<code>estimator</code>	None (decision tree)	the base model (formerly <code>base_estimator</code>)
<code>n_estimators</code>	10	number of base models
<code>max_samples</code>	None (all observations)	observations per model: a share (0.25) or a count (50)
<code>bootstrap</code>	True	observations with replacement (bagging) or without (pasting)
<code>max_features</code>	1.0 (all features)	features per model: a share or a count
<code>bootstrap_features</code>	False	features with replacement or without
<code>oob_score</code>	False	score the ensemble on out-of-bag observations (<code>oob_score_</code>)
<code>n_jobs</code>	None (1 core)	cores to use; -1 means all

- On the moons data, bagging lifts one tree from 0.856 to 0.912; with only two features, feature sampling hurts (0.648).
- Bagging helps unstable models (trees); here it does not help KNN or the SVM.
- On 10,000 observations, bagging lifts one tree from 0.927 to 0.945; a grid search finds 0.952.
- The OOB score (0.943) estimates test accuracy (0.945) without a test set.
- Try bagging and pasting, start `max_samples` at 0.25 to 0.5, sample features only when there are many, and tune with a grid search.

7. Sources

Built from

- CampusX, “Bagging Ensemble | Part 2 | Bagging Classifiers”, YouTube, https://www.youtube.com/watch?v=-1T54G_E-ys
- StatQuest with Josh Starmer, “StatQuest: Random Forests Part 1 - Building, Using and Evaluating”, YouTube, https://www.youtube.com/watch?v=J4Wdy0Wc_xQ (the out-of-bag procedure: each left-out observation is run through the trees built without it; Figure 4 is redrawn from this idea with our own data)

Other references

- Breiman, L. (1996). “Bagging Predictors”. *Machine Learning* 24(2), 123–140, sections 1 and 6.3.
- Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The Elements of Statistical Learning* (ESL), 2nd ed. Springer, section 15.2, equation 15.1.

- scikit-learn developers. Release notes, version 1.2 (December 2022), `base_estimator` renamed to `estimator`. [scikit-learn.org, whats_new/v1.2](https://scikit-learn.org/whats_new/v1.2).
- scikit-learn developers. API reference, `sklearn.ensemble.BaggingClassifier`, versions 1.3 (“`base_estimator ...` will be removed in 1.4”) and 1.4 (no `base_estimator`). [scikit-learn.org, modules/generated/sklearn.ensemble.BaggingClassifier](https://scikit-learn.org/modules/generated/sklearn.ensemble.BaggingClassifier).

8. Key terms

Term	Meaning
<code>BaggingClassifier</code> (G-253)	scikit-learn class for bagging, pasting, random subspaces and random patches in classification
<code>estimator</code>	The base model that bagging copies (formerly <code>base_estimator</code>)
<code>n_estimators</code>	The number of base models in an ensemble
Observation	One record of the data: one row of the data table
Feature	An input variable: one column of the data table
Target	The output we predict
Unstable model (G-2156)	A model whose fit changes a lot when the training data changes a little; bagging helps it most
Out-of-bag (OOB) evaluation, OOB score (G-1411)	Testing a bagging model by predicting each training observation with only the base models that never saw it
<code>max_samples</code>	The number or share of observations each base model gets
<code>bootstrap</code> (G-320)	<code>BaggingClassifier</code> setting: draw observations with replacement (True, bagging) or without (False, pasting)
<code>bootstrap_features</code> (G-321)	<code>BaggingClassifier</code> setting: draw features with replacement or without
<code>estimators_samples__</code>	The row numbers (observations) each trained base model was given
<code>estimators_features__</code>	The column numbers (features) each trained base model was given
<code>verbose</code>	scikit-learn setting that prints progress messages during training