

Feature Importance with Decision Trees and Random Forests

Contents

1. Overview	2
2. What feature importance is for	3
2.1 Feature selection	3
2.2 Interpretability	3
3. Feature importance in code	4
4. How a decision tree computes feature importance	4
4.1 The decrease at one node	4
4.2 From node decreases to column importance	5
4.3 Worked example: 5 rows	5
4.4 Worked example: 15 rows	6
5. Feature importance in a random forest	7
6. The weakness: columns with many unique values	8
7. Permutation importance	9
8. Summary	10
Sources	10
9. Key terms	10

Where this fits: Pipeline map steps 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** Overfitting ([Note 91](#)); Decision trees ([Note 100](#)); Bagging ([Note 107](#)); Hyperparameter tuning ([Note 111](#)); Grid and random search ([Note 112](#)); OOB score ([Note 113](#)).
- **Leads to:** Balanced random forest ([Note 133](#)).
- **Compare with:** Bagging ([Note 107](#)); Dropout ([Note 1024](#)).

1. Overview

Key point: A trained random forest gives every input column a score, its feature importance: the share of all the impurity reduction its splits achieved, averaged over the trees. The scores add up to 1.




Figure 1: MNIST: (a) one handwritten digit; (b) the random forest’s feature importance of each of the 784 pixels. Darker means more important

Tree-based algorithms (decision trees, random forests, and later bagging and boosting ensembles such as AdaBoost, gradient boosting and XGBoost) can all tell us how much each column helped. Figure 1 shows this for handwritten digits: the forest relies on the pixels in the middle of the image and ignores the border.

For a single tree, `feature_importances_` was introduced in the [regression trees Note](#), section 7.3. This Note covers why feature importance is useful, how a tree computes it step by step,

how a random forest combines its trees, and the method's main weakness, with permutation importance as the alternative.

The Notebook (`notebook.ipynb`) runs every example.

2. What feature importance is for

Key point: It tells us which columns a model relies on: useful for dropping weak columns (feature selection) and for explaining the model's decisions.

Feature importance is a number for each input column that says how much the model relied on it. Two uses stand out.

2.1 Feature selection

Key point: Keep the columns with high importance and drop the rest.

Too many columns hurt: training slows down, and useless columns can mislead the model. **Feature selection** keeps only the useful columns (the [feature engineering Note](#)). Feature importance is one way to do it: compute the importance of every column, keep the important ones, drop the others.

Take MNIST (the [PCA on MNIST Note](#)): 42,000 images of handwritten digits, each $28 \times 28 = 784$ pixels, stored as one row with 784 pixel columns plus a label. Digits are written in the middle of the image, so the pixels near the border are almost always blank. They cannot help tell a 0 from a 4; the central pixels can.

We expect a model to rank the central pixels high and the border pixels near zero. That is exactly what Figure 1b shows.

2.2 Interpretability

Key point: Importances help explain to people why a model decides as it does.

Suppose a bank's model rejects a loan application. The customer will ask why. If the model reports which columns matter most, say income, existing debt and credit history, those are the main factors behind its decisions, and the bank can explain them.

A model whose decisions can be explained in this way is **interpretable**. Feature importance gives an overall view: which columns matter to the model in general, not why one particular application was rejected.

3. Feature importance in code

Key point: After fit, `feature_importances_` holds one number per column; on MNIST, 784 numbers that add up to 1.

Python: The importance of every MNIST pixel.

```
rf = RandomForestClassifier(random_state=42, n_jobs=-1)
rf.fit(X, y)                                # 42,000 images, 784 pixels
imp = rf.feature_importances_                # shape (784,), sums to 1
px.imshow(imp.reshape(28, 28))              # back to the image grid
reshape(28, 28) turns the 784 numbers back into a 28 × 28 grid, row by row, the
way the pixels were unrolled. px.imshow (Plotly Express) draws a grid of numbers as a
heatmap.
```

Figure 1b is that heatmap. The darkest pixels, the most important, sit near the centre; the largest single importance is 0.0086. **122 pixels have importance exactly 0:** they lie along the border, where every image is blank, so no tree ever splits on them.

The forest was trained on all 42,000 images, with no test set: here we only want to see what it relies on, not to measure its accuracy.

4. How a decision tree computes feature importance

Key point: Each split’s weighted impurity decrease is credited to the column it splits on; a column’s importance is its total, divided by the total over all columns.

The scikit-learn documentation defines a column’s importance as “the (normalized) total reduction of the criterion brought by that feature”, also called **Gini importance** or **mean decrease in impurity (MDI)**. The steps make this concrete.

4.1 The decrease at one node

Key point: The same weighted decrease Δ that `min_impurity_decrease` uses.

Each split lowers the impurity. Its **weighted impurity decrease** Δ was defined, with a worked example, in the [decision tree hyperparameters Note](#), section 4.8:

$$\Delta = \frac{N_t}{N} \left(G_t - \frac{N_L}{N_t} G_L - \frac{N_R}{N_t} G_R \right)$$

Here N is the number of training rows, N_t the rows reaching the node, N_L and N_R the rows in its left and right children, and G the Gini impurities. The factor N_t/N makes a split near the root, which handles many rows, count more than a split deep down.

4.2 From node decreases to column importance

Key point: Add up the Δ of every node that splits on the column, then divide by the Δ of all nodes.

1. **In words:** a column's importance is its share of the tree's total impurity decrease.
2. **Formula:**

$$\text{importance}(j) = \frac{\text{sum of } \Delta \text{ over the nodes that split on column } j}{\text{sum of } \Delta \text{ over all split nodes}}$$

3. **Example:** see section 4.3: column 0 gets $0.30/0.48 = 0.625$.

Dividing by the total **normalizes** the importances: they always add up to 1.

4.3 Worked example: 5 rows

Key point: The root, on column 1, decreases impurity by 0.18; the second split, on column 0, by 0.30; so column 0 has importance 0.625 and column 1 0.375.




Figure 2: A tree trained on 5 rows and 2 columns, with the impurity decrease of each split. The square brackets give the number of rows of class 0 and class 1

We make a tiny dataset of 5 rows and 2 columns (column 0 and column 1) with `make_classification` and train a fully grown tree. Figure 2 shows it: two split nodes and three pure leaves.

The root splits on column 1. It holds all 5 rows (3 of class 0, 2 of class 1), Gini 0.48. Its left child has 1 row (Gini 0), its right child 4 rows (Gini 0.375):

$$\Delta_{\text{root}} = \frac{5}{5} \left(0.48 - \frac{4}{5}(0.375) - \frac{1}{5}(0) \right) = 0.48 - 0.30 = 0.18$$

The second node splits on column 0. It holds 4 of the 5 rows, Gini 0.375, and both children are pure:

$$\Delta_{\text{node 2}} = \frac{4}{5} \left(0.375 - \frac{3}{4}(0) - \frac{1}{4}(0) \right) = 0.8 \times 0.375 = 0.30$$

The total decrease is $0.18 + 0.30 = 0.48$, so:

- column 0: $0.30/0.48 = 0.625$;
- column 1: $0.18/0.48 = 0.375$.

These are exactly the values of `tree.feature_importances_`. Column 0 wins although column 1 makes the first split: the second split removes more impurity.

4.4 Worked example: 15 rows

Key point: When a column splits several nodes, its decreases are added: column 0 splits two nodes ($0.119 + 0.089$) and gets 0.417.

With 15 rows, the tree has three split nodes: the root on column 1, and two lower nodes on column 0.

Node	Column	Rows	Gini	Δ
root	1	15	0.498	0.290
2	0	9	0.346	0.119
4	0	3	0.444	0.089

Each Δ comes from the formula of section 4.1 (leaves and pure children contribute 0, so they drop out):

$$\Delta_{\text{root}} = \frac{15}{15} \left(0.498 - \frac{9}{15}(0.346) \right) = 0.290$$

$$\Delta_2 = \frac{9}{15} \left(0.346 - \frac{3}{9}(0.444) \right) = 0.119$$

$$\Delta_4 = \frac{3}{15}(0.444) = 0.089$$

The total is $0.290 + 0.119 + 0.089 = 0.498$, so column 0 gets $(0.119 + 0.089)/0.498 = 0.417$ and column 1 gets $0.290/0.498 = 0.583$. `feature_importances_` gives 0.4167 and 0.5833; rounding each step to two digits gives slightly different numbers, such as 0.39 and 0.61, so we keep three.

5. Feature importance in a random forest

Key point: Every tree computes its own importances; the forest takes their mean.

A random forest adds nothing new: each tree computes its importances as in section 4, and the forest averages them.

1. **In words:** a column's importance in the forest is the mean of its importance in each tree.
2. **Formula:** for a forest of T trees,

$$\text{importance}_{\text{forest}}(j) = \frac{1}{T} \sum_{t=1}^T \text{importance}_t(j)$$

3. **Example:** a forest of 2 trees on the 5-row data. Tree 1 gives (0.625, 0.375); tree 2 splits only on column 1 and gives (0, 1):

$$\text{column 0 : } \frac{0.625 + 0}{2} = 0.3125 \quad \text{column 1 : } \frac{0.375 + 1}{2} = 0.6875$$

`rf.feature_importances_` gives exactly (0.3125, 0.6875).

Python: Each tree's importances, and their mean.

```
rf2 = RandomForestClassifier(n_estimators=2, random_state=3)
rf2.fit(X5, y5)
for t in rf2.estimators_:
    print(t.feature_importances_) # [0.625 0.375], [0. 1.]
rf2.feature_importances_        # [0.3125 0.6875]
```

A single tree's importances can change a lot when the data changes slightly (high variance); the forest's mean over many trees is much more stable. In the Notebook, retraining on 20 resamples of the data of section 6, the forest's importances spread about a third as much as one tree's: the same variance reduction that makes bagging work.

Extra: Two details of scikit-learn's version. Each tree computes its decreases on its own bootstrap sample, not on the full training set. And a tree that never split (only a root) is left out of the mean; the result is divided by its total again so it still sums to 1. (scikit-learn source, `ensemble/_forest.py`)

6. The weakness: columns with many unique values

Key point: Impurity-based importance favours columns with many different values: a column of pure noise with 1,000 unique values gets 10% of the importance.

The scikit-learn documentation warns that impurity-based feature importances “can be misleading for high cardinality features (many unique values)”. High cardinality (the [pandas profiling Note](#)) here means a column with very many different values, numeric or categorical. A test shows why.




Figure 3: Two useless columns added to four real ones: random_id (1,000 unique values) and random_coin (2 values). (a) Impurity-based importance; (b) permutation importance on the test set

We make a dataset of 1,000 rows with four real columns, x1 to x4 (52 to 75 unique values each), and add two columns of noise that have nothing to do with the target:

- **random_id:** a different number for every row, 1,000 unique values;
- **random_coin:** 0 or 1 at random, 2 unique values.

A random forest trained on 700 rows scores 0.87 on the other 300. Its impurity-based importances (Figure 3a) give **random_id 0.100**, more than half as much as the real column x4, while **random_coin** gets only 0.015.

The reason: importances are computed from the **training** data, where trees are grown until every leaf is pure. A column with 1,000 different values offers about 1,000 possible thresholds, so deep in a tree it can almost always separate the last few rows by chance. Each such split lowers the training impurity and earns the column credit, though it means nothing on new data. A column with 2 values offers one threshold, so it has far fewer chances to fit noise. The bias towards features with many values is well known (Strobl et al., 2007; scikit-learn User Guide §5.2).

Extra: A test of this explanation. If `random_id` earns its credit from the deepest splits, forcing bigger leaves should shrink it. The Notebook changes only `min_samples_leaf`: with leaves of at least 1, 5, 20 and 50 rows, `random_id`'s importance falls from 0.100 to 0.051, 0.032 and 0.021, while the test accuracy barely moves (0.870 to 0.847).

7. Permutation importance

Key point: Shuffle one column in the test set and measure how much the score drops; a column the model truly uses causes a big drop, a useless column none.

For data with high-cardinality columns, scikit-learn recommends **permutation importance**, `sklearn.inspection.permutation_importance`, instead. Figure 3b shows it for the same forest: the four real columns keep clear importances (0.055 to 0.210), and both noise columns fall to about 0 (`random_id` -0.003, `random_coin` 0.004).

Python: Permutation importance on the test set.

```
from sklearn.inspection import permutation_importance

perm = permutation_importance(rf, X_test, y_test,
                              n_repeats=20, random_state=0)

perm.importances_mean      # one number per column
perm.importances_std       # their spread over the 20 shuffles
```

Extra: How permutation importance works.

1. Score the trained model on the test set: accuracy 0.870.
2. Shuffle one column's values among the test rows, which breaks its link with the target while keeping its values. Score again: with `x1` shuffled, the accuracy drops to 0.720.
3. The drop, here 0.15 for this one shuffle, is that column's importance. Repeat with new shuffles (`n_repeats`) and average; then do the same for every column.

A noise column like `random_id` was never truly used, so shuffling it changes almost nothing; its importance is near 0 and can even come out slightly negative by chance. Because it is measured on unseen data, permutation importance cannot be fooled by splits that only fit the training set (scikit-learn User Guide §5.2).

It has costs: it needs a test set and many extra predictions, so it is slower. And when two columns are strongly correlated, shuffling one barely hurts the model, because the other still carries the same information, so both can look unimportant (scikit-learn User Guide §5.2.3).

If the data has no high-cardinality columns, the impurity-based importance of `feature_importances_` works well and is free: it is computed during training.

8. Summary

	Impurity-based (MDI)	Permutation
Where Measures	<code>feature_importances_</code> share of the impurity decrease from each column's splits	<code>permutation_importance</code> drop in score when a column is shuffled
Computed on	the training data, during <code>fit</code>	a test set, after training
Cost	free	many extra predictions
Noise column with 1,000 values	0.100 (misleading)	-0.003 (correct)

- Feature importance scores how much a model relies on each column; it helps with feature selection and interpretability.
- In a tree, each split's weighted impurity decrease Δ is credited to its column; a column's importance is its share of the total (5 rows: $0.30 / 0.48 = 0.625$).
- A random forest averages its trees' importances, which makes them more stable.
- Impurity-based importance favours high-cardinality columns; permutation importance on a test set does not.

Sources

- Strobl, C., Boulesteix, A.-L., Zeileis, A. and Hothorn, T. (2007). Bias in random forest variable importance measures. *BMC Bioinformatics* 8: 25.
- scikit-learn developers. User Guide, section 5.2, "Permutation feature importance" (version 1.9). https://scikit-learn.org/stable/modules/permutation_importance.html
- scikit-learn source code, `sklearn/ensemble/_forest.py`, property `feature_importances_` (version 1.9).

9. Key terms

Term	Meaning
Mean decrease in impurity (MDI)	Impurity-based feature importance: a column's share of the total weighted impurity decrease of its splits, averaged over the trees; also called Gini importance
Weighted impurity decrease	A split's impurity drop, weighted by the share of rows reaching the node: the <code>min_impurity_decrease</code>
Interpretability	How well people can understand why a model makes its decisions
Permutation importance	The drop in a model's test score when one column's values are shuffled

Term	Meaning
permutation_importance	scikit-learn function (in <code>sklearn.inspection</code>) that computes permutation importance
Normalized importances	Importances divided by their total, so they add up to 1