


Feature Importance with Decision Trees and Random Forests

Contents

1. Overview	2
2. What feature importance is for	3
2.1 Feature selection	3
2.2 Interpretability	3
3. Feature importance in code	4
4. How a decision tree computes feature importance	4
4.1 The decrease at one node	4
4.2 From node decreases to feature importance	5
4.3 Worked example: 5 observations	5
4.4 Worked example: 15 observations	6
5. Feature importance in a random forest	8
6. The weakness: features with many unique values	9
7. Permutation importance	11
8. Summary	13
9. Sources	13
10. Key terms	13

Where this fits: Pipeline map steps 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** Overfitting ([Note 91](#)); Decision trees ([Note 100](#)); Bagging ([Note 107](#)); Hyperparameter tuning ([Note 111](#)); Grid and random search ([Note 112](#)); OOB score ([Note 113](#)).
- **Leads to:** Balanced random forest ([Note 133](#)).
- **Compare with:** Bagging ([Note 107](#)); Dropout ([Note 1024](#)).

1. Overview

Key point: A trained random forest gives every **feature** (input variable, one column of the data table) a score, its feature importance: the share of all the impurity reduction its splits achieved, averaged over the trees. The scores add up to 1.




Figure 1: MNIST: (a) one handwritten digit; (b) the random forest's feature importance of each of the 784 pixels. Darker means more important

Tree-based algorithms (decision trees, random forests, and later bagging and boosting ensembles such as AdaBoost, gradient boosting and XGBoost) can all tell us how much each feature helped in predicting the **target** (the output we predict) from the **observations** (records, one row of the table each). Figure 1 shows this for handwritten digits: the forest relies on the pixels in the middle of the image and ignores the border.

For a single tree, `feature_importances_` was introduced in the [regression trees Note](#), section 7.3. This Note covers:

- why feature importance is useful (section 2);
- how a tree computes it, step by step (section 4);
- how a random forest combines its trees (section 5);
- the method's main weakness (section 6), with permutation importance as the alternative (section 7).

The Notebook (`notebook.ipynb`) runs every example.

2. What feature importance is for

Key point: Feature importance tells us which features a model relies on: useful for dropping weak features (feature selection) and for explaining the model's decisions.

Feature importance (G-764) is a number for each input feature that says how much the model relied on it. Two uses stand out.

2.1 Feature selection

Key point: Keep the features with high importance and drop the rest.

Too many features hurt: training slows down, and useless features can mislead the model. **Feature selection** (G-768) keeps only the useful features (the [feature engineering Note](#)). Feature importance is one way to do it: compute the importance of every feature, keep the important ones, drop the others.

Take MNIST (the [PCA on MNIST Note](#)): 42,000 images of handwritten digits, each $28 \times 28 = 784$ pixels, stored as one row of 784 pixel columns plus a label. Digits are written in the middle of the image, so the pixels near the border are almost always blank. Border pixels cannot help tell a 0 from a 4; the central pixels can.

We expect a model to rank the central pixels high and the border pixels near zero. Figure 1b shows exactly that pattern.

2.2 Interpretability

Key point: Importances help explain to people why a model decides as it does.

Suppose a bank's model rejects a loan application. The customer will ask why. If the model reports which features matter most, say income, existing debt and credit history, those are the main factors behind its decisions, and the bank can explain them.

A model whose decisions can be explained in this way is **interpretable**; the property is **interpretability** (G-965). Feature importance gives an overall view: which features matter to the model in general, not why one particular application was rejected.

3. Feature importance in code

Key point: After fit, `feature_importances_` holds one number per feature; on MNIST, 784 numbers that add up to 1.

Python: The importance of every MNIST pixel.

```
rf = RandomForestClassifier(random_state=42, n_jobs=-1)
rf.fit(X, y)                                # 42,000 images, 784 pixels
imp = rf.feature_importances_               # shape (784,), sums to 1
px.imshow(imp.reshape(28, 28))             # back to the image grid
reshape(28, 28) turns the 784 numbers back into a 28 × 28 grid, row by row, the
way the pixels were unrolled. px.imshow (Plotly Express) draws a grid of numbers as a
heatmap.
```

Figure 1b is that heatmap. The darkest pixels, the most important, sit near the centre; the largest single importance is 0.0086. **122 pixels have importance exactly 0:** they lie along the border, where every image is blank, so no tree ever splits on them.

The forest was trained on all 42,000 images, with no test set: here we only want to see what it relies on, not to measure its accuracy.

4. How a decision tree computes feature importance

Key point: Each split’s weighted impurity decrease is credited to the feature it splits on; a feature’s importance is its total, divided by the total over all features.

The scikit-learn documentation defines a feature’s importance as “the (normalized) total reduction of the criterion brought by that feature”, also called **Gini importance** or **mean decrease in impurity (MDI)** (G-1196). The steps make this concrete.

4.1 The decrease at one node

Key point: The same weighted decrease Δ that `min_impurity_decrease` uses.

Every split makes its node’s observations less mixed, and it gets credit for how much it cleaned up, counted in proportion to how many observations reached it. A split near the root handles all the observations, so its clean-up counts in full; a split deep down handles only a few, so its clean-up counts for little.

The measure of “mixed” is the impurity (**Gini impurity**, G-847), and the credit is the split’s **weighted impurity decrease** (G-2115) Δ , defined, with a worked example, in the [decision tree hyperparameters Note](#), section 4.8:

$$\Delta = \frac{N_t}{N} \left(G_t - \frac{N_L}{N_t} G_L - \frac{N_R}{N_t} G_R \right)$$

Here N is the number of training observations, N_t the observations reaching the node, N_L

and N_R the observations in its left and right children, and G the Gini impurities. The factor N_t/N makes a split near the root, which handles many observations, count more than a split deep down.

4.2 From node decreases to feature importance

Key point: Add up the Δ of every node that splits on the feature, then divide by the Δ of all nodes.

1. **In words:** a feature's importance is its share of the tree's total impurity decrease.
2. **Formula:**

$$\text{importance}(j) = \frac{\text{sum of } \Delta \text{ over the nodes that split on column } j}{\text{sum of } \Delta \text{ over all split nodes}}$$

3. **Example:** see section 4.3: feature 0 gets $0.30/0.48 = 0.625$.

Dividing by the total **normalizes** the importances: they always add up to 1.

4.3 Worked example: 5 observations

Key point: The root, on feature 1, decreases impurity by 0.18; the second split, on feature 0, by 0.30; so feature 0 has importance 0.625 and feature 1 0.375.




Figure 2: A tree trained on 5 observations and 2 features, with the impurity decrease of each split. The square brackets give the number of observations of class 0 and class 1

We make a tiny dataset of 5 observations and 2 features (feature 0 and feature 1) with `make_classification` and train a fully grown tree. Figure 2 shows it: two split nodes and three pure leaves.

The root splits on feature 1. The root holds all 5 observations (3 of class 0, 2 of class 1), Gini 0.48. Its left child has 1 observation (Gini 0), its right child 4 observations (Gini 0.375):

$$\Delta_{\text{root}} = \frac{5}{5} \left(0.48 - \frac{4}{5}(0.375) - \frac{1}{5}(0) \right) = 0.48 - 0.30 = 0.18$$

The second node splits on feature 0. The node holds 4 of the 5 observations, Gini 0.375, and both children are pure:

$$\Delta_{\text{node 2}} = \frac{4}{5} \left(0.375 - \frac{3}{4}(0) - \frac{1}{4}(0) \right) = 0.8 \times 0.375 = 0.30$$

The total decrease is $0.18 + 0.30 = 0.48$, so:

- feature 0: $0.30/0.48 = 0.625$;
- feature 1: $0.18/0.48 = 0.375$.

These are exactly the values of `tree.feature_importances_`. Feature 0 wins although feature 1 makes the first split: the second split removes more impurity.

4.4 Worked example: 15 observations

Key point: When a feature splits several nodes, its decreases are added: feature 0 splits two nodes ($0.119 + 0.089$) and gets 0.417.

With 15 observations, the tree has three split nodes: the root on feature 1, and two lower nodes on feature 0.

Figure 3 adds up the credit split by split. Watch the blue bar of feature 0 grow twice, once for each of its two nodes, before both bars are divided by the total.




Figure 3: Feature importance built up on the 15-observation tree: each split's decrease is added to the bar of the feature it splits on; the two totals are then divided by their sum, 0.498

Node	Feature	Observations	Gini	Δ
root	1	15	0.498	0.290
2	0	9	0.346	0.119
4	0	3	0.444	0.089

Each Δ comes from the formula of section 4.1 (leaves and pure children contribute 0, so they drop out):

$$\Delta_{\text{root}} = \frac{15}{15} \left(0.498 - \frac{9}{15}(0.346) \right) = 0.290$$

$$\Delta_2 = \frac{9}{15} \left(0.346 - \frac{3}{9}(0.444) \right) = 0.119$$

$$\Delta_4 = \frac{3}{15}(0.444) = 0.089$$

The total is $0.290 + 0.119 + 0.089 = 0.498$, so feature 0 gets $(0.119 + 0.089)/0.498 = 0.417$ and feature 1 gets $0.290/0.498 = 0.583$. `feature_importances_` gives 0.4167 and 0.5833; rounding each step to two digits gives slightly different numbers, such as 0.39 and 0.61, so we keep three.

5. Feature importance in a random forest

Key point: Every tree computes its own importances; the forest takes their mean.

A random forest adds nothing new: each tree computes its importances as in section 4, and the forest averages them.

1. **In words:** a feature's importance in the forest is the mean of its importance in each tree.
2. **Formula:** for a forest of T trees,

$$\text{importance}_{\text{forest}}(j) = \frac{1}{T} \sum_{t=1}^T \text{importance}_t(j)$$

3. **Example:** a forest of 2 trees on the 5-observation data. Tree 1 gives (0.625, 0.375); tree 2 splits only on feature 1 and gives (0, 1):

$$\text{column 0 : } \frac{0.625 + 0}{2} = 0.3125 \quad \text{column 1 : } \frac{0.375 + 1}{2} = 0.6875$$

`rf.feature_importances_` gives exactly (0.3125, 0.6875).

Python: Each tree's importances, and their mean.

```
rf2 = RandomForestClassifier(n_estimators=2, random_state=3)
rf2.fit(X5, y5)
for t in rf2.estimators_:
    print(t.feature_importances_) # [0.625 0.375], [0. 1.]
rf2.feature_importances_        # [0.3125 0.6875]
```

A single tree's importances can change a lot when the data changes slightly (high variance); the forest's mean over many trees is much more stable. The Notebook tests this on the data of section 6. It draws 20 bootstrap resamples of the 700 training observations and, on each one, trains one tree and one forest. Figure 4 plots every resample's importances. Watch the red dots of the single tree spread widely along each feature's row, while the blue dots of the forest sit in tight clusters. Measured by the standard deviation, averaged over the six features, the forest's importances spread about a quarter as much as one tree's (0.011 against 0.043): the same variance reduction that makes bagging work.




Figure 4: Feature importances from 20 bootstrap resamples of the training data. Left: one decision tree per resample. Right: one random forest of 100 trees per resample. Each dot is one resample

Extra: Two details of scikit-learn’s version. Each tree computes its decreases on its own bootstrap sample, not on the full training set. And a tree that never split (only a root) is left out of the mean; the result is divided by its total again so it still sums to 1. (scikit-learn source, `ensemble/_forest.py`)

6. The weakness: features with many unique values

Key point: Impurity-based importance favours features with many different values: a feature of pure noise with 1,000 unique values gets 10% of the importance.

The scikit-learn documentation warns that impurity-based feature importances “can be misleading for high cardinality features (many unique values)”. High cardinality (the [pandas profiling Note](#)) here means a feature with very many different values, numeric or categorical. A test shows why.




Figure 5: Two useless features added to four real ones: `random_id` (1,000 unique values) and `random_coin` (2 values). (a) Impurity-based importance; (b) permutation importance on the test set

We make a dataset of 1,000 observations with four real features, `x1` to `x4` (52 to 75 unique values each), and add two features of noise that have nothing to do with the target:

- `random_id`: a different number for every observation, 1,000 unique values;
- `random_coin`: 0 or 1 at random, 2 unique values.

A random forest trained on 700 observations scores 0.87 on the other 300. Its impurity-based importances (Figure 5a) give `random_id` **0.100**, more than half as much as the real feature `x4`, while `random_coin` gets only 0.015.

The reason: importances are computed from the **training** data, where trees are grown until every leaf is pure. A feature with 1,000 different values offers about 1,000 possible thresholds, so deep in a tree it can almost always separate the last few observations by chance. Each such split lowers the training impurity and earns the feature credit, though it means nothing on new data. A feature with 2 values offers one threshold, so it has far fewer chances to fit noise. The bias towards features with many values is well known (Strobl et al., 2007; scikit-learn User Guide §5.2).

Extra: A test of this explanation. If `random_id` earns its credit from the deepest splits, forcing bigger leaves should shrink it. The Notebook changes only `min_samples_leaf`: with leaves of at least 1, 5, 20 and 50 observations, `random_id`'s importance falls from 0.100 to 0.051, 0.032 and 0.021, while the test accuracy barely moves (0.870 to 0.847).

7. Permutation importance

Key point: Shuffle one feature in the test set and measure how much the score drops; a feature the model truly uses causes a big drop, a useless feature none.

For data with high-cardinality features, scikit-learn recommends **permutation importance** (G-1490), `sklearn.inspection.permutation_importance`, instead. Figure 5b shows it for the same forest: the four real features keep clear importances (0.055 to 0.210), and both noise features fall to about 0 (`random_id` -0.003, `random_coin` 0.004).

Python: Permutation importance on the test set.

```
from sklearn.inspection import permutation_importance

perm = permutation_importance(rf, X_test, y_test,
                             n_repeats=20, random_state=0)
perm.importances_mean      # one number per column
perm.importances_std       # their spread over the 20 shuffles
```

Extra: How permutation importance works.

1. Score the trained model on the test set: accuracy 0.870.
2. Shuffle one feature's values among the test observations, which breaks its link with the target while keeping its values. Score again: with `x1` shuffled, the accuracy drops to 0.720.
3. The drop, here 0.15 for this one shuffle, is that feature's importance. Repeat with new shuffles (`n_repeats`) and average; then do the same for every feature.

A noise feature like `random_id` was never truly used, so shuffling it changes almost nothing; its importance is near 0 and can even come out slightly negative by chance. Because it is measured on unseen data, permutation importance cannot be fooled by splits that only fit the training set (scikit-learn User Guide §5.2).

Permutation importance has costs: it needs a test set and many extra predictions, so it is slower. And when two features are strongly correlated, shuffling one barely hurts the model, because the other still carries the same information, so both can look unimportant (scikit-learn User Guide §5.2.3).

Figure 6 animates the procedure of the Extra box. Each dot is the test accuracy after one shuffle of one feature. Watch the dots of `x1` land far left of the dashed line at 0.870, a drop of 0.210 on average, while the dots of `random_id` and `random_coin` stay on the line: shuffling a feature the model never truly used changes nothing.




Figure 6: Permutation importance step by step: each feature is shuffled 20 times in the 300 test observations. Dots: test accuracy after each shuffle; dashed: accuracy with no shuffle (0.870); right: the mean drop, the feature’s permutation importance

If the data has no high-cardinality features, the impurity-based importance of `feature_importances_` is a quick first look that costs nothing: it is computed during training.

8. Summary

	Impurity-based (MDI)	Permutation
Where Measures	<code>feature_importances_</code> share of the impurity decrease from each feature's splits	<code>permutation_importance</code> drop in score when a feature is shuffled
Computed on	the training data, during <code>fit</code>	a test set, after training
Cost	free	many extra predictions
Noise feature with 1,000 values	0.100 (misleading)	-0.003 (correct)

- Feature importance scores how much a model relies on each feature; it helps with feature selection and interpretability.
- In a tree, each split's weighted impurity decrease Δ is credited to its feature; a feature's importance is its share of the total (5 observations: $0.30 / 0.48 = 0.625$).
- A random forest averages its trees' importances, which makes them more stable.
- Impurity-based importance favours high-cardinality features; permutation importance on a test set does not.

9. Sources

Built from

- CampusX, "Feature Importance using Random Forest and Decision Trees | How is Feature Importance calculated", YouTube, <https://www.youtube.com/watch?v=R47JAob1xBY>

Other references

- Strobl, C., Boulesteix, A.-L., Zeileis, A. and Hothorn, T. (2007). Bias in random forest variable importance measures. *BMC Bioinformatics* 8: 25.
- scikit-learn developers. User Guide, §5.2 "Permutation feature importance" (version 1.9). scikit-learn.org/stable/modules/permutation_importance.html
- scikit-learn source code, `sklearn/ensemble/_forest.py`, property `feature_importances_` (version 1.9).

10. Key terms

Term	Meaning
Mean decrease in impurity (MDI)	Impurity-based feature importance: a feature's share of the total weighted impurity decrease of its splits, averaged over the trees; also called Gini importance

Term	Meaning
Weighted impurity decrease	A split's impurity drop, weighted by the share of observations reaching the node: the Δ of <code>min_impurity_decrease</code>
Interpretability	How well people can understand why a model makes its decisions
Permutation importance	The drop in a model's test score when one feature's values are shuffled
<code>permutation_importance</code>	scikit-learn function (in <code>sklearn.inspection</code>) that computes permutation importance
Normalized importances	Importances divided by their total, so they add up to 1