

AdaBoost Step by Step

Contents

1. Overview	2
2. The toy data	3
3. Step 1: every observation gets the same weight	4
4. Step 2: fit a decision stump	4
5. Step 3: the error is a sum of weights	4
6. Step 4: the say of the stump, alpha	5
6.1 What shape alpha should have	5
6.2 The formula	6
7. Step 5: update the observation weights	7
8. Step 6: normalise the weights	8
9. Step 7: upsampling, a new dataset drawn by weight	9
10. Repeat, then vote	10
11. Summary	12
12. Sources	12
13. Key terms	12

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Decision trees ([Note 100](#)); Boosting ([Note 101](#)); Hyperparameter tuning ([Note 111](#)).
- **Compare with:** Gradient boosting ([Note 120](#)).

1. Overview

Key point: One AdaBoost stage, in seven steps: give the observations weights, fit a stump, add up the weights it got wrong, turn that error into a say alpha, raise the weights of the mistakes and lower the rest, normalise, and draw a new dataset by weight.




Figure 1: One stage of AdaBoost; the new dataset starts the next stage

We met the idea in the [AdaBoost intuition Note](#): stumps trained one after another, each focusing on the previous one's mistakes, combined by a weighted vote. The intuition Note left two questions open:

- how the say α of each stump is computed;
- how the mistakes are made “more important” for the next stump.

The present Note answers both on a toy dataset of 5 observations, following Figure 1 step by step. The [AdaBoost from scratch Note](#) runs the same steps in Python.

2. The toy data

Key point: 5 observations, two features x_1 and x_2 , and a target y of $+1$ or -1 .

Each **observation** is one record, one row of the table below. Each **feature** is an input variable, one column (x_1 , x_2). The **target** y is the output we predict.

Row	x_1	x_2	y
1	3	7	$+1$
2	2	9	$+1$
3	1	9	-1
4	9	8	-1
5	7	4	-1

The output has two classes, so this is classification. The number of observations is $n = 5$. As in the intuition Note, the classes are written $+1$ and -1 .

Figure 2 places the 5 observations in the plane of the two features. Watch where the two classes sit: the $+1$ observations 1 and 2 lie at the left, next to the -1 observation 3, so no single cut on x_1 or x_2 puts every observation on its class's side.




Figure 2: The toy data: each observation at its (x_1, x_2) , numbered as in the table; blue $y = +1$, orange $y = -1$

3. Step 1: every observation gets the same weight

Key point: At the start every observation has weight $1/n$, so all observations are equally important and the weights add up to 1.

AdaBoost gives each observation a **sample weight** (G-1730): a number saying how important that observation is. At the start all observations are equally important.

1. **In words:** each observation's starting weight is 1 divided by the number of observations.
2. **Formula:**

$$w_i = \frac{1}{n}$$

3. **Example:** with $n = 5$, every observation gets $w_i = 1/5 = 0.2$, and $0.2 \times 5 = 1$.

The weights always add up to 1. We add them to the table as a new column, *weight*.

4. Step 2: fit a decision stump

Key point: Train a decision stump on the data and record its prediction for every observation.

We train a **decision stump** (G-559): a decision tree with `max_depth=1` (the intuition Note, section 2.2). The stump tries every cut on x_1 and on x_2 , for example “ $x_1 > 5$ ” or “ $x_2 < 10$ ”, and keeps the one with the largest information gain (or the largest drop in Gini impurity: the [decision trees Note](#) compares the two). Call the result **model 1**.

We then pass the training observations through model 1 and write its predictions in a new column. Suppose they are:

Row	y	weight	prediction	correct?
1	+1	0.2	+1	yes
2	+1	0.2	-1	no
3	-1	0.2	+1	no
4	-1	0.2	-1	yes
5	-1	0.2	-1	yes

Model 1 gets observations 2 and 3 wrong. The predictions are an assumed example: the steps that follow are the same for any stump.

5. Step 3: the error is a sum of weights

Key point: A stump's error is the total weight of the observations it misclassified, not simply the share of observations it got wrong.

1. **In words:** add up the weights of the misclassified observations. The total is the **weighted error** (G-2114).

2. Formula:

$$\text{error} = \sum_{i: \hat{y}_i \neq y_i} w_i$$

The sum runs over the observations where the prediction $\hat{y}_i$ differs from the true class y_i .

3. **Example:** observations 2 and 3 are wrong, each with weight 0.2:

$$\text{error} = 0.2 + 0.2 = 0.4$$

In this first stage all weights are equal, so the error equals the share of observations wrong, 2 out of 5. In later stages the weights differ, and a mistake on a heavy observation costs more than a mistake on a light one.

6. Step 4: the say of the stump, alpha

Key point: Alpha, half the log of $(1 - \text{error})/\text{error}$, is large and positive for a small error, 0 at error 0.5, and large and negative for an error near 1.

The **model weight** α (**alpha**, G-192) is the stump's say in the final vote. The say must depend on the error: a model that makes many mistakes should have a small say, and a model that makes few a large one.

6.1 What shape alpha should have

Key point: A model that is always wrong is as useful as one that is always right: we just flip its answer. The useless model is the one that is right half the time.

Take three models:

- **model A** is always right: error 0, accuracy 100%;
- **model B** is always wrong: error 1, accuracy 0%;
- **model C** is right half the time: error 0.5, accuracy 50%.

(The error is 1 minus the accuracy: 90% accuracy means an error of 0.1.)

A is clearly the most reliable. Between B and C, B is the more reliable. Think of two friends: one lies every time, the other lies half the time. Whatever the first one says, the opposite is true, so we can trust that friend completely by flipping every answer. The second one is no help at all: we never know when to believe them.

So we want a function of the error that:

- is large and positive when the error is near 0 (trust the model);
- is 0 when the error is 0.5 (ignore the model);
- is large and negative when the error is near 1 (trust the model, but flip its vote).

6.2 The formula

Key point: Half the natural log of the ratio of correct weight to wrong weight.




Figure 3: Alpha against the error: model A (error near 0) gets a large positive say, model C (error 0.5) none, model B (error near 1) a large negative say; red: the stump of this Note; green: two stumps with fewer mistakes

The standard AdaBoost formula (Schapire 2013, Algorithm 1), plotted in Figure 3, has exactly that shape.

1. **In words:** divide the weight the stump got right by the weight it got wrong, take the natural logarithm, and halve it.
2. **Formula:**

$$\alpha = \frac{1}{2} \ln \left(\frac{1 - \text{error}}{\text{error}} \right)$$

$\ln$ is the natural logarithm, the log to base $e \approx 2.718$ (logs were introduced in the [log loss Note](#), section 4).

3. **Example:** model 1 has error 0.4:

$$\alpha_1 = \frac{1}{2} \ln \left(\frac{1 - 0.4}{0.4} \right) = \frac{1}{2} \ln(1.5) = \frac{1}{2} \times 0.405 = 0.20$$

So model 1's say in the final vote is $\alpha_1 = 0.20$, a small say, since it got 40% wrong (Figure 3, red point). Fewer mistakes earn a larger say: a stump with error 0.3 gets $\alpha = 0.42$ and a stump with error 0.1 gets $\alpha = 1.10$ (Figure 3, green points; these two stumps appear in the [AdaBoost from scratch Note](#)). The checks: an error of 0.5 gives $0.5 \times \ln 1 = 0$, and an error of 0.98 gives $0.5 \times \ln(0.02/0.98) = -1.95$, the mirror image of an error of 0.02.

Extra: A negative alpha flips the stump’s vote in the final sum, which is the “believe the liar backwards” idea. In practice a stump that is worse than guessing is rarely kept: scikit-learn stops adding stumps when a new one’s error reaches 0.5 or more on two classes (scikit-learn source; the [AdaBoost hyperparameters Note](#)).

7. Step 5: update the observation weights

Key point: Multiply the weight of each misclassified observation by e^α (it grows) and of each correct observation by $e^{-\alpha}$ (it shrinks).

Now we tell the next stump about the mistakes, by **boosting** the weights of the misclassified observations and lowering the rest: the **weight update** (G-2110). Raising these weights is where boosting gets its name.

1. **In words:** a misclassified observation’s weight is multiplied by e to the power alpha; a correctly classified observation’s weight by e to the power minus alpha.
2. **Formula:**

$$w_i^{\text{new}} = \begin{cases} w_i e^\alpha & \text{if observation } i \text{ was misclassified} \\ w_i e^{-\alpha} & \text{if observation } i \text{ was classified correctly} \end{cases}$$

3. **Example:** with $w_i = 0.2$ and $\alpha_1 = 0.2027$:

$$\text{misclassified: } 0.2 \times e^{0.2027} = 0.2 \times 1.2247 = 0.2449$$

$$\text{correct: } 0.2 \times e^{-0.2027} = 0.2 \times 0.8165 = 0.1633$$

Observations 2 and 3 rise from 0.2 to about 0.24; observations 1, 4 and 5 fall to about 0.16. Why the exponential is the right choice is shown in the [AdaBoost from scratch Note](#).

Figure 4 runs steps 1 to 7 on these weights. Watch the two red bars, the mistakes, grow above the dotted line at 0.2 while the green bars shrink below it. Its last frames show step 7 (section 9).



Figure 4: The 5 sample weights through one stage of the worked example: equal at 0.2, observations 2 and 3 misclassified (error 0.4, α 0.2027), multiplied by e^α or $e^{-\alpha}$, normalised, then laid end to end on the line from 0 to 1, where five random numbers pick the new dataset

Extra: With classes +1 and -1, both cases fit in one formula: $w_i^{\text{new}} = w_i e^{-\alpha y_i h(x_i)}$. If the stump is right, $y_i h(x_i) = (+1)(+1)$ or $(-1)(-1) = +1$, giving $e^{-\alpha}$. If it is wrong, $y_i h(x_i) = -1$, giving e^α . The single formula is another reason AdaBoost uses +1 and -1.

8. Step 6: normalise the weights

Key point: Divide every new weight by their total, so the weights add up to 1 again.

After the update the weights no longer add up to 1.

1. **In words:** divide each weight by the sum of all the weights.
2. **Formula:**

$$w_i \leftarrow \frac{w_i^{\text{new}}}{\sum_{j=1}^n w_j^{\text{new}}}$$

3. **Example:** the sum is $2 \times 0.2449 + 3 \times 0.1633 = 0.4898 + 0.4899 = 0.9798$, so

$$\text{misclassified: } \frac{0.2449}{0.9798} = 0.25 \quad \text{correct: } \frac{0.1633}{0.9798} = 0.1667$$

Check: $2 \times 0.25 + 3 \times 0.1667 = 1$. Figure 4 shows the normalised weights at the end of step 6. The two mistakes now carry half of the total weight between them, against 40% before.

Extra: Here the misclassified observations end up with exactly half the total weight. The half-and-half split holds after every normalised AdaBoost update. Before normalising, the mistakes weigh $\text{error} \cdot e^\alpha$ in total and the correct observations $(1 - \text{error}) \cdot e^{-\alpha}$. With

$e^\alpha = \sqrt{(1 - \text{error})/\text{error}}$ from step 4, both totals equal $\sqrt{\text{error}(1 - \text{error})}$: in our example $\sqrt{0.4 \times 0.6} = 0.4899$, the two equal halves of the sum above. So the old stump, judged on the new weights, has an error of exactly 0.5 and would get $\alpha = 0$: repeating it adds nothing to the vote, and the next stump only earns a say by doing better on the reweighted observations.

9. Step 7: upsampling, a new dataset drawn by weight

Key point: Lay the weights end to end on the line from 0 to 1, draw n random numbers, and pick the observation whose range each number lands in. Heavy observations get picked more often.

The new weights are passed on through the data itself. We build a new dataset of the same size, $n = 5$ observations, in which heavy observations appear more often. Drawing a dataset this way is called **upsampling** (resampling by weight) (G-2063).




Figure 5: Upsampling: each observation owns a stretch of the line from 0 to 1 as long as its weight; five random numbers pick observations 1, 3, 3, 3 and 4

1. **Make ranges.** Each observation owns a stretch of the line from 0 to 1, as long as its weight. Observation 1 owns 0 to 0.167, observation 2 owns 0.167 to 0.417, observation 3 owns 0.417 to 0.667, observation 4 owns 0.667 to 0.833, and observation 5 owns 0.833 to 1 (Figure 5). Each boundary is the running total, the **cumulative sum** (G-519), of the weights.
2. **Draw random numbers.** Draw 5 random numbers between 0 and 1, say 0.13, 0.43, 0.62, 0.50 and 0.80.
3. **Pick observations.** Each number picks the observation whose range it falls in: 0.13 picks observation 1; 0.43, 0.62 and 0.50 all pick observation 3; 0.80 picks observation 4.

The new dataset is observations 1, 3, 3, 3 and 4. Observation 3, a mistake, appears three times; observations 2 and 5 do not appear at all. Another draw could give, for example, observations 1, 3, 2, 2 and 5. Either way, observations with larger weights own longer stretches of the line and are picked more often.

The last frames of Figure 4 play the draw: each random number lands as a dart on one observation's stretch.

The next stump trains on this new dataset, so it pays most attention to the observations the first stump got wrong. In the new dataset every observation starts with the same weight, $1/n$, again. Nothing is lost by the reset: a heavy observation is now present as several copies, and the copies carry its importance.

Extra: Upsampling is one way to make a model respect weights. The other is to hand the weights straight to the learner: scikit-learn's decision trees accept a `sample_weight` argument and count each observation in proportion to its weight. scikit-learn's AdaBoost works this way, with no random draws (scikit-learn source; the [AdaBoost from scratch Note](#), section 10).

10. Repeat, then vote

Key point: The new dataset starts the next stage; after T stages we have T stumps and T alphas, combined by the weighted vote.

With the new dataset, everything repeats (Figure 1):

1. every observation of the new dataset gets weight $1/n$ again;
2. a new stump is trained on it, model 2;
3. its error and its say α_2 are computed;
4. the weights are updated and normalised, and a new dataset is drawn for model 3.

We repeat this for as many stumps as we want, T . At the end we have $\alpha_1, \alpha_2, \dots, \alpha_T$ and stumps $h_1, \dots, h_T$, and the prediction is the weighted vote of the intuition Note, section 5:

$$H(x) = \text{sign} \left(\alpha_1 h_1(x) + \alpha_2 h_2(x) + \dots + \alpha_T h_T(x) \right)$$

Figure 6 runs three real stages on our 5 observations. Here the stumps are fitted by scikit-learn, so their mistakes differ from the assumed example of section 4, and the weights go straight to each stump as sample weights instead of through upsampling (the Extra in section 9). Watch observation 3 swell to half of all the weight after stump 1 misses it, the next stumps cut around it, and the weighted vote of 3 stumps get all 5 observations right, even though no single stump can.




Figure 6: Three AdaBoost stages on the toy data. Dot size and label: the observation's weight. Line and shading: the stump's cut and the class on each side; red rings: its mistakes. Right: each stump's alpha. Last frame: the sign of the weighted vote.

11. Summary

Step	What we do	Our numbers (stage 1)
1	weight every observation $1/n$	0.2 each
2	fit a decision stump	observations 2 and 3 wrong
3	error = sum of the weights of the mistakes	0.4
4	$\alpha = 0.5 \ln((1 - \text{error})/\text{error})$	0.20
5	mistakes $\times e^\alpha$, correct $\times e^{-\alpha}$	0.2449 and 0.1633
6	divide by the sum	0.25 and 0.1667
7	draw n observations by weight	observations 1, 3, 3, 3, 4

- The error is a weighted error: the total weight of the misclassified observations.
- Alpha is large for a small error, 0 at error 0.5, negative above 0.5.
- The weights of the mistakes grow, the others shrink, and all are rescaled to add up to 1.
- Upsampling turns weights into a new dataset in which heavy observations appear more often.

12. Sources

Built from

- CampusX, “AdaBoost - A Step by Step Explanation”, YouTube, <https://www.youtube.com/watch?v=F...>
- StatQuest with Josh Starmer, “AdaBoost, Clearly Explained”, YouTube, <https://www.youtube.com/watch?v=LsK-xG1cLYA> (the same steps on a second example; equal weights again after the new dataset is drawn)

Other references

- Schapire, R. E. (2013). Explaining AdaBoost. In *Empirical Inference*, Springer, pp. 37–52. (Algorithm 1: weights start at $1/m$, $\alpha_t = \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$, weights multiplied by $e^{-\alpha_t y_i h_t(x_i)}$ and normalised.) schapire.net/papers/explaining-adaboost.pdf
- scikit-learn developers. Source file `sklearn/ensemble/_weight_boosting.py` (version 1.9), method `_boost` of `AdaBoostClassifier`: fits each stump with `sample_weight` and stops when the error reaches $1 - 1/K$, which is 0.5 for two classes. github.com/scikit-learn/scikit-learn

13. Key terms

Term	Meaning
Sample weight	A number attached to each observation saying how important it is; AdaBoost starts every observation at $1/n$

Term	Meaning
Weighted error	The total sample weight of the observations a model misclassifies
Weight update	Multiplying misclassified observations' weights by e^α and correct observations' weights by $e^{-\alpha}$
Normalisation (of weights)	Dividing every weight by their sum so they add up to 1
Upsampling (resampling by weight)	Drawing a new dataset in which each observation is picked with probability equal to its weight
Cumulative sum	The running total of a list of numbers; it turns weights into ranges on the line from 0 to 1