

Statistics Roadmap for Machine Learning

Contents

1. Overview	2
2. Why ML needs statistics	3
3. The four modules	3
3.1 Descriptive statistics	3
3.2 Probability distributions	4
3.3 Inferential statistics	4
3.4 Miscellaneous topics	4
4. How to study the roadmap	5
5. Summary	6
6. Key terms	6

Where this fits: Pipeline map steps 0 (Foundations), 3 (Understand data). See the [Course map](#).




- **Builds on:** Exploratory data analysis ([Note 19](#)).
- **Leads to:** Measures of central tendency ([Note 221](#)); Variance ([Note 222](#)); Correlation ([Note 231](#)); Normal distribution ([Note 240](#)); Probability mass function (PMF) ([Note 241](#)); Uniform distribution ([Note 241](#)).

1. Overview

Key point: Statistics for ML splits into four modules: descriptive statistics, probability distributions, inferential statistics, and a few miscellaneous tools.

Statistics for ML: four modules



each module builds on the one before it;
the miscellaneous topics are studied when an algorithm needs them

Figure 1: The four modules of statistics for ML

Figure 1 shows the whole map. Each module builds on the one before it. This Note says what each module covers, why ML needs it, and where each topic is taught.

2. Why ML needs statistics

Key point: Statistics is how we summarise data and draw conclusions from it, and many ML algorithms are built on it.

Statistics turns raw data into summaries and decisions. In data work it shows up in three places:

- **Understanding data:** every exploratory analysis is built from statistical summaries and graphs.
- **Algorithms:** linear regression, logistic regression and Naive Bayes all come out of statistics.
- **Decisions:** checking whether a new model, medicine or website design is really better needs a statistical test.

3. The four modules

Key point: Descriptive statistics summarises the data we have; probability distributions describe the shapes data can take; inferential statistics draws conclusions about data we do not have.

3.1 Descriptive statistics

Key point: Summarise the data we have, with numbers and graphs.

This module describes data that is already in our hands. Its topics, and where each is taught:

Topic	Where it is taught
Population and sample, types of data	What is statistics Note
Central tendency (mean, median, mode, ...)	Measures of central tendency Note
Dispersion (range, variance, standard deviation, ...)	Measures of dispersion Note
Quantiles, percentiles, box plots	Percentiles and box plots Note
Skewness	Univariate analysis Note
Kurtosis	Pandas Profiling Note
Univariate analysis	Univariate analysis Note
Bivariate and multivariate analysis	Bivariate and multivariate analysis Note , frequency tables and graphs Note
Covariance and correlation	Covariance and correlation Note

3.2 Probability distributions

Key point: A probability distribution describes how likely each value of a column is; the hypothesis tests of the next module depend on it.

This module studies the shapes data can take:

- **Random variables**, and the functions that describe them: the PMF (for counts), the PDF (for measurements) and the CDF (for running totals of probability).
- **Kernel density estimation**, the smooth curve over a histogram, met in the density plot of the [univariate analysis Note](#).
- **2D density plots**, the same idea for two columns at once.
- **Named distributions**: normal, uniform, Bernoulli, binomial, log-normal and others. The normal distribution's 68-95-99.7 rule appears in the [z-score outliers Note](#).

The basic rules of probability are in the [conditional probability Note](#) and the Notes after it.

3.3 Inferential statistics

Key point: Use a sample to say something about the whole population, and say how sure we are.

This module draws conclusions about a population from a sample:

- **Central limit theorem**: why averages of samples behave predictably.
- **Confidence intervals**: a range that most likely holds the true population value.
- **Hypothesis testing**: checking a claim about a population with a sample. The common tests are the z-test, t-test, chi-square test and ANOVA, and the module ends with when to use which.

3.4 Miscellaneous topics

Key point: A few tools do not fit one module; we study them when an algorithm needs them.

Topic	Where it is taught
Chebyshev's inequality	a later maths Note
Q-Q plot	Function transformer Note
Sampling techniques	a later maths Note; sampling bias in the challenges in ML Note
Resampling (bootstrap)	Bagging intuition Note
Statistical moments	skewness in the univariate analysis Note
Bayesian statistics	Bayes' theorem Note

4. How to study the roadmap

Key point: About 60 hours, or a month at 2 to 2.5 hours a day; some topics can wait until an algorithm needs them.

The whole roadmap takes roughly 60 hours. At 2 to 2.5 hours a day, that is about one month. Two habits help:

- **Know why each topic matters.** For every topic, note where ML or data analysis uses it.
- **Defer some topics.** A few topics, such as Chebyshev's inequality or Bayesian statistics, can wait until an algorithm needs them. The rest are worth covering in order, since each builds on the last.

Extra: Good resources for this roadmap:

- **Books:** *An Introduction to Statistical Learning* (James, Witten, Hastie, Tibshirani; statistics through ML), *Practical Statistics for Data Scientists* (Bruce, Bruce, Gedeck; practical, with Python and R), *Think Stats* (Downey; short, Python-based).
- **Videos:** StatQuest's *Statistics Fundamentals* series (Josh Starmer).
- **Blog:** Machine Learning Mastery (Jason Brownlee) has many short statistics tutorials.
- **Revision:** a one-page statistics cheat sheet, once the topics are learned.

5. Summary

Module	Question it answers	Main topics
Descriptive statistics	What does the data we have look like?	central tendency, dispersion, percentiles, graphs, correlation
Probability distributions	What shapes can data take?	PMF, PDF, CDF, normal and other distributions
Inferential statistics	What can a sample tell us about the population?	central limit theorem, confidence intervals, hypothesis tests
Miscellaneous	Extra tools	Q-Q plot, sampling, resampling, moments, Bayes

- Descriptive statistics comes first; inferential statistics depends on probability distributions.
- About 60 hours covers the whole roadmap.

6. Key terms

Term	Meaning
Probability distribution	A description of how likely each value of a variable is