


Frequency Tables and Graphs by Data Type

Contents

1. Overview	2
2. Frequency tables for a categorical column	2
2.1 Frequency distribution table	3
2.2 Relative frequency	3
2.3 Cumulative frequency	4
3. Histograms for a numerical column	4
3.1 The shapes of a histogram	5
4. Graphs for two columns	6
4.1 Categorical and numerical: any aggregate per category	6
4.2 Categorical and numerical: a contingency table of bins	6
5. Graphs for more than two columns	7
5.1 The 3D scatter plot	8
5.2 The facet grid	8
6. Summary	10
7. Key terms	10

Where this fits: Pipeline map step 3 (Understand data). See the [Course map](#).



- **Builds on:** Exploratory data analysis ([Note 19](#)).
- **Leads to:** Correlation ([Note 231](#)); Z-score outlier method ([Note 251](#)); Skewness ([Note 252](#)).
- **Compare with:** Inferential statistics ([Note 220](#)); Cumulative distribution function (CDF) ([Note 241](#)).

1. Overview

Key point: A categorical column is summarised by a frequency table and its relative and cumulative versions; a numerical column by a histogram; two or more columns by graphs chosen from their types.




Figure 1: A frequency table drawn three ways: frequency as a bar chart, relative frequency as a pie chart, cumulative frequency as a line chart

Descriptive statistics is about summarising data and presenting it in a meaningful way. Figure 1 shows the three tables we can build from one categorical column, each with the chart that suits it. This Note covers those tables, the shapes a histogram can take, and the graphs for two and for more than two columns.

2. Frequency tables for a categorical column

Key point: Count each category (frequency), turn the counts into shares (relative frequency), and add them up as we go (cumulative frequency).

2.1 Frequency distribution table

Key point: A frequency distribution table lists each category with the number of times it occurs.

We survey 200 people about their favourite type of vacation. Each answer is one of six categories: beach, city, adventure, nature, cruise or other. The data has one row per person, with their name and their answer.

A **frequency distribution table** summarises how many times each value occurs in a dataset. For the survey, we count each category:

Vacation	Frequency	Relative frequency	Cumulative frequency
Beach	60	0.300	60
City	40	0.200	100
Adventure	30	0.150	130
Nature	35	0.175	165
Cruise	20	0.100	185
Other	15	0.075	200
Total	200	1.000	

The frequencies are drawn as a bar chart, one bar per category (Figure 1, left). This is the count plot of the [univariate analysis Note](#) (section 4), and pandas builds the table with `value_counts()`.

2.2 Relative frequency

Key point: Relative frequency is a category's share of all the data: its frequency divided by the total.

The **relative frequency** of a category is its proportion, or percentage, of the data.

1. **In words:** divide the category's count by the total number of rows.
2. **Formula:** for a category with frequency f out of n rows,

$$\text{relative frequency} = \frac{f}{n}$$

3. **Example:** 60 of the 200 people chose the beach, so

$$\text{relative frequency}_{\text{Beach}} = \frac{60}{200} = 0.30 = 30\%$$

The relative frequencies always add up to 1 (100%). That makes them the natural input for a pie chart, whose slices show each category's share (Figure 1, middle).

2.3 Cumulative frequency

Key point: Cumulative frequency is the running total of the frequencies, so the last value is always the total.

The **cumulative frequency** of a category is its frequency plus the frequencies of all the categories before it.

1. **In words:** walk down the table, keeping a running total of the frequencies.
2. **Formula:** for categories in the table's order with frequencies $f_1, f_2, \dots$, the cumulative frequency of the k -th category is

$$F_k = f_1 + f_2 + \dots + f_k$$

3. **Example:** Beach gives 60; City adds 40, giving 100; Adventure adds 30, giving 130; then 165, 185 and finally 200, the total.

A line chart shows the running total (Figure 1, right). The same running total of the relative frequencies, the **cumulative relative frequency**, climbs from 0.30 to 0.50, and so on, to 1.

Cumulative frequency is most meaningful when the categories have an order, such as ordinal data or the bins of a histogram: “how many people scored 60 or less?”. For unordered categories it depends on the order we list them in.

Python: The three tables with pandas.

```
freq = df["vacation"].value_counts(sort=False)
rel = freq / freq.sum()           # relative frequency
cum = freq.cumsum()               # cumulative frequency
value_counts(normalize=True) gives the relative frequencies directly. cumsum() is the
running total. The Notebook (notebook.ipynb) builds all three tables and charts.
```

3. Histograms for a numerical column

Key point: A numerical column has no categories, so we make our own by cutting its range into bins; the histogram is the frequency table of those bins.

A numerical column such as age has no categories to count. So we create them: we cut the range into bins (or buckets), such as 0 to 10, 10 to 20, 20 to 30, and put each person in their bin. The result is a frequency distribution table on bins instead of categories, and its chart is the histogram (see the [univariate analysis Note](#), section 6).

A histogram looks like a bar chart with one difference: its bars touch. The categories of a bar chart are separate; the bins of a histogram are continuous ranges, one ending where the next begins.

The number of bins changes the picture, as section 6.1 of that Note shows. Too few bins (say, two bins of 0 to 30 and 30 to 100) hide the shape; too many give thin, noisy bars. We try a few and keep the clearest.

3.1 The shapes of a histogram

Key point: Histograms come in a few typical shapes: symmetric, bimodal, skewed to the right or left, uniform, or no clear pattern.

Figure 2 shows six typical shapes.




Figure 2: Six shapes a histogram can take

- **Symmetric:** most values in the middle, fewer and fewer towards both sides.
- **Bimodal:** two separate peaks, two groups of values where points are dense. With three peaks it is trimodal. Two peaks often mean two kinds of data mixed together, such as the heights of children and adults.
- **Right or left skew:** a long tail on one side; the shapes and the skewness number are taught in section 10 of the [univariate analysis Note](#).
- **Uniform:** every bin holds about the same number of values. Too few bins also make data look uniform.
- **No pattern:** the bars jump up and down. Usually there are too many bins for the amount of data, here 30 bins for 60 values.

Named shapes such as the normal and uniform distributions come in the probability distribution Notes.

Extra: Remember the direction of skew by the tail, not by the hump. In a right-skewed histogram the hump is on the left and the long tail points right.

4. Graphs for two columns

Key point: Two columns are categorical + categorical (contingency table), numerical + numerical (scatter plot) or categorical + numerical (an aggregate per category, or a contingency table of bins).

Bivariate analysis studies two columns together; the [bivariate and multivariate analysis Note](#) (Figure 1) maps each pair of column types to its plots. In short:

- **Categorical + categorical:** a **contingency table**, also called a crosstab: the counts for every pair of categories (that Note, section 7). On the Titanic, a contingency table of class against survival shows 372 of the 549 passengers who died were in third class. A side-by-side or stacked bar chart draws it.
- **Numerical + numerical:** a scatter plot (that Note, section 3). Its pattern can show a positive relation, a negative one, or none, which the [covariance and correlation Note](#) measures.
- **Categorical + numerical:** two options, below.

4.1 Categorical and numerical: any aggregate per category

Key point: A bar per category whose height is any summary of the numerical column: mean, median, maximum, standard deviation.

A bar chart of a categorical column against a numerical one does not show counts. Each bar shows an **aggregate**, one summary number computed from the numerical values in that category. The bar plot of that Note (section 4) uses the mean, but any summary works:

Sex	Mean age	Median age	Maximum age
Female	27.9	27	63
Male	30.7	29	80

The mean bars compare typical ages; the maximum bars show that the oldest man (80) was much older than the oldest woman (63).

4.2 Categorical and numerical: a contingency table of bins

Key point: Cutting the numerical column into bins turns it into a categorical one, so a contingency table works.

We can also bin the numerical column, as for a histogram, and then count every pair of age band and category:

Age band	Female	Male
0 to 10	31	33

Age band	Female	Male
10 to 20	46	69
20 to 30	81	149
30 to 40	55	100
40 to 50	31	55
50 to 60	14	28
60 to 80	3	19

The cell “20 to 30, male” means 149 passengers were men aged over 20 and up to 30. The table shows the men outnumbered the women in every band, most of all among the older passengers (19 against 3 over 60).

Python: Binning, then a contingency table.

```
bands = pd.cut(df["Age"], bins=[0, 10, 20, 30, 40, 50, 60, 80])
pd.crosstab(bands, df["Sex"])
pd.cut puts each age in its band; a band such as (20, 30] includes 30 but not 20.
```

Extra: A crosstab only counts pairs of categories. A pivot table can also aggregate a third, numerical column for each pair: for example, the mean age of the male passengers in first class. Spreadsheets and pandas (`pivot_table`) both offer it; the [bivariate and multivariate analysis Note](#) (section 11) uses one.

5. Graphs for more than two columns

Key point: A 3D scatter plot, colour (hue), facet grids, pair plots and bubble charts each fit three or more columns into one figure.

Multivariate analysis studies more than two columns at once. Most of its graphs are bivariate graphs with an extra column added:

- **Hue:** colour shows a categorical column, on a scatter plot, bar plot or box plot. A bar plot of mean age by sex with `hue = class` shows three columns. See section 3.2 of the [bivariate and multivariate analysis Note](#).
- **Bubble chart:** a scatter plot whose dot size shows a third numerical column, such as countries by GDP and population with bubble size for literacy rate (the “size” setting of the same section).
- **Pair plot:** a scatter plot for every pair of numerical columns, with a histogram on the diagonal (section 9 of that Note).
- **3D scatter plot** and **facet grid:** below.

5.1 The 3D scatter plot

Key point: Three numerical columns on three axes; colour can add a fourth.

A **3D scatter plot** places each row as a point in three dimensions, one numerical column per axis. Figure 3 plots the tips data: total bill, party size and tip, with lunch and dinner by colour.




Figure 3: 3D scatter plot of total bill, party size and tip

Bigger bills and bigger parties go with bigger tips: the points rise towards the far corner. On paper a 3D plot is hard to read; in the Notebook it can be rotated with the mouse, which is where it is useful.

5.2 The facet grid

Key point: The same plot repeated side by side, one panel per category of another column.

A **facet grid** splits the data by a categorical column and draws the same plot for each part, side by side, on shared axes. Figure 4 draws bill against tip once for lunch and once for dinner, with smokers in orange.




Figure 4: Bill against tip, one panel per meal time, smokers by colour

Figure 4 shows four columns at once: bill, tip, meal time (the panels) and smoker (the colour). Dinner has more customers and more of the large bills.

Python: 3D scatter plot and facet grid.

```
import plotly.express as px
```

```
px.scatter_3d(tips, x="total_bill", y="size", z="tip", color="time")
px.scatter(tips, x="total_bill", y="tip", color="smoker",
           facet_col="time")
```

In seaborn, `sns.relplot(data=tips, x="total_bill", y="tip", hue="smoker", col="time")` draws the facet grid.

Extra: A joint plot (a scatter plot with a histogram of each column along its edges) looks busy, but it shows only two columns. It is a bivariate graph, not a multivariate one.

6. Summary

Columns	Types	Table	Graph
One	categorical	frequency, relative, cumulative frequency	bar chart, pie chart, line chart
One	numerical	frequency table of bins	histogram
Two	categorical + categorical	contingency table (crosstab)	side-by-side or stacked bars
Two	numerical + numerical	none	scatter plot
Two	categorical + numerical	aggregate per category, or crosstab of bins	bar chart of an aggregate, box plots
Three or more	any	pivot table	hue, bubble chart, pair plot, 3D scatter, facet grid

- Relative frequency = frequency / total; cumulative frequency = running total.
- Histogram bars touch because bins are continuous ranges.
- Histogram shapes: symmetric, bimodal, right skew, left skew, uniform, no pattern.

7. Key terms

Term	Meaning
Frequency distribution table	A table of each value or category with the number of times it occurs
Relative frequency	A category's share of all the data: its frequency divided by the total
Cumulative frequency	The running total of the frequencies, up to and including a category
Cumulative relative frequency	The running total of the relative frequencies; ends at 1
Contingency table	A table of counts for every pair of categories of two columns; another name for a crosstab
Aggregate	One summary number (mean, median, maximum, ...) computed from a group of values
3D scatter plot	A scatter plot of three numerical columns on three axes
Facet grid	The same plot repeated side by side, one panel per category of another column