

Random Variables and Probability Distributions

Contents

1. Overview	2
2. Random variables	3
2.1 Algebra variables and random variables	3
2.2 Random experiments	3
2.3 Discrete and continuous random variables	4
3. Probability distributions as tables	4
3.1 The sum of two dice	5
4. From a table to a function	6
5. Discrete and continuous graphs	6
5.1 PMF, PDF and CDF	7
6. Famous distributions	7
7. Why distributions matter	8
8. Parameters of a distribution	9
9. Summary	11
10. Key terms	11

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



- **Builds on:** Discrete and continuous data ([Note 220](#)).
- **Leads to:** Probability mass function (PMF) ([Note 241](#)); Probability density function (PDF) ([Note 242](#)).

1. Overview

Key point: A random variable turns the outcomes of a random experiment into numbers; a probability distribution says how likely each number is, and it comes as a PMF, a PDF or a CDF.

The family of probability distribution functions




Figure 1: The family of probability distribution functions

Figure 1 is the map of this Note and the three after it. A **probability distribution function** is a formula that links every outcome to its probability. It has two kinds, one for each kind of random variable, and both lead to a third function, the CDF.

This Note builds the base: random variables, distributions as tables, why we want a formula instead, the famous distributions, and their parameters. The three functions themselves are taught in the [PMF and discrete CDF Note](#), the [PDF and continuous CDF Note](#) and the [density estimation Note](#).

2. Random variables

Key point: An algebra variable holds one unknown value; a random variable can take several values, and which one appears is decided by chance.

2.1 Algebra variables and random variables

Key point: We solve for an algebra variable; we cannot solve for a random variable, only describe how likely each of its values is.

In algebra, a variable stands for one unknown value. In $x + 5 = 10$ we simplify and get $x = 5$: there is exactly one answer.

Statistics and probability use a different kind of variable. A **random variable** is the set of possible numerical values of a random experiment. It differs from an algebra variable in two ways:

- **It holds many values.** A die can show any of 1 to 6.
- **The value that appears is random.** We cannot solve for it; we can only say how likely each value is.

2.2 Random experiments

Key point: A random experiment has an outcome we cannot predict; the random variable writes its outcomes as numbers.

A **random experiment** is any experiment whose outcome is random: tossing a coin, rolling a die, picking a student and asking their CGPA. Figure 2 turns three such experiments into random variables.

A random variable turns outcomes into numbers




Figure 2: Three random experiments and their random variables

- **Coin toss:** writing head as 1 and tail as 0 gives $X = \{1, 0\}$, assuming a fair (unbiased) coin.
- **Die roll:** $Y = \{1, 2, 3, 4, 5, 6\}$.
- **CGPA of a random student:** Z can be any number from 0 to 10.

By convention, a random variable gets a **capital letter** (X, Y) and an algebra variable a small one (x). A particular value of X is written with the small letter: “ $X = x$ ” reads “the random variable X takes the value x ”. The set of all possible outcomes is the sample space (see the [conditional probability Note](#)).

Extra: Strictly, a random variable is a **function**: it assigns a number to each outcome of the sample space. Head $\rightarrow 1$ and tail $\rightarrow 0$ is such a function. “The set of possible values” describes what the function can output. Thinking of it as a function explains why the same experiment can have several random variables: rolling two dice can give “the sum”, “the larger face” or “the number of sixes”.

2.3 Discrete and continuous random variables

Key point: A discrete random variable takes separate values (a die never shows 1.5); a continuous one can take any value in a range (a CGPA can be 8.527).

Random variables come in the same two kinds as numerical data (see the [what is statistics Note](#), section 6.2):

- A **discrete random variable** takes separate values. The coin gives 0 or 1, the die 1 to 6; a value of 1.5 or 1.6 can never appear.
- A **continuous random variable** can take any value in a range. A CGPA can be 8.5, 8.52, 8.527 or 6.532: every decimal between 0 and 10 is possible.

The difference decides which function describes the variable (Figure 1): a PMF for a discrete variable, a PDF for a continuous one.

3. Probability distributions as tables

Key point: A probability distribution lists every possible outcome of a random variable with its probability, like a frequency table with probabilities in place of counts.

A **probability distribution** is a list of all the possible outcomes of a random variable, each with its probability. It looks like the frequency table of the [frequency tables Note](#), with probabilities instead of counts.

Coin toss. Two outcomes, equally likely:

x	1 (head)	0 (tail)
$P(X = x)$	1/2	1/2

One die. Six outcomes, all **equiprobable** (equally likely):

x	1	2	3	4	5	6
$P(X = x)$	1/6	1/6	1/6	1/6	1/6	1/6

3.1 The sum of two dice

Key point: Two dice give 36 equally likely pairs; counting the pairs for each sum gives probabilities from 1/36 (sums 2 and 12) up to 6/36 (sum 7).

When we roll two dice and add the faces, the sum can be anything from $1+1=2$ to $6+6=12$. The 36 pairs of faces are equally likely; the [conditional probability Note](#) (Figure 2) draws them as a 6 by 6 grid with the sum in each cell. Here the sums are not equally likely, because some sums can be made in more ways than others.

1. **In words:** the probability of a sum is the number of pairs that give it, divided by the 36 pairs.
2. **Formula:**

$$P(X = x) = \frac{\text{number of pairs with sum } x}{36}$$

3. **Example:** the sum 7 comes from (1, 6), (2, 5), (3, 4), (4, 3), (5, 2) and (6, 1), six pairs:

$$P(X = 7) = \frac{6}{36} \approx 0.167$$

The sum 2 comes only from (1, 1), so $P(X = 2) = 1/36 \approx 0.028$.

Counting every sum gives the whole distribution:

Sum x	Pairs	$P(X = x)$
2	1	1/36
3	2	2/36
4	3	3/36
5	4	4/36
6	5	5/36
7	6	6/36
8	5	5/36
9	4	4/36
10	3	3/36
11	2	2/36
12	1	1/36

The most likely sum is 7; the least likely are 2 and 12. The counts rise by one up to 7 and fall by one after it.

4. From a table to a function

Key point: A table fails when there are too many outcomes or infinitely many; a formula $y = f(x)$ from outcome to probability works in both cases, and it can be graphed.

A table works for a coin or two dice. It breaks down in two situations:

- **Too many outcomes.** With 10 dice the sum runs from 10 to 60: 51 sums, built from $6^{10} = 60,466,176$ equally likely combinations. Writing that table by hand is tedious.
- **A continuous random variable.** A CGPA can be any number between 0 and 10, so the outcomes cannot even be listed.

The way out is a function. Call the outcome x and its probability y . Instead of a table, we look for a formula $y = f(x)$ that **models** the relationship between them. A formula has two benefits:

1. **Any outcome at once:** we plug in any x and the formula returns its y .
2. **A graph:** a formula can be plotted, and the graph shows the shape of the distribution at a glance.

This formula is the **probability distribution function**: a mathematical function that links every possible outcome of a random variable to its probability. For one die it is $f(x) = 1/6$ for x in 1 to 6, and 0 for any other x .

Extra: “Probability distribution” (the table, or the idea) and “probability distribution function” (the formula) are used interchangeably in most books. Both mean the description of how likely each value is; the table and the formula are two ways of writing it down.

5. Discrete and continuous graphs

Key point: The graph of a discrete distribution has values only at separate points; the graph of a continuous one is an unbroken curve.

The two kinds of random variable give two kinds of graph (Figure 3):

- **Discrete:** a die has values only at 1, 2, ..., 6. At 1.5 there is nothing, because 1.5 is not in the sample space. The graph is a set of separate bars or sticks with gaps between them.
- **Continuous:** a CGPA has a value at every point from 0 to 10, so the graph is an unbroken curve. We can pick any x and read off its y .




Figure 3: The distribution of one die (separate values) and of CGPA (an unbroken curve)

The two graphs also differ in what the height means. For the die, the height is a probability. For the CGPA curve the y axis is labelled “probability density”, which is not a probability; the [PDF and continuous CDF Note](#) explains why.

5.1 PMF, PDF and CDF

Key point: For a discrete variable the function is a PMF (mass); for a continuous one it is a PDF (density); either can be turned into a CDF (cumulative).

The probability distribution function is the umbrella term for three functions (Figure 1):

- **Probability mass function (PMF):** the probability distribution function of a **discrete** random variable, such as the die.
- **Probability density function (PDF):** the probability distribution function of a **continuous** random variable, such as the CGPA.
- **Cumulative distribution function (CDF):** built from either of them; it gives the probability that the variable is **at most** x .

“PDF” is ambiguous: both “probability **distribution** function” and “probability **density** function” shorten to it. In these Notes, **PDF always means the density function**; the distribution function is always written in full.

6. Famous distributions

Key point: Data from very different fields often follows one of a few common shapes; these famous distributions have names, known formulas and known parameters.

When people in many fields (science, medicine, engineering, economics) plotted the distributions of their data, the same few shapes kept appearing. These are the **famous probability**

distributions. Figure 4 shows eight of them: discrete ones as bars, continuous ones as curves.




Figure 4: Eight famous distributions: four discrete (top) and four continuous (bottom), with their parameters

- **Discrete:** Bernoulli, binomial, Poisson, discrete uniform.
- **Continuous:** normal, uniform, log-normal, exponential.
- **Others** we will meet later: beta, chi-square, Pareto.

When we plot the distribution of a real column, there is a good chance it resembles one of these. The normal distribution is the most famous; the [z-score outliers Note](#) already used its 68-95-99.7 rule. The binomial distribution appeared in the [voting ensemble Note](#). Each named distribution gets its own Note later.

7. Why distributions matter

Key point: A distribution shows the shape of the data, and if that shape matches a famous distribution, everything already known about it applies to our data.

The graph of a distribution is only x (outcomes) against y (their probabilities), yet it gives us two things.

1. The shape of the data. The graph shows where the values lie, which are common and which are rare:

- If most marks in a class (out of 10) sit around 5 to 6, the class has many weak students.
- If most marks sit around 8, with few below 7, most of the class is strong.
- If a country's salary distribution has a big mass of high earners, it is probably a developed country; if most of the mass sits at low salaries, poverty is likely widespread.

2. Ready-made knowledge. If our data's shape matches a famous distribution, we can apply everything known about that distribution to our data. The normal distribution has been studied for centuries and its mathematics is fully worked out. Knowing that a column is normal, we immediately know, for example, that about 95% of its values lie within two standard deviations of the mean.

8. Parameters of a distribution

Key point: Parameters are the tuning knobs of a distribution: numbers such as μ and σ that set its location, scale and shape.

Every famous distribution comes with a few numbers written next to its name in Figure 4: μ and σ , a and b , n and p , λ . These are its **parameters**: numerical values that determine the shape, location and scale of the distribution. Changing them changes the graph, like turning a knob.

Figure 5 turns the two knobs of the normal distribution:

- μ , **the location (mean)**: moves the curve left or right without changing its shape.
- σ , **the scale (standard deviation)**: widens or narrows the curve. A wider curve is lower, because the total area under it stays the same (see the [PDF and continuous CDF Note](#)).




Figure 5: Changing the mean moves the normal curve; changing the standard deviation widens it

Different distributions have different sets of parameters. Whenever we study a distribution, we study what each of its parameters does.

This is the same word as in the [what is statistics Note](#) (section 4.2), where a parameter is a number that describes the population. The two meanings meet: if a population is normal, its mean μ and standard deviation σ are exactly the knobs of its normal curve. Estimating those knobs from a sample is the idea behind the [density estimation Note](#).

Extra: The normal distribution's scale parameter is σ ; its square, σ^2 , is the variance (see the [measures of dispersion Note](#)). Some books list the parameters as (μ, σ) , others as (μ, σ^2) : the same curve, written two ways.

9. Summary

Idea	Meaning	Example
Random variable	The numerical outcomes of a random experiment	die: $Y = \{1, \dots, 6\}$
Discrete random variable	Takes separate values	coin, die, sum of two dice
Continuous random variable	Takes any value in a range	CGPA from 0 to 10
Probability distribution	Every outcome with its probability	two dice: $P(X = 7) = 6/36$
Probability distribution function	A formula $y = f(x)$ from outcome to probability	$f(x) = 1/6$ for $x = 1, \dots, 6$
PMF / PDF / CDF	Discrete / continuous / cumulative	die / CGPA / $P(X \leq x)$
Parameters	Knobs that set location, scale and shape	normal: μ, σ

- Random variables get capital letters; their values get small letters.
- A table cannot list too many outcomes or a continuous range; a function can, and it can be plotted.
- In these Notes, PDF means probability density function.
- A distribution shows the shape of the data; matching a famous one lets us reuse its known mathematics.

10. Key terms

Term	Meaning
Random experiment	An experiment whose outcome cannot be predicted, such as a coin toss
Random variable	The possible numerical outcomes of a random experiment; strictly, a function from outcomes to numbers
Discrete random variable	A random variable that takes separate values, such as a die's face
Continuous random variable	A random variable that can take any value in a range, such as a CGPA
Equiprobable	Equally likely
Probability distribution	A list of every possible outcome of a random variable with its probability
Probability distribution function	A formula $y = f(x)$ giving the probability of each outcome; the umbrella term for PMF and PDF
Probability mass function (PMF)	The probability distribution function of a discrete random variable

Term	Meaning
Cumulative distribution function (CDF)	The function giving $P(X \leq x)$, the probability of a value at most x
Famous probability distributions	Common named shapes such as normal, uniform, binomial and Poisson
Parameters (of a distribution)	The numbers, such as μ and σ , that set a distribution's location, scale and shape