

Linear Algebra Roadmap for Machine Learning

Contents

1. Overview	2
2. Why ML needs linear algebra	3
3. How to read the roadmap	3
4. The eight modules	4
4.1 Scalars	4
4.2 Vectors	5
4.3 Matrices	5
4.4 Tensors	6
4.5 Eigenvalues and eigenvectors	6
4.6 Matrix factorisation	6
4.7 Advanced topics	6
4.8 Tools and libraries	7
5. Resources	7
6. Summary	9
7. Sources	9
8. Key terms	9

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



1. Overview

Key point: Linear algebra for ML splits into eight modules: scalars, vectors, matrices, tensors, eigenvalues and eigenvectors, matrix factorisation, advanced topics, and tools.

Linear algebra for ML: eight modules




Figure 1: The eight modules of linear algebra for ML

Figure 1 shows the whole map. This Note says what each module covers, why ML needs it, which topics can wait, and where each topic is taught.

2. Why ML needs linear algebra

Key point: Statistics and linear algebra are the two pillars of ML: statistics is how we see the data, linear algebra does the heavy lifting.

The first pillar, statistics, has its own [roadmap Note](#). If statistics is the eyes of an ML engineer, linear algebra is the hands and feet: it stores the data and does every calculation on it.

Why ML needs it (high dimensions, data as numbers, GPU speed) is explained in section 3 of the [vectors and feature vectors Note](#).

3. How to read the roadmap

Key point: Each topic is marked by importance, and topics needed only by a few algorithms are marked “later”.




Figure 2: The marks on the roadmap: three levels of importance, and “later” for topics only particular algorithms need

Figure 2 is the key to the marks; everything above the dashed line is worth covering in one go.

The roadmap marks every topic in three ways:

- **Importance.** A *very important* topic appears almost everywhere in ML and DL; an *important* topic appears in several places; a *normal* topic is still worth learning.
- **“Later”.** A few specialised topics are needed only by particular algorithms. We can skip them now and learn them when we reach that algorithm. Most of them sit in the

last modules.

- **Why it matters.** Every topic comes with a one-line description and where ML and DL use it. For example, vectors represent data points in ML, and features, weights and biases in DL.

Everything not marked “later” is worth covering in one go before moving deep into ML and DL.

4. The eight modules

Key point: Vectors and matrices are the core; tensors, eigenvectors and factorisations build on them; a few advanced topics wait until an algorithm needs them.




Figure 3: A preview of six modules, each picture taken from the Note that teaches it: a vector as a point (Note 360), cosine similarity (Note 362), matrices as transformations (Note 500), tensors (Note 11), eigenvectors (Note 530) and the eigen-decomposition in PCA (Note 48)

Figure 3 shows what each module looks like once we study it; the subsection numbers match the captions.

4.1 Scalars

Key point: A scalar is a plain number.

Scalars are ordinary numbers such as 2, 5 or -6, and the arithmetic on them. The [tensors Note](#) places them as 0D tensors. They need no separate study.

4.2 Vectors

Key point: Vectors are where linear algebra starts, and where ML data starts: every observation of a dataset is a vector.

Each **observation** (one record, one row of the data table) is a vector. This module tries to cover everything about vectors that ML uses:

Topic	Why ML needs it	Where it is taught
What vectors are, components, dimension	every data point is a vector	Vectors and feature vectors Note
Types of vectors (row, column, ...)	data and parameters are written as rows or columns	Vectors and feature vectors Note
Distance from origin (magnitude)	the length of a vector	Magnitude, distance and scalar operations Note
Euclidean distance	KNN, K-means, recommender systems	KNN imputer Note , magnitude Note
Operations with a scalar	shifting and scaling data, mean centring	Magnitude, distance and scalar operations Note
Operations between two vectors, dot product	similarity, projection, matrix multiplication	Dot product and cosine similarity Note
Equation of a line in n dimensions	linear models: linear and logistic regression, SVM	Equation of a hyperplane Note
Vector norms	regularisation	norm in the SVM maths Note ; penalties in the Ridge and Lasso Notes
Vector spaces	the setting for everything above	a later maths Note

4.3 Matrices

Key point: Matrices come in two parts: the school-level mechanics, then the intuition of what a matrix does.

The first part is the mechanics taught in school mathematics (classes 11 and 12): what matrices are, types of matrices, operations with a scalar and with a vector, the inverse and the determinant.

The second part is conceptual: the rank of a matrix, column space, change of basis, solving a system of linear equations, linear transformations, and the dot product seen as the engine of matrix multiplication. **Linear transformations** are marked very important.

So far, the inverse appears in the [normal equation](#) of linear regression, and the idea of a matrix as a transformation in the [PCA step by step Note](#) (section 4.1). The rest of this module comes in a later maths Note.

4.4 Tensors

Key point: Tensors are super important for deep learning, and for representing any kind of data.

What a tensor is, and how tables, text, images and videos become tensors, is taught in the [tensors Note](#).

4.5 Eigenvalues and eigenvectors

Key point: Eigenvectors are the directions a matrix only stretches; PCA is built on them.

Anyone who has studied PCA has met eigenvalues and eigenvectors. They are taught in the [PCA step by step Note](#) (section 4), and in more depth in a later maths Note.

4.6 Matrix factorisation

Key point: Factorising a matrix splits it into simpler matrices; LU, QR, eigen-decomposition and SVD are the ones ML meets.

Matrix factorisation (or **decomposition**) writes one matrix as a product of simpler ones. Four techniques are marked important: LU decomposition, QR decomposition, eigen-decomposition and SVD (singular value decomposition). Eigen-decomposition appears in the [PCA step by step Note](#); the others come in a later maths Note.

Extra: Where these show up. scikit-learn's Ridge can solve its equation with SVD (`solver="svd"`, see the [Ridge gradient descent Note](#)). Least-squares problems can be solved through a QR factorisation (Trefethen and Bau 1997, Lecture 11, Algorithm 11.2), and recommender systems use SVD-like factorisations of the user-item rating matrix to predict missing ratings (Koren, Bell and Volinsky 2009).

4.7 Advanced topics

Key point: Quadratic forms and the pseudo-inverse are needed only by particular algorithms, so they can wait.

- **Quadratic forms** are used heavily when solving SVMs (see the [SVM maths Note](#)).
- The **Moore-Penrose pseudo-inverse** gives an inverse for matrices that are not square, or that have no ordinary inverse. The [multiple linear regression maths Note](#) mentions why libraries need it.

Both come in a later maths Note.

4.8 Tools and libraries

Key point: NumPy is how linear algebra is done in Python; SciPy adds more.

All linear algebra in the ML and DL world runs through **NumPy**, so a strong command of it pays off everywhere. **SciPy**, built on NumPy, adds more routines; it is used less, but we should know it exists.

Python: NumPy's linear algebra lives in `np.linalg`.

```
import numpy as np

v = np.array([3, 4])
np.linalg.norm(v)           # 5.0: length of v
A = np.array([[2, 1], [1, 3]])
np.linalg.inv(A)            # inverse of A
np.linalg.eig(A)            # eigenvalues, eigenvectors
```

5. Resources

Key point: The 3Blue1Brown series for intuition, Gilbert Strang for depth, and chapter 2 of the *Deep Learning* book as a summary.

Lesson series:

- Sanderson, G., *Essence of Linear Algebra*, 3Blue1Brown, [3blue1brown.com](https://www.3blue1brown.com). An animated series that builds the geometric intuition for vectors, matrices and transformations. Each lesson is dense, so it pays to go through it more than once.
- Strang, G., *18.06 Linear Algebra*, MIT OpenCourseWare, ocw.mit.edu. About 35 lectures by a renowned professor. The course covers far more than ML needs, but it is the place for depth.

Books:

- Gilbert Strang, *Introduction to Linear Algebra*. The book behind his lectures.
- NCERT mathematics textbooks (classes 11 and 12). Light on intuition, but they cover the school-level topics well.
- Goodfellow, Bengio and Courville, *Deep Learning*, chapter 2. A compact summary of all the linear algebra deep learning needs; worth reading twice.

Online courses:

- Khan Academy, *Linear Algebra*.
- Coursera, *Mathematics for Machine Learning: Linear Algebra* (Imperial College London).

The vectors Notes that follow this one cover the first module in depth; the rest of the matrix and eigen topics come in later maths Notes.

Extra: One more free book worth knowing: *Mathematics for Machine Learning* by Deisenroth, Faisal and Ong (mml-book.github.io). Its chapters 2 to 4 (Linear Algebra, Analytic Geometry, Matrix Decompositions) cover the vectors, matrices, eigenvalues and factorisations of this roadmap, including SVD in section 4.5.

6. Summary

Module	What it covers	Status
Scalars	plain numbers	no study needed
Vectors	components, distance, operations, dot product, hyperplanes, norms	core
Matrices	mechanics, then rank, spaces, linear transformations	core
Tensors	data of any shape	core for DL
Eigenvalues and eigenvectors	directions a matrix only stretches	needed for PCA
Matrix factorisation	LU, QR, eigen-decomposition, SVD	important
Advanced topics	quadratic forms, pseudo-inverse	later
Tools	NumPy, SciPy	throughout

- Statistics and linear algebra are the two pillars of ML and DL.
- Learn the unmarked topics in one go; leave the “later” topics until an algorithm needs them.
- For every topic, know where ML uses it.

7. Sources

Built from

- CampusX, “Linear Algebra Roadmap for Machine Learning and Deep Learning | Complete Guide | CampusX”, YouTube, <https://www.youtube.com/watch?v=rIsCKVyh4dI>

Other references

- Trefethen, L. N. and Bau, D. (1997). *Numerical Linear Algebra*. SIAM. Lecture 11, “Least squares problems”, Algorithm 11.2.
- Koren, Y., Bell, R. and Volinsky, C. (2009). “Matrix factorization techniques for recommender systems.” *IEEE Computer* 42(8), 30–37.
- Deisenroth, M. P., Faisal, A. A. and Ong, C. S. (2020). *Mathematics for Machine Learning*. Cambridge University Press. Chapters 2–4; §4.5.

8. Key terms

Term	Meaning
Matrix factorisation (decomposition)	Writing a matrix as a product of simpler matrices

Term	Meaning
SVD (singular value decomposition)	A factorisation that works for any matrix, square or not
Quadratic form	An expression like $x^T Ax$: a sum of squared and cross terms of a vector's components
Moore-Penrose pseudo-inverse	A generalised inverse for matrices that are not square or have no inverse
NumPy	Python's library for arrays and linear algebra
SciPy	A library built on NumPy with more scientific routines