

The Poisson Distribution

Contents

1. Overview	2
2. What the Poisson distribution describes	3
2.1 Counts in an interval	3
2.2 Notation	3
2.3 The values a Poisson variable can take	4
3. The Poisson PMF	4
3.1 The example: a spike in questions	4
3.2 Two pieces of the formula	4
3.3 Exactly 7 questions	4
3.4 Scaling the rate to the interval	5
4. The graph of the Poisson distribution	6
5. Mean and variance	7
6. Probability of a range of counts	8
7. The Poisson distribution as a limit of the binomial	9
8. When a count is Poisson, and where it is used	10
9. Summary	11
10. Sources	11
11. Key terms	11

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



- **Builds on:** Probability mass function (PMF) ([Note 241](#)); Expected value and variance of a random variable ([Note 332](#)).
- **Compare with:** Bernoulli and binomial distributions ([Note 270](#)).

1. Overview

Key point: The Poisson distribution gives the probability of each count of events (0, 1, 2, ...) in a fixed interval of time or space, when we know only the average count λ .




Figure 1: Events on a time line, counted per day; the average count is the rate λ

Figure 1 shows the kind of data the Poisson distribution describes. Events (here, questions from students) arrive at random moments; we cut time into equal intervals (days) and count the events in each. The counts vary from day to day, but their average, the **rate** λ , stays about the same.

The Poisson distribution was named in one line in the [PDF Note](#): a discrete distribution of counts, with parameter λ . This Note teaches it in full:

- what kind of question it answers, and its notation;
- its PMF, with a worked example;
- the shape of its graph and how λ moves it;
- its mean and variance, which are both λ ;
- probabilities of a range of counts;
- its link to the binomial distribution, and where it is used.

2. What the Poisson distribution describes

Key point: A Poisson variable counts how often an event happens in one interval; instead of a success probability, we need the average count per interval.

2.1 Counts in an interval

Key point: The binomial distribution needs a number of trials and a probability; the Poisson distribution needs only a rate.

The binomial distribution counts successes in a fixed number n of trials, each with success probability p (see the [Bernoulli and binomial Note](#)). Many counts have no trials to count:

- the questions students send in one day;
- the flashes of a firefly in 10 seconds;
- customers entering a shop in one hour;
- typing errors on one page of a book.

For these we know how **often** the event happens on average, its frequency in a standard interval of time or distance. The **Poisson distribution** turns that average into a probability for every possible count.

For example, a firefly lights up 3 times in 10 seconds on average. The Poisson distribution answers questions such as: how likely is it that it lights up 8 times in 20 seconds?

2.2 Notation

Key point: $Y \sim \text{Po}(\lambda)$ reads “ Y follows a Poisson distribution with rate λ ”.

A distribution is often written as a short name with its parameters in brackets. The symbol $\sim$ reads “follows”:

Notation	Read as	Parameters
$X \sim \text{Bern}(p)$	Bernoulli with success probability p	p
$X \sim B(n, p)$	binomial with n trials and success probability p	n, p
$Y \sim \text{Po}(\lambda)$	Poisson with rate λ	λ

For example, $X \sim B(10, 0.6)$ means 10 trials with success probability 0.6 each, and $Y \sim \text{Po}(4)$ means a count that is 4 on average. The Poisson distribution has a single parameter, λ (lambda), the average number of events per interval.

2.3 The values a Poisson variable can take

Key point: The count starts at 0 and has no upper limit.

A count can never be negative, so the possible values start at 0. Unlike the binomial count, which stops at n , a Poisson count has no cap: in principle the students could send 50 questions in one day. Such large counts are possible, only extremely unlikely.

3. The Poisson PMF

Key point: $P(Y = y) = \frac{\lambda^y e^{-\lambda}}{y!}$, for $y = 0, 1, 2, \dots$

3.1 The example: a spike in questions

Key point: Students usually ask 4 questions a day; yesterday they asked 7. How likely was exactly 7?

Suppose we run an online help forum on probability. Students usually ask around 4 questions per day, but yesterday they asked 7. To judge whether this spike is unusual, we want the probability of exactly 7 questions in one day.

The three ingredients:

- the interval is one working day;
- the average count per interval is 4, so $\lambda = 4$;
- the count we ask about is 7, so $y = 7$.

3.2 Two pieces of the formula

Key point: $e \approx 2.718$ is a fixed number, and $e^{-\lambda}$ means $1/e^\lambda$.

The formula uses two pieces of notation:

- **Euler's number** e , also called Napier's constant, is a fixed number, $e \approx 2.71828$. It turns up across mathematics, physics and nature; here we only need its value.
- A **negative power** means one over the positive power: $a^{-n} = 1/a^n$. So $e^{-4} = 1/e^4 = 1/54.6 = 0.0183$.

The factorial $y!$ is the product $1 \times 2 \times \dots \times y$, with $0! = 1$ (as in the binomial coefficient of the [Bernoulli and binomial Note](#)).

3.3 Exactly 7 questions

Key point: $P(Y = 7) = 0.060$: about one day in 17 brings exactly 7 questions.

1. **In words:** raise the rate to the power of the count, multiply by e to the power of minus the rate, and divide by the factorial of the count.

2. Formula:

$$P(Y = y) = \frac{\lambda^y e^{-\lambda}}{y!}, \quad y = 0, 1, 2, \dots$$

3. **Example:** with $\lambda = 4$ and $y = 7$:

$$P(Y = 7) = \frac{4^7 e^{-4}}{7!} = \frac{16384 \times 0.0183}{5040} = \frac{300.1}{5040} = 0.0595$$

So there was only about a 6% chance of exactly 7 questions. The bar at 7 in the left panel of Figure 2 has this height.

Extra: A common slip is to read e^{-4} as 0.183. The correct value is 0.0183; with 0.183 the answer would come out as 0.595, ten times too large. The final answer 0.06 needs 0.0183.

Python: `stats.poisson` gives the PMF in one call.

```
import math
from scipy import stats

4 ** 7 * math.exp(-4) / math.factorial(7)    # 0.0595, by hand
# scipy: pmf(y, lambda)
stats.poisson.pmf(7, 4)                      # 0.0595
```




Figure 2: The Poisson PMF with $\lambda = 4$: the bar for exactly 7 questions (left) and the bars for 7 or more (right)

3.4 Scaling the rate to the interval

Key point: λ belongs to one interval length; for a longer or shorter interval, scale it in proportion.

The firefly flashes 3 times in 10 seconds on average, and we ask about 20 seconds. Twice the interval means twice the average, so $\lambda = 3 \times 20/10 = 6$.

1. **In words:** the probability of 8 flashes when 6 are expected.
2. **Formula:**

$$P(Y = 8) = \frac{6^8 e^{-6}}{8!}$$

3. **Example:**

$$P(Y = 8) = \frac{1679616 \times 0.002479}{40320} = \frac{4163.4}{40320} = 0.103$$

Extra: Using $\lambda = 3$ (the rate for 10 seconds) with a 20-second question is a common mistake. The rate and the interval must always match.

4. The graph of the Poisson distribution

Key point: One bar per count, starting at 0 and running on to the right without end; the bars are tallest near λ .

The graph of a Poisson distribution plots each count y against its probability. Figure 2 shows it for $\lambda = 4$:

- it starts at 0, since no event can happen a negative number of times;
- it peaks at 3 and 4 (both have probability 0.195);
- it has a tail to the right that never ends, though after about 12 the bars are too small to see.

The bars add to 1, as in every PMF (see the [PMF Note](#)).




Figure 3: Poisson PMFs for three rates; the dashed line marks the mean λ

Figure 3 shows how λ changes the shape:

- $\lambda = 1$: most days have 0 or 1 events; the distribution is strongly right-skewed (see the [skewness Note](#)).
- $\lambda = 4$: the peak moves right and the bars spread out; a mild right skew remains.

- $\lambda = 10$: the peak sits at 9 and 10, the spread is wider, and the shape is close to a symmetric bell.

Extra: The skewness of a Poisson distribution is $1/\sqrt{\lambda}$ (NIST Handbook §1.3.6.6.19): 1 for $\lambda = 1$, 0.5 for $\lambda = 4$, 0.32 for $\lambda = 10$. It shrinks towards 0 as λ grows, which matches Figure 3: the shape gets closer to a symmetric bell, a normal distribution with mean λ and variance λ (see the [normal distribution Note](#)).

5. Mean and variance

Key point: A Poisson variable has mean λ and variance λ : one number fixes both its centre and its spread.

The expected value is the sum of every value times its probability (see the [expected value Note](#)). Applied to the Poisson PMF, this long sum simplifies to λ . The variance, from the same shortcut formula $E[Y^2] - (E[Y])^2$, is also λ .

1. **In words:** the average count is the rate, and so is the variance; the standard deviation is the square root of the rate.
2. **Formula:**

$$E[Y] = \lambda, \quad \text{Var}(Y) = \lambda, \quad \sigma = \sqrt{\lambda}$$

3. **Example:** for the questions, $\lambda = 4$: on average 4 questions per day, variance 4, standard deviation $\sqrt{4} = 2$. So 7 questions is $(7 - 4)/2 = 1.5$ standard deviations above the mean.

A simulation of 100,000 days with `rng.poisson(lam=4)` gives an average of 3.994 and a variance of 3.987, both close to 4.

Extra: Why the mean is λ . Writing out the sum and cancelling one y against $y!$:

$$E[Y] = \sum_{y=0}^{\infty} y \frac{\lambda^y e^{-\lambda}}{y!} = \lambda \sum_{y=1}^{\infty} \frac{\lambda^{y-1} e^{-\lambda}}{(y-1)!} = \lambda \times 1 = \lambda$$

The remaining sum is the Poisson PMF summed over all counts, which is 1. The same trick, cancelling $y(y-1)$ against $y!$, gives $E[Y(Y-1)] = \lambda^2$. Then $E[Y^2] = E[Y(Y-1)] + E[Y] = \lambda^2 + \lambda$, and $\text{Var}(Y) = \lambda^2 + \lambda - \lambda^2 = \lambda$.

Extra: Mean equal to variance is a quick check on real count data. A variance clearly larger than the mean is called **overdispersion**. The notebook tests two causes by breaking one Poisson condition at a time while keeping the mean near 4:

Simulated days (100,000)	Mean	Variance	$P(Y \geq 12)$ seen	Poisson with the same mean says
constant rate, independent events	3.99	3.99	0.0008	0.0009
rate varies from day to day	3.99	11.99	0.0382	0.0009
events come in pairs (not independent)	4.01	7.96	0.0163	0.0009

Each broken condition raises the variance above the mean, and the Poisson model then underestimates days with 12 or more events by a factor of about 20 to 40. With both conditions kept, mean, variance and tail all match.

6. Probability of a range of counts

Key point: For “at least”, “at most” or “between”, add the probabilities of the single counts in the range.

The counts 0, 1, 2, ... are separate outcomes: one day cannot have both exactly 6 and exactly 7 questions. So the probability of a range is the **sum** of the bars in it, as for any discrete distribution (see the [PMF Note](#)).

1. **In words:** the probability of 7 or more questions is everything except 0 to 6.
2. **Formula:**

$$P(Y \geq 7) = 1 - P(Y \leq 6) = 1 - \sum_{y=0}^6 \frac{4^y e^{-4}}{y!}$$

3. **Example:** the bars for 0 to 6 are 0.0183, 0.0733, 0.1465, 0.1954, 0.1954, 0.1563 and 0.1042:

$$P(Y \leq 6) = 0.8893, \quad P(Y \geq 7) = 1 - 0.8893 = 0.111$$

The right panel of Figure 2 shows these red bars. A day with 7 or more questions happens about once in 9 days: unusual, but not rare.

More examples with $\lambda = 4$:

Question	Probability
no questions at all, $P(Y = 0)$	0.018
at most 2, $P(Y \leq 2)$	0.238
between 2 and 6, $P(2 \leq Y \leq 6)$	0.798


```

Python: cdf adds the bars up to a count; sf adds the bars above it.
questions = stats.poisson(4)
questions.sf(6) # 0.111: P(Y >= 7)
questions.cdf(2) # 0.238: P(Y <= 2)
questions.cdf(6) - questions.cdf(1) # 0.798: P(2 <= Y <= 6)

```

7. The Poisson distribution as a limit of the binomial

Key point: A binomial count with many trials and a small success probability is close to a Poisson count with $\lambda = np$.

Extra: This whole section is extra material.

A website has 1000 visitors a day, and each visitor buys with probability 0.004. The number of buyers is binomial, $B(1000, 0.004)$, with mean $np = 4$. The count is also, very nearly, $Po(4)$.




Figure 4: A day cut into n moments, each with a question with probability $4/n$ (top); the binomial PMF $B(n, 4/n)$ (bars) settles onto the Poisson PMF with $\lambda = 4$ (dots) as n grows from 5 to 1000

Figure 4 keeps $np = 4$ and lets n grow. The top strip shows one simulated day cut into n moments, with a dot where a question arrived. Watch the bars: at $n = 5$ they are bunched around 4, and as the moments get finer they spread out and sit on the Poisson dots.

n	p	Largest gap between the two PMFs
10	0.4	0.055
40	0.1	0.011
1000	0.004	0.0004

With 10 trials the binomial is narrower than the Poisson: its variance $np(1-p) = 2.4$ is below 4. With 1000 trials $1-p$ is almost 1, the variance $np(1-p) = 3.98$ is almost λ , and the bars sit on the dots.

The binomial limit is where the Poisson distribution comes from. Cut a day into many tiny moments; in each, a question arrives or not, with a tiny probability. The count is binomial with huge n and tiny p , and in the limit it becomes Poisson. The approximation needs n large and p small (Ross §4.7). For example, at $n = 20$ and $p = 0.05$ the notebook finds a largest gap of about 0.01, and the gap grows when p is larger.

8. When a count is Poisson, and where it is used

Key point: Events must happen one at a time, independently, at a constant average rate; then counts of arrivals, calls, clicks or defects are Poisson.

Extra: This whole section is extra material.

A count follows a Poisson distribution when these three conditions hold (Ross §4.7):

1. **Events are independent.** One question does not make the next more or less likely.
2. **The rate is constant.** The average count is the same for every interval of the same length; an exam week with more questions breaks this.
3. **Events happen one at a time.** Two events never happen at exactly the same moment.

Typical Poisson counts:

- customers arriving at a shop or a website per hour;
- calls to a support centre per minute;
- defects per metre of cable, typos per page;
- goals in a football match (Maher 1982);
- rare-event counts, such as accidents at a crossing per month.

In machine learning, a **target** (the output we predict) that is a count, such as bike rentals per hour or insurance claims per year, is often modelled with **Poisson regression**. The model predicts λ for each **observation** (one record, a row of the data table) from its **features** (the input variables); scikit-learn has it as **PoissonRegressor** (scikit-learn docs). Counts of words in a document also appear in text models.

9. Summary

	Binomial	Poisson
Counts	successes in n trials	events in one interval
Notation	$B(n, p)$	$Po(\lambda)$
Parameters	n, p	λ
Values	0 to n	0, 1, 2, ... without limit
PMF	$\binom{n}{y} p^y (1-p)^{n-y}$	$\lambda^y e^{-\lambda} / y!$
Mean, variance	$np, np(1-p)$	λ, λ
Example	2 likes out of 3 viewers, $p = 0.5$: 3/8	7 questions when 4 are usual: 0.060

- The Poisson distribution needs only the average count per interval, λ ; scale λ when the interval changes.
- Its graph starts at 0, peaks near λ and has an endless right tail; a larger λ moves it right and spreads it out.
- Mean and variance are both λ .
- A range of counts has the sum of the single-count probabilities.
- A binomial with large n and small p is approximately Poisson with $\lambda = np$.

10. Sources

Built from

- 365 Data Science, “Data Science & Statistics Tutorial: The Poisson Distribution”, YouTube, <https://www.youtube.com/watch?v=BbLfV0wOeyc>

Other references

- Maher, M. J. (1982). “Modelling Association Football Scores”. *Statistica Neerlandica* 36(3).
- NIST/SEMATECH (2012). *e-Handbook of Statistical Methods*. itl.nist.gov/div898/handbook. Section 1.3.6.6.19, Poisson distribution.
- Ross, S. (2010). *A First Course in Probability*, 8th ed. Pearson. Section 4.7, the Poisson random variable.
- scikit-learn documentation. `sklearn.linear_model.PoissonRegressor`.

11. Key terms

Term	Meaning
Poisson distribution	The distribution of the number of events in a fixed interval, given their average count λ

Term	Meaning
Rate λ	The average number of events per interval; the only parameter of the Poisson distribution
$Y \sim \text{Po}(\lambda)$	Notation: Y follows a Poisson distribution with rate λ ; likewise $\text{Bern}(p)$ and $B(n, p)$
Euler's number e	A fixed number, about 2.71828, that appears in the Poisson PMF
Negative power	$a^{-n} = 1/a^n$
Overdispersion	Count data whose variance is clearly larger than its mean, a sign that a Poisson model does not fit
Poisson regression	A model that predicts the rate λ of a count target from the features