

One-way ANOVA

Contents

1. Overview	2
2. Why not several t-tests?	3
3. The hypotheses	3
4. Splitting the variation	4
4.1 Total sum of squares	4
4.2 Between-group sum of squares	5
4.3 Within-group sum of squares	5
5. The F statistic	5
5.1 Mean squares	5
5.2 F	6
6. The F distribution and the decision	6
7. Assumptions	8
8. Case study: age by class on the Titanic	8
8.1 The test	8
8.2 Checking the assumptions	9
8.3 Which classes differ? Tukey's test	10
9. ANOVA in machine learning	10
10. Summary	11
11. Key terms	11

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



- **Builds on:** Variance ([Note 222](#)); Hypothesis testing: null and alternative ([Note 291](#)).
- **Compare with:** T-tests: one-sample, two-sample, paired ([Note 570](#)).

1. Overview

Key point: One-way ANOVA tests whether three or more group means are equal by comparing two variances: how far the group means lie from each other, and how far the values lie from their own group mean.




Figure 1: Same group means, different spread: the gaps between means are clear on the left and lost in the noise on the right

ANOVA (**analysis of variance**) was named in the [what is statistics Note](#) as the test that compares the means of several groups. In the [choosing a hypothesis test Note](#) it is the test for one numerical column against a categorical column with three or more categories.

Figure 1 shows the whole idea. Both panels have the same group means, 5, 7 and 9. On the left the values sit close to their group mean, so the gaps between the means stand out. On the right the values are spread widely, and the same gaps could easily be chance. ANOVA turns this comparison into one number, the F statistic.

This Note covers:

- why we do not run a t-test for every pair of groups;
- the split of the total variation into between-group and within-group parts;
- the F statistic, the F distribution and a worked example;

- the assumptions, a Titanic case study, and the follow-up question “which groups differ?”.

Extra: All of this Note is extra material beyond the “which test when” guide; the guide only names ANOVA.

2. Why not several t-tests?

Key point: Each t-test at $\alpha = 0.05$ has a 5% false-alarm risk; three of them together have a 14% risk, ten of them 40%.

A two-sample t-test compares two group means (see the [two-sample and paired t-tests Note](#)). With three groups A, B and C, we could run three t-tests: A against B, A against C, B against C.

The trouble is the Type I error (see the [errors, power and tails Note](#)). Each test, run at $\alpha = 0.05$, wrongly rejects a true H_0 5% of the time. Over several tests these risks add up.

1. **In words:** the chance of at least one false alarm is one minus the chance that every test stays quiet.
2. **Formula:** for m independent tests,

$$P(\text{at least one Type I error}) = 1 - (1 - \alpha)^m$$

3. **Example:** three groups need $m = 3$ tests:

$$1 - 0.95^3 = 1 - 0.857 = 0.143$$

Five groups need $m = 10$ tests: $1 - 0.95^{10} = 0.40$.

So with five groups whose true means are all equal, we would “find” a difference 40% of the time. ANOVA asks one question about all groups at once, with a single 5% risk.

3. The hypotheses

Key point: H_0 : all group means are equal; H_1 : at least one group mean differs from the others.

With k groups:

$$H_0 : \mu_1 = \mu_2 = \cdots = \mu_k, \quad H_1 : \text{at least one } \mu_i \text{ differs}$$

H_1 is not “all means differ”. Rejecting H_0 tells us that some difference exists, not which groups differ; section 8 answers that.

“One-way” means one categorical column defines the groups. The worked example uses $k = 3$ sections of a class, with the marks (out of 10) of 3 students each:

Section A	Section B	Section C
4, 5, 6	6, 7, 8	8, 9, 10
mean 5	mean 7	mean 9

The **grand mean**, the mean of all $N = 9$ values, is 7. These are the left-hand groups of Figure 1.

4. Splitting the variation

Key point: Total variation = variation between the group means + variation within the groups: $SST = SSB + SSW$.




Figure 2: The variation split into its two parts, and the steps from there to F

Variance is built from squared distances from a mean (see the [measures of dispersion Note](#)). ANOVA works with the **sums of squares** before dividing, and splits them in two, as Figure 2 shows.

4.1 Total sum of squares

Key point: SST measures how far all values lie from the grand mean.

1. **In words:** square each value's distance from the grand mean and add over all values.
2. **Formula:**

$$SST = \sum_{\text{all values}} (x - \bar{x})^2$$

where $\bar{x}$ is the grand mean.

3. **Example:** distances from 7 are $-3, -2, -1; -1, 0, 1; 1, 2, 3$:

$$SST = 9 + 4 + 1 + 1 + 0 + 1 + 1 + 4 + 9 = 30$$

4.2 Between-group sum of squares

Key point: SSB measures how far the group means lie from the grand mean, each counted once per value in its group.

1. **In words:** square each group mean's distance from the grand mean, multiply by the group size, and add over groups.
2. **Formula:**

$$SSB = \sum_{\text{groups}} n_i(\bar{x}_i - \bar{x})^2$$

3. **Example:**

$$SSB = 3(5 - 7)^2 + 3(7 - 7)^2 + 3(9 - 7)^2 = 12 + 0 + 12 = 24$$

4.3 Within-group sum of squares

Key point: SSW measures how far the values lie from their own group mean: the noise inside the groups.

1. **In words:** square each value's distance from its own group mean and add over all values.
2. **Formula:**

$$SSW = \sum_{\text{groups}} \sum_{\text{values in group}} (x - \bar{x}_i)^2$$

3. **Example:** each group gives $1 + 0 + 1 = 2$, so

$$SSW = 2 + 2 + 2 = 6$$

The check: $SSB + SSW = 24 + 6 = 30 = SST$. Most of the variation here lies between the groups.

5. The F statistic

Key point: $F = MSB/MSW$: the between-group variance divided by the within-group variance; a large F means the group means differ by more than the noise explains.

5.1 Mean squares

Key point: Dividing each sum of squares by its degrees of freedom turns it into a variance, a **mean square**.

The degrees of freedom are:

- between groups: $k - 1$ (here $3 - 1 = 2$), since k group means vary around one grand mean;
- within groups: $N - k$ (here $9 - 3 = 6$), since each group loses one degree of freedom to its own mean, as in Bessel's correction (see the [measures of dispersion Note](#)).

The two add to $N - 1 = 8$, the degrees of freedom of SST .

5.2 F

1. **In words:** divide each sum of squares by its degrees of freedom, then divide the between-group mean square by the within-group one.

2. **Formula:**

$$MSB = \frac{SSB}{k - 1}, \quad MSW = \frac{SSW}{N - k}, \quad F = \frac{MSB}{MSW}$$

3. **Example:**

$$MSB = \frac{24}{2} = 12, \quad MSW = \frac{6}{6} = 1, \quad F = \frac{12}{1} = 12$$

If H_0 is true, both mean squares estimate the same noise variance, and F is close to 1. Here the between-group variance is 12 times the within-group variance.

For the right-hand groups of Figure 1 (1, 5, 9; 3, 7, 11; 5, 9, 13) the means and so $SSB = 24$ are the same, but $SSW = 96$. Then $MSW = 96/6 = 16$ and $F = 12/16 = 0.75$: no evidence of a difference.

Extra: The results are usually set out as an **ANOVA table**:

Source	SS	df	MS	F	p
between groups	24	2	12	12	0.008
within groups	6	6	1		
total	30	8			

6. The F distribution and the decision

Key point: Under H_0 , F follows an F distribution with $k - 1$ and $N - k$ degrees of freedom; the p-value is its right-tail area beyond our F .




Figure 3: The F distribution with 2 and 6 degrees of freedom: the whole curve (left) and its tail (right)

The **F distribution** is the distribution of a ratio of two independent variances. Like the chi-square distribution (see the [chi-square tests Note](#)), it is never negative and is skewed to the right. It has two parameters: the degrees of freedom of the top and of the bottom variance.

Only a large F counts against H_0 (group means further apart than noise explains), so the test is always right-tailed. Figure 3 shows the F distribution with 2 and 6 degrees of freedom:

- the 5% critical value is 5.14;
- our $F = 12$ lies far beyond it;
- the tail area beyond 12 is $p = 0.008$.

Since $0.008 \leq 0.05$, we reject H_0 : the three sections do not all have the same mean mark.

Python: `f_oneway` takes one array per group.

```
from scipy import stats
```

```
stats.f_oneway([4, 5, 6], [6, 7, 8], [8, 9, 10])
# statistic = 12.0, pvalue = 0.008
stats.f.ppf(0.95, 2, 6)      # 5.14, the critical value
```

Extra: With exactly two groups, ANOVA and the equal-variance two-sample t-test are the same test: $F = t^2$ and the p-values match. For sections A and B alone, $t^2 = 6$ and $F = 6$.

7. Assumptions

Key point: Independent observations, roughly normal values in each group, and equal variances across groups.

ANOVA makes the same assumptions as the two-sample t-test (see the [two-sample and paired t-tests Note](#)), for every group:

1. **Independence.** The groups are separate (no one is in two groups), and the observations do not influence each other.
2. **Normality.** The values in each group are roughly normal. With about 30 or more per group, the central limit theorem makes this less critical. Check with a Q-Q plot or the Shapiro-Wilk test (see the [one-sample t-test Note](#)).
3. **Equal variances.** The groups have about the same variance; Levene's test checks this, with H_0 : "the variances are equal".

When the assumptions fail:

- unequal variances: **Welch's ANOVA**, or scipy's `stats alexandergovern`, drops the equal-variance assumption;
- strongly non-normal small groups: the **Kruskal-Wallis test** (`stats.kruskal`) compares the groups by ranks instead of means.

8. Case study: age by class on the Titanic

Key point: Mean ages of 38.2, 29.9 and 25.1 years in the three classes give $F = 57.4$ and $p \approx 10^{-23}$; Tukey's test shows that every pair of classes differs.

8.1 The test

Key point: Did the mean age of passengers differ between first, second and third class?

The Kaggle Titanic training file has 714 passengers with a known age.




Figure 4: Age of Titanic passengers by class: every passenger as a dot, the class mean and one standard deviation either side in black

Class	Passengers	Mean age	Standard deviation
1	186	38.2	14.8
2	173	29.9	14.0
3	355	25.1	12.5

Figure 4 shows a clear downward step in the means, with much overlap between the classes.

1. **Hypotheses.** $H_0 : \mu_1 = \mu_2 = \mu_3$; H_1 : at least one class mean differs.
2. **Significance level.** $\alpha = 0.05$.
3. **Statistic.** $SSB = 20,930$ with 2 df, $SSW = 129,527$ with $714 - 3 = 711$ df:

$$MSB = \frac{20,930}{2} = 10,465, \quad MSW = \frac{129,527}{711} = 182.2, \quad F = \frac{10,465}{182.2} = 57.4$$

4. **P-value.** $p = 7.5 \times 10^{-24}$, far below 0.05; the 5% critical value is only 3.01.
5. **Decide.** Reject H_0 : mean age differed between the classes.

8.2 Checking the assumptions

Key point: Levene's test rejects equal variances ($p = 0.004$), but the tests that drop that assumption agree: the classes differ.

- **Normality:** Shapiro-Wilk rejects normality for classes 2 and 3. With 173 and 355 passengers, the central limit theorem protects the means.

- **Equal variances:** Levene gives $p = 0.004$: the variances (standard deviations 14.8, 14.0, 12.5) differ.

Both checks raise doubts, so we confirm with the tests of section 7. Alexander-Govern gives $p = 4 \times 10^{-21}$ and Kruskal-Wallis gives $p = 1 \times 10^{-21}$. The conclusion stands.

8.3 Which classes differ? Tukey's test

Key point: A post-hoc test compares every pair of groups after ANOVA rejects, while keeping the overall false-alarm risk at 5%.

ANOVA said that at least one mean differs. A **post-hoc test** finds which. The usual choice is **Tukey's HSD** (honestly significant difference), which compares every pair with the overall Type I error held at α :

Pair	Difference in mean age	p-value
class 1 – class 2	8.4 years	2×10^{-8}
class 1 – class 3	13.1 years	below 10^{-9}
class 2 – class 3	4.7 years	0.0005

All three pairs differ: first-class passengers were the oldest, third-class the youngest.

Python: ANOVA, the assumption checks and Tukey's test on the Titanic ages.

```
ages = [titanic.loc[titanic["Pclass"] == c, "Age"].dropna()
        for c in (1, 2, 3)]
stats.f_oneway(*ages)           # F = 57.4, p = 7.5e-24
stats.levene(*ages)             # p = 0.004
stats.alexandergovern(*ages)    # p = 4.4e-21
stats.tukey_hsd(*ages)          # every pair
```

9. ANOVA in machine learning

Key point: `f_classif` scores each numerical feature by its ANOVA F against the class labels, for feature selection.

`SelectKBest` can score features with `f_classif` instead of `chi2` (see the [pipelines Note](#) and the [errors, power and tails Note](#)). For each numerical feature it runs a one-way ANOVA with the target classes as groups. A feature whose mean differs strongly between the classes gets a large F, and is kept.

10. Summary

	Several t-tests	One-way ANOVA
Question	is this pair different?	do any of the k means differ?
Type I risk, 3 groups	14%	5%
Statistic	t per pair	$F = MSB/MSW$
Distribution	t with $n_1 + n_2 - 2$ df	F with $k - 1$ and $N - k$ df
Follow-up	none needed	Tukey's HSD to find which pairs differ

- $SST = SSB + SSW$: total variation splits into variation between and within groups.
- F compares the between-group variance with the within-group variance; near 1 under H_0 , large when the means differ.
- The test is right-tailed; reject H_0 when $p \leq \alpha$.
- Assumptions: independence, normality in each group, equal variances; Welch's ANOVA or Kruskal-Wallis when they fail.
- Rejecting H_0 says some mean differs; Tukey's HSD says which.

11. Key terms

Term	Meaning
One-way ANOVA	A test of whether three or more group means are equal, with the groups defined by one categorical column
Grand mean	The mean of all values from all groups together
Sum of squares (SST, SSB, SSW)	Total, between-group and within-group squared distances; $SST = SSB + SSW$
Mean square (MSB, MSW)	A sum of squares divided by its degrees of freedom: a variance
F statistic	MSB/MSW : between-group variance over within-group variance
F distribution	The distribution of a ratio of two variances; two degrees-of-freedom parameters, right-skewed
ANOVA table	The table of SS, df, MS, F and p for each source of variation
Familywise error rate	The probability of at least one Type I error over several tests: $1 - (1 - \alpha)^m$
Post-hoc test	A test run after ANOVA rejects, to find which groups differ

Term	Meaning
Tukey's HSD	A post-hoc test comparing every pair of groups with the overall Type I error held at α
Welch's ANOVA	A version of ANOVA that does not assume equal variances
Kruskal-Wallis test	A rank-based alternative to one-way ANOVA that does not assume normality