

Deep Learning: What It Covers and What to Know First

Contents

1. Overview	2
2. The four families of networks	2
2.1 Artificial neural networks	3
2.2 CNNs, RNNs and transformers	4
2.3 The software	4
3. What to know first	4
3.1 Python and the basics of ML	6
3.2 Linear algebra	6
3.3 Derivatives	6
4. Summary	7
5. Sources	7
6. Key terms	7

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



- **Builds on:** Machine learning ([Note 2](#)).
- **Leads to:** Types of neural networks ([Note 1003](#)).
- **Compare with:** Machine learning ([Note 2](#)).

1. Overview

Key point: The Deep Learning Notes move through four families of neural networks: ANNs, CNNs, RNNs and transformers. Before starting, we need Python, the basics of ML, some linear algebra and derivatives.




Figure 1: The four families of networks in the Deep Learning Notes, the data each one suits, and what to know first

Figure 1 shows the plan. Each family builds on the one before it, and each suits a different kind of data. The first family, the **artificial neural network (ANN)** (G-216), takes the most space, because every later network reuses its ideas.

2. The four families of networks

Key point: ANNs come first and take about a third of the Notes; CNNs, RNNs and transformers each reuse the ANN machinery for one kind of data.




Figure 2: One picture from each family, taken from the Notes ahead: the perceptron (Note 1004), learned convolution filters (Note 1042), a recurrent layer (Note 1056) and an attention map (Note 1069)

Figure 2 previews what each family looks like once we reach it.

2.1 Artificial neural networks

Key point: From one artificial neuron to a full network: how it predicts, how it is trained, and how to make it train well.

The ANN part runs in this order:

1. **Foundations:** what deep learning is, how it differs from ML, the main types of network, their history and applications.
2. **The perceptron** (G-1486): a single artificial **neuron** (G-1318), how it predicts, how it is trained, and the problem that a single neuron cannot solve.
3. **The multi-layer perceptron (MLP)** (G-1270): many neurons in **layers** (G-1056), the names of its **weights** (G-2106) and **biases** (G-287), and how it makes a prediction.
4. **Training:** **loss functions** (G-1130), **backpropagation** (G-247) and **gradient descent** (G-862).
5. **First projects** in the **Keras** (G-1003) library: one classification and one regression problem.
6. **Improving a network:** vanishing and exploding gradients, dropout, regularisation, activation functions, weight initialisation, batch normalisation, optimisers and hyperparameter tuning.

2.2 CNNs, RNNs and transformers

Key point: CNNs for images, RNNs for sequences, transformers for language.

- **Convolutional neural networks (CNNs)** (G-484) work best on images: convolution, pooling, pretrained models and transfer learning.
- **Recurrent neural networks (RNNs)** (G-1647) work on sequences such as text, speech and time series: RNNs, LSTMs and GRUs.
- **Transformers** (G-2007) replace recurrence with **attention** (G-226; Vaswani et al. 2017) and are the basis of language models such as GPT-3 (Brown et al. 2020): attention, self-attention and the full encoder-decoder transformer.

Extra: Generative networks (GANs, autoencoders), object detection and image segmentation are not covered in these Notes. The [types of neural networks Note](#) describes GANs and autoencoders in a paragraph each.

2.3 The software

Key point: Code uses TensorFlow with Keras, the pair most used in industry; PyTorch is the common choice in research.

Most deep learning code in these Notes uses **TensorFlow** (G-1959) with **Keras**, its built-in high-level interface. The other major library, **PyTorch** (G-1595), is used more in research than in industry. Where a Note has code, its Notebook rebuilds it for current library versions.

3. What to know first

Key point: Four things: basic Python, the overall flow of an ML project, vectors and matrices, and derivatives.




Figure 3: Where each prerequisite shows up in one layer of a network, $\sigma(W\mathbf{x} + \mathbf{b})$

Figure 3 ties the four prerequisites to one formula:

- **linear algebra** (G-1090) builds $W\mathbf{x} + \mathbf{b}$;
- logistic regression supplies the sigmoid;
- derivatives train the weights;
- Python runs the code.

The Deep Learning Notes assume four things. Each row of the table links to the Notes that teach it:

Prerequisite	How much	Where it is taught
Python	read and write simple code; NumPy and pandas	Python boxes from the toy project Note onward; setup in the environment Note
ML basics	how a model is trained, how data is prepared, how a model is served	AI vs ML vs DL , ML development life cycle , toy project
Linear algebra	vectors, dot product, matrices and matrix multiplication	vectors , dot product , linear transformations and matrices , matrix multiplication
Derivatives	slopes, partial derivatives, the chain rule	derivatives , partial derivatives and gradients , gradient descent

3.1 Python and the basics of ML

Key point: We do not need every ML algorithm, only the overall flow: prepare data, train, evaluate, serve.

Every Notebook is written in Python, so we need to read loops, functions and NumPy arrays without trouble. From ML we need the overall flow of a project rather than every algorithm. The [toy project Note](#) walks through that flow once, from loading data to a website.

A few ML Notes come back again and again in the Deep Learning Notes, so they are worth reading first:

- **Logistic regression** (G-1120): the [perceptron trick](#), the [sigmoid function](#) and the [log loss](#). A single neuron with a **sigmoid** (G-1798) activation computes $\sigma(\mathbf{w}^T \mathbf{x} + b)$, the same formula as logistic regression.
- **Gradient descent**: the [gradient descent Note](#) and its [stochastic](#) variant. Every network is trained this way.
- **Tensors** (G-1957): the [tensors Note](#). TensorFlow stores inputs, weights and outputs as tensors (TensorFlow guide, Tensors).

3.2 Linear algebra

Key point: A layer of a network is a matrix multiplication plus a shift, so matrices and the dot product are the core of every prediction.

A neural network spends almost all its time multiplying matrices; deep learning libraries are built on linear algebra. One neuron computes a **dot product** (G-634) of its inputs and weights. A whole layer computes a **matrix product** (G-1179) $W\mathbf{x} + \mathbf{b}$, as previewed in section 7.3 of the [linear transformations and matrices Note](#).

3.3 Derivatives

Key point: Training adjusts every weight against the slope of the loss, and the chain rule carries that slope back through the layers.

Training is like walking downhill in fog: we cannot see the valley, but we can feel which way the ground slopes under our feet and step that way. The slope is a **derivative** (G-595), one per weight. The [derivatives Note](#) and the [partial derivatives and gradients Note](#) teach everything we need.

Extra: Training a network is calculus. Backpropagation applies the **chain rule** (G-371) layer by layer; section 10 of the [Jacobian Note](#) previews it.

4. Summary

Family	Data it suits	Main idea
ANN	tables of numbers	neurons in layers, trained by backpropagation
CNN	images	small filters slid over the image
RNN	text, speech, time series	a memory passed from step to step
Transformer	text (language models)	attention between all positions at once

- The ANN part comes first and is the longest; the later families reuse it.
- Code uses TensorFlow with Keras; PyTorch is the research favourite.
- Before starting: Python, the flow of an ML project, vectors and matrices, and derivatives.
- Read the logistic regression and gradient descent Notes first: a sigmoid neuron computes the same formula as logistic regression, and every network is trained with gradient descent.

5. Sources

Built from

- CampusX, “100 Days of Deep Learning | Course Announcement”, YouTube, https://www.youtube.com/watch?v=2dH_qjc9mFg

Other references

- Vaswani et al., “Attention Is All You Need”, NeurIPS 2017.
- Brown et al., “Language Models are Few-Shot Learners”, NeurIPS 2020.
- TensorFlow guide, “Introduction to Tensors”, [tensorflow.org](https://www.tensorflow.org/).
- Linux Foundation, “Meta Transitions PyTorch to the Linux Foundation” (press release, 12 September 2022), [linuxfoundation.org](https://www.linuxfoundation.org/), press releases.

6. Key terms

Term	Meaning
TensorFlow	Google’s deep learning library
Keras	The high-level interface built into TensorFlow for defining and training networks
PyTorch	Deep learning library first built at Meta (Facebook), run by the PyTorch Foundation since 2022 (Linux Foundation 2022); most used in research