

What Is Deep Learning? Deep Learning vs Machine Learning

Contents

1. Overview	2
2. Two definitions of deep learning	3
2.1 The simple definition: ML with neural networks	3
2.2 The parts of a neural network	3
2.3 The technical definition: representation learning	4
2.4 Layers extract features of rising complexity	5
3. Why deep learning became famous	6
3.1 Applicability	6
3.2 Performance	7
4. Five differences between DL and ML	8
4.1 Data	8
4.2 Hardware	8
4.3 Training time and prediction time	9
4.4 Feature selection	9
4.5 Interpretability	10
4.6 DL does not replace ML	10
5. Why deep learning took off after 2010	11
5.1 Datasets	11
5.2 Hardware	12
5.3 Frameworks	13
5.4 Architectures and transfer learning	14
5.5 Community	14
6. Summary	15
7. Sources	15
8. Key terms	16

Where this fits: Pipeline map step 0 (Foundations). See the [Course map](#).



- **Builds on:** Machine learning ([Note 2](#)); Linear transformations and matrices ([Note 500](#)); Matrix multiplication as composition ([Note 510](#)).
- **Leads to:** Types of neural networks ([Note 1003](#)); Multi-layer perceptron (MLP) ([Note 1003](#)); Convolutional neural network (CNN) ([Note 1040](#)).
- **Compare with:** Machine learning ([Note 2](#)); Feature engineering ([Note 23](#)).

1. Overview

Key point: Deep learning is the part of ML that uses neural networks: layers of simple units that learn their own features from raw data. Deep learning needs more data, hardware and time than ML, and it is harder to explain, but on images, text and speech it is far stronger.

The [AI vs ML vs DL Note](#) already placed **deep learning (DL)** (G-568) inside ML and showed its two big advantages: it learns features by itself, and it keeps improving with more data. This Note goes further in three directions:

1. What a neural network looks like, and the technical definition of DL built on **representation learning** (G-1670).
2. Five practical differences between DL and ML.
3. Why DL took off only after about 2012, although its ideas are much older.




Figure 1: Who finds the features. In ML we choose them (size, colour); in DL the layers learn them, from edges to shapes to a face

Figure 1 shows the idea that runs through this Note: in DL the features come from the network's layers, not from us.

2. Two definitions of deep learning

Key point: The simple definition: DL is ML with neural networks, inspired by the brain. The technical one: DL learns its own representations, layer by layer, from raw data.

2.1 The simple definition: ML with neural networks

Key point: ML finds the input-output relationship with statistical techniques; DL finds it with a neural network.

Both ML and DL do the same job: in **supervised learning** (G-1919), they find the relationship between the inputs and the output. Most ML algorithms do it with statistical techniques. DL does it with a **neural network** (G-1316), a structure loosely inspired by the brain (see section 5.1 of the [AI vs ML vs DL Note](#)).

The reasoning behind the brain as a model is simple: to build intelligent machines, copy the most intelligent thing we know. The copy is very loose, as the [perceptron Note](#) shows.

2.2 The parts of a neural network

Key point: Neurons are arranged in layers: an input layer, one or more hidden layers and an output layer. Each connection carries a weight. Many hidden layers make a network “deep”.




Figure 2: The parts of an artificial neural network: input layer, two hidden layers, output layer

Figure 2 shows the simplest kind of network, the **artificial neural network (ANN)** (G-216):

- Each circle is a **neuron** (G-1318), also called a **perceptron** (G-1486): the network’s basic unit.
- Each line connecting two neurons carries a **weight** (G-2106), a number the network learns.
- Neurons in one column form a **layer** (G-1056). The **input layer** (G-952) takes the data, one node per **feature** (G-772; an input variable, one column of the data table). The **output layer** (G-1424) gives the prediction.
- Every layer in between is a **hidden layer** (G-890). We can add as many as we like.

A network with many hidden layers is called **deep**, and this is where the name “deep learning” comes from.

The ANN is only one type of network. Three others:

- **convolutional neural networks (CNNs)** (G-484) work best on images;
- **recurrent neural networks (RNNs)** (G-1647) work on speech and text;
- **generative adversarial networks (GANs)** (G-840) create new images and text.

The [types of neural networks Note](#) describes each one.

2.3 The technical definition: representation learning

Key point: Representation learning means the algorithm finds the useful features in raw data by itself, instead of us engineering them by hand.

Start with a task that looks easy: read a handwritten 3. Figure 3 shows three 3s written by three people.




Figure 3: Three handwritten 3s from the MNIST digits, each 28 by 28 pixels, and (in red) the pixels that are ink in all three

We see the same digit three times. The pixels do not agree:

1. Each image is a grid of $28 \times 28 = 784$ pixels. The three 3s put ink on 137, 100 and 100 of them.

2. Only 34 pixels are ink in all three images (the red panel).
3. So a hand-written rule such as “these pixels must be dark” fits one 3 and fails on the next.

A rule on raw pixels does not work. We need descriptions that stay the same when the handwriting changes, such as “a curve open to the left, twice”. Such a description is a **feature**, and a set of features that describes the data well is a **representation**. Writing these features by hand for every digit, face or word is not practical, so we let the algorithm find them.

A more technical definition reads: deep learning is part of the broader family of ML methods based on artificial neural networks **with representation learning**. Its algorithms use multiple layers to progressively extract higher-level features from the raw input.

Representation learning (also called **feature learning**) is a set of techniques that let a system discover, from raw data, the features it needs for a task. Representation learning replaces manual [feature engineering](#) (G-761): the machine both learns the features and uses them.

Take a dog-vs-cat classifier:

- **With ML**, we design the features ourselves: size (dogs are usually bigger), colour, and other physical attributes. Then we give these features to the algorithm.
- **With DL**, we give the network the image itself, pixel by pixel, in the right format. The network extracts the features on its own.

2.4 Layers extract features of rising complexity

Key point: Early layers find edges, middle layers find shapes, late layers find whole objects such as faces.

The second half of the technical definition says how the layers share the work. In a dog-vs-cat or face classifier:

1. The first hidden layers detect **primitive features**: edges.
2. The next layers combine edges into **shapes**.
3. The deepest layers combine shapes into **high-level concepts**: a face, a digit, a letter.

How can a neuron detect an edge at all? Its weights are positive on a short strip of pixels and negative around the strip, so its weighted sum is large only when that strip is bright and its surroundings are dark. Section 4.5 of the [MLP intuition Note](#) shows one such neuron.

The output layer then gives the answer. Section 5.3 of the [AI vs ML vs DL Note](#) animates the same idea for a handwritten 7.

We can watch representation learning happen. Figure 4 trains a small network on the MNIST handwritten digits (LeCun et al. 1998) and looks inside it.




Figure 4: Representation learning on MNIST. Each dot is one test image of a 3, 5 or 8. First frame: the raw pixels; later frames: the 32 numbers of the network’s last hidden layer after 0 to 10 epochs. Each view is squeezed to two dimensions with PCA.

1. **The network.** It takes the 784 pixels of a digit, passes them through hidden layers of 64 and 32 neurons, and outputs one of 10 digits.
2. **Raw pixels.** Squeezed to two dimensions, the 3s, 5s and 8s overlap: only 56 percent of them have a nearest neighbour of the same digit.
3. **Before training (epoch 0).** The hidden layer holds random features, and the digits are just as mixed (48 percent).
4. **Training.** After each **epoch** (G-696), one pass over the 60,000 training images, the hidden layer’s 32 numbers group the digits better: 64, 69 and 78 percent after 1, 2 and 5 epochs.
5. **After 10 epochs.** The three digits form three clear groups (87 percent), and the network labels 96.9 percent of the test images correctly.

Nobody told the network which features to compute. It changed its hidden layer, step by step, into numbers that keep similar digits together: the representation it learned.

3. Why deep learning became famous

Key point: Two reasons: it applies to a huge range of problems, and in many of them it gives the best results ever recorded, sometimes better than human experts.

3.1 Applicability

Key point: The same family of methods works on images, speech, text, medicine, science and games.

DL is used in many fields:

- computer vision;
- speech recognition;
- **natural language processing (NLP)** (G-1305);
- **machine translation** (G-1141);
- bioinformatics;
- drug design;
- medical image analysis;
- climate science;
- material inspection;
- board-game programs.

Few methods reach that many fields.




Figure 5: The fields where deep learning is used, as listed in section 3.1

Figure 5 shows the reach: one family of methods, ten very different fields.

3.2 Performance

Key point: In many fields DL holds the state-of-the-art result, and in some it beats human experts.

In most of these fields, the best results today come from DL. In March 2016 the program AlphaGo, built on deep networks, beat Lee Sedol, one of the world’s best players of the board game Go, by 4 games to 1 (Silver et al. 2017). Self-driving cars, new drug candidates and new chemistry research are further examples.

4. Five differences between DL and ML

Key point: DL needs more data, stronger hardware and longer training, but it predicts fast, learns its own features, and is hard to interpret.




Figure 6: The five differences of sections 4.1 to 4.5 as a trade: what DL costs and what it gives

Figure 6 is the balance this section weighs: four costs on the left, two gains on the right.

4.1 Data

Key point: With little data ML wins; with lots of data DL keeps improving while ML levels off.

DL is **data hungry** (G-534): its results become reliable only with a lot of data. Figure 8 of the [AI vs ML vs DL Note](#) shows the classic curves. With little data, ML performs better; beyond some amount, ML stagnates while DL keeps rising.

4.2 Hardware

Key point: ML trains on an ordinary CPU; DL needs a GPU, because it multiplies very large matrices.

A neural network does huge numbers of matrix multiplications. A [GPU](#) (G-856) with plenty of memory does them in parallel, while a CPU does them slowly. So ML runs on cheap hardware, and DL needs costly hardware.

Figure 7 measures the difference on one laptop. Each point is the time to multiply two square matrices of random numbers, the operation a network repeats millions of times.




Figure 7: Time to multiply two n by n matrices on a laptop CPU and a laptop GPU, with TensorFlow. Labels: how many times faster the GPU is.

- **Small matrices** (256×256): the GPU is 2.7 times faster.
- **Large matrices** (8192×8192): the CPU takes 1.6 seconds and the GPU 0.13 seconds, 12.5 times faster.

The bigger the matrices, the more of the work the GPU can do at the same time, and the larger its lead.

4.3 Training time and prediction time

Key point: DL trains slowly (days to months on big data) but predicts fast; ML trains in minutes to hours, and its prediction speed depends on the algorithm.

- **Training time:** a DL model on a large dataset can train for weeks, and some research models for months. Most ML models train in minutes, at most hours.
- **Prediction time:** a trained network predicts quickly, because a prediction is only a fixed series of matrix products (the [forward propagation Note](#) shows it). In ML it varies: [KNN](#), for example, predicts slowly because it compares the new point with every stored **observation** (G-1374; one record, one row of the data table).

4.4 Feature selection

Key point: In ML we build the features with domain experts; in DL the network extracts them.

Suppose we predict whether a student will be placed from the content of their resume:

- **With ML**, we extract features first: 12th marks, 10th marks, number of achievements,

number of courses, quality of college. Choosing them needs domain experts, such as HR staff or teachers.

- **With DL**, we give the whole resume, in a suitable form, to the network. The network extracts the features behind the scenes and predicts.

Letting the network extract the features is representation learning (section 2.3), and it is one of DL's biggest practical benefits.

4.5 Interpretability

Key point: A trained network is a black box: it cannot say why it gave an answer. Linear models and decision trees can.

Interpretability (G-965) is how well people can understand why a model makes its decisions. The features a network learns are internal numbers that no one chose, so we cannot say what each one means. A trained network is a **black box** (G-312): it gives an answer without the reasons.

Interpretability matters wherever we must justify a decision. Suppose a social network bans users based on their comments, using a DL model. A banned user asks why, and we have no answer.

ML models are often much easier to explain:

- **Logistic regression** (G-1120) on CGPA and IQ learns two weights, w_1 and w_2 . The larger weight marks the more important input (on inputs of similar scale), so we can tell a student “your CGPA is too low”.
- **A decision tree** (G-561) is a flowchart of questions, so it shows exactly why a point got its class (see the [decision trees Note](#)).

Extra: Researchers have built tools that explain single predictions of any model, including networks, such as LIME and SHAP, and heat maps of the image regions a CNN used. LIME, for example, fits a simple model that imitates the network near one input (Ribeiro et al. 2016, §3). So these tools explain the network from outside; its own weights stay unreadable.

4.6 DL does not replace ML

Key point: Use a needle to sew and a sword to fight: pick ML or DL to fit the problem.

	Machine learning	Deep learning
Data	works with little data, levels off	needs a lot, keeps improving
Hardware	CPU	GPU (or TPU)
Training time	minutes to hours	hours to months
Prediction time	depends on the algorithm (KNN is slow)	fast

	Machine learning	Deep learning
Features	engineered by us	learned by the network
Interpretability	often high (linear models, trees)	low: a black box

Since DL is so strong, why not use it everywhere? Because on small or tabular data (section 4.1; Grinsztajn et al. 2022), or when decisions must be explained (section 4.5), ML is the better tool. An old saying puts it well: where a needle is needed, a sword is no use. Section 6 of the [AI vs ML vs DL Note](#) gives the rule of thumb for choosing.

5. Why deep learning took off after 2010

Key point: Five forces: large public datasets, faster hardware, easy frameworks, ready-made architectures, and a large community.

The core ideas of neural networks are decades old (the [history section](#) tells the story), yet DL became famous only around 2012. Figure 8 shows the five forces behind the change.




Figure 8: The five forces that made deep learning practical

5.1 Datasets

Key point: Smartphones and cheap internet created huge amounts of data; big companies labelled some of it and made it public.

Around 2010 two revolutions met: smartphones, and much cheaper mobile internet (in India, Jio's free and then very cheap data from 2016). With billions of people on social media apps, the amount of data generated every year began to grow exponentially.

Raw data alone is not enough. To train a dog-vs-cat classifier we need photos **labelled** “dog” or “cat”, and a photo uploaded to a social network carries no such label. Companies such as Microsoft, Google and Facebook paid to label large datasets, and then released them as **public datasets** (G-1590) that anyone may use, because open data speeds up research for everyone:

Data type	Public dataset	Contents
Images	Microsoft COCO	images with a labelled box around every object; used for object detection
Video	YouTube-8M	about 6.1 million labelled online video clips
Text	SQuAD	about 150,000 questions and answers on Wikipedia articles
Audio	Google AudioSet	about 2 million labelled sound clips in over 600 categories

Thousands more are available today, for example on Kaggle. Without data there would be no deep learning, so datasets are the biggest single reason.

5.2 Hardware

Key point: Moore's law made chips steadily faster and cheaper, and GPUs turned out to do a network's matrix maths in parallel, 10 to 20 times faster than a CPU.

Moore's law (G-1261), stated by Intel co-founder Gordon Moore, says that the number of transistors on a microchip doubles about every two years. Electronics therefore get faster and cheaper every year: phones went from about 128 MB of memory in the mid-2000s to 12 GB today.

Around 2010, researchers realised that a network's matrix multiplications can run in parallel, just like the graphics work a GPU was built for. NVIDIA's **CUDA** (G-514), a platform for programming GPUs, made this practical, and training on GPUs became the norm. A GPU typically cuts training time by a factor of 10 to 20 compared with a CPU.

Extra: CUDA 1.0 came out in 2007 (NVIDIA 2007). Deep networks were trained on GPUs by 2009 (Raina et al. 2009), and the 2012 ImageNet winner described in the [history section](#) was trained on two GPUs (Krizhevsky et al. 2012, §3.2).

Once the speed-up was clear, chips designed for DL followed:

- **FPGA** (G-800; field-programmable gate array): a reprogrammable chip that is fast and low-power, but expensive. Microsoft runs much of its search engine’s AI on FPGAs.
- **ASIC** (G-217; application-specific integrated circuit): a chip custom-made for one job, costly to design but cheap in large numbers. Examples: Google’s **TPU** (G-1996; tensor processing unit), the small **Edge TPU** for drones, smartwatches and smart glasses, and the **NPU** (G-1359; neural processing unit) in phones.

Which hardware to use depends on the job:

Situation	Hardware
Learning, small projects	CPU
Training large networks	GPU (e.g. NVIDIA RTX series) or TPU
DL inside a phone app	mobile CPU or GPU, digital signal processors, NPU
DL on a smartwatch, drone or smart glasses	Edge TPU or NPU

5.3 Frameworks

Key point: Libraries do the hard maths behind the scenes, as scikit-learn does for ML: TensorFlow with Keras in industry, PyTorch in research.

Writing a network’s training code from scratch takes longer than the problem it solves. Deep learning needed libraries that handle this code, just as scikit-learn does for ML. Two families dominate:

- **TensorFlow (Google)** (G-1959): Google built an internal framework, DistBelief, in 2011, and released TensorFlow publicly in 2015. TensorFlow was powerful but hard to use, so **Keras** (G-1003), a simpler library running on top of it, became popular. Since TensorFlow 2.0 (2019), Keras is built in.
- **PyTorch (Facebook, now Meta)** (G-1595): released in 2016, it became the favourite of researchers. Facebook’s Caffe2, a library for running models on servers, was merged into it in 2018.

Today TensorFlow with Keras is used more in industry, and PyTorch more in research. Converting a model from one to the other is awkward, so drag-and-drop tools appeared that build a network in a browser and export code for either: Google’s AutoML, Microsoft’s Custom Vision and Apple’s Create ML.

Extra: Since Keras 3 (2023), Keras can again run on several back ends: JAX, TensorFlow or PyTorch (Keras 3 docs).

5.4 Architectures and transfer learning

Key point: Researchers publish trained state-of-the-art networks; we download one and apply it to our problem instead of designing our own.

A network's **architecture** (G-209) is how its nodes are connected by weights: how many nodes, what kind, and which connections. Different problems need different architectures, and finding a good one takes many experiments, which cost time, effort and money.

Instead, we can reuse an architecture that researchers already designed and trained on a big dataset. Downloading such a trained network and applying it to our own problem is called **transfer learning** (G-2005), taught in detail later. Well-known ready-made architectures:

Task	Architecture
Image classification	ResNet
Text classification, other language tasks	BERT (a transformer)
Image segmentation	U-Net
Image-to-image translation	pix2pix
Object detection	YOLO (You Only Look Once)
Speech generation	WaveNet

These models reach the best known accuracy, in some tasks above human level, and transfer learning made adopting DL far easier.

5.5 Community

Key point: People who believed in neural networks through decades of failure, plus today's engineers, teachers, students and communities like Kaggle.

None of the above would exist without people. Researchers worked on neural networks from the 1960s and failed several times before succeeding around 2012. Today engineers, teachers, students and communities such as Kaggle keep the field moving.

6. Summary

Reason DL took off	What changed
Datasets	smartphones and cheap internet; public labelled datasets (COCO, YouTube-8M, SQuAD, AudioSet)
Hardware	Moore's law; GPUs with CUDA; FPGAs, TPUs, Edge TPUs, NPUs
Frameworks	TensorFlow with Keras (industry), PyTorch (research)
Architectures	ready-made trained networks (ResNet, BERT, YOLO) and transfer learning
Community	researchers, engineers, teachers, students, Kaggle

- DL is ML with neural networks: neurons in an input layer, hidden layers and an output layer, joined by weights.
- Many hidden layers make a network deep.
- Representation learning: the network learns its own features, edges first and whole objects last.
- DL needs more data, a GPU and longer training, predicts fast, and is a black box.
- ML stays the better choice for small or tabular data and when decisions must be explained.

7. Sources

Built from

- CampusX, “What is Deep Learning? Deep Learning Vs Machine Learning | Complete Deep Learning Course”, YouTube, <https://www.youtube.com/watch?v=fHF22Wxuyw4>
- Sanderson, G. (3Blue1Brown), “But what is a neural network? | Deep learning chapter 1”, 2017, <https://www.youtube.com/watch?v=aircAruvnKk> (the three different 3s; an edge neuron's weights)

Other references

- LeCun, Y., Bottou, L., Bengio, Y. and Haffner, P. (1998). Gradient-Based Learning Applied to Document Recognition. *Proceedings of the IEEE* 86(11). (The MNIST digits.)
- Silver et al., “Mastering the game of Go without human knowledge”, *Nature*, 2017 (AlphaGo against Lee Sedol, March 2016).
- NVIDIA, “CUDA Toolkit Archive” (CUDA Toolkit 1.0, June 2007), developer.nvidia.com.
- Raina, Madhavan and Ng, “Large-scale Deep Unsupervised Learning using Graphics Processors”, ICML 2009.

- Krizhevsky, Sutskever and Hinton, “ImageNet Classification with Deep Convolutional Neural Networks”, NeurIPS 2012.
- Keras documentation, “Introducing Keras 3.0”, keras.io.
- Ribeiro, Singh and Guestrin, “Why Should I Trust You? Explaining the Predictions of Any Classifier”, KDD 2016 (LIME). Lundberg and Lee, “A Unified Approach to Interpreting Model Predictions”, NeurIPS 2017 (SHAP).
- Grinsztajn, Oyallon and Varoquaux, “Why do tree-based models still outperform deep learning on typical tabular data?”, NeurIPS 2022 (Datasets and Benchmarks).

8. Key terms

Term	Meaning
Artificial neural network (ANN)	The simplest neural network: neurons in layers, each layer connected to the next by weights
Neuron (in a network)	One unit of a neural network; in an ANN, a perceptron
Weight (in a network)	A learned number on a connection between two neurons
Input layer	The first layer, with one node per feature
Output layer	The last layer, which gives the prediction
Hidden layer	Any layer between the input and output layers
Epoch	One pass of training over all the training data
Deep network	A neural network with many hidden layers
Representation	A set of features that describes the data; a good one stays the same when unimportant details (such as handwriting) change
Representation learning (feature learning)	Letting the algorithm discover useful features from raw data, instead of engineering them by hand
Data hungry	Needing a lot of data before results become reliable
Public dataset	A labelled dataset released for anyone to use
Moore’s law	The number of transistors on a chip doubles about every two years
CUDA	NVIDIA’s platform for programming GPUs
FPGA	A reprogrammable chip: fast and low-power, but expensive
ASIC	A chip custom-made for one job, such as the TPU
Edge TPU	A small Google chip for running networks on drones, watches and glasses

Term	Meaning
NPU	A neural processing unit: a chip in phones that speeds up networks
Architecture (of a network)	How a network's nodes are connected: how many, of what kind, and which connections
Transfer learning	Reusing a network trained by others on a big dataset for our own problem