

Forward Propagation: How a Neural Network Predicts

Contents

1. Overview	2
2. The network and the row	3
3. What one node computes	3
4. Layer 1 as one matrix product	4
4.1 The weight matrix	4
4.2 The product	4
5. Layers 2 and 3	5
5.1 Layer 2	5
5.2 Layer 3: the prediction	5
6. The whole network in one formula	6
7. Summary	7
8. Sources	7
9. Key terms	7

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Matrix multiplication as composition ([Note 510](#)); MLP notation and parameter count ([Note 1008](#)).
- **Leads to:** Loss functions in deep learning ([Note 1014](#)); Backpropagation ([Note 1015](#)).

1. Overview

Key point: To predict, a network passes one row of data forward, layer by layer. Each layer is one matrix product, one added bias and one sigmoid, however many nodes it has.

Forward propagation is the computation that turns one row of inputs into a prediction, moving from the input layer through every hidden layer to the output layer. It is how a trained network predicts. It is also the first half of every training step: the network predicts, measures its error, and then **backpropagation**, the training algorithm of later Notes, sends that error backwards to update the weights (previewed in section 10 of the [Jacobian Note](#)).




Figure 1: Forward propagation through the 4-3-2-1 network: the values of each layer appear in its nodes, with that layer's matrix equation

Figure 1 shows the whole computation for one student. The rest of this Note works through it, layer by layer, with every number.

2. The network and the row

Key point: The 4-3-2-1 network has 26 parameters. We feed it one student: CGPA 7.2, IQ 72, 10th marks 69, 12th marks 81.

We use the 4-3-2-1 network and the notation of the [MLP notation Note](#): 4 inputs, hidden layers of 3 and 2 nodes, and 1 output, with $15 + 8 + 3 = 26$ weights and biases. Every node is a perceptron with a sigmoid activation, so every output lies between 0 and 1.

The data has four input columns and the output column *placed*. One row enters the input layer:

CGPA	IQ	10th marks	12th marks
7.2	72	69	81

Before it enters, we scale the row into the range 0 to 1: CGPA divided by 10, the others by 100. So the input is $a^0 = (0.72, 0.72, 0.69, 0.81)$.

Extra: Why scale first? With the raw values, the first hidden layer's sums are 20.1, -13.7 and 12.5, and the sigmoid turns them into 1.000, 0.000 and 1.000. Every node is pinned at an extreme. The sigmoid's slope there, $\sigma(z)(1 - \sigma(z))$ (see the [sigmoid derivative Note](#)), is only 1.9×10^{-9} , 1.1×10^{-6} and 3.7×10^{-6} , so a small change to a weight hardly changes the node's output. Backpropagation multiplies every gradient through a node by this slope, so these weights would barely move in training. The [data scaling Note](#) standardizes the inputs instead (see the [standardization Note](#)); dividing by 10 and 100 keeps the numbers here easy to follow.

The weights are not trained here: we set them by hand to one decimal, so that every step can be checked with a calculator. Training would find better values; forward propagation works the same either way.

3. What one node computes

Key point: Each node takes the outputs of the previous layer, multiplies each by its weight, adds its bias and applies the sigmoid.

Every node does what a single perceptron does (see the [perceptron Note](#)), with the sigmoid as its activation. For node 1 of layer 1:

1. **In words:** multiply each input by the weight on its connection to this node, add them up, add the node's bias, and pass the total through the sigmoid.

2. **Formula:**

$$O_{11} = \sigma(W_{11}^1 x_1 + W_{21}^1 x_2 + W_{31}^1 x_3 + W_{41}^1 x_4 + b_{11})$$

3. **Example:** with weights 0.2, 0.4, -0.5, 0.3 and bias 0.1,

$$z = 0.2 \times 0.72 + 0.4 \times 0.72 - 0.5 \times 0.69 + 0.3 \times 0.81 + 0.1 = 0.430$$

$$O_{11} = \sigma(0.430) = \frac{1}{1 + e^{-0.430}} = 0.606$$

Doing this node by node works, but a large network has thousands of nodes. Linear algebra does a whole layer in one step.

4. Layer 1 as one matrix product

Key point: Put the layer's weights in a matrix with one row per input node and one column per receiving node; then $z^1 = W^{1\top}a^0 + b^1$ gives all three sums at once.

4.1 The weight matrix

Key point: Row i , column j of W^1 holds W_{ij}^1 : from input node i to node j .

We collect the 12 weights entering layer 1 in a 4×3 matrix. Row i holds the weights leaving input node i ; column j holds the weights entering node j :

$$W^1 = \begin{bmatrix} W_{11}^1 & W_{12}^1 & W_{13}^1 \\ W_{21}^1 & W_{22}^1 & W_{23}^1 \\ W_{31}^1 & W_{32}^1 & W_{33}^1 \\ W_{41}^1 & W_{42}^1 & W_{43}^1 \end{bmatrix} = \begin{bmatrix} 0.2 & -0.3 & 0.5 \\ 0.4 & 0.1 & -0.2 \\ -0.5 & 0.2 & 0.1 \\ 0.3 & -0.4 & 0.2 \end{bmatrix}, \quad b^1 = \begin{bmatrix} 0.1 \\ -0.1 \\ 0.2 \end{bmatrix}$$

The first column, $(0.2, 0.4, -0.5, 0.3)$, is exactly the four weights node 1 used in Section 3.

4.2 The product

Key point: Transpose to 3×4 , multiply by the 4×1 input, add the 3×1 bias: a 3×1 result, one number per node.

To get one sum per node, each column of W^1 must meet the input. Transposing turns the columns into rows, and then the product does exactly that (a matrix times a vector, see the [linear transformations Note](#), section 7.3).

1. **In words:** multiply the transposed weight matrix by the input vector, add the bias vector, then apply the sigmoid to every entry.
2. **Formula:**

$$z^1 = W^{1\top}a^0 + b^1, \quad a^1 = \sigma(z^1)$$

The shapes follow the shape rule of the [matrix multiplication Note](#) (section 7.1):
 $(3 \times 4)(4 \times 1) + (3 \times 1) = 3 \times 1$.

3. **Example:**

$$z^1 = \begin{bmatrix} 0.2 & 0.4 & -0.5 & 0.3 \\ -0.3 & 0.1 & 0.2 & -0.4 \\ 0.5 & -0.2 & 0.1 & 0.2 \end{bmatrix} \begin{bmatrix} 0.72 \\ 0.72 \\ 0.69 \\ 0.81 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.1 \\ 0.2 \end{bmatrix} = \begin{bmatrix} 0.330 \\ -0.330 \\ 0.447 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.1 \\ 0.2 \end{bmatrix} = \begin{bmatrix} 0.430 \\ -0.430 \\ 0.647 \end{bmatrix}$$

$$a^1 = \sigma(z^1) = \begin{bmatrix} 0.606 \\ 0.394 \\ 0.656 \end{bmatrix} = \begin{bmatrix} O_{11} \\ O_{12} \\ O_{13} \end{bmatrix}$$

The first entry, 0.606, is the O_{11} we computed by hand in Section 3. The vector a^1 , the outputs of layer 1, is called the **activation** of layer 1. It is the input to layer 2.

5. Layers 2 and 3

Key point: The same step repeats: the previous layer's activation goes in, this layer's activation comes out, until the output layer gives $\hat{y}$.

5.1 Layer 2

Key point: $(2 \times 3)(3 \times 1) + (2 \times 1) = 2 \times 1$.

Layer 2 has 6 weights, from 3 nodes to 2, in a 3×2 matrix:

$$W^2 = \begin{bmatrix} 0.6 & -0.4 \\ -0.2 & 0.5 \\ 0.3 & 0.7 \end{bmatrix}, \quad b^2 = \begin{bmatrix} 0.1 \\ -0.2 \end{bmatrix}$$

$$z^2 = W^{2T}a^1 + b^2 = \begin{bmatrix} 0.6 & -0.2 & 0.3 \\ -0.4 & 0.5 & 0.7 \end{bmatrix} \begin{bmatrix} 0.606 \\ 0.394 \\ 0.656 \end{bmatrix} + \begin{bmatrix} 0.1 \\ -0.2 \end{bmatrix} = \begin{bmatrix} 0.582 \\ 0.214 \end{bmatrix}, \quad a^2 = \sigma(z^2) = \begin{bmatrix} 0.641 \\ 0.553 \end{bmatrix}$$

These are O_{21} and O_{22} .

5.2 Layer 3: the prediction

Key point: $(1 \times 2)(2 \times 1) + (1 \times 1) = 1 \times 1$: a single number, the probability of placement.

The output layer has 2 weights and 1 bias:

$$W^3 = \begin{bmatrix} 0.8 \\ -0.6 \end{bmatrix}, \quad b^3 = 0.2$$

$$z^3 = W^{3T}a^2 + b^3 = 0.8 \times 0.641 - 0.6 \times 0.553 + 0.2 = 0.381, \quad a^3 = \sigma(0.381) = 0.594$$

So $\hat{y} = O_{31} = 0.594$: the network gives this student a 59.4% probability of being placed. Since it is above 0.5, it predicts *placed*. Four input numbers became one prediction, using nothing but matrix products, additions and sigmoids.

6. The whole network in one formula

Key point: $\hat{y} = \sigma(W^{3\top} \sigma(W^{2\top} \sigma(W^{1\top} a^0 + b^1) + b^2) + b^3)$. More layers only make the chain longer.

Each layer applies the same rule to the activation of the layer before it:

$$a^k = \sigma(W^{k\top} a^{k-1} + b^k), \quad k = 1, 2, 3$$

Substituting each activation into the next writes the whole network as one nested formula:

$$\hat{y} = a^3 = \sigma\left(W^{3\top} \sigma\left(W^{2\top} \underbrace{\sigma(W^{1\top} a^0 + b^1)}_{a^1} + b^2\right) + b^3\right)$$

$\underbrace{\hspace{10em}}_{a^2}$

This is what a neural network is, as a function: a chain of matrix products, each followed by a bias and an activation. However large the architecture, prediction stays this organised. The sigmoids between the matrices are essential: without them, the three matrices would collapse into one (see the [matrix multiplication Note](#), section 7.2), and the network would be no more powerful than one perceptron.

Python: Forward propagation with NumPy.

```
import numpy as np
```

```
def sigmoid(z):  
    return 1 / (1 + np.exp(-z))
```

```
a = np.array([0.72, 0.72, 0.69, 0.81]) # a0  
for W, b in [(W1, b1), (W2, b2), (W3, b3)]:  
    a = sigmoid(W.T @ a + b)           # one layer  
print(a)                               # [0.594]
```

$W1$ is the 4×3 array of Section 4.1 and $b1$ the bias of length 3; $W.T$ is the transpose and $@$ the matrix product. The same three-line loop works for any number of layers.

Extra: Libraries predict many rows at once. Stacking n rows into an $n \times 4$ matrix X (a 2D tensor, see the [tensors Note](#)), the first layer becomes $\sigma(XW^1 + b^1)$: an $(n \times 4)(4 \times 3) = n \times 3$ matrix, one row of activations per student. With rows instead of columns, no transpose is needed (as in section 7.1 of the [linear transformations Note](#)). scikit-learn's `MLPClassifier` and Keras both store each layer's weights with shape (nodes in, nodes out), exactly like our W^1 , and use this row form: a Keras `Dense` layer computes `activation(dot(input, kernel) + bias)` (Keras docs, `Dense`). The Notebook loads our hand-set weights into both and gets the same 0.594.

7. Summary

Layer	Weights	Computation	Product shapes	Activation
1	4×3	$\sigma(W^{1\top}a^0 + b^1)$	$(3 \times 4)(4 \times 1)$	$(0.606, 0.394, 0.656)$
2	3×2	$\sigma(W^{2\top}a^1 + b^2)$	$(2 \times 3)(3 \times 1)$	$(0.641, 0.553)$
3	2×1	$\sigma(W^{3\top}a^2 + b^3)$	$(1 \times 2)(2 \times 1)$	$\hat{y} = 0.594$

- Forward propagation: one row moves from the input layer to the output, layer by layer.
- Each layer: weighted sums ($W^\top a$), plus bias, then sigmoid.
- W^k has one row per node of layer $k - 1$ and one column per node of layer k .
- The whole network is the nested formula $\sigma(W^{3\top}\sigma(W^{2\top}\sigma(W^{1\top}a^0 + b^1) + b^2) + b^3)$.
- Backpropagation, next, uses this forward pass to compute the error and update the 26 parameters.

8. Sources

- Keras documentation, `keras.layers.Dense`.

9. Key terms

Term	Meaning
Forward propagation	Passing one row of inputs through the network, layer by layer, to get the prediction
Backpropagation	The training algorithm that sends the prediction error backwards to update the weights (taught in later Notes)
Activation (a^k)	The vector of outputs of layer k ; a^0 is the input row
Weight matrix (W^k)	All weights entering layer k : one row per node of layer $k - 1$, one column per node of layer k
Bias vector (b^k)	The biases of all nodes of layer k
Weighted input (z^k)	$W^{k\top}a^{k-1} + b^k$: a layer's sums before the activation