


Customer Churn Prediction with an ANN in Keras

Contents

1. Overview	2
2. The churn data	3
3. Preparing the inputs	4
3.1 Dropping and encoding	4
3.2 Splitting and scaling	5
4. Building a network in Keras	5
4.1 The first architecture	6
4.2 Counting the parameters with summary	7
5. Compiling and training	7
5.1 Compile: loss and optimizer	7
5.2 Fit: epochs	8
5.3 Where the weights are stored	8
6. Predicting and measuring accuracy	8
6.1 From probabilities to classes	9
6.2 Why 79% means nothing here	9
7. Improving the network	10
7.1 Four changes	10
7.2 The second architecture	11
7.3 Tracking accuracy and a validation set	11
8. Training curves	12
9. Summary	14
10. Sources	14
11. Key terms	14

Where this fits: Pipeline map steps 8 (Model), 9 (Evaluate). See the [Course map](#).



- **Builds on:** Standardization ([Note 24](#)); One-hot encoding ([Note 27](#)); Log loss (binary cross entropy) ([Note 73](#)); Confusion matrix ([Note 76](#)); Softmax regression ([Note 79](#)); Overfitting ([Note 91](#)).
- **Leads to:** ANN for regression ([Note 1013](#)); Vanishing gradient ([Note 1018](#)); Early stopping ([Note 1021](#)).
- **Compare with:** K-nearest neighbours ([Note 91](#)); ANN for regression ([Note 1013](#)).

1. Overview

Key point: We build, train and improve our first neural network in Keras, on real bank data: which customers will leave the bank?

The Notes so far built the multi-layer perceptron (MLP) on paper: how it is organised (see the [MLP notation Note](#)), why it can draw curved boundaries (see the [MLP intuition Note](#)) and how it predicts (see the [forward propagation Note](#)). How it *learns* its weights is backpropagation, the subject of later Notes. Before that theory, three small projects show a network being built, trained and used in code:

Project	Problem type	This Note?
Customer churn	Binary classification	yes
Handwritten digits (MNIST)	Multi-class classification	MNIST Note
Graduate admissions	Regression	graduate admission Note

The goal is not the best possible model. It is to see every step of a Keras project once, so that the training theory has something concrete to attach to.




Figure 1: The steps of a Keras project. The dashed arrow is the improvement loop.

Figure 1 shows the steps. Preparing the data is the same as in the [toy project Note](#); the four Keras steps (build, compile, fit, predict) are new. The Notebook ([notebook.ipynb](#)) runs every step.

2. The churn data

Key point: 10,000 bank customers, 13 possible inputs and one output, **Exited:** 1 if the customer left the bank.

Churn means customers leaving a service (see the [framing an ML problem Note](#)). Banks, phone companies and streaming services all try to predict it, so that they can act before a customer goes. Our data holds 10,000 customers of one bank, 14 columns each:

Column	Meaning
RowNumber, CustomerId, Surname	Identifiers
CreditScore	How creditworthy the customer is
Geography	France, Germany or Spain
Gender	Male or Female
Age	Age in years
Tenure	Years with the bank
Balance	Account balance
NumOfProducts	Number of bank products used (cards, deposits, funds, ...)
HasCrCard	1 if the customer has a credit card
IsActiveMember	1 if the customer uses the account actively

Column	Meaning
EstimatedSalary	The bank's estimate of the yearly salary
Exited	Output: 1 if the customer left, 0 if not

Extra: The file is `Churn_Modelling.csv` (10,000 rows, 685 KB), published on Kaggle as a bank churn dataset; the copy in `data/` is the full file.

Python: Loading and checking the data.

```
import pandas as pd

df = pd.read_csv("data/Churn_Modelling.csv")
df.shape           # (10000, 14)
df.info()          # no missing values
df.duplicated().sum() # 0 duplicated rows
df["Exited"].value_counts() # 0: 7963, 1: 2037
```

The checks find no missing values and no duplicated rows. About 8,000 customers stayed and 2,000 left, so the classes are imbalanced, 4 to 1. That matters for reading accuracy, as Section 6 shows (see also the [imbalanced data Note](#)).

The two text columns are balanced enough: France 5,014, Germany 2,509, Spain 2,477; male 5,457, female 4,543.

3. Preparing the inputs

Key point: Drop the three identifier columns, one-hot encode Geography and Gender, split 80/20, then standardize.

3.1 Dropping and encoding

Key point: Identifiers carry no pattern; text columns become 0/1 columns.

RowNumber, CustomerId and Surname identify a customer but say nothing about whether they will leave, so we drop them. A careful project would first check every other column with exploratory analysis; here we keep the remaining ten.

A network only takes numbers, so Geography and Gender are one-hot encoded with `get_dummies` and `drop_first=True` (see the [one-hot encoding Note](#), section 6). Geography becomes two columns, `Geography_Germany` and `Geography_Spain` (France is 0, 0), and Gender becomes one, `Gender_Male` (female is 0). That leaves 11 input columns.

Python: Dropping and encoding.

```
df = df.drop(columns=["RowNumber", "CustomerId", "Surname"])
df = pd.get_dummies(df, columns=["Geography", "Gender"],
```

```
drop_first=True, dtype=int)
```

Since pandas 2, `get_dummies` writes True/False; `dtype=int` gives 0/1 instead.

3.2 Splitting and scaling

Key point: 8,000 customers to train on, 2,000 to test; every column standardized with the training set's mean and spread.

We separate X (11 columns) from y (`Exited`) and split 80/20 with `train_test_split`, as in the [toy project Note](#). This leaves 8,000 training rows and 2,000 test rows.

The columns live on very different scales: Balance and EstimatedSalary run to six digits, NumOfProducts is 1 to 4. With such inputs the weights take much longer to settle during training (the forward propagation Note shows what raw inputs do to a sigmoid, in section 2). So before a network sees any data we **standardize** every column to mean 0, standard deviation 1 (see the [standardization Note](#)).

Python: Split, then scale.

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
X = df.drop(columns="Exited")
y = df["Exited"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=1)
```

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

4. Building a network in Keras

Key point: A `Sequential` model is a stack of layers; each `Dense` layer is a row of perceptrons connected to every node of the layer before.

Keras is the high-level interface of TensorFlow for defining and training networks (see the [deep learning scope Note](#)). Keras has two ways to build a model:

- **Sequential model:** layers stacked one after the other, each feeding the next. Every MLP is of this kind.
- **Functional (non-sequential) model:** layers wired in any pattern, for example with branches. It comes in later Notes.

4.1 The first architecture

Key point: 11 inputs, one hidden layer of 3 sigmoid nodes, one sigmoid output node.

We start small: an input layer of 11 nodes, one per column; a hidden layer of 3 perceptrons; and an output layer of 1 perceptron. All of them use the sigmoid activation, so the output is the probability that the customer leaves. Figure 2(a) shows this 11-3-1 network.




Figure 2: The two churn networks. Every input reaches every node of the next layer; the numbers are the trainable parameters of each layer.

A **dense layer** (also **fully connected layer**) is a layer in which every node receives the output of every node in the layer before. Every layer of an MLP is dense, so in Keras each one is a `Dense`.

Python: Building the 11-3-1 network.

```
import keras
from keras import layers

model = keras.Sequential([
    keras.Input(shape=(11,)),           # 11 input columns
    layers.Dense(3, activation="sigmoid"), # hidden layer
    layers.Dense(1, activation="sigmoid"), # output layer
])
```

`Dense(3, ...)` is a layer of 3 nodes. `keras.Input(shape=(11,))` tells the first layer how many inputs each row has; Keras works out the inputs of every later layer by itself.

Extra: Older code passes the input size to the first layer, `Dense(3, activation="sigmoid", input_dim=11)`, and imports from `tensorflow.keras`. Keras 3 still builds such a model but prints a warning asking for an `Input(shape)` object as the first layer instead. So we write `keras.Input` as the first item and `import keras`

directly.

4.2 Counting the parameters with summary

Key point: `model.summary()` lists each layer and its trainable parameters: $36 + 4 = 40$.

`model.summary()` prints one line per layer. Keras does not list the input layer, because it has no parameters:

Layer	Output shape	Param #
dense (Dense)	(None, 3)	36
dense_1 (Dense)	(None, 1)	4
Total		40

The counts follow the rule of the [MLP notation Note](#): one weight per pair of connected nodes, one bias per node. Into the hidden layer, $11 \times 3 + 3 = 36$; into the output, $3 \times 1 + 1 = 4$. Training must find these 40 numbers. **None** in the output shape stands for the number of rows, which can be anything.

5. Compiling and training

Key point: `compile` chooses the loss and the optimizer; `fit` runs the training for a set number of epochs and prints the loss after each.

5.1 Compile: loss and optimizer

Key point: Two classes and a sigmoid output: the loss is binary cross-entropy. The optimizer is Adam.

Compiling tells Keras how the model will be trained. Two settings are needed:

- **Loss:** the number that training makes smaller. For two classes with a sigmoid output it is binary cross-entropy, also called log loss (see the [log loss Note](#) and the [perceptron loss Note](#), section 8).
- **Optimizer:** the method that updates the weights to reduce the loss, a variant of gradient descent. We use **Adam**, which usually works well; optimizers are taught in later Notes.

Python: Compiling.

```
model.compile(loss="binary_crossentropy", optimizer="adam")
```

5.2 Fit: epochs

Key point: One epoch is one pass over all 8,000 training customers; after 10 epochs the loss fell from 0.71 to 0.43.

`fit` does the training. We give it the scaled training inputs, the training outputs and the number of **epochs**: how many times the network goes through the whole training set (see the [gradient descent Note](#)). We choose 10.

Python: Training for 10 epochs.

```
history1 = model.fit(X_train_scaled, y_train, epochs=10)
```

Keras prints one line per epoch. The loss falls quickly at first, then more slowly:

Epoch	1	2	3	5	10
Loss	0.713	0.569	0.502	0.455	0.430

Extra: Each epoch line also shows 250/250. Keras does not update the weights once per epoch: by default it updates them after every **batch** of 32 rows (Keras docs, `Model.fit`), which is mini-batch gradient descent (see the [mini-batch gradient descent Note](#)). 8,000 rows / 32 = 250 updates per epoch. The `batch_size` argument of `fit` changes this.

5.3 Where the weights are stored

Key point: Each layer keeps its weight matrix and its bias vector; `get_weights()` returns them.

After training, the 40 parameters live inside the model, one layer at a time. `model.layers[0]` is the hidden layer, `model.layers[1]` the output layer.

Python: Reading the trained weights.

```
W1, b1 = model.layers[0].get_weights()
W1.shape, b1.shape      # (11, 3), (3,)
W2, b2 = model.layers[1].get_weights()
W2.ravel(), b2          # [-0.734 -1.314 0.572], [-0.571]
```

The weight matrix has one row per input node and one column per receiving node, exactly the W^1 of the [forward propagation Note](#). Its $11 \times 3 = 33$ weights plus 3 biases are the 36 of the summary.

6. Predicting and measuring accuracy

Key point: `predict` gives probabilities; a threshold of 0.5 turns them into 0 or 1, and `accuracy` compares those with the true labels.

6.1 From probabilities to classes

Key point: Above 0.5 means “will leave” (1); 0.5 or below means “stays” (0).

`model.predict` runs forward propagation on every test row. Because the output node uses a sigmoid, each answer is a probability between 0 and 1, not a class: the first five test customers get 0.128, 0.146, 0.135, 0.113 and 0.102.

To get a class we pick a **threshold**: a probability above it means 1, at or below it means 0. We use 0.5.

Python: Predicting and thresholding.

```
import numpy as np
from sklearn.metrics import accuracy_score

y_log = model.predict(X_test_scaled)           # probabilities
y_pred = np.where(y_log > 0.5, 1, 0)           # classes
accuracy_score(y_test, y_pred)                 # 0.7925
np.where(condition, a, b) gives a where the condition holds and b elsewhere.
```

Extra: 0.5 is only the default. The best threshold for a problem can be found from the ROC curve (see the [ROC and AUC Note](#)); a bank that would rather contact too many customers than miss leavers might choose a lower one.

6.2 Why 79% means nothing here

Key point: 79.25% is exactly the share of customers who stay: the first network predicts “stays” for everyone.

The first network scores 79.25% accuracy on the 2,000 test customers. But 1,585 of them, 79.25%, stayed. A “model” that answers 0 for every customer scores the same.




Figure 3: Confusion matrices on the 2,000 test customers. The first network never predicts “leaves”; the second finds 189 of the 415 leavers.

Figure 3 (left) confirms it: the first network predicted “leaves” for nobody. After 10 epochs its probabilities run from 0.069 to 0.498: even the most likely leaver falls just short of the 0.5 threshold. This is the trap of accuracy on imbalanced data, described in section 6 of the [accuracy and confusion matrix Note](#). The confusion matrix shows it at once; accuracy alone hides it.

Extra: The exact result depends on the random starting weights. With four other seeds, the same network flags 35 to 128 customers as leaving and scores 79.9% to 80.7% (Notebook). Either way, a network this small, trained for 10 epochs, has barely learned the pattern: trained for 100 epochs instead, it reaches 83.4%.

7. Improving the network

Key point: Four things to try: more epochs, ReLU in the hidden layers, more nodes per hidden layer, more hidden layers.

7.1 Four changes

Key point: Each change gives the network more time or more capacity; too much capacity leads to overfitting.

A first network is rarely the best one. We experiment with:

1. **More epochs:** more passes over the data, more time to find good weights. We go from 10 to 100.
2. **ReLU in the hidden layers:** hidden layers with the ReLU activation usually train better than sigmoid ones (see the [MLP intuition Note](#), section 5; Goodfellow et al. 2016,

§6.3). The output stays sigmoid, since we still want a probability. Activation functions get their own Notes later.

3. **More nodes per hidden layer:** 11 instead of 3.
4. **More hidden layers:** two instead of one.

Changes 3 and 4 add parameters, so the network can fit more complex patterns. Too many, and it starts to memorise the training rows instead of learning the pattern: **overfitting** (see the [challenges in ML Note](#) and the [bias-variance Note](#)). Finding the right size takes experiments.

7.2 The second architecture

Key point: 11-11-11-1 with ReLU: $132 + 132 + 12 = 276$ parameters.

Figure 2(b) shows the new network. Counting its parameters:

1. **In words:** for each layer, weights from every node before plus one bias per node.
2. **Formula:** $n_{l-1} n_l + n_l$ per layer, added up.
3. **Example:**

$$(11 \times 11 + 11) + (11 \times 11 + 11) + (11 \times 1 + 1) = 132 + 132 + 12 = 276$$

`model.summary()` prints the same 276.

Python: The second network.

```
model2 = keras.Sequential([
    keras.Input(shape=(11,)),
    layers.Dense(11, activation="relu"),
    layers.Dense(11, activation="relu"),
    layers.Dense(1, activation="sigmoid"),
])
```

7.3 Tracking accuracy and a validation set

Key point: `metrics=["accuracy"]` reports accuracy every epoch; `validation_split=0.2` holds back 20% of the training rows and reports the loss and accuracy on them too.

Two additions make training easier to watch:

- **`metrics=["accuracy"]` in `compile`:** Keras reports the accuracy after each epoch, next to the loss. A **metric** is only reported; training does not minimise it.
- **`validation_split=0.2` in `fit`:** Keras sets aside 20% of the 8,000 training rows as a **validation set** (see the [OOB score Note](#)). It trains on the other 6,400 and, after each epoch, measures the loss and accuracy on the 1,600 it never trains on.

Python: Compiling and training with both.

```
model2.compile(loss="binary_crossentropy", optimizer="adam",
```

```

        metrics=["accuracy"])
history = model2.fit(X_train_scaled, y_train, epochs=100,
                    validation_split=0.2)

```

Each epoch line now has four numbers. For the last epoch:

loss	accuracy	val_loss	val_accuracy
0.325	0.864	0.359	0.853

We want the loss to fall and the accuracy to rise on **both** sets. If the training accuracy keeps rising while the validation accuracy stalls or falls, the network is overfitting. Here training (86.4%) is a little ahead of validation (85.3%): a slight gap.

Extra: `validation_split` takes the **last** 20% of the rows, not a random 20% (Keras docs, `Model.fit`). Our rows were already shuffled by `train_test_split`, so that is fine; on data sorted by date or by class, shuffle first or pass `validation_data=(X_val, y_val)` instead.

On the test set the second network scores **86.45%**, against 79.25%. Figure 3 (right) shows what changed: it now finds 189 of the 415 customers who left, while wrongly flagging only 45 who stayed.

8. Training curves

Key point: `fit` returns a History object: the loss and accuracy of every epoch, on both sets. Plotting them shows how training went and whether it overfits.

`fit` returns a **History** object. Its `.history` attribute is a dictionary with one list per quantity, one value per epoch: `loss`, `accuracy`, `val_loss` and `val_accuracy`. Plotting these lists against the epoch number gives the **training curves** (also **learning curves**) in Figure 4.




Figure 4: Training curves of the second network: training (blue) and validation (orange) loss and accuracy over 100 epochs

Reading Figure 4:

- **Loss (left):** both losses fall fast for about 10 epochs, then slowly. The training loss keeps falling to the end; the validation loss is lowest around epoch 40 (0.355) and then creeps up to 0.359.
- **Accuracy (right):** both rise together at first; after about 20 epochs training accuracy stays roughly 1 point above validation accuracy.

The gap between the two curves measures overfitting. Here it is small, but it is opening: more epochs would not help. Techniques that close such a gap (regularization, dropout, stopping training early) come in later Notes.

Python: Plotting the curves with Plotly.

```
import plotly.graph_objects as go

h = history.history
fig = go.Figure()
fig.add_scatter(y=h["loss"], name="training loss")
fig.add_scatter(y=h["val_loss"], name="validation loss")
fig.show()
```

9. Summary

Step	First network	Second network
Architecture	11-3-1	11-11-11-1
Hidden activation	sigmoid	ReLU
Parameters	40	276
Epochs	10	100, with 20% validation
Test accuracy	79.25% (predicts “stays” for all)	86.45%

- Prepare as usual: drop identifiers, one-hot encode, split, standardize.
- Keras: `Sequential` + `Dense` layers, `summary`, `compile` (loss, optimizer, metrics), `fit` (epochs, validation split), `predict`.
- Binary classification: one sigmoid output node, binary cross-entropy loss, threshold 0.5 on the probability.
- On imbalanced data, compare accuracy with always predicting the majority class, and look at the confusion matrix.
- Improve by changing epochs, activation, nodes and layers; watch the training curves for a gap that signals overfitting.

10. Sources

- Keras documentation, `Model.fit` (`batch_size` defaults to 32; `validation_split` takes the last samples, before shuffling).
- Goodfellow, Bengio and Courville, *Deep Learning*, MIT Press, 2016, §6.3 (rectified linear units as the default hidden unit).

11. Key terms

Term	Meaning
Keras workflow	Build, compile, fit, predict: the four steps of every Keras model
Sequential model	A Keras model whose layers form one stack, each feeding the next
Dense (fully connected) layer	A layer whose every node receives the output of every node in the layer before
<code>keras.Input</code>	The first item of a Sequential model; gives the shape of one input row
<code>model.summary()</code>	Prints each layer’s output shape and number of trainable parameters
Compile	Choosing the loss, the optimizer and the metrics before training
Adam	The optimizer used in these projects; optimizers are taught later

Term	Meaning
<code>fit</code>	Trains the model on given inputs and outputs for a number of epochs
Batch	The rows used for one weight update; Keras uses 32 by default
<code>get_weights()</code>	Returns a layer's weight matrix and bias vector
Threshold (classification)	The probability above which a prediction counts as class 1
Metric	A score reported during training, such as accuracy, that training does not minimise
<code>validation_split</code>	The share of the training rows Keras holds back as a validation set
History object	What <code>fit</code> returns: the loss and metrics of every epoch
Training curves (learning curves)	Loss or accuracy plotted against the epoch, for the training and validation sets