

Scaling the Inputs of a Neural Network

Contents

1. Overview	2
2. Prerequisites	2
3. The problem in Keras	3
3.1 The data and the network	3
3.2 What happens	3
4. Why unscaled inputs break training	4
4.1 The weight of a large input gets a large gradient	4
4.2 The shape of the loss	5
5. The fix: scale the inputs	5
5.1 Standardization or normalization	5
5.2 The same network on standardized inputs	6
5.3 When to scale	6
6. Summary	7
7. Sources	7
8. Key terms	7

Where this fits: Pipeline map step 5 (Engineer features). See the [Course map](#).



- **Builds on:** Train-test split ([Note 13](#)); Feature engineering ([Note 23](#)); Feature transformation ([Note 23](#)); Normalization ([Note 25](#)); Descriptive statistics ([Note 230](#)); Gradient descent ([Note 1020](#)).
- **Compare with:** Normalization ([Note 25](#)); Decision trees ([Note 100](#)).

1. Overview

Key point: When input columns live on very different scales, a network trains slowly or not at all. Bringing every column to the same scale before training fixes it.

Suppose a network predicts whether a user buys a product from two columns: age, in tens, and salary, in tens of thousands. Fed the raw columns, the network in this Note never gets past guessing. Fed the same columns standardized, it reaches 94% accuracy within 20 epochs.




Figure 1: The same network on raw and on standardized inputs. Left: validation accuracy per epoch. Right: the average size of the gradients of the age weights and the salary weights at the start of training (log scale).

Figure 1 shows both runs and the reason. This Note shows the problem in Keras (section 3), explains why it happens (section 4), and fixes it (section 5). The Notebook ([notebook.ipynb](#)) runs every step.

2. Prerequisites

- The [standardization Note](#) and the [normalization Note](#): the two ways to put columns on one scale.

- The [gradient descent Note](#), section 9: what unscaled inputs do to the shape of the loss.
- The [customer churn Note](#): a Keras network and its validation accuracy.

3. The problem in Keras

Key point: On raw age and salary, the validation accuracy jumps between 35% and 65% for 100 epochs: the network only ever predicts one class.

3.1 The data and the network

Key point: 400 users, two inputs on very different scales, one 0/1 output; one hidden layer of 128 ReLU nodes.

The data is `Social_Network_Ads.csv`, also used in the [standardization Note](#). It holds 400 users of a social network, and we keep three columns:

Column	Range	Meaning
Age	18 to 60	age in years
EstimatedSalary	15,000 to 150,000	estimated yearly salary
Purchased	0 or 1	output: 1 if the user bought the advertised product

Age is in tens and salary in tens of thousands: the second column is about 2,000 times larger. We split 80/20, leaving 320 users for training and 80 for validation.

The network has an input layer of 2 nodes, one hidden layer of 128 ReLU nodes and one sigmoid output node, since this is binary classification. It is compiled with Adam, binary cross-entropy and accuracy, and trained for 100 epochs.

Python: The network, trained on the raw columns.

```
model = keras.Sequential([
    keras.Input(shape=(2,)),
    keras.layers.Dense(128, activation="relu"),
    keras.layers.Dense(1, activation="sigmoid"),
])
model.compile(optimizer="adam", loss="binary_crossentropy",
              metrics=["accuracy"])
history = model.fit(X_train, y_train, epochs=100,
                    validation_data=(X_test, y_test))
```

3.2 What happens

Key point: The accuracy never converges, and the loss stays in the tens.

Figure 1 (left, red) shows the validation accuracy. It jumps up and down between about 35% and 65% for all 100 epochs and never settles.

The two values give the game away. Of the 80 validation users, 65% did not buy and 35% did. An accuracy of 65% means the network predicted “did not buy” for everyone; 35% means it predicted “bought” for everyone. It flips between the two and never learns to separate the users.

The training loss tells the same story. Binary cross-entropy of a coin toss is 0.69, yet the loss starts at 3,618 and still swings between about 20 and 170 in the last 10 epochs.

4. Why unscaled inputs break training

Key point: The gradient of a weight is proportional to its input. The salary weights get gradients about 2,000 times larger than the age weights, so training attends to salary and ignores age.

4.1 The weight of a large input gets a large gradient

Key point: For a first-layer weight, gradient = (error signal of its node) $\times$ (its input).

Call the two inputs x_1 (age) and x_2 (salary), and their weights into one hidden node w_1 and w_2 . Backpropagation updates both with $w_{\text{new}} = w_{\text{old}} - \eta \partial L / \partial w$.

1. **In words:** the node computes $z = w_1 x_1 + w_2 x_2 + b$. By the chain rule, the gradient of each weight is the node’s error signal $\delta = \partial L / \partial z$ times that weight’s input.
2. **Formula:**

$$\frac{\partial L}{\partial w_1} = \delta \cdot x_1, \quad \frac{\partial L}{\partial w_2} = \delta \cdot x_2$$

3. **Example:** a user aged 40 earning 80,000, and $\delta = 0.01$:

$$\frac{\partial L}{\partial w_1} = 0.01 \times 40 = 0.4, \quad \frac{\partial L}{\partial w_2} = 0.01 \times 80,000 = 800$$

The salary weight gets a gradient 2,000 times larger.

So salary takes over. With plain gradient descent, nearly all the change would go into w_2 . Our network uses Adam, which moves every weight by about the same small amount (Kingma and Ba 2015, §2.1), but salary still takes over, because the same small step moves z 2,000 times further through salary than through age. Think of two volume knobs, one normal and one extremely sensitive: the same small twist barely changes the first, but makes the second jump from silent to full blast.

In numbers: in the first epoch both kinds of weight moved by about 0.004, which shifts z by about $0.004 \times 70,000 = 280$ through salary and only $0.004 \times 38 = 0.15$ through age. So z swings wildly from step to step, the predictions flip between all 0 and all 1, and training is unstable, as in Figure 1.

Extra: The Notebook tests this by scaling one feature at a time. Standardizing only salary gives 85% to 86% validation accuracy over the last 10 epochs; standardizing only age leaves the accuracy jumping between 35% and 85%. A 100 times smaller learning rate on the raw features does not help either: the network then predicts “did not buy” for everyone (65%).

The Notebook measures this on the real network before any training (Figure 1, right). With raw inputs, the salary weights’ gradients average 1,393 and the age weights’ 0.73: about 1,900 times smaller, close to the ratio of the average salary to the average age ($69,742 / 38$). After standardizing, both are around 0.01, within a factor of 2.

4.2 The shape of the loss

Key point: Unscaled inputs make the loss a long, narrow valley that gradient descent zigzags through; scaled inputs make it round.

The same effect has a geometric picture, taught in the [gradient descent Note](#), section 9. With inputs on very different scales, the contours of the loss are long, narrow ellipses: steep in one direction, flat in the other. Gradient descent oscillates across the valley before it gets anywhere. With inputs on the same scale, the contours are nearly circles, and gradient descent heads straight for the minimum.

5. The fix: scale the inputs

Key point: Standardize (or normalize) every input column, fitting the scaler on the training rows only. The same network then reaches 94%.

5.1 Standardization or normalization

Key point: Either works; standardization is the usual choice, normalization when a column’s minimum and maximum are known.

Both techniques put the columns on one scale (see the [feature scaling Notes](#)):

- Standardization subtracts the mean and divides by the standard deviation: every column gets mean 0 and standard deviation 1.
- Normalization (min-max scaling) subtracts the minimum and divides by the range: every column lands between 0 and 1 (see the [normalization Note](#)).

Which to use follows the rules of the [normalization Note](#), section 10. In short: normalize when the minimum and maximum are known in advance, such as a CGPA from 0 to 10; otherwise standardize, especially when the column is close to normally distributed. Salary has no known maximum, so we standardize.

Extra: Standardized values are not confined to -1 to $+1$. They have standard deviation 1, so most values fall between about -2 and $+2$, and outliers further out. Here age runs

from -1.95 to 2.17 and salary from -1.61 to 2.32 . Only min-max scaling guarantees a fixed range.

5.2 The same network on standardized inputs

Key point: 80% validation accuracy after 2 epochs, 94% after 20, and stable from there.

Python: Standardizing, then training the same network.

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

history = model.fit(X_train_scaled, y_train, epochs=100,
                    validation_data=(X_test_scaled, y_test))
```

The scaler learns the mean and standard deviation from the training rows only, then applies them to both sets (see the [standardization Note](#), section 6.3).

Plotted as a scatter, the scaled data looks exactly like the raw data; only the numbers on the axes change (see the [standardization Note](#), section 7.1). Nothing else changes either: same network, same starting weights, same 100 epochs.

The result is Figure 1 (left, blue). The validation accuracy climbs steadily: 80% after 2 epochs, 94% after 20, and it stays there. The training loss falls smoothly from 0.67 to 0.24.

5.3 When to scale

Key point: Always: scaling never hurts a network and often helps a lot. Columns already on the same scale can skip it.

Scaling is a standard pre-processing step whenever data goes into a neural network. It does no harm, and in cases like this one it decides whether the network learns at all. The only exception is data whose columns are already on the same scale, such as the pixels of an image, which all run from 0 to 255 and are simply divided by 255 (see the [MNIST Note](#), section 3).

6. Summary

	Raw inputs	Standardized inputs
Age, salary range	18 to 60, 15,000 to 150,000	about -2 to $+2$ each
Mean gradient, age weights	0.73	0.016
Mean gradient, salary weights	1,393	0.009
Validation accuracy	jumps between 35% and 65%	94% from epoch 20
Training loss after 100 epochs	55	0.24

- A weight's gradient is proportional to its input, so large inputs dominate the updates.
- Geometrically, unscaled inputs stretch the loss into a narrow valley that gradient descent zigzags through.
- Standardize (or normalize) every input, fitting the scaler on the training data only.
- Make it a habit: scale before any data reaches a network.

7. Sources

- Kingma and Ba, “Adam: A Method for Stochastic Optimization”, ICLR 2015, §2.1 (step size about the learning rate, invariant to the scale of the gradient).

8. Key terms

Term	Meaning
Normalizing inputs	Bringing every input column of a network to the same scale before training, by standardization or normalization