

Nesterov Accelerated Gradient (NAG)

Contents

1. Overview	2
2. Prerequisites	2
3. Momentum's problem: oscillations	3
4. Two pushes at once, or one after the other	3
5. The update rule	4
6. Why NAG overshoots less	5
7. The weakness: less momentum to escape a dip	6
8. NAG on real data: MNIST	7
9. Momentum and NAG in Keras	8
10. Summary	9
11. Sources	9
12. Key terms	9

1. Overview

Key point: NAG is momentum with one change: it first takes the momentum jump, then measures the gradient at the point where that jump lands (the look-ahead point), and corrects from there. Seeing the slope ahead lets it brake before the minimum, so it overshoots less than momentum.

Nesterov accelerated gradient (NAG; Nesterov 1983, adapted to deep learning by Sutskever et al. 2013) is a small upgrade of [momentum](#). Momentum's speed is its strength, but near the minimum the same speed makes it overshoot and swing. NAG damps those swings.




Figure 1: Momentum and NAG along the same narrow valley, both with learning rate 0.01 and $\beta = 0.9$. Momentum swings across the valley and past the minimum; NAG's swings are much smaller

Figure 1 shows the difference. On this valley momentum needs 59 steps to bring the loss below 0.01; NAG needs 25.

2. Prerequisites

- The [SGD with momentum Note](#): the velocity $v_t = \beta v_{t-1} + \eta \nabla L(w_t)$, and overshooting.
- The [EWMA Note](#): the decay factor β .

3. Momentum's problem: oscillations

Key point: Momentum reaches the neighbourhood of the minimum fast, then wastes many steps swinging past it. A smaller β reduces the swings, but also the speed.

Momentum covers the early part of the road quickly: in the valley of Figure 1 it is close to the minimum within about 20 steps. Then it crosses the minimum, comes back, crosses again, and only settles once its velocity has died away.

Lowering the decay factor β tames the swings, at the cost of some speed. On the more complex, non-convex losses of neural networks the oscillations cost even more time. NAG offers another fix: keep the speed, and reduce the oscillations.

4. Two pushes at once, or one after the other

Key point: Momentum adds the momentum jump and the gradient step, both computed at the current point. NAG first makes the momentum jump, then computes the gradient where it lands.

Write momentum's update in one line by putting v_t into $w_{t+1} = w_t - v_t$:

$$w_{t+1} = w_t - \beta v_{t-1} - \eta \nabla L(w_t)$$

Each step has two parts: the push of the past velocity, βv_{t-1} , and the push of the current gradient, $\eta \nabla L(w_t)$. Plain gradient descent has only the second. Momentum computes both at the current point w_t and applies them together.

NAG changes only **where the gradient is computed**. NAG first applies the momentum part alone, which takes it to a **look-ahead point**. It computes the gradient there and then moves on, or back, according to that gradient (Goodfellow et al. 2016, §8.3.3).




Figure 2: One step near the minimum. Left, momentum: the momentum jump and the gradient step are both computed at w_t , where the slope still points right, so the step overshoots. Right, NAG: jump first, then measure the slope at the look-ahead point, which points back, so the step is corrected

Figure 2 shows one step near the minimum. Hinton’s lecture slides put the difference in one line: first make a big jump in the direction of the previous accumulated gradient, then measure the gradient where you end up and make a correction; it is better to correct a mistake after you have made it (Hinton 2012, lecture 6c).

5. The update rule

Key point: Look-ahead $w_{\text{la}} = w_t - \beta v_{t-1}$; velocity $v_t = \beta v_{t-1} + \eta \nabla L(w_{\text{la}})$; update $w_{t+1} = w_t - v_t$. Only the point where the gradient is taken changes.

1. **In words:** find where the momentum jump alone would take us. Take the gradient there. Build the velocity from the old velocity and that gradient, and move from the current point.
2. **Formula:**

$$\begin{aligned}
 w_{\text{la}} &= w_t - \beta v_{t-1} \\
 v_t &= \beta v_{t-1} + \eta \nabla L(w_{\text{la}}) \\
 w_{t+1} &= w_t - v_t
 \end{aligned}$$

3. **Example:** the same start as for momentum: $L(w) = w^2/2$ (gradient w), $w_0 = -10$, $\eta = 0.1$, $\beta = 0.9$, $v_0 = 0$.
 - Step 1: the look-ahead is $w_0 - 0.9 \times 0 = -10$, so $v_1 = 0.1 \times (-10) = -1$ and $w_1 = -9$, the same as momentum.
 - Step 2: the look-ahead is $-9 - 0.9 \times (-1) = -8.1$. Its gradient is -8.1 , not -9 , so

$$v_2 = 0.9 \times (-1) + 0.1 \times (-8.1) = -1.71, \quad w_2 = -9 - (-1.71) = -7.29$$

Momentum reached -7.2 . NAG's step is a little shorter, because the slope at the look-ahead point is gentler (Notebook).

Putting v_t into the last line gives the whole step: $w_{t+1} = w_t - \beta v_{t-1} - \eta \nabla L(w_{\text{la}})$. It is the distance to the look-ahead point plus a gradient step taken from the look-ahead point.

6. Why NAG overshoots less

Key point: When momentum is about to carry the weights past the minimum, the slope at the look-ahead point already points back. NAG feels that and brakes one step earlier.

Follow a ball rolling down towards the minimum on a curve.

Momentum. At every point it computes the momentum and the slope at that same point. Far from the minimum, both push forward and the ball speeds up. Past the minimum the slope turns and pushes back, but the built-up momentum dominates, so the ball keeps going: it overshoots, makes a U-turn, comes back, makes a smaller U-turn, and so on until it stops.

NAG. Near the minimum, the momentum jump alone already lands past the minimum. The slope there points back, so the gradient step pulls the ball back towards the minimum, and it stops much nearer to it. NAG makes smaller U-turns and settles sooner. NAG has a sense of where it is going, so it slows down before the hill slopes up again (Ruder 2016, §4.2).

If the momentum jump is a poor move that increases the loss, the gradient at the look-ahead point points back towards w_t more strongly than the gradient at w_t itself, which gives a larger and more timely correction to the velocity. On an oblong bowl, momentum swings strongly across the steep direction while NAG avoids these swings almost entirely (Sutskever et al. 2013, §2.1).




Figure 3: The weight over the steps. Left: the bowl $L = w^2/2$ from $w = -10$, $\eta = 0.1$, $\beta = 0.9$; NAG (purple) swings less far past the minimum and settles sooner than momentum (orange). Right: a curve with a local minimum from $w = -3$, $\eta = 0.05$, $\beta = 0.9$; momentum rolls on to the global minimum, NAG stays in the local one

Figure 3 (left) and Figure 1 measure the effect (Notebook):

	Momentum	NAG
Bowl: furthest past the minimum	6.04	3.55
Bowl: stays within 0.1 of the minimum from step	81	42
Valley: steps until the loss is below 0.01	59	25

Extra: On convex problems with exact gradients, Nesterov’s method provably converges faster than plain gradient descent: the excess loss falls like $1/k^2$ after k steps instead of $1/k$ (Nesterov 1983; Goodfellow et al. 2016, §8.3.3). With noisy mini-batch gradients that guarantee is lost (Goodfellow et al. 2016, §8.3.3); the benefit that remains in practice is the damping seen here.

7. The weakness: less momentum to escape a dip

Key point: The damping that removes overshooting also removes some of the extra push. Where momentum rolls out of a small local minimum, NAG can stay stuck in it.

Damping the oscillations has a possible downside. On a loss with a small dip on the way to a deeper one, momentum may build up enough speed to roll over the bump and out of the dip; NAG brakes earlier, so it may not gain that speed and can settle in the local minimum.

Figure 3 (right) shows exactly this on the curve $L(w) = (w^2 - 4)^2/8 - 0.6w$: with the same settings, momentum ends in the global minimum near $w = 2.14$ and NAG in the local one near -1.83 . Over 12 settings (two starts, three learning rates, two values of β), momentum

reached the global minimum in 8 and NAG in 4, and NAG never escaped where momentum did not (Notebook). On such losses another optimizer may do better.

8. NAG on real data: MNIST

Key point: On MNIST with $\beta = 0.9$, momentum and NAG train almost identically. With $\beta = 0.99$, where momentum's swings matter more, NAG's training is steadier and ends slightly lower.

We repeat the MNIST experiment of the [momentum Note](#): 10,000 training images, each with 784 pixel **features** (input variables) and the digit as **target** (the output we predict), the 10,000 test images for validation, hidden layers of 128 and 64 ReLU nodes, learning rate 0.01, batch size 64, 20 epochs, 3 seeds. Only **nesterov** changes, at $\beta = 0.9$ and at $\beta = 0.99$.




Figure 4: Training loss on MNIST per epoch, mean of 3 seeds. Left: $\beta = 0.9$, the two curves overlap. Right: $\beta = 0.99$, NAG (purple) mostly below momentum (orange)

Figure 4 and the Notebook give, as means over the 3 seeds:

	$\beta = 0.9$: momentum	$\beta = 0.9$: NAG	$\beta = 0.99$: momentum	$\beta = 0.99$: NAG
Training loss, epoch 5	0.177	0.175	0.122	0.108
Training loss, epoch 20	0.021	0.020	0.036	0.033
Validation accuracy, epoch 20	0.950	0.949	0.944	0.946
Epochs where the loss went up	0	0	4	1

With $\beta = 0.9$ the two are practically the same: the gradient at the look-ahead point is almost the gradient at the current point. With $\beta = 0.99$ momentum's velocity is ten times longer-lived, its training loss went up from one epoch to the next in 4 epochs, against 1 for NAG, and NAG ends with a slightly lower loss and slightly higher validation accuracy. The differences are small and vary from seed to seed: in one of the three seeds momentum ended lower (Notebook).

Sutskever et al. (2013, §2.1) found the same pattern: NAG changes the velocity in a quicker and more responsive way, which lets it behave more stably than momentum in many situations, especially for higher values of β .

9. Momentum and NAG in Keras

Key point: Keras' SGD class covers all three: `momentum=0` is plain SGD, `momentum=0.9` is momentum, and `momentum=0.9, nesterov=True` is NAG.

Python: Three optimizers from one class.

```
import keras
sgd = keras.optimizers.SGD(learning_rate=0.01, momentum=0.0)
mom = keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
nag = keras.optimizers.SGD(learning_rate=0.01, momentum=0.9,
                           nesterov=True)
```

`momentum` is the decay factor β , usually 0.9 or 0.5; `nesterov` switches the look-ahead on.

Extra: Keras never computes the gradient at a separate look-ahead point. Its weights are the look-ahead point itself: write $u_t = w_t - \beta v_{t-1}$. Putting $v_t = \beta v_{t-1} + \eta \nabla L(u_t)$ into $u_{t+1} = w_{t+1} - \beta v_t = w_t - v_t - \beta v_t$ gives

$$u_{t+1} = u_t - \beta v_t - \eta \nabla L(u_t)$$

a step that needs only the gradient at the stored weights. With $m = -v$ this is Keras' rule $m \leftarrow \beta m - \eta g$, $w \leftarrow w + \beta m - \eta g$ (Keras SGD documentation). The Notebook checks it: Keras' weights $-10, -8.1, -5.75, -3.27, \dots$ are exactly our look-ahead points.

10. Summary

	Momentum	NAG
Gradient taken at	the current point w_t	the look-ahead point $w_t - \beta v_{t-1}$
Near the minimum	overshoots, large swings	brakes earlier, small swings
Bowl $L = w^2/2$: furthest past the minimum	6.04	3.55
Valley: steps to loss < 0.01	59	25
Small local dip	more push to roll out	can stay stuck
Keras	<code>SGD(momentum=0.9)</code>	<code>SGD(momentum=0.9, nesterov=True)</code>

- NAG is momentum with the gradient measured after the momentum jump, at the look-ahead point.
- Seeing the slope ahead lets NAG correct the velocity before it overshoots, so it oscillates less and often converges faster.
- The same damping can leave it in a small local minimum that momentum would roll out of.
- On MNIST it matches momentum at $\beta = 0.9$ and is slightly steadier at $\beta = 0.99$.

11. Sources

- Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*. MIT Press. §8.3.3 Nesterov momentum.
- Hinton, G. (2012). *Neural Networks for Machine Learning*, Coursera, lecture 6c: The momentum method (slides).
- Nesterov, Y. (1983). A method for solving the convex programming problem with convergence rate $O(1/k^2)$. *Soviet Mathematics Doklady* 27, 372–376.
- Ruder, S. (2016). An overview of gradient descent optimization algorithms. arXiv:1609.04747. §4.2 Nesterov accelerated gradient.
- Sutskever, I., Martens, J., Dahl, G. and Hinton, G. (2013). On the importance of initialization and momentum in deep learning. ICML 2013. §2.1.
- Keras API documentation: SGD optimizer (`momentum`, `nesterov`), keras.io/api/optimizers/sgd.

12. Key terms

Term	Meaning
Nesterov accelerated gradient (NAG)	Momentum that computes the gradient at the look-ahead point instead of the current point
Look-ahead point	Where the momentum jump alone would take the weights: $w_t - \beta v_{t-1}$

Term	Meaning
Oscillation	Swinging back and forth past the minimum before settling
Damping	Reducing the size of the oscillations
Feature	An input variable, such as one pixel of an image
Target	The output we predict, such as the digit