

What Is Self-Attention

Contents

1. Overview	2
2. Prerequisites	3
3. Turning words into numbers	3
4. Static embeddings hold an average meaning	4
4.1 The average, worked out	4
4.2 The same effect in real text	5
5. Static and contextual embeddings	6
6. Self-attention: static in, contextual out	6
7. Summary	9
8. Sources	9
9. Key terms	10

Where this fits: Pipeline map steps 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** Word embeddings ([Note 1057](#)); Attention mechanism ([Note 1069](#)).
- **Leads to:** Scaled dot-product attention ([Note 1074](#)); Positional encoding ([Note 1078](#)); Masked self-attention ([Note 1081](#)).
- **Compare with:** Cross-attention ([Note 1082](#)).

1. Overview

Key point: A word embedding gives each word one fixed vector, learned from the average way the word is used. The word “bank” gets the same vector in “money bank grows” and in “river bank flows”. **Self-attention** is a mechanism that takes the static embeddings of all the words of a sentence and returns new, **contextual embeddings**, in which each word’s vector depends on the words around it.

Every NLP application, whether sentiment analysis, named entity recognition or machine translation, starts from words, and computers work with numbers. The first step is always to turn words into numbers. Word embeddings do this well, but they have one weakness: they are **static**, the same in every sentence. Self-attention is the transformer’s answer to that weakness, and it is the core of the transformer (the [introduction to transformers Note](#)).

This Note explains **what** self-attention does and why it is needed. The next Note, the [self-attention step by step Note](#), explains **how** it works inside.




Figure 1: Self-attention as a box: the same static vector e_{bank} goes in for both sentences, but a different contextual vector y_{bank} comes out

2. Prerequisites

- [One-hot encoding](#): a category as a vector of 0s with one 1.
- [Vectors and feature vectors](#), section 6.2: bag of words.
- [RNN sentiment analysis](#), section 6: word embeddings and Keras' **Embedding** layer.
- [Dot product and cosine similarity](#): how similar two vectors are.

3. Turning words into numbers

Key point: One-hot encoding and bag of words turn words into numbers but carry no meaning. Word embeddings carry meaning: similar words get similar vectors.

Turning text into numbers is called **vectorization**. The early methods count words:

- **One-hot encoding** lists the **vocabulary** (the distinct words) and writes each word as a vector with a 1 at its own place. For the vocabulary mat, cat, rat: mat is $[1, 0, 0]$, cat is $[0, 1, 0]$, rat is $[0, 0, 1]$. The sentence “mat cat mat” becomes $[1, 0, 0]$, $[0, 1, 0]$, $[1, 0, 0]$.
- **Bag of words** counts how often each vocabulary word appears in a sentence: “mat cat mat” becomes $[2, 1, 0]$ (two mats, one cat, no rat) and “cat rat rat” becomes $[0, 1, 2]$.
- **TF-IDF** improves bag of words by weighting each count by how rare the word is across documents (SLP3 ch. 5).

None of these says anything about meaning: in one-hot space, every pair of words is equally far apart. **Word embeddings** fix this (the [RNN sentiment analysis Note](#), section 6). A neural network reads a large text collection, such as all of Wikipedia, sees how each word is

used, and represents every word as a dense vector of n numbers, where n is typically 64, 256 or 512. Words used in similar ways get similar vectors.

A useful picture: each embedding is an arrow, and **each direction in the space is one aspect of meaning**, not each number. King and queen point in similar directions, and cricketer points elsewhere. The step from “man” to “woman” is roughly the same arrow as the step from “king” to “queen”, so that arrow can be read as a direction for gender, and the dot product of a word’s vector with it measures how much of that aspect the word carries (Sanderson 2024, Ch 5). The single numbers of an embedding have no labels and rarely a clean meaning of their own; the meaning lives in directions, which mix many numbers. The [meaning as direction](#) Note shows this on real word vectors. For this Note one fact is enough: similar words lie close together.

4. Static embeddings hold an average meaning

Key point: An embedding is trained once on a whole corpus, so a word’s vector reflects how the word is used **on average** in that corpus. Whatever the sentence, the word gets that one vector.

4.1 The average, worked out

Key point: If “apple” is a fruit in 9 of 10 sentences, its vector leans 9 to 1 towards fruit.

Suppose we train embeddings with two dimensions, which we read as **taste** and **technology**. Sentences such as “an apple a day keeps the doctor away”, “apple is healthy” and “apple is better than orange” pull the vector of “apple” towards taste; “apple makes great phones” pulls it towards technology.

1. **In words:** each use of the word pulls its vector towards the meaning of that use. The final vector is the average pull over the whole corpus.
2. **Formula:** with n_f fruit sentences pulling towards $(1, 0)$ and n_c company sentences pulling towards $(0, 1)$,

$$e_{\text{apple}} = \frac{n_f (1, 0) + n_c (0, 1)}{n_f + n_c}$$

3. **Example:** corpus A has 9,000 fruit sentences and 1,000 company sentences:

$$e_{\text{apple}} = \frac{9000 (1, 0) + 1000 (0, 1)}{10000} = (0.9, 0.1)$$

Corpus B, with the numbers swapped, gives $(0.1, 0.9)$ (Figure 2).




Figure 2: The same word gets a different static vector from a different corpus: its vector is the average meaning in that corpus

The formula is a simplification: real embedding methods do not literally average. But they are built from counts of which words appear near which, summed over every occurrence of the word (SLP3 ch. 5), so the vector is fixed by the mix of uses in the corpus.

4.2 The same effect in real text

Key point: Embeddings learned from 116,830 real English sentences put “bank” next to shop and supermarket: in that text, a bank is always a building where money is kept.

The Notebook learns 50-number embeddings from the 116,830 distinct English sentences of the Keras English–French translation dataset. It counts which words appear within 4 words of each other and compresses the counts with the SVD, the method of latent semantic analysis (the [SVD in machine learning Note](#), section 3).

Extra: The counts are first weighted by **positive pointwise mutual information** (PPMI), which scores how much more often two words appear together than chance would give, and sets negative scores to 0. The SVD then keeps the 50 strongest directions. Word2vec, the best-known embedding method, can be seen as implicitly working with such a PPMI-weighted count matrix (Levy and Goldberg 2014, cited in SLP3 ch. 5).

The nearest neighbours of a few words, by cosine similarity:

Word	Six nearest words
bank	shop, bookstore, supermarket, hall, convenience, hospital
river	lake, sea, across, flows, village, bridge

Word	Six nearest words
apple	pie, fruit, apples, baked, banana, orange
phone	telephone, cell, passport, address, number, card

“Bank” appears in 95 sentences of this corpus, and none of them mentions a river (“he works at a bank”, “is there a bank near here?”). So its static vector is the money kind of bank: its cosine similarity is 0.40 with “money” and only 0.24 with “river”. “Apple” is a fruit here, because the corpus never mentions the company.

5. Static and contextual embeddings

Key point: A static embedding gives a word the same vector in every sentence. A **contextual embedding** changes the word’s vector according to the other words of the sentence.

Static embeddings are trained once and then used again and again, in every sentence. That is the problem. Suppose we translate into Hindi the sentence

“Apple launched a new phone while I was eating an orange.”

With static embeddings, “apple” enters the translation model with its average vector, (0.9, 0.1) in corpus A: a fruit. In this sentence, it is the company. What we want is an embedding that notices “launched” and “phone”, raises the technology part and lowers the taste part. It should also not be misled by “orange”, which belongs to a different part of the sentence.

The same holds for two shorter phrases, “money bank grows” and “river bank flows”. The word “bank” means a financial institution in the first and the side of a river in the second, but a static embedding gives it exactly the same vector in both. In the Notebook, Keras’ **Embedding** layer, loaded with the embeddings of section 4.2, returns identical numbers for “bank” in the two phrases: an embedding layer is a lookup table, and the lookup ignores the neighbours. SLP3 (ch. 7) makes the same point: with static embeddings, a word’s representation “is always the same vector irrespective of the context”.

An embedding that changes with the sentence is a **contextual embedding** (SLP3 ch. 7).

6. Self-attention: static in, contextual out

Key point: Self-attention is a function: the static embeddings of all the words of a sentence go in, one contextual embedding per word comes out. In the Notebook, “bank” moves towards “river” in “river bank flows” and towards “money” in “money bank grows”.

Self-attention is a mechanism that turns static embeddings into contextual embeddings. Think of it as a box (Figure 1). The embeddings of all the words of a sentence go in together, $e_{\text{river}}, e_{\text{bank}}, e_{\text{flows}}$. Some calculations happen inside, and one new vector per word comes out, $y_{\text{river}}, y_{\text{bank}}, y_{\text{flows}}$. Each output is a **contextual embedding**: it takes into account

how the word is used in this sentence. The original paper describes self-attention as “an attention mechanism relating different positions of a single sequence in order to compute a representation of the sequence” (Vaswani et al. 2017, section 2). The transformer then uses these contextual embeddings in its later layers.

The Notebook passes the two phrases through the simplest form of self-attention, the one the next Note builds from scratch, with the real embeddings of section 4.2. The cosine similarity of the vector for “bank” with “money” and with “river”:

Vector for “bank”	Similarity to “money”	Similarity to “river”
static, in any sentence	0.40	0.24
after self-attention, “money bank grows”	0.68	0.13
after self-attention, “river bank flows”	0.21	0.71

The static vector leans a little towards money, the average meaning of the corpus. After self-attention, the vector for “bank” in “river bank flows” is now much closer to “river” than to “money”, although the corpus never once used “bank” in that sense. The neighbours alone changed it.




Figure 3: Real word vectors, flattened to 2-D. The static “bank” (grey square) sits with the money words. After self-attention (red diamonds), the “bank” of “river bank flows” has moved towards the river words; the “bank” of “money bank grows” stays on the money side and moves towards “grows”

Figure 3 shows the same thing on a map. The map keeps only 2 of the 50 directions, so distances on it are approximate; the table gives the exact similarities.

How the box computes its output, and how it learns to do so for a particular task, is the subject of the [self-attention step by step Note](#).

7. Summary

	Static embedding	Contextual embedding
Vector for a word	one, fixed after training	different in every sentence
Learned from	the average use of the word in the training corpus	the word and its neighbours in the current sentence
“bank” in “river bank flows”	the money meaning (similarity 0.24 to “river”)	moved to the river meaning (0.71)
Produced by	an embedding (Word2Vec, GloVe, Keras Embedding)	self-attention

- Every NLP task starts by turning words into numbers: one-hot, bag of words, TF-IDF, then word embeddings, which capture meaning.
- A static embedding holds the average meaning of a word in its training corpus.
- A word with several meanings needs a vector that depends on the sentence: a contextual embedding.
- Self-attention takes the static embeddings of a sentence and returns contextual embeddings, one per word.

8. Sources

Built from

- CampusX, “What is Self Attention | Transformers Part 2 | CampusX”, YouTube, <https://www.youtube.com/watch?v=XnGGmvpDLA0>
- Sanderson, G. (3Blue1Brown), “Transformers, the tech behind LLMs | Deep Learning Chapter 5”, 2024, [3blue1brown.com/lessons/gpt](https://www.youtube.com/watch?v=wjZofJX0v4M), <https://www.youtube.com/watch?v=wjZofJX0v4M>. 13:33–17:44 (directions in embedding space carry meaning; woman – man $\approx$ queen – king; the dot product measures how much of a direction a vector has).

Other references

- Jurafsky, D. and Martin, J. H. *Speech and Language Processing*, 3rd ed. draft (19 August 2026), web.stanford.edu/~jurafsky/slp3. Chapter 5 (static embeddings; tf-idf; word2vec and PPMI, citing Levy and Goldberg 2014; SVD on word co-occurrence counts); chapter 7, section 7.1 (static embeddings are the same “irrespective of the context”; contextual embeddings; the “bank” of a river).
- Vaswani, A. et al. (2017). Attention Is All You Need. *NeurIPS 2017*. arXiv:1706.03762. Section 2 (definition of self-attention).
- Levy, O. and Goldberg, Y. (2014). Neural Word Embedding as Implicit Matrix Factorization. *NeurIPS 2014*. Cited in SLP3 ch. 5.
- Tatoeba (tatoeba.org) English–French sentence pairs, prepared by many-things.org/anki (file of 7 April 2019) and distributed for the Keras examples at storage.googleapis.com/download.tensorflow.org/data/fra-eng.zip. Licence CC BY 2.0.

9. Key terms

Term	Meaning
Vectorization	Turning words or text into vectors of numbers
Vocabulary	The list of distinct words in the data
Bag of words	A sentence as the count of each vocabulary word in it
Word embedding	A dense vector of real numbers for each word, learned so that words used similarly get similar vectors
Static embedding	An embedding that gives a word the same vector in every sentence
Average meaning	What a static embedding stores: the word's meaning averaged over all its uses in the training corpus
Contextual embedding	A vector for a word that depends on the other words of the sentence it appears in
Self-attention	A mechanism that takes the embeddings of all the words of a sequence and returns a contextual embedding for each
PPMI	Positive pointwise mutual information: a score of how much more often two words appear together than by chance, with negative scores set to 0