

The Transformer, End to End

Contents

1. Overview	2
2. Prerequisites	3
3. The chain of problems that led to the transformer	4
4. One sentence through the model	5
4.1 The encoder	6
4.2 The decoder	7
4.3 Parameter count	8
5. Training and inference	9
6. Why each part is there	10
7. How the paper trained the transformer	12
7.1 Data and batching	13
7.2 Hardware and training time	13
7.3 The optimizer and the warm-up learning rate	14
7.4 Regularisation: residual dropout and label smoothing	15
8. What the paper found	17
8.1 Translation results (Table 2)	17
8.2 What changing each part does (Table 3)	18
8.3 English constituency parsing	21
9. Where the idea went next	22
10. Reading the paper	23
11. Summary	25
12. Sources	25
13. Key terms	26

Where this fits: Pipeline map steps 8 (Model), 10 (Tune). See the [Course map](#).



- **Builds on:** Overfitting ([Note 1026](#)); Adam ([Note 1038](#)); Transformer encoder ([Note 1080](#)); Transformer decoder ([Note 1083](#)); Transformer inference (autoregressive decoding, KV cache, beam search) ([Note 1084](#)).

1. Overview

Key point: The transformer is an encoder and a decoder built from attention alone. The encoder reads the whole input sentence at once and gives every word a context-aware vector. The decoder writes the output one word at a time; at each step it looks at the words it has written (masked self-attention) and at the input (cross-attention). Every part fixes a problem of the models that came before it. The paper “Attention Is All You Need” (Vaswani et al. 2017) adds a training recipe on top: a warm-up learning rate, dropout and label smoothing. With that recipe, the big model beat every earlier translation model in the paper’s comparison, at a fraction of their training cost.

The Notes from 1060 to 1084 built the transformer one part at a time. This Note puts the parts back together into one story:

1. **Why** the transformer exists: the chain of problems from RNNs to attention (section 3).
2. **How** one sentence flows through the whole model, with every shape (sections 4 and 5).
3. **Why each part is there**, with the evidence our Notes measured (section 6).
4. **What only the paper teaches:** how it was trained (section 7), what it achieved and what its experiments say about each part (section 8), and where the idea went next (section 9).

Figure 1 is the map of the whole concept. Each box carries the number of the Note that teaches it.




Figure 1: The transformer, base model. Grey numbers: the Note that teaches each part. Blue: the shapes for “we’re friends .” → “nous sommes amis .”

The earlier Notes teach each part in depth. Here each part gets a one-line recap and a link.

2. Prerequisites

- [Problems with RNNs](#): vanishing gradients through time.
- [Encoder–decoder](#) and [attention](#): the context vector, and a fresh one per output word.
- [Self-attention step by step](#) and [scaled dot-product attention](#): queries, keys, values, $\text{softmax}(QK^T / \sqrt{d_k})V$.
- [Multi-head attention](#), [positional encoding](#), [layer normalisation](#).
- [Transformer encoder](#), [masked self-attention](#), [cross-attention](#), [transformer decoder](#), [transformer inference](#).
- [Adam](#), [dropout](#) and the [loss functions Note](#) (categorical cross-entropy), for section 7.

3. The chain of problems that led to the transformer

Key point: Each model in the chain fixed the main problem of the one before and left a new one. RNNs forget long-range information; LSTMs remember better but read one word at a time; the encoder–decoder squeezes a sentence into one vector; attention removes that bottleneck but still runs on an LSTM; self-attention drops the LSTM, so all words are processed at once.




Figure 2: From the RNN to the transformer. Red: the problem each stage left, which the next stage fixed

1. **RNNs read word by word, and forget.** An RNN reads one word per step and carries a hidden state forward. During training the gradient is multiplied by the same recurrent weights once per step, so it vanishes over long distances. On IMDB reviews, a simple RNN fell from 0.736 accuracy to 0.488, guessing level, when 80 blank steps separated the review from the prediction (the [problems with RNNs Note](#)).
2. **LSTMs and GRUs remember better, but stay sequential.** An LSTM adds a protected long-term memory, the cell state (the [LSTM Note](#)); a GRU does the same with one memory and two gates (the [GRU Note](#)). Both still compute step t from step $t - 1$. The paper names the cost: “This inherently sequential nature precludes parallelization within training examples” (Vaswani et al. 2017, §1, p. 2).
3. **The encoder–decoder squeezes a sentence into one vector.** Translation needs an output of a different length from the input. An LSTM encoder reads the input into one **context vector** (G-461), and an LSTM decoder writes the output from it (the [encoder–decoder Note](#)). Long sentences do not fit: in the [attention Note](#), without attention, BLEU fell from 14.3 on the shortest sentences to 5.9 on the longest.
4. **Attention removes the bottleneck.** The decoder gets a fresh context vector at every step, a weighted mix of all encoder states. On the same data, test BLEU rose from 9.8 to 25.7 (the [attention Note](#)). Luong’s dot-product score was both better and faster than Bahdanau’s small network: BLEU 31.6 against 25.1, and 22 against 51 seconds per epoch (the [Bahdanau and Luong Note](#)). But the encoder and decoder were

still LSTMs.

5. **Self-attention drops the recurrence.** If attention can relate any two words directly, the LSTM is no longer needed. In **self-attention** (G-1764) the words of one sentence attend to each other, all at the same time (the [what is self-attention](#) and [why “self”](#) Notes). On a GPU, an LSTM step’s time grew 3.9 times from 16 to 256 words, while a self-attention step stayed near 1 millisecond (the [introduction to transformers](#) Note). The paper’s conclusion states the result: “the first sequence transduction model based entirely on attention, replacing the recurrent layers most commonly used in encoder-decoder architectures with multi-headed self-attention” (Vaswani et al. 2017, §7, p. 10).

Self-attention alone brings new problems: it ignores word order, it has one point of view, and in the decoder it would see future words. The rest of the transformer fixes these (section 6).

4. One sentence through the model

Key point: We follow “we’re friends .” → “nous sommes amis .” through the base model ($d_{\text{model}} = 512$, $h = 8$ heads, $d_{\text{ff}} = 2048$, $N = 6$ blocks per side). Inside the model every word is a 512-number vector from start to end; only the feed-forward hidden layer (2048) and the output layer (one score per vocabulary word) change the width.

Before the shapes, watch the whole flow once in a model that really translates (Figure 3). It is the small trained transformer of the [transformer inference](#) Note: 2 encoder and 2 decoder blocks, 128 numbers per word, 4 heads, trained on 40,000 English–French pairs. Each English word becomes a strip of 128 coloured numbers; the encoder blocks mix the words (the dots are the self-attention weights) and change every strip; then the decoder runs once per French word. At each step, watch three things: the masked self-attention looks only at the words already written, the green lines show how much the newest position reads each English word, and the bars give the 5 most likely next words. The most likely one is appended, and the loop runs again until `<end>` (the predict, append, repeat loop that Sanderson 2024, Ch 5, animates for GPT). The model writes “nous sommes amies .”, reading “we’re” for “nous” (0.93 of its cross-attention weight in the last block) and “friends” for “amies” (0.85).




Figure 3: One sentence through the trained transformer of the transformer inference Note. Left: each English word as a strip of 128 numbers (blue positive, red negative), after the embedding and after each encoder block. Right: the decoder, one run per French word: masked self-attention of the newest word (purple dots), cross-attention to the English words (green line widths), the 5 most likely next words, and the chosen word appended

The example pair is the one of the [transformer decoder Note](#). The Notebook builds a base-size transformer in Keras with random weights and passes the pair through it. With random weights the outputs mean nothing, but the shapes and the parameter counts are those of the real model. The vocabulary is one shared list of about 37,000 tokens, as in the paper’s English–German setup (Vaswani et al. 2017, §5.1, p. 7).

4.1 The encoder

Key point: Tokens $\rightarrow$ embeddings $\times \sqrt{512} \rightarrow$ + positional encoding $\rightarrow$ 6 encoder blocks. Out comes H_{enc} : one context-aware 512-number vector per input token.

Step	What it does	Why it is there	Shape	Taught in
Tokens	“we’re”, “friends”, “.”	the model reads tokens, not letters	3 tokens	1080, §4

Step	What it does	Why it is there	Shape	Taught in
Embedding $\times \sqrt{d_{\text{model}}}$	looks up a learned vector per token, multiplied by $\sqrt{512} \approx 22.6$	words become numbers	3×512	1083 , §5.2
+ positional encoding	adds a sine-cosine vector per position	attention alone ignores order	3×512	1078
Multi-head self-attention	every word mixes in the other words, 8 heads of 64	context: “bank” next to “river” differs from “bank” next to “money”	3×512 (weights: $8 \times 3 \times 3$)	1073–1077
Add & norm	LayerNorm($x + \text{Sublayer}(x)$)	residual path and stable scale	3×512	1079 , 1080
Feed-forward	$512 \rightarrow 2048$ (ReLU) $\rightarrow 512$, each word separately	a non-linear transformation per word	3×2048 , then 3×512	1080 , §5.3
Add & norm	as above	as above	3×512	1080
Repeat 6 times	same design, own weights	depth	$H_{\text{enc}}: 3 \times 512$	1080 , §6

4.2 The decoder

Key point: The target sentence, shifted right, gets the same embedding and positional encoding. Each of the 6 decoder blocks runs masked self-attention, then cross-attention to H_{enc} , then a feed-forward network, each with add & norm. A linear layer and a softmax turn each position’s vector into a probability for every word.

Step	What it does	Why it is there	Shape	Taught in
Shift right	<start> nous sommes amis .	position i holds the word before the one it predicts	5 positions	1083 , §5.1
Embedding, positions	as in the encoder	as in the encoder	5×512	1083 , §5.2
Masked self-attention	each position sees only itself and earlier positions	at prediction time the future words do not exist yet	5×512 (weights: $8 \times 5 \times 5$)	1081

Step	What it does	Why it is there	Shape	Taught in
Cross-attention	queries from the French side, keys and values from H_{enc}	the only place the decoder reads the input	5×512 (weights: $8 \times 5 \times 3$)	1082
Feed-forward	as in the encoder	as in the encoder	5×512	1083 , §6.5
Add & norm; 6 blocks	after each of the 3 sub-layers	as in the encoder	5×512	1083 , §7
Linear	one score (logit) per vocabulary word	turn a vector into word scores	$5 \times 37,000$	1083 , §8
Softmax	scores $\rightarrow$ probabilities	a probability per word, summing to 1	$5 \times 37,000$	1083 , §8

In the Notebook every row of the $5 \times 37,000$ output sums to 1. The paper shares one weight matrix between the two embedding layers and the final linear layer (Vaswani et al. 2017, §3.4, p. 5); the Notebook does the same.

4.3 Parameter count

Key point: One encoder block has 3,152,384 parameters and one decoder block 4,204,032; the 12 blocks and the shared embedding give about 63 million, close to the paper’s 65 million.

Part	Parameters (Notebook, <code>count_params()</code>)
One encoder block	3,152,384
One decoder block	4,204,032
6 encoder + 6 decoder blocks	44,138,496
Shared embedding, $37,000 \times 512$	18,944,000
Total	63,082,496




Figure 4: Where the base model’s 63,082,496 parameters sit: the 6 encoder blocks, the 6 decoder blocks and the shared embedding

In Figure 4, the decoder blocks hold the largest share, because each one has a cross-attention that the encoder blocks lack. These are the counts of the [encoder](#) and [decoder](#) Notes. The paper lists 65 million for the base model (Vaswani et al. 2017, Table 3, p. 9); the [transformer decoder Note](#), section 7, traces the gap to the exact vocabulary size, which the paper does not give.

5. Training and inference

Key point: Training knows the whole target sentence, so all positions are computed in one parallel pass and the loss compares every position with its correct next word. Inference does not know the target, so the decoder runs once per new word, feeding each chosen word back in.

Training. The training data are sentence pairs. Each pair is one **observation** (G-1374) (one record of the data): an English sentence and its French translation, the **target** (G-1949) (the output the model must learn to produce).

1. **Teacher forcing.** The decoder input is the correct target, shifted right, not the model’s own guesses (the [encoder–decoder Note](#), section 5.2).
2. **One pass.** Because all inputs are known in advance, all 5 positions are computed at once. The causal mask keeps the pass honest: in the [masked self-attention Note](#), the masked pass gave the same outputs as feeding the prefixes one at a time, and was 17 times faster at 512 positions.
3. **Loss.** Each position is a classification over the vocabulary, scored with categorical cross-entropy and averaged over positions (the [transformer decoder Note](#), section 9). The paper changes the target of this loss slightly with label smoothing (section 7.4).




Figure 5: Left: in training, the whole shifted target goes into the decoder at once and all five positions are predicted in one pass. Right: in inference, the decoder runs once per new word, and each chosen word is appended to the next input

Figure 5 sets the two side by side: one decoder call for the whole sentence in training, one call per word in inference.

Inference. The encoder runs once. The decoder starts from <start>, picks a word, appends it, and runs again, until it writes <end> (the [transformer inference Note](#)). The mask stays on, so earlier positions never change. Two refinements:

- **KV cache:** the keys and values of earlier positions are stored and reused, so each step computes them only for the new word (the [transformer inference Note](#), Extra; SLP3 §7.8).
- **Beam search:** several candidate sentences are kept at each step. The paper used 4 candidates and a length penalty $\alpha = 0.6$, and set the maximum output length to the input length + 50 (Vaswani et al. 2017, §6.1, p. 8).

6. Why each part is there

Key point: Every part of the transformer answers a specific problem, and our Notes measured most of these problems directly.

Part	Problem it solves	Evidence	Note
Attention	one context vector forgets long sentences	test BLEU 9.8 without attention, 25.7 with	1069
Dot-product scores	an additive scoring network is slow	Luong dot: BLEU 31.6, 22 s per epoch; Bahdanau: 25.1, 51 s	1070

Part	Problem it solves	Evidence	Note
Self-attention	an LSTM reads one word at a time	LSTM time grew 3.9 times from 16 to 256 words; self-attention stayed near 1 ms	1071
Learned W_Q , W_K , W_V	plain dot products of embeddings learn nothing	IMDB accuracy 0.84 with learned matrices, 0.69 without	1073
Divide by $\sqrt{d_k}$	long vectors give huge scores and a saturated softmax	unscaled, at $d = 1,024$: largest weight 0.97 on average, gradient 7.6 times smaller	1074
Multiple heads	one head gives one point of view	same 1,050,624 parameters as one big head; trained BERT heads attend differently; paper: 1 head is 0.9 BLEU worse	1077
Positional encoding	self-attention ignores order	shuffling the input only shuffled the output	1078
Residual connections	deep stacks lose each word's information	without them, one untrained block gave all 30 words of a review the same vector (similarity 1.00, against 0.09 with them)	1080
Layer normalisation	batch normalisation mixes padding into the statistics	72 percent of the positions in a real IMDB batch were padding	1079
Feed-forward network	attention only mixes words; it has no per-word ReLU	changing “you” changed only the output of “you”; the network holds two-thirds of a block's parameters	1080

Part	Problem it solves	Evidence	Note
Causal mask	in the parallel training pass, a word could see the future	same outputs as step-by-step decoding; 17 times faster at 512 positions	1081
Cross-attention	the decoder must read the input	in a trained model, most French words put their largest weight on the English word they translate	1082
Mask at inference	inputs must look like the training inputs	removing it dropped BLEU from 41.5 to 38.9	1084




Figure 6: Five parts of the transformer and the effect measured in the Note that teaches each one: the metric without the part (red) and with it (green). For the residual connections, the metric is the similarity between words after one block, where lower is better

Figure 6 draws five rows of the table as before-and-after bars, each on its own scale. Every part moves its metric in the right direction.

The whole model works: a small transformer trained for about 4 minutes on 40,000 sentence pairs reached a test BLEU of 41.4, against 13.1 for the LSTM encoder–decoder on the same test sentences (the [transformer inference Note](#)). The two models also differ in size and training time, so that gap is not a measure of the architecture alone.

7. How the paper trained the transformer

The parts above are the architecture. The paper also gives a training recipe, which no earlier Note teaches. This section and the next cover it, with the paper’s own numbers.

7.1 Data and batching

Key point: About 4.5 million English–German sentence pairs and 36 million English–French sentences, split into pieces of words, in batches of about 25,000 source and 25,000 target tokens.

- **English–German:** WMT 2014, “about 4.5 million sentence pairs”, encoded with byte-pair encoding into “a shared source-target vocabulary of about 37000 tokens” (Vaswani et al. 2017, §5.1, p. 7).
- **English–French:** WMT 2014, “36M sentences”, split into a 32,000 word-piece vocabulary (§5.1).
- **Batches:** sentence pairs of similar length are grouped together; each batch holds about 25,000 source tokens and 25,000 target tokens (§5.1).

Byte-pair encoding splits words into smaller pieces (the [transformer encoder Note](#), section 4).

7.2 Hardware and training time

Key point: One machine with 8 NVIDIA P100 GPUs: 12 hours for the base model, 3.5 days for the big model.

Model	Steps	Time per step	Total time
Base	100,000	about 0.4 s	12 hours
Big	300,000	1.0 s	3.5 days

(Vaswani et al. 2017, §5.2, p. 7.) The Notebook checks the arithmetic: $100,000 \times 0.4 \text{ s}$ is 11.1 hours, and $300,000 \times 1.0 \text{ s}$ is 3.47 days.

The **big** model is wider: $d_{\text{model}} = 1024$, $d_{\text{ff}} = 4096$, 16 heads, dropout 0.3, 213 million parameters, against 65 million for the base model (Table 3, bottom row, p. 9).

Training cost in FLOPs. A **FLOP** (G-790) is one floating-point operation, one addition or multiplication of decimal numbers. The paper estimates a model’s training cost by “multiplying the training time, the number of GPUs used, and an estimate of the sustained single-precision floating-point capacity of each GPU”, with 9.5 trillion operations per second (9.5 TFLOPS) for a P100 (§6.1 and footnote 5, p. 8). For the base model:

$$12 \times 3600 \text{ s} \times 8 \text{ GPUs} \times 9.5 \times 10^{12} \text{ FLOPs per second} = 3.28 \times 10^{18} \text{ FLOPs}$$

and for the big model $3.5 \times 86400 \times 8 \times 9.5 \times 10^{12} = 2.30 \times 10^{19}$ (Notebook). These are the $3.3 \cdot 10^{18}$ and $2.3 \cdot 10^{19}$ of Table 2.

7.3 The optimizer and the warm-up learning rate

Key point: The learning rate rises in a straight line for the first 4,000 steps, peaks at about 7×10^{-4} , then falls in proportion to $1/\sqrt{\text{step}}$. Starting small protects the first updates, when the gradients near the output are large.

The optimizer is Adam (the [Adam Note](#)) with $\beta_1 = 0.9$, $\beta_2 = 0.98$ and $\epsilon = 10^{-9}$ (Vaswani et al. 2017, §5.3, p. 7). The learning rate is not fixed. It changes with the training step:

1. **In words:** for the first 4,000 steps, raise the learning rate in proportion to the step number; after that, lower it in proportion to one over the square root of the step number.
2. **Formula** (§5.3, eq. 3):

$$\text{lr} = d_{\text{model}}^{-0.5} \cdot \min(\text{step}^{-0.5}, \text{step} \cdot \text{warmup}^{-1.5}), \quad \text{warmup} = 4000$$

The two terms inside the min are equal when $\text{step} = \text{warmup}$, so the peak is at step 4,000, with value $(d_{\text{model}} \cdot \text{warmup})^{-0.5} = (512 \times 4000)^{-0.5}$.

3. **Example** with $d_{\text{model}} = 512$ (Notebook):

Step	1	1,000	2,000	4,000	16,000	100,000
Learning rate	1.7×10^{-7}	1.75×10^{-4}	3.49×10^{-4}	6.99×10^{-4}	3.49×10^{-4}	1.40×10^{-4}

Doubling the step from 1,000 to 2,000 doubles the rate (the straight-line rise). Multiplying the step by 4, from 4,000 to 16,000, halves it (the inverse square root). By step 100,000, the end of base training, the rate is one fifth of the peak.




Figure 7: The learning rate of the base model over its 100,000 training steps

Why start small. The [improving a neural network Note](#), section 3.4, met warm-up as a way to train with large batches. For the transformer, Xiong et al. (2020) give a specific reason. In the original transformer, with layer normalisation after each residual addition, “the expected gradients of the parameters near the output layer are large. Therefore, using a large learning rate on those gradients makes the training unstable. The warm-up stage is practically helpful for avoiding this problem” (Xiong et al. 2020, abstract). They also show that moving the layer normalisation inside the residual branch, the **Pre-LN** transformer (G-1545), makes the warm-up unnecessary.

7.4 Regularisation: residual dropout and label smoothing

Key point: Dropout with rate 0.1 on every sub-layer’s output and on the embeddings, and label smoothing with $\varepsilon = 0.1$: the target puts 0.9 plus a small share on the correct word and spreads the rest over all words, so the model is never pushed to be 100 percent sure.

Residual dropout. Dropout (the [dropout Note](#)) is applied “to the output of each sub-layer, before it is added to the sub-layer input and normalized”, and to the sums of the embeddings and positional encodings, with rate $P_{\text{drop}} = 0.1$ for the base model (Vaswani et al. 2017, §5.4, p. 8).

Label smoothing. With plain cross-entropy, the target for each position is **one-hot** (G-1380): probability 1 on the correct word, 0 on all others. The loss $-\ln p_{\text{correct}}$ keeps falling as p_{correct} approaches 1, so training always pushes the model towards complete certainty. Szegedy et al. (2016, §7), who introduced label smoothing, name the two risks of that push: overfitting, because full probability on every training label “is not guaranteed to generalize”, and a model that becomes “too confident about its predictions”. Label smoothing (G-1031) mixes the one-hot target with a uniform distribution:

1. **In words:** keep $1 - \varepsilon$ of the target on the correct word, and share ε equally among all K words of the vocabulary, the correct one included.
2. **Formula** (Szegedy et al. 2016, §7):

$$q'(k) = (1 - \varepsilon) \delta_{k,y} + \frac{\varepsilon}{K}$$

where y is the correct word and $\delta_{k,y}$ is 1 when $k = y$ and 0 otherwise. The loss is the cross-entropy against q' : $L = -\sum_k q'(k) \ln p(k)$.

3. **Example:** a vocabulary of $K = 4$ words, nous, sommes, amis, <end>, the correct word “nous”, and the paper’s $\varepsilon = 0.1$ (Notebook):

$$q' = [0.9 + 0.025, 0.025, 0.025, 0.025] = [0.925, 0.025, 0.025, 0.025]$$

Three predictions, scored both ways (Notebook):

Prediction p (nous, sommes, amis, <end>)	Loss with one-hot target	Loss with smoothed target
[0.609, 0.224, 0.136, 0.030] (logits 2, 1, 0.5, -1)	0.495	0.633
[0.925, 0.025, 0.025, 0.025] (equal to q')	0.078	0.349
[0.999, 0.0003, 0.0003, 0.0003] (over-confident)	0.001	0.601




Figure 8: The two losses as the probability of the correct word “nous” grows from 0.30 to 0.9999, with the rest shared evenly. The one-hot loss keeps falling towards certainty; the smoothed loss is lowest at 0.925 and rises again

With the one-hot target, the over-confident prediction wins (Figure 8, grey curve). With the smoothed target, the best prediction is q' itself: scanning the probability of “nous” from 0.30 to 0.9999, with the rest shared evenly, the smoothed loss is lowest at exactly 0.925 (Notebook), and at 0.9999 it rises to 0.773. Over-confidence is now penalised.

The paper reports the trade-off: label smoothing “hurts perplexity, as the model learns to be more unsure, but improves accuracy and BLEU score” (Vaswani et al. 2017, §5.4, p. 8).

Perplexity (G-1492) is the exponential of the mean cross-entropy per token; lower means the model gives the correct tokens higher probability (SLP3 §7.7.1). Table 3, rows (D), shows both effects on the English–German development set:

ε_{ls}	0.0	0.1 (base)	0.2
Perplexity	4.67	4.92	5.47

ε_{ls}	0.0	0.1 (base)	0.2
BLEU	25.3	25.8	25.7

Perplexity is best without smoothing, because the model is then free to be fully confident; BLEU is best with $\varepsilon = 0.1$.

8. What the paper found

8.1 Translation results (Table 2)

Key point: The big transformer set new records on both language pairs: 28.4 BLEU on English–German, over 2 BLEU above the best earlier models and ensembles, and 41.8 on English–French. Even the base model beat every earlier English–German model, at the lowest training cost in the table.

The test sets are newstest2014. BLEU (the [history of LLMs Note](#), section 4) measures how many word sequences of a translation match a human reference. An **ensemble** (G-690) combines the predictions of several trained models (the [ensemble learning Note](#)).




Figure 9: Table 2 of the paper: BLEU against training cost. Orange: the transformer. Hollow circles: ensembles. ByteNet has no listed cost and is left out

Model	BLEU EN–DE	BLEU EN–FR	Cost EN–DE (FLOPs)	Cost EN–FR (FLOPs)
ByteNet	23.75			
Deep-Att + PosUnk		39.2		$1.0 \cdot 10^{20}$
GNMT + RL	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$

Model	BLEU EN-DE	BLEU EN-FR	Cost EN-DE (FLOPs)	Cost EN-FR (FLOPs)
ConvS2S	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base)	27.3	38.1	$3.3 \cdot 10^{18}$	(same)
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	(same)

(Vaswani et al. 2017, Table 2, p. 8; the paper gives one cost per transformer for both language pairs.) From these numbers (Notebook):

- **English-German:** the big model beats the best earlier result, the ConvS2S ensemble, by 2.04 BLEU, at 3.3 times less training cost. The base model beats every earlier model and ensemble at 2.9 times less cost than the cheapest of them with a listed cost, ConvS2S.
- **English-French:** the big model beats every single model, at about a fifth (0.19) of the cost of MoE. The three ensembles score just below it but cost 35 to 52 times more.

The text of §6.1 gives the English-French score of the big model as 41.0, while the abstract and Table 2 give 41.8; we use the table. The paper also averaged the last 5 saved copies of the weights (checkpoints) for the base model and the last 20 for the big model (§6.1, p. 8).

8.2 What changing each part does (Table 3)

Key point: Table 3 changes one setting of the base model at a time and measures English-German BLEU on the development set. Eight heads beat one; shorter keys hurt; bigger models score higher; dropout and label smoothing help; learned and sinusoidal positions tie.

Table 3 is the paper’s own evidence on its design choices (Vaswani et al. 2017, Table 3, p. 9; development set newstest2013, beam search but no checkpoint averaging). The base model scores 25.8 BLEU with 65 million parameters.

Row	What changes	BLEU (dev)	What it shows	Note
(A)	heads h : 1, 4, 16, 32 (same total size)	24.9, 25.5, 25.8, 25.4	one head is 0.9 BLEU worse; “quality also drops off with too many heads”	1077
(B)	key size d_k : 16, 32	25.1, 25.4	“reducing the attention key size d_k hurts model quality”	here
(C)	blocks N : 2, 4, 8	23.7, 25.3, 25.5	6 blocks score best of those tried	1080
(C)	d_{model} : 256, 1024	24.5 (28 M), 26.0 (168 M)	wider vectors score higher	here
(C)	d_{ff} : 1024, 4096	25.4 (53 M), 26.2 (90 M)	a wider feed-forward layer scores higher	here
(D)	dropout: 0.0, 0.2	24.6, 25.5	“dropout is very helpful in avoiding over-fitting”	1024
(D)	label smoothing: 0.0, 0.2	25.3, 25.7	section 7.4	here
(E)	learned positions instead of sinusoids	25.7	“nearly identical results”	1078
big	$d_{\text{model}} = 1024$, $d_{\text{ff}} = 4096$, 16 heads, dropout 0.3, 300K steps	26.4 (213 M)	the best row	here




Figure 10: Table 3 as bars around the base model’s 25.8 BLEU: blue rows beat the base model, red rows fall below it

Figure 10 draws the table. The longest red bar is $N = 2$ blocks; the blue bars are all wider or bigger models.

On rows (B), the paper adds: “This suggests that determining compatibility is not easy and that a more sophisticated compatibility function than dot product may be beneficial” (§6.2, p. 9). On rows (C) and (D): “as expected, bigger models are better” (§6.2).


Table 3 changes sizes and settings; it never removes a part such as the residual connections, the layer normalisation or the mask. The evidence for those parts is the table of section 6.

8.3 English constituency parsing

Key point: To test whether the transformer works beyond translation, the paper trained it to turn sentences into grammar trees. A 4-block transformer with little tuning scored 91.3 F1 trained on 40,000 sentences and 92.7 with extra data, better than all earlier models except one.

Constituency parsing (G-453) finds the grammatical structure of a sentence: which words form a noun phrase, a verb phrase, and so on, as a tree. Vinyals et al. (2015, §2.2 and Figure 2) turned the tree into a sequence by writing it out in depth-first order, so a sequence-to-sequence model can produce it:

“John has a dog .” $\rightarrow$ (S (NP NNP)NP (VP VBZ (NP DT NN)NP)VP .)S



(S (NP NNP)NP (VP VBZ (NP DT NN)NP)VP .)S

the tree written out depth-first: a sequence the decoder can produce

Figure 11: The parse tree of “John has a dog .” and the same tree written out depth-first as one sequence, the form a sequence-to-sequence model produces

Figure 11 shows the tree and its sequence form. The output is “subject to strong structural constraints and is significantly longer than the input” (Vaswani et al. 2017, §6.3, p. 9).

- **Model:** 4 blocks, $d_{\text{model}} = 1024$; the other settings as the English–German base model, with only “a small number of experiments” to choose dropout, learning rates and beam size (§6.3).
- **Data:** the Wall Street Journal part of the Penn Treebank, “about 40K training sentences”; and a semi-supervised setting with about 17 million extra sentences (§6.3).
- **Score:** F1 (the [precision, recall and F1 Note](#)) on section 23 of the WSJ, computed with the standard EVALB tool that compares the predicted tree with the human one (Vinyals et al. 2015, §3.2).

Training data	Best earlier result in the same setting (Table 4)	Transformer, 4 blocks
WSJ only (40K sentences)	91.7 (Dyer et al. 2016)	91.3
Semi-supervised	92.1 (McClosky et al. 2006; Vinyals and Kaiser et al. 2014)	92.7

(Vaswani et al. 2017, Table 4, p. 10.) The paper concludes that “despite the lack of task-specific tuning our model performs surprisingly well, yielding better results than all previously reported models with the exception of the Recurrent Neural Network Grammar”, and that unlike RNN sequence-to-sequence models it outperforms the BerkeleyParser “even when training only on the WSJ training set of 40K sentences” (§6.3, p. 10).

9. Where the idea went next

Key point: The encoder alone became BERT, the decoder alone became GPT, and scaling GPT-style decoders up gave the large language models.

The paper ends with plans: “to extend the Transformer to problems involving input and output modalities other than text” and to make generation “less sequential” (Vaswani et al. 2017, §7, p. 10). What followed is the story of the [history of LLMs Note](#), section 8:

- **Encoder only: BERT** (Google, October 2018), pre-trained to predict hidden words from the words on both sides.
- **Decoder only: GPT** (OpenAI, June 2018), pre-trained to predict the next word.
- **Scale:** GPT grew from 117 million parameters to 175 billion in GPT-3, and ChatGPT was built on a GPT model (the [history of LLMs Note](#), sections 8 and 9).




Figure 12: From the 2017 transformer to ChatGPT: the encoder alone became BERT, the decoder alone became GPT, and scaling GPT up led to GPT-3 and ChatGPT

Figure 12 puts the three bullets on one line of descent.

The [introduction to transformers Note](#), section 6, follows the same architecture into images, proteins and code.

10. Reading the paper

Key point: Every section of “Attention Is All You Need” maps onto a Note. With the Notes behind us, the paper can be read from start to end.

Paper section	Pages	What it says	Note
Abstract	1	attention only; 28.4 and 41.8 BLEU; 3.5 days on 8 GPUs	this Note, §8.1
1 Introduction	2	RNNs are sequential; attention used “in conjunction with a recurrent network”	1071 , this Note §3
2 Background	2	convolutional models; self-attention defined	1072 , 1076
3.1 Encoder and decoder stacks	3	$N = 6$; sub-layers; LayerNorm($x + \text{Sublayer}(x)$); mask	1080 , 1083 , 1079
3.2.1 Scaled dot-product attention	4	eq. 1; why divide by $\sqrt{d_k}$	1073 , 1074 , 1075
3.2.2 Multi-head attention	4–5	$h = 8$, $d_k = d_v = 64$	1077
3.2.3 Applications of attention	5	encoder self-, masked decoder self-, encoder–decoder attention	1076 , 1081 , 1082
3.3 Position-wise feed-forward	5	eq. 2; $d_{\text{ff}} = 2048$	1080
3.4 Embeddings and softmax	5	shared weights; $\times \sqrt{d_{\text{model}}}$	1083
3.5 Positional encoding	6	sines and cosines	1078
4 Why self-attention, Table 1	6–7	complexity, sequential operations, path length	1071 , section 5
5.1–5.2 Data, hardware	7	4.5M pairs; 8 P100 GPUs; 12 hours, 3.5 days	this Note, §7.1–7.2
5.3 Optimizer	7	Adam; warm-up schedule, eq. 3	1038 , this Note §7.3
5.4 Regularisation	8	residual dropout; label smoothing	1024 , this Note §7.4

Paper section	Pages	What it says	Note
6.1 Machine translation, Table 2	8	records at a fraction of the cost; beam search	this Note §8.1, 1084
6.2 Model variations, Table 3	9	one change at a time	this Note §8.2
6.3 Constituency parsing, Table 4	9–10	the transformer beyond translation	this Note §8.3
7 Conclusion	10	other modalities; less sequential generation	this Note §9

Table 1 in brief: per layer, self-attention costs $O(n^2 \cdot d)$ with $O(1)$ sequential operations and a maximum path length of $O(1)$ between any two positions; a recurrent layer costs $O(n \cdot d^2)$ with $O(n)$ sequential operations and path length $O(n)$ (Vaswani et al. 2017, Table 1, p. 6). Here n is the sentence length and d the vector size. “The shorter these paths between any combination of positions in the input and output sequences, the easier it is to learn long-range dependencies” (§4, p. 6).

11. Summary

Stage	The transformer’s answer
Words to numbers	embeddings $\times \sqrt{512}$, plus sine–cosine positions
Context	multi-head self-attention: 8 heads of 64, scaled by $\sqrt{d_k}$
Stability in depth	residual connection, then layer normalisation, after every sub-layer
Per-word processing	feed-forward $512 \rightarrow 2048 \rightarrow 512$
Writing the output	masked self-attention, cross-attention to H_{enc} , linear, softmax
Training	teacher forcing, one parallel pass, cross-entropy with label smoothing 0.1
Optimiser	Adam; learning rate up for 4,000 steps, then down as $1/\sqrt{\text{step}}$
Inference	one word per step, mask on, KV cache, beam search of 4

- The transformer is the end of a chain: RNN $\rightarrow$ LSTM $\rightarrow$ encoder–decoder $\rightarrow$ attention $\rightarrow$ self-attention, each fixing the last one’s main problem.
- Every vector inside the base model has 512 numbers; the model has about 63–65 million parameters.
- The warm-up keeps the first updates small, when the gradients near the output are large.
- Label smoothing makes the target 0.925 on the correct word (for 4 words): perplexity gets worse, BLEU better.
- Trained 12 hours on 8 GPUs, the base model beat every earlier English–German model; the big model set records on both language pairs at a fraction of the earlier models’ cost.

12. Sources

Built from

- Sanderson, G. (3Blue1Brown), “Transformers, the tech behind LLMs | Deep Learning Chapter 5”, 2024, 3blue1brown.com/lessons/gpt, <https://www.youtube.com/watch?v=wjZofJX0v4M>. 1:31–5:42 (data flowing through embeddings, attention and MLP blocks to a next-word distribution; predict, append, repeat).
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł. and Polosukhin, I. (2017). Attention Is All You Need. *NeurIPS 2017*. arXiv:1706.03762 (v7). Abstract; §1 (sequential computation); §3.1–3.5; Table 1 and §4; §5.1 (data, 4.5M pairs, 37,000 tokens, 36M sentences, 25,000-token batches); §5.2 (8 P100 GPUs, 0.4 s and 1.0 s per step, 12 hours and 3.5 days); §5.3 (Adam settings, eq. 3, 4,000 warm-up steps); §5.4 (residual dropout 0.1, label smoothing 0.1, the perplexity quote); §6.1 and

footnote 5 (results, checkpoint averaging, beam search, FLOPs estimate, 9.5 TFLOPS for a P100); Table 2; §6.2 and Table 3; §6.3 and Table 4; §7.

- The earlier Notes of this collection, 1060 to 1084, whose parts this Note puts back together (linked in its text).

Other references

- Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J. and Wojna, Z. (2016). Rethinking the Inception Architecture for Computer Vision. *CVPR 2016*. arXiv:1512.00567. §7 (label smoothing, $q'(k) = (1 - \epsilon)\delta_{k,y} + \epsilon/K$; over-fitting and over-confidence).
- Xiong, R. et al. (2020). On Layer Normalization in the Transformer Architecture. *ICML 2020*. arXiv:2002.04745. Abstract (why the warm-up is needed in the Post-LN transformer; Pre-LN without warm-up).
- Vinyals, O., Kaiser, Ł., Koo, T., Petrov, S., Sutskever, I. and Hinton, G. (2015). Grammar as a Foreign Language. *NeurIPS 2015*. arXiv:1412.7449. §2.2 and Figure 2 (linearising a parse tree); §3.2 (EVALB, F1).
- Jurafsky, D. and Martin, J. H. *Speech and Language Processing*, 3rd ed. draft (19 August 2026), ch. 7, §7.7.1 (perplexity as the exponential of the mean cross-entropy) and §7.8 (KV cache).

13. Key terms

Term	Meaning
Transformer	An encoder–decoder built from attention, feed-forward layers and layer normalisation, with no RNN
Observation	One record of the data; here, one sentence pair
Target	The output the model must learn to produce; here, the French sentence and, at each position, the correct next word
Base model	The paper’s standard transformer: $N = 6$, $d_{\text{model}} = 512$, 8 heads, $d_{\text{ff}} = 2048$, 65 million parameters
Big model	The paper’s larger transformer: $d_{\text{model}} = 1024$, $d_{\text{ff}} = 4096$, 16 heads, 213 million parameters
Learning-rate warm-up	Starting training with a small learning rate and raising it over the first steps; 4,000 steps in the paper
Inverse square root decay	Lowering the learning rate in proportion to $1/\sqrt{\text{step}}$ after the warm-up
Residual dropout	Dropout applied to each sub-layer’s output before the residual addition

Term	Meaning
Label smoothing	Replacing the one-hot target by $(1 - \varepsilon)$ on the correct word plus ε/K on every word
One-hot target	A target with probability 1 on the correct word and 0 on all others
Perplexity	The exponential of the mean cross-entropy per token; lower is better
FLOP	One floating-point operation; training cost is counted in FLOPs
Ensemble	Several trained models whose predictions are combined
Checkpoint averaging	Averaging the weights of the last few saved copies of a model
Beam search	Decoding that keeps several candidate sentences at each step
Constituency parsing	Finding the grammatical tree of a sentence: its phrases and how they nest
Pre-LN transformer	A transformer variant with layer normalisation inside the residual branch, which trains without warm-up