

The Problem with the Perceptron: Non-Linear Data

Contents

1. Overview	2
2. Prerequisites	2
3. Three tiny datasets: AND, OR and XOR	3
4. What the perceptron learns	3
4.1 Another way to see it: one input	5
5. The same failure on larger data	6
6. The way forward	7
7. Summary	9
8. Sources	9
9. Key terms	9

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Perceptron ([Note DL-004](#)).
- **Leads to:** Multi-layer perceptron (MLP) ([Note DL-009](#)).

1. Overview

Key point: A perceptron has one straight decision boundary. It learns AND and OR easily, but no line can separate XOR, so on XOR (and on any data that needs a curved decision boundary) it fails, however long it trains.




Figure 1: scikit-learn's Perceptron on AND, OR and XOR

The [perceptron Note](#) showed that a perceptron is a line, plane or hyperplane. This Note shows in code what that means for data that no line can split. Figure 1 is the whole story: two easy tables and one impossible one.

The XOR weakness stalled perceptron research for years (see the [history of deep learning Note](#)), and the reason **multi-layer perceptrons** (G-1270) were needed.

2. Prerequisites

- The [perceptron Note](#): a perceptron is a linear **decision boundary** (G-555).
- The [perceptron loss Note](#): how a perceptron is trained.

3. Three tiny datasets: AND, OR and XOR

Key point: Each table has two binary inputs and four rows. AND outputs 1 only when both inputs are 1; OR when at least one is; XOR when exactly one is.

The three datasets are the logic functions **AND**, **OR** (G-199) and **XOR** (G-2134; exclusive or, defined in the [history section](#)). Each has inputs $x_1, x_2 \in \{0, 1\}$, so there are only four possible rows. The table uses three terms:

- each row is one **observation** (G-1374; one record of the table);
- x_1 and x_2 are its two **features** (G-772; input variables);
- the AND, OR or XOR column is the **target** (G-1949; the output we predict).

x_1	x_2	AND	OR	XOR
0	0	0	0	0
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0

Plotted, the four observations are the corners of a square (Figure 1):

- for AND, the single 1 sits in one corner;
- for OR, the single 0 sits in one corner;
- for XOR, the 1s sit on opposite corners, and so do the 0s.

Python: Building the tables and training one perceptron per table.

```
import numpy as np
from sklearn.linear_model import Perceptron

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y_and = np.array([0, 0, 0, 1])
y_or = np.array([0, 1, 1, 1])
y_xor = np.array([0, 1, 1, 0])

for y in (y_and, y_or, y_xor):
    p = Perceptron(random_state=0).fit(X, y)
    print(p.coef_, p.intercept_, p.score(X, y))
```

4. What the perceptron learns

Key point: AND and OR get a separating line and 100% accuracy. XOR ends with all weights 0, no line at all, and 50% accuracy.

The three perceptrons give:

Table	w_1	w_2	b	Line	Accuracy
AND	2	2	-2	$x_1 + x_2 = 1$	100%
OR	2	2	-1	$x_1 + x_2 = 0.5$	100%
XOR	0	0	0	none	50%

For AND and OR, one line puts the 1s on one side and the 0s on the other (Figure 1, left and middle). For XOR, the Notebook prints the weights after each **epoch** (G-696; one full pass over the training data): they are back at $(0, 0, 0)$ every time. The update for a misclassified observation adds $y_i(x_{i1}, x_{i2}, 1)$ to (w_1, w_2, b) , with labels ± 1 (see the [perceptron loss Note](#), section 7.2). For the four XOR observations these add up to zero:

$$-(0, 0, 1) + (0, 1, 1) + (1, 0, 1) - (1, 1, 1) = (0, 0, 0)$$

So when all four observations are misclassified in turn, as here, each epoch's updates cancel and the epoch ends where it began. Training stops with all weights at 0: $z = 0$ everywhere, the whole plane gets one class, and only half the observations are right (Figure 1, right).

More epochs would not help. Whatever the perceptron does, its decision boundary is always one straight line.

Figure 2 tries every line. A line turns through 340 degrees, and at each angle it takes the position that classifies the most observations correctly.




Figure 2: A line turning through every angle on the four XOR observations, always at its best position. The green side predicts 1; a black ring marks a wrong observation. The count never reaches 4 of 4. Last frame: two lines, $x_1 + x_2 = 0.5$ and $x_1 + x_2 = 1.5$, with the 1s between them

Watch the black ring: at every angle some corner is on the wrong side, and the count never passes 3 of 4. The Extra below proves that no line can do better.

Extra: Why no line can separate XOR. A line works if $b < 0$ for $(0, 0)$, $w_2 + b \geq 0$ for $(0, 1)$, $w_1 + b \geq 0$ for $(1, 0)$, and $w_1 + w_2 + b < 0$ for $(1, 1)$. Adding the middle two gives $w_1 + w_2 + 2b \geq 0$, so $w_1 + w_2 + b \geq -b$. Since $b < 0$, $-b > 0$, so $w_1 + w_2 + b > 0$, which contradicts the last condition. No w_1, w_2, b can satisfy all four.

4.1 Another way to see it: one input

Key point: With one input, XOR becomes the pattern 0, 1, 0. A straight line can rise or fall, but it cannot rise and then fall.

The same failure shows with a single input. Suppose a drug is tested at three doses:

- a **low** dose does not work (0);
- a **medium** dose works (1);
- a **high** dose does not work (0).

We fit a straight line to the nine patients and predict “works” wherever the line is at 0.5 or above. Figure 3 turns the line through every slope and, at each slope, puts it at its best height.




Figure 3: A drug that works only at a medium dose. A straight line is turned through every slope and placed at its best height; a black ring marks a patient predicted wrong. The line never gets more than 2 of the 3 dose groups right. Idea after StatQuest, “The Essential Main Ideas of Neural Networks”.

1. A **rising line** is low on the left and high on the right. It gets the low and medium doses right, and wrongly says the high dose works.
2. A **falling line** makes the mirror mistake on the low dose.
3. A **flat line** gives every dose the same answer, so the medium dose is wrong.

At every slope the count stays at 2 of 3. The target goes up and then down, and a straight line changes direction never. The data needs a curve that bends; the [MLP intuition Note](#) builds such a curve from two neurons for this same drug.

5. The same failure on larger data

Key point: One sigmoid neuron separates two blobs perfectly but stays near 50% on XOR-shaped and circular data.

The **TensorFlow Playground** (G-1958; playground.tensorflow.org) is a website that trains small neural networks in the browser, with no code. Removing every **hidden layer** (G-890) leaves a single neuron (with a **sigmoid** (G-1798) activation, that is **logistic regression** (G-1120); see the [perceptron loss Note](#), section 8). On two separate clusters it finds a line almost at once; on XOR-shaped data its error stays high no matter how many epochs it runs.

The Notebook repeats this experiment in scikit-learn with `LogisticRegression`, which is exactly one sigmoid neuron:




Figure 4: One sigmoid neuron on linearly separable, XOR-shaped and circular data

- **Two blobs** (Figure 4, left): a line separates them, accuracy 100%.
- **XOR quadrants** (middle): class 1 where x_1 and x_2 have the same sign. A line cuts across all four quadrants, accuracy 52%.
- **Circles** (right): one class inside a ring of the other. A line cannot enclose anything, accuracy 50%.

So the problem is not the four-row table: any data whose classes need a bent or closed decision boundary defeats a single neuron. Such data is called [non-linear data](#) (G-1335): no line, plane or hyperplane separates its classes.

Extra: Two ways out exist.

1. **Add a feature by hand.** Feed x_1x_2 as a third input, like the [polynomial features](#) (G-1513) of linear regression (the Playground offers this input too). With it, one neuron reaches 99% on the XOR quadrants in the Notebook. On the four table observations, $x_1 + x_2 - 2x_1x_2$ equals XOR exactly, so $z = x_1 + x_2 - 2x_1x_2 - 0.5$ is a plane in the three inputs x_1, x_2, x_1x_2 that puts the 1s at $z = 0.5$ and the 0s at $z = -0.5$.

2. Add hidden layers, so the network builds such features itself (next section).

6. The way forward

Key point: Combining several perceptrons in layers gives curved decision boundaries: the multi-layer perceptron.

A single perceptron is a linear model, so it can only capture linear relationships. The fix is to use several perceptrons and feed their outputs into another one. The last frame of Figure 2 shows the idea with two lines, one per hidden perceptron: $x_1 + x_2 = 0.5$ (the OR line of section 4) and $x_1 + x_2 = 1.5$. The sum $x_1 + x_2$ is 0 for (0, 0), 1 for (0, 1) and (1, 0), and 2 for (1, 1), so exactly the two 1s fall between the lines. The [MLP intuition Note](#) shows how two lines combined give a curved decision boundary, and how such a network solves XOR on the Playground.

Figure 5 shows why the second perceptron layer then has an easy job. Each hidden perceptron outputs 0 or 1, so each observation gets a new pair of numbers, (h_1, h_2) .




Figure 5: The four XOR observations moved by two hidden perceptrons, $h_1 = \text{step}(x_1 + x_2 - 0.5)$ and $h_2 = \text{step}(x_1 + x_2 - 1.5)$. Left: the input square with the two hidden lines. Right: the new positions, where one line separates the classes.

1. h_1 is the OR line. It outputs 1 for every corner except (0, 0).
2. h_2 is the AND line. It outputs 1 only for (1, 1).
3. The corners move. (0, 0) goes to (0, 0), both (0, 1) and (1, 0) go to (1, 0), and (1, 1) goes to (1, 1). The two green observations now sit at the same spot.
4. One line is enough. In the new positions, the line $h_1 - h_2 = 0.5$ has both green observations on one side and both red ones on the other. An output perceptron with this line classifies all four correctly.

The hidden layer has changed how the data is represented, into a form where the classes are linearly separable. Finding such a form from the data is **representation learning** (G-1670), and it is what the hidden layers of a multi-layer perceptron do.

7. Summary

Data	Separable by a line?	Single perceptron or neuron
AND	yes	100%
OR	yes	100%
XOR (4 observations)	no	50%, all weights end at 0
Two blobs	yes	100%
XOR quadrants	no	52%
Circles	no	50%

- A perceptron can only have one straight decision boundary.
- XOR puts each class on opposite corners of a square: no line separates them, which a short proof confirms.
- More training never fixes this; more neurons, in layers, do.

8. Sources

Built from

- CampusX, “Problem with Perceptron”, YouTube, <https://www.youtube.com/watch?v=Jp44b27VnOg>
- StatQuest with Josh Starmer, “The Essential Main Ideas of Neural Networks”, YouTube, <https://www.youtube.com/watch?v=CqOfi41LfDw> (the three-dose example of section 4.1)

9. Key terms

Term	Meaning
AND, OR	Logic functions: 1 when both inputs are 1 (AND), or when at least one is (OR)
TensorFlow Playground	A website that trains small neural networks in the browser and shows their boundaries