


Handwritten Digit Classification (MNIST) with an ANN in Keras

Contents

1. Overview	2
2. The MNIST data	2
3. Scaling the pixels	3
4. The network	4
4.1 Flatten: from an image to a row	5
4.2 The architecture	6
4.3 Counting the parameters	7
5. Compiling and training	8
5.1 The loss for many classes	8
5.2 Training	8
6. Predicting with argmax	10
7. A bigger network and its training curves	11
7.1 The second network	11
7.2 Reading the curves	12
8. Where the network goes wrong	13
9. What the network really learned	14
9.1 The hidden nodes as pictures	14
9.2 Noise in, confident answer out	15
10. Predicting a single image	16
11. Summary	17
12. Sources	17
13. Key terms	17

Where this fits: Pipeline map steps 8 (Model), 9 (Evaluate). See the [Course map](#).



- **Builds on:** Overfitting ([Note ML-007](#)); Standardization ([Note ML-009](#)); Imbalanced data ([Note ML-009](#)); One-hot encoding ([Note ML-010](#)); Log loss (binary cross entropy) ([Note ML-072](#)); Confusion matrix ([Note ML-075](#)).
- **Leads to:** ANN for regression ([Note DL-013](#)); Vanishing gradient ([Note DL-018](#)); Early stopping ([Note DL-021](#)); Keras Tuner ([Note DL-039](#)); Image classification with a CNN (cats vs dogs) ([Note DL-049](#)); Keras functional API ([Note DL-054](#)).
- **Compare with:** K-nearest neighbours ([Note ML-085](#)); ANN for regression ([Note DL-013](#)).

1. Overview

Key point: The same Keras steps as for churn, now with 10 classes: a Flatten layer turns each image into 784 inputs, and an output layer of 10 softmax nodes gives one probability per digit.

A person reads a sloppy handwritten 3 at a glance, in any handwriting. Writing down rules that do the same, pixel by pixel, is very hard, so we let a network learn the rules from examples.

The [customer churn network](#) was built for **binary** classification: one output node, two classes. Here the task is **multi-class** classification: an image shows one handwritten digit, and the network must say which of the 10 digits, 0 to 9, it is.

The Keras workflow is unchanged: [build](#), [compile](#), [fit](#), [predict](#). Four things are new:

1. **Input:** an image, not a row of a table; a **Flatten** layer turns it into one row.
2. **Output layer:** 10 nodes with the **softmax** activation (a function that turns 10 scores into 10 probabilities that add up to 1; section 4.2).
3. **Loss:** sparse categorical cross-entropy (the loss for many classes, with labels kept as plain digits; section 5.1).
4. **Prediction:** the digit with the highest probability, found with **argmax** (G-212).

After training, section 9 opens the network up to see what it learned. The goal is to see how a network handles more than two classes, not to reach the best accuracy. The Notebook (DL-012-mnist-ann.ipynb) runs every step.

2. The MNIST data

Key point: 70,000 images of handwritten digits, 28×28 pixels each, with values 0 to 255. Keras ships them already split: 60,000 for training, 10,000 for testing.

MNIST (G-1249) is a set of 70,000 grey images of handwritten digits, each $28 \times 28 = 784$ pixels (see [pixels as columns](#) and [the MNIST data](#)). MNIST is so widely used that Keras includes it.

Python: Loading MNIST from Keras.

```
import keras

(X_train, y_train), (X_test, y_test) = \
    keras.datasets.mnist.load_data()
X_train.shape    # (60000, 28, 28)
X_test.shape     # (10000, 28, 28)
y_train[:3]      # [5 0 4]
load_data() downloads the data once and returns it already split into training and test
sets, so no train_test_split is needed.
```

`X_train` is a 3D array, a 3D [tensor](#) (a table of numbers with any number of axes; here 3 axes): 60,000 images, each a 28×28 grid of numbers. Each number is one pixel's brightness, from 0 (blank) to 255 (full ink), stored row by row. Each image is one **observation** (G-1374) (one record, one row of the data table), and each pixel is one **feature** (G-772) (an input variable, one column of the table): 784 features per image. `y_train` holds the **label** (G-1032) of each image, the **target** (G-1949) (the output we predict): the digit the image shows.




Figure 1: The first 10 training images and their labels

Figure 1 draws the first 10 training images, each as a grid of grey squares: dark where the pixel value is large (ink), light where it is 0. The title above each image is its label. The first three are a 5, a 0 and a 4, matching `y_train[:3]`.

Python: Drawing one image with Plotly.

```
import plotly.express as px

px.imshow(X_train[0], color_continuous_scale="gray_r")
```

3. Scaling the pixels

Key point: Divide every pixel by 255, the largest possible value, so that all inputs lie between 0 and 1.

As with the churn data, a network trains faster when its inputs share a small range. Pixels already share one range, 0 to 255, so we only shrink it:

1. **In words:** divide each pixel by the largest possible value, 255.
2. **Formula:**

$$x_{\text{scaled}} = \frac{x}{255}$$

3. **Example:** three pixels.

$$255/255 = 1$$

$$0/255 = 0$$

$$51/255 = 0.2$$

So a full-ink pixel (255) becomes 1, a blank pixel (0) stays 0, and a light-grey pixel (51) becomes 0.2.




Figure 2: The first training image, and the 8×8 patch in the red box as stored (0 to 255) and after dividing by 255 (0 to 1)

In Figure 2, compare the two grids cell by cell: every number keeps its place and its shading, only the scale changes, so 255 becomes 1.00 and 126 becomes 0.49. Dividing by 255 is [min-max scaling](#) with a known minimum of 0 and maximum of 255.

Python: Scaling the pixels.

```
X_train = X_train / 255
```

```
X_test = X_test / 255
```

4. The network

Key point: Flatten (784 inputs), a hidden layer of 128 ReLU nodes, an output layer of 10 softmax nodes: 101,770 trainable parameters.

4.1 Flatten: from an image to a row


Key point: A Flatten layer lays the 28 rows of pixels side by side, giving one row of 784 numbers.

A **Dense layer** (G-583; a layer in which every node is connected to every input of the layer) takes a flat list of numbers, but each image is a 28×28 grid. **Flatten** reshapes any multi-dimensional input into one dimension: first the 28 pixels of row 1, then the 28 of row 2, and so on to row 28. The count of numbers:


$$28 \times 28 = 784$$

Flatten does exactly the same as [turning each image into one table row](#) before PCA, but inside the network.

Flatten: 28 rows of 28 pixels become one row of 784



Flatten: 28 rows of 28 pixels become one row of 784



Flatten: 28 rows of 28 pixels become one row of 784




Figure 3: Flatten on the first training image: the 28 rows leave the image one after another and line up into one strip of 784 values. In the strip, each pixel is a thin bar with its grey level

In Figure 3, watch the orange box move down the image: each row it marks is appended to the right end of the strip, and the counter grows by 28. After the last row the strip holds all 784 values, and each value goes to its own input node. An input node does nothing more than hold that one number, the brightness of its pixel.

Flatten only rearranges numbers, so it has no weights and nothing to train.

4.2 The architecture

Key point: One input per pixel, 128 hidden nodes, one output node per digit.




Figure 4: The first MNIST network: the image is flattened into 784 inputs, which feed 128 ReLU nodes and then 10 softmax nodes, one per digit

Figure 4 shows the network:

- **Input layer:** 784 nodes, one per pixel.
- **Hidden layer:** 128 nodes with the **ReLU** activation (G-1668; a node outputs its weighted sum z if z is positive and 0 otherwise, so $\text{ReLU}(3) = 3$ and $\text{ReLU}(-2) = 0$), the usual default for hidden layers (Goodfellow et al. 2016, §6.3; see [ReLU](#) among the activation functions).
- **Output layer:** 10 nodes, one per class.

With more than two classes the output layer has **one node per class**, and the class whose node gives the highest probability is the prediction (see [more nodes in the output layer](#)). The

10 nodes use **softmax**, which turns their 10 scores into 10 probabilities that add up to 1 (see [the softmax function](#)). So for every image the network gives $P(0)$, $P(1)$, $\dots$, $P(9)$.

Python: Building the network.

```
from keras import layers
```

```
model = keras.Sequential([
    keras.Input(shape=(28, 28)),          # one image
    layers.Flatten(),                     # 28 x 28 -> 784
    layers.Dense(128, activation="relu"),  # hidden layer
    layers.Dense(10, activation="softmax"), # one node per digit
])
```

The input shape is now (28, 28), one image. The Dense layers work out their own number of inputs from the layer before.

4.3 Counting the parameters

Key point: 100,480 parameters into the hidden layer and 1,290 into the output layer: over 100,000 numbers to learn.

`model.summary()` lists:

Layer	Output shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100,480
dense_1 (Dense)	(None, 10)	1,290
Total		101,770

With the [counting rule for trainable parameters](#) (weights plus one bias per node):

Hidden layer (784 inputs, 128 nodes):

$$784 \times 128 = 100,352 \text{ weights}$$

$$100,352 + 128 = 100,480$$

Output layer (128 inputs, 10 nodes):

$$128 \times 10 = 1,280 \text{ weights}$$

$$1,280 + 10 = 1,290$$

The 784-128-10 network is the one previewed in the Extra box of that section. Over 100,000 weights and biases is already far more than the 276 of the churn network, though small by deep learning standards.

5. Compiling and training

Key point: Sparse categorical cross-entropy takes the labels as plain digits 0 to 9; categorical cross-entropy needs them one-hot encoded. Otherwise they are the same loss.

5.1 The loss for many classes

Key point: Both losses are categorical cross-entropy; “sparse” only means the labels stay as integers.

With a softmax output, the loss is **categorical cross-entropy** (G-349): the average of $-\log$ of the probability the network gives the true class (see [one loss for all classes](#) and [one perceptron, many models](#)). For one image of a 5, the loss is small when $P(5)$ is high and large when $P(5)$ is low:

$$P(5) = 0.9 : -\ln 0.9 = 0.105$$

$$P(5) = 0.1 : -\ln 0.1 = 2.30$$

Keras offers the loss in two forms, which differ only in how the labels are written:

Loss	Labels	Label of a “5”
categorical cross-entropy	one-hot vectors	[0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
sparse categorical cross-entropy	plain integers	5

Our labels are integers 0 to 9, so the sparse form saves a one-hot encoding step (see [how one-hot encoding works](#)). The loss values are identical (see [categorical cross-entropy](#) among the loss functions).

Python: Compiling, and the one-hot alternative.

```
model.compile(loss="sparse_categorical_crossentropy",  
              optimizer="adam")
```

only needed for categorical_crossentropy:

```
keras.utils.to_categorical(5, 10) # [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

The optimizer is again Adam (see [Adam](#) among the optimizers).

5.2 Training

Key point: 10 epochs with 20% of the training images held back for validation.

An **epoch** (G-696) is one pass of the network over all its training images; a **batch** (G-263) is the small group of images used for one weight update ([fit and epochs](#)). The **validation split** holds back part of the training data to measure the loss on images the network does not train on ([tracking accuracy and a validation set](#)).

Python: Training.

```
history1 = model.fit(X_train, y_train, epochs=10,  
                    validation_split=0.2)
```

Keras trains on 80% of the images and validates on the other 20%:

$$60,000 \times 0.8 = 48,000 \text{ training}$$

$$60,000 \times 0.2 = 12,000 \text{ validation}$$

`fit` was given no `batch_size`, so Keras uses its default of 32 images per batch (Keras docs, `Model.fit`). The number of weight updates per epoch:

$$48,000/32 = 1,500$$

The output shows this as 1500/1500. The training loss falls from 0.280 after the first epoch to 0.014 after the tenth. The validation loss is lowest after epoch 4 (0.088) and has risen to 0.099 by epoch 10: a first sign of overfitting.




Figure 5: Training and validation loss of the first network over its 10 epochs

Figure 5 plots the loss (the average error, lower is better) against the epoch number, one line for the training images and one for the validation images. Watch the gap open: the training loss keeps falling to 0.014, while the validation loss flattens after epoch 4 and then creeps up.

6. Predicting with argmax

Key point: `predict` gives 10 probabilities per image; the position of the largest one is the predicted digit.

`model.predict(X_test)` returns an array of shape (10000, 10): for each test image, the probability of each digit. For the first test image, $P(7)$ is almost exactly 1 and the other nine are almost 0.

To turn the 10 probabilities into one digit we take the position of the largest: **argmax**. Since position k holds $P(k)$, the position *is* the digit.

1. **In words:** the predicted class is the index of the largest probability.
2. **Formula:**

$$\hat{y} = \arg \max_k P(k), \quad k = 0, 1, \dots, 9$$

3. **Example:** if an image gets $P(1) = 0.03$, $P(2) = 0.90$, $P(7) = 0.05$ and 0.02 shared among the rest, the largest is at position 2, so $\hat{y} = 2$.




Figure 6: The worked example as a bar chart: one probability per digit, and `argmax` picks the position of the tallest bar

Figure 6 draws the example: `argmax` ignores how much bigger the tallest bar is and returns only its position, 2.

Python: From probabilities to digits.

```
from sklearn.metrics import accuracy_score
```

```
y_prob = model.predict(X_test)      # shape (10000, 10)
y_pred = y_prob.argmax(axis=1)      # [7 2 1 ...]
accuracy_score(y_test, y_pred)      # 0.9769
axis=1 takes the argmax across each row, one answer per image.
```

The first network gets **97.69%** of the 10,000 test images right. For comparison, KNN on all 784 pixels reached 96.8% in [KNN on all 784 features](#), on a different split of MNIST. KNN

compares each new image with every stored image (33,600 there), while the network only passes the image once through its two layers. A plain network with one hidden layer and no tuning already does better.

Extra: Image data is the home ground of the **convolutional neural network (CNN)** (G-484), which looks at small patches of pixels instead of treating every pixel as a separate input. CNNs reach over 99% on MNIST: LeNet-5 had a test error of 0.95% (LeCun et al. 1998). CNNs are taught in [what makes a network a CNN](#).

7. A bigger network and its training curves

Key point: An extra hidden layer and 25 epochs make the training accuracy reach 99.7%, but the test accuracy barely moves (97.72%): the network overfits.

7.1 The second network

Key point: 784-128-32-10, with accuracy tracked: 104,938 parameters.

To try to improve, we add a second hidden layer of 32 ReLU nodes, train for 25 epochs and track accuracy with `metrics=["accuracy"]` (see [tracking accuracy and a validation set](#)).

Parameters per layer (inputs times nodes, plus one bias per node):

$$784 \times 128 + 128 = 100,480$$

$$128 \times 32 + 32 = 4,128$$

$$32 \times 10 + 10 = 330$$

$$100,480 + 4,128 + 330 = 104,938$$

Python: The second network.

```
model2 = keras.Sequential([
    keras.Input(shape=(28, 28)),
    layers.Flatten(),
    layers.Dense(128, activation="relu"),
    layers.Dense(32, activation="relu"),
    layers.Dense(10, activation="softmax"),
])
model2.compile(loss="sparse_categorical_crossentropy",
               optimizer="adam", metrics=["accuracy"])
history = model2.fit(X_train, y_train, epochs=25,
                    validation_split=0.2)
```

The test accuracy is **97.72%**, almost the same as the first network's 97.69%. More capacity and more epochs did not help.

7.2 Reading the curves

Key point: The training loss keeps falling while the validation loss climbs after epoch 3: a clear case of overfitting.




Figure 7: Training curves of the second network. The circle marks the lowest validation loss, after epoch 3.

Figure 7 has two panels, both with the epoch on the horizontal axis. Left: the loss; right: the accuracy (the share of images classified correctly). Each has a training line and a validation line. The two panels explain the result:

- **Loss (left):** the training loss falls to about 0.01. The validation loss is lowest after epoch 3 (0.095), then rises, to 0.155 by epoch 25.
- **Accuracy (right):** training accuracy climbs to 99.7%, validation accuracy stays around 97 to 97.6%.

After epoch 3 the network keeps getting better on the images it trains on and worse on images it has not seen. The network is memorising the training set: **overfitting** (G-1429) (see [variance](#)). A student who memorises last year's answer sheet scores full marks on that sheet and still fails a new paper; the validation images are the new paper. The churn network showed a small gap; here it is large. [Regularization, dropout and stopping training at the right epoch](#) fix this.

Extra: If an accuracy curve is requested but `metrics=["accuracy"]` was not given to `compile`, `history.history` has no `accuracy` key and the plot fails with a `KeyError`. The model must be compiled again with the metric and retrained.

8. Where the network goes wrong

Key point: 228 of the 10,000 test images are misclassified; the most common mistake is a 4 read as a 9.




Figure 8: Confusion matrix of the second network on the 10,000 test images. Rows: actual digit; columns: predicted digit.

Figure 8 is the 10×10 confusion matrix (a table of counts; see [more than two classes](#)). Each row is the actual digit, each column the predicted digit, and each cell counts the test images with that pair, shaded darker for larger counts. The diagonal holds the correct predictions. The largest off-diagonal cells are the most common mistakes:

Actual	Predicted	Count
4	9	35
7	2	9
2	7	7
8	0	7

These pairs look alike when handwritten: a 4 with a closed top resembles a 9, a 2 with a straight base resembles a 7. Figure 9 shows five of the misclassified test images, each with its

actual and predicted digit; several are hard to read even for a person.

Extra: Are the confused digits close in raw pixels? A nearest-centroid classifier, which compares each test image only with the average image of each digit, makes the network's top mistake too: 4 read as 9 is the 2nd most common of its 79 different (actual, predicted) error pairs. But 7 read as 2 is only 30th and 2 read as 7 only 36th, so not every confusion comes from two digits with similar average images (Notebook).




Figure 9: Five test images the second network gets wrong

9. What the network really learned

Key point: High accuracy does not mean the network sees digits the way we do. Its hidden nodes do not look for clean strokes, and it gives a confident answer even to random noise.

9.1 The hidden nodes as pictures

Key point: Each hidden node has one weight per pixel, so its 784 weights can be drawn as a 28×28 picture of what the node responds to.

A natural hope is that each hidden node learns one piece of a digit, such as a stroke or a loop, and that the output layer combines the pieces. We can check this hope on the first network. Each of its 128 hidden nodes has 784 incoming weights, one per pixel. Drawn in the shape of the image, the weights show which pixels raise the node's sum (positive weight, blue) and which lower it (negative weight, red).




Figure 10: The incoming weights of the first 16 hidden nodes of the first network, each drawn as a 28×28 picture: blue is a positive weight, red a negative one. Idea after 3Blue1Brown, “Gradient descent, how neural networks learn”

Figure 10 draws each weight as one square, at the position of its pixel. In the figure, look for strokes and loops. A few nodes show a short streak, but most pictures are blotchy patches of blue and red with no clear shape. The network reaches 97.69% with detectors that a person cannot read.

9.2 Noise in, confident answer out

Key point: Fed random noise, the network still names a digit, with a probability near 1. Softmax has no way to answer “none of these”.

The second check feeds the network an image that shows no digit at all: every pixel is a random number between 0 and 1.




Figure 11: Left: an image of random noise. Right: the 10 probabilities the first network gives it

Figure 11 shows the noise image (left) and, on the right, one bar per digit giving the network's probability for it. The network says the noise is a 3, with probability 0.997. The noise image of Figure 11 is not a lucky pick: over 1,000 noise images the median top probability is 0.984, and 72% of the images get a top probability above 0.9. For real test images the median top probability rounds to 1.000, so the network is nearly as sure about noise as about real digits. It calls 771 of the 1,000 noise images a 3 and 193 a 5 (Notebook).

Two things cause this behaviour:

1. **Softmax must pick.** The 10 probabilities always add up to 1 (section 4.2), so the network cannot give a low score to every digit.
2. **The training data held only digits.** Every training image was a clear, centred digit, and the loss rewarded confident answers. Nothing taught the network what a non-digit looks like.

So a high test accuracy shows that the network separates the 10 digits well. It does not show that the network knows what a digit is. Networks built for images, the CNNs (see [what makes a network a CNN](#)), look at small patches of pixels and raise the accuracy, but they also end in a softmax layer, so they too must pick one of the 10 digits for a noise image.

10. Predicting a single image

Key point: `predict` expects a batch of images, so one image is reshaped from (28, 28) to (1, 28, 28) first.

Keras models always predict a batch of observations. A single image has shape (28, 28); `reshape(1, 28, 28)` turns it into a batch of one.

Python: Predicting one test image.

```
p = model2.predict(X_test[0].reshape(1, 28, 28))
p.argmax(axis=1)      # [7]: the first test image is a 7
```

The first two test images, a 7 and a 2, are both predicted correctly, each with a probability that rounds to 1.

11. Summary

	Churn (binary)	MNIST (multi-class)
Input	11 features, standardized	28×28 image, divided by 255, flattened
Output layer	1 node, sigmoid	10 nodes, softmax
Loss	binary cross-entropy	sparse categorical cross-entropy
Prediction	probability > 0.5	argmax of 10 probabilities
Test accuracy	86.45%	97.69% (first), 97.72% (second)

- MNIST: 60,000 training and 10,000 test images of 28×28 pixels, loaded with `keras.datasets.mnist.load_data()`.
- Divide pixels by 255; a Flatten layer turns each image into 784 inputs.
- Multi-class output: one softmax node per class; predict with `argmax(axis=1)`.
- Sparse categorical cross-entropy for integer labels, categorical cross-entropy for one-hot labels.
- A bigger network trained longer overfitted: training loss down, validation loss up after epoch 3.
- The hidden nodes' weights show no clear strokes, and random noise gets a confident answer (a 3 with probability 0.997): accuracy alone does not show what a network learned.

12. Sources

Built from

- CampusX, “Handwritten Digit Classification using ANN | MNIST Dataset”, YouTube, <https://www.youtube.com/watch?v=3xPT2Pk0Jds>
- 3Blue1Brown, “But what is a neural network? | Deep learning chapter 1”, YouTube, <https://www.youtube.com/watch?v=aircAruvnKk> (the task and the input nodes holding pixel values)
- 3Blue1Brown, “Gradient descent, how neural networks learn | Deep Learning Chapter 2”, YouTube, <https://www.youtube.com/watch?v=IHZwWFHwa-w> (section 9: the hidden weights as pictures, and the noise test)

Other references

- LeCun, Bottou, Bengio and Haffner, “Gradient-Based Learning Applied to Document Recognition”, *Proceedings of the IEEE*, 1998.
- Goodfellow, Bengio and Courville, *Deep Learning*, MIT Press, 2016, §6.3 (rectified linear units as the default hidden unit).

13. Key terms

Term	Meaning
MNIST	70,000 images of handwritten digits, 28×28 pixels, labelled 0 to 9
<code>keras.datasets.mnist</code>	Keras's built-in copy of MNIST, already split 60,000 / 10,000
Label	The target of one observation: its true class, here the digit an image shows
Flatten layer	A layer that reshapes a multi-dimensional input into one dimension; no parameters
Softmax output layer	An output layer with one node per class whose outputs are probabilities adding up to 1
Sparse categorical cross-entropy	Categorical cross-entropy for labels written as integers
<code>to_categorical</code>	Keras function that one-hot encodes integer labels
<code>argmax</code>	The position of the largest value; on 10 class probabilities, the predicted class
Convolutional neural network (CNN)	A network built for images that looks at small patches of pixels; taught later