

Loss Functions in Deep Learning

Contents

1. Overview	2
2. Prerequisites	3
3. Why the loss function matters	3
3.1 The loss in machine learning	3
3.2 The loss inside a neural network	4
4. Loss function versus cost function	5
5. Mean squared error	6
5.1 The formula	6
5.2 Far observations push harder	6
5.3 Advantages and disadvantages	7
5.4 MSE in a Keras network	8
6. Mean absolute error	8
7. Huber loss	9
7.1 When neither MSE nor MAE fits	9
7.2 The formula	10
7.3 The three losses on data with outliers	10
8. Binary cross-entropy	12
8.1 Why not the squared error?	13
8.2 Advantages and disadvantages	14
9. Categorical cross-entropy	15
9.1 The setup	15
9.2 Two students	17
10. Sparse categorical cross-entropy	17
11. Choosing a loss	18
12. Summary	20
13. Sources	20
14. Key terms	21

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).




- **Builds on:** One-hot encoding ([Note ML-010](#)); Softmax regression ([Note ML-078](#)); Maximum likelihood estimation (MLE) ([Note MA-070](#)); Forward propagation ([Note DL-010](#)).
- **Leads to:** Backpropagation ([Note DL-015](#)).
- **Compare with:** Regression metrics ([Note ML-051](#)); Hinge loss and soft margin ([Note ML-088](#)); Entropy, information gain and Gini ([Note ML-091](#)).

1. Overview

Key point: A loss function scores how wrong the network is on one **observation** (one record, one row of the data table). Training changes the weights and biases until that score is as small as possible. Which loss we use depends on the problem, and it fixes the activation of the output layer.

A [loss function](#) measures how wrong a model's predictions are. The loss is a function of the model's parameters: change any weight or bias, and the loss changes (see the [convex and non-convex cost functions Note](#), section 2). A large loss means the model is doing badly; a small one means it is doing well.



dashed: the losses of this Note

Figure 1: The main loss functions of deep learning, grouped by the problem they solve

Figure 1 groups the loss functions by problem. This Note covers the dashed part: three losses

for regression and three for classification, with the output layer and Keras setting each one needs.

2. Prerequisites

- [Forward propagation](#): how a network turns an observation into a prediction.
- MSE and MAE as metrics: the [regression metrics Note](#).
- Binary cross-entropy: the [log loss Note](#). Categorical cross-entropy and softmax: the [softmax regression Note](#).

3. Why the loss function matters

Key point: We cannot improve what we cannot measure. The loss is the number training watches: every weight update is decided by it.

3.1 The loss in machine learning

Key point: Start with a random line, measure its loss, move the line, measure again; stop when the loss is smallest.

In simple linear regression, the loss is the mean squared error, and $\hat{y}_i = mx_i + b$. So the loss is a function of m and b only. Gradient descent starts from a random line, computes its loss, and changes m and b to lower it, again and again (see the [gradient descent Note](#)).

The loss is the algorithm's eyes. Without it, the algorithm has no way to know which direction is better.




Figure 2: Gradient descent on the 30 points of Figure 6 that follow $y = 2x + 1$: the line moves (left) while its MSE loss falls (right)

Figure 2 runs this loop from $m = 0$, $b = 0$ with learning rate 0.01. Watch the loss fall from 156 to 1.45 in two steps, then creep down to 0.12 by step 500 as the line settles on the points.

3.2 The loss inside a neural network

Key point: Each observation goes forward through the network, its loss is computed, and the weights and biases are adjusted to lower that loss. Repeat for every observation, for many epochs.

A network is trained in the same spirit. Take the student data (CGPA, IQ, package in LPA):

1. Start with random weights and biases.
2. Feed the first student forward (7.1 and 83 into the input nodes) and get a prediction, say $\hat{y} = 3.7$.
3. Put y and $\hat{y}$ into the loss function, here the mean squared error, because the output is a number.
4. Use that loss to adjust every weight and bias with gradient descent.
5. Take the next student and repeat; go over the whole dataset many times (epochs).

The weights and biases that give the smallest loss are the trained network. How step 4 is done, for every weight at once, is **backpropagation**, taught in the [backpropagation Note](#).

4. Loss function versus cost function

Key point: The loss is computed on one observation; the cost is the average loss over a batch or the whole training set. Many people use the two words for the same thing.

Strictly, the two words mean different things:

- **Loss function** (G-706, also called **error function**): the error on a single training observation, e.g. $(y_i - \hat{y}_i)^2$.
- **Cost function** (G-492): the average of the losses over a batch of observations, or over the whole training set, e.g. $\frac{1}{n} \sum (y_i - \hat{y}_i)^2$.

Four students show the difference:

Student	y (LPA)	$\hat{y}$	$y - \hat{y}$	Loss $(y - \hat{y})^2$
1	6.3	6.1	0.2	0.04
2	4.1	4.0	0.1	0.01
3	3.5	3.7	-0.2	0.04
4	7.2	7.0	0.2	0.04

1. **In words:** compute the loss of each observation, then average them.
2. **Formula:**

$$J = \frac{1}{n} \sum_{i=1}^n L_i$$

3. **Example:**

$$J = \frac{0.04 + 0.01 + 0.04 + 0.04}{4} = 0.0325$$




Figure 3: Each bar is one student's loss; the dashed line is the cost, their average

Figure 3 shows the difference at a glance: four losses, one per student, and one cost line through their middle.

Extra: Other Notes, such as the [convex and non-convex cost functions Note](#), use “cost function” loosely as another name for the loss. Libraries also blur the line: by default Keras averages the losses over the batch and reports that average as the loss (Keras docs, Losses).

5. Mean squared error

Key point: MSE squares each error: far-off observations get much bigger weight updates. Use it for regression without outliers, with a linear output node.

5.1 The formula

Key point: Loss $(y - \hat{y})^2$; cost $\frac{1}{n} \sum (y_i - \hat{y}_i)^2$. Also called squared loss or L2 loss.

The **mean squared error** (MSE, G-1201) is defined in the [regression metrics Note](#), section 3. As a training loss it is also called the **squared loss** or **L2 loss**.

The square is there so that errors cannot cancel. In the table of Section 4, the raw errors add up to $0.2 + 0.1 - 0.2 + 0.2 = 0.3$: student 3’s negative error hides part of the others. Squaring makes every error count as positive.

5.2 Far observations push harder

Key point: The slope of the loss is $-2(y - \hat{y})$. An observation with an error of 42 pushes the weights 42 times harder than an observation with an error of 1.

Squaring has a second effect: errors grow quadratically. An error of 1 costs 1, an error of 2 costs 4, an error of 4 costs 16.

What moves the weights is the slope of the loss, written $\partial L / \partial \hat{y}$ (read it as: how much L changes when $\hat{y}$ alone changes a little; backpropagation multiplies every weight’s update by it):

1. **In words:** for MSE, the slope is -2 times the error, so the push grows in proportion to the error.
2. **Formula:**

$$\frac{\partial}{\partial \hat{y}} (y - \hat{y})^2 = -2(y - \hat{y})$$

3. **Example:** a student with package 8 predicted as 7 has slope $-2 \times 1 = -2$. A student with a package of 50 lakhs predicted as 8 has loss $42^2 = 1764$ and slope $-2 \times 42 = -84$: a 42 times bigger push.




Figure 4: Left: the loss of one observation against its error. Right: its slope, which sets how hard the observation pushes the weights

Figure 4 (blue) shows both: the loss is a parabola, and its slope a straight line through 0. Observations far from the prediction drive the weight updates.

5.3 Advantages and disadvantages

Key point: Easy to read and smooth everywhere, but in squared units and pulled hard by outliers.

Advantages:

- **Easy to interpret:** the formula says exactly how the error is measured.
- **Differentiable everywhere:** the parabola has no corners, so gradient descent can always compute a slope.
- **One minimum** for a linear model such as linear regression (see section 5.1 of the [convex and non-convex cost functions Note](#)).

Disadvantages:

- **Squared units:** a loss of 5 means 5 LPA², not 5 LPA. Its square root, the RMSE, is back in LPA.
- **Not robust to outliers:** the 50-lakh student of Section 5.2 dominates the updates, and the network bends towards it, away from the other students (Figure 6 in Section 7).

Extra: “One minimum” holds for a linear model only. As a function of a network’s weights, even the MSE has several minima, because the network sits between the loss and the weights (see section 5.2 of the [convex and non-convex cost functions Note](#)). This is true for every loss in this Note.

5.4 MSE in a Keras network

Key point: The output node must be linear, so it can output any number. Hidden layers can use any activation.

With MSE, the output layer needs a linear activation (no activation at all), because a package can be any number. A sigmoid output would trap every prediction between 0 and 1. The hidden layers can use any activation, such as ReLU or sigmoid.

Python: A regression network with MSE.

```
import keras

model = keras.Sequential([
    keras.Input(shape=(2,)),
    keras.layers.Dense(3, activation="relu"),
    keras.layers.Dense(1, activation="linear"), # any number
])
model.compile(optimizer="sgd", loss="mse")
loss="mse" is short for "mean_squared_error". activation="linear" is the default
of Dense, so it can also be left out.
```

6. Mean absolute error

Key point: MAE takes the absolute error instead of the square. Its units are those of y and outliers pull less, but its slope jumps at 0.

The **mean absolute error** (MAE, G-1194) is defined in the [regression metrics Note](#), section 2. As a training loss it is also called **L1 loss**: loss $|y_i - \hat{y}_i|$ for one observation, cost $\frac{1}{n} \sum |y_i - \hat{y}_i|$ over n observations.

The only change from MSE is the absolute value in place of the square. The MAE's slope (Figure 4, orange) is -1 or $+1$, whatever the size of the error. The 50-lakh student now pushes exactly as hard as a student whose prediction is off by 1.




Figure 5: One point is lifted from the pattern to 42 above it; the line is refitted with MSE (blue) and MAE (orange) each time

Figure 5 drags one point away from the 30 points of Figure 6 that follow $y = 2x + 1$. Watch the MSE line tilt towards the point, rising by 3.7 at $x = 9$, while the MAE line does not move at all.

Advantages:

- **Intuitive:** the average size of the error.
- **Same units as y :** an MAE of 2 means the predictions are off by 2 LPA on average.
- **Robust to outliers:** no square, so far observations do not dominate.

Disadvantage: the loss has a corner at an error of 0, where it has no derivative, and its slope jumps from -1 to $+1$ there. Gradient descent then needs a special rule at that point (a **subgradient**, see section 7.1 of the [perceptron loss Note](#)), which makes the computation more awkward.

So: no outliers, use MSE; outliers in the data, use MAE. In Keras, `loss="mae"`, again with a linear output node.

7. Huber loss

Key point: Huber loss is MSE for small errors and MAE for large ones. The threshold δ decides which errors count as outliers.

7.1 When neither MSE nor MAE fits

Key point: With many outliers, MSE follows them and MAE ignores them; Huber loss sits in between.

Suppose 25% of the points are **outliers** (G-1420): points far from the pattern the rest follow. So many outliers are no longer rare mistakes: they are part of the data. MSE gives them a huge pull; MAE treats them like any other point. **Huber loss** (G-905; Huber 1964) combines the two:

- if an observation's error is small ($|y - \hat{y}| \leq \delta$), it behaves like MSE;
- if the error is large, it behaves like MAE.

7.2 The formula

Key point: $\frac{1}{2}e^2$ up to $|e| = \delta$, then a straight line $\delta(|e| - \frac{1}{2}\delta)$ that continues it smoothly.

1. **In words:** with error $e = y - \hat{y}$, use half the squared error while the error is within δ ; beyond it, use a straight line that rises by δ per unit of error.
2. **Formula:**

$$L_{\delta}(e) = \begin{cases} \frac{1}{2}e^2 & |e| \leq \delta \\ \delta(|e| - \frac{1}{2}\delta) & |e| > \delta \end{cases}$$

3. **Example:** with $\delta = 1$, an error of 0.5 costs $\frac{1}{2} \times 0.5^2 = 0.125$; an error of 3 costs $1 \times (3 - 0.5) = 2.5$; the 50-lakh student's error of 42 costs $1 \times (42 - 0.5) = 41.5$, not 1764.

The two pieces meet at $|e| = \delta$ with the same value and the same slope, so the loss is smooth everywhere (Figure 4, green). The Huber slope is the error itself for small errors and $\pm\delta$ beyond, so no observation can push harder than δ .

δ is a hyperparameter: we try several values and keep the one that works best. In Keras: `loss=keras.losses.Huber(delta=1.0)`, which uses this same formula (Keras docs, Huber).

Extra: The $\frac{1}{2}$ only makes the slope come out as e instead of $2e$. The factor changes the size of every update by the same factor, so it plays the same role as the learning rate.

7.3 The three losses on data with outliers

Key point: On data with 25% outliers, MSE lifts the line by about 2; MAE and Huber stay on the main points.




Figure 6: A line fitted with each loss to 40 points, 10 of them outliers

Figure 6 fits a straight line to 40 points: 30 follow $y = 2x + 1$, and 10 (25%) sit 8 units above it. Each loss gives a different line:

Loss	Fitted line
MSE	$\hat{y} = 1.88x + 3.55$
MAE	$\hat{y} = 1.96x + 1.31$
Huber, $\delta = 1$	$\hat{y} = 1.95x + 1.48$

MSE has moved the line up by about $0.25 \times 8 = 2$, between the two groups. MAE and Huber stay with the 75% of points that follow the pattern. With $\delta = 1$, Huber treats every outlier error of about 8 like MAE, so its line is close to MAE's. A larger δ moves it towards MSE's.




Figure 7: The Huber line (green) refitted on the 40 points of Figure 6 as δ grows from 0.25 to 16, with the MSE and MAE lines dashed

In Figure 7, watch the green line leave the orange MAE line and climb to the blue MSE line as δ grows: $\hat{y} = 1.96x + 1.29$ at $\delta = 0.25$, $1.89x + 2.78$ at $\delta = 4$, and $1.88x + 3.55$, exactly the MSE line, from $\delta = 6$ on. Once δ is larger than every error, no observation is treated as an outlier, and Huber loss is the squared error.

Extra: Here Huber does not land halfway between MSE and MAE: it behaves like MSE only for errors below δ . Huber's real gain is that it is smooth (no corner at 0, unlike MAE) while still resisting outliers.

8. Binary cross-entropy

Key point: For two classes. One sigmoid output node; loss $-y \log \hat{y} - (1 - y) \log(1 - \hat{y})$.

When the output is a probability, the loss should be near 0 if the network gives the true class a high probability, and very large if it gives the true class a probability near 0. Minus the log of that probability behaves this way: $-\log 1 = 0$, $-\log 0.5 = 0.69$ and $-\log 0.01 = 4.6$.

Binary cross-entropy (G-303), also called **log loss**, is this loss for two classes, and the loss of logistic regression. The formula and where it comes from (maximum likelihood) are in the [log loss Note](#):

$$L = -y \log \hat{y} - (1 - y) \log(1 - \hat{y}), \quad J = -\frac{1}{n} \sum_{i=1}^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

We use it only when the output has two classes, such as placed (1) or not placed (0). In a network:

- the **output layer has one node with a sigmoid activation**, so $\hat{y}$ is a probability between 0 and 1;
- the hidden layers can use any activation.

Two students, with natural logs:

1. Student 1 is placed ($y = 1$) and the network gives $\hat{y} = 0.73$. Only the first term is on: $L = -\log 0.73 = 0.315$.
2. Student 2 is not placed ($y = 0$) and $\hat{y} = 0.25$. Only the second term is on: $L = -\log(1 - 0.25) = -\log 0.75 = 0.288$.




Figure 8: The loss for each label as the predicted probability changes, with the two students marked

Figure 8 draws the two halves of the formula. Watch each curve shoot up as the prediction moves towards the wrong label: a confident wrong answer costs far more than a hesitant one.

After each loss, gradient descent updates the weights, then the next student comes.

8.1 Why not the squared error?

Key point: For a badly wrong probability, the slope of cross-entropy is huge and the slope of the squared error is small. The slope is the push on the weights, so cross-entropy corrects bad predictions much faster.

The squared error $(1 - p)^2$ could also score a probability p given to the true class: it is 0 when $p = 1$ and largest when $p = 0$. The difference lies in the slope, which sets how hard the observation pushes the weights (section 5.2).




Figure 9: The probability p given to the true class slides from 0.99 to 0.01. Red: cross-entropy and its tangent. Blue: the squared error and its tangent. The title prints both slopes. Idea after StatQuest, “Neural Networks Part 6: Cross Entropy”

In Figure 9, watch the two dotted tangent lines as p falls:

1. **In words:** the slope of $-\log p$ is $-1/p$, and the slope of $(1-p)^2$ is $-2(1-p)$.
2. **Formula:**

$$\frac{d}{dp}(-\log p) = -\frac{1}{p}, \quad \frac{d}{dp}(1-p)^2 = -2(1-p)$$
3. **Example:** for student 1, $p = 0.73$: the slopes are $-1/0.73 = -1.37$ and $-2 \times 0.27 = -0.54$. For a badly wrong prediction, $p = 0.01$: the slopes are $-1/0.01 = -100$ and $-2 \times 0.99 = -1.98$.

The squared error’s slope can never be steeper than -2 , however wrong the prediction. The cross-entropy’s slope keeps growing as p approaches 0. So with cross-entropy, the worse the prediction, the larger the correction.

8.2 Advantages and disadvantages

- **Advantage:** it is differentiable, so gradient descent applies directly.
- **Disadvantage:** it is less intuitive than MSE or MAE: the number has no everyday unit.

Extra: The log here is the natural log, as in Keras.

Python: A binary classifier.

```
model = keras.Sequential([
    keras.Input(shape=(2,)),
    keras.layers.Dense(3, activation="relu"),
    keras.layers.Dense(1, activation="sigmoid"),
])
model.compile(optimizer="sgd", loss="binary_crossentropy")
```

9. Categorical cross-entropy

Key point: For more than two classes. One softmax output node per class; one-hot labels; only the log probability of the true class counts.

9.1 The setup

Key point: k classes means k output nodes with softmax, and one-hot encoded labels.

Suppose the output has three classes: placed *yes*, *no* or *maybe*. **Categorical cross-entropy** (G-349) is the loss of softmax regression (see section 4.2 of the [softmax regression Note](#)). For one observation with k classes:

The symbol $\sum_{j=1}^k$ means “add the terms for $j = 1, 2, \dots, k$ ”; here $k = 3$, so it adds three terms:

$$L = - \sum_{j=1}^k y_j \log \hat{y}_j = -y_1 \log \hat{y}_1 - y_2 \log \hat{y}_2 - y_3 \log \hat{y}_3$$

and the cost over n observations is $J = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k y_{ij} \log \hat{y}_{ij}$.

The network changes in two ways:

- the **output layer has one node per class** (here 3: yes, no, maybe);
- they all use the **softmax** activation, $\hat{y}_j = e^{z_j} / (e^{z_1} + e^{z_2} + e^{z_3})$, so the three outputs lie between 0 and 1 and add up to 1 (see the [softmax regression Note](#), section 2).

Why softmax, and not simply the largest raw output? Take three raw outputs of the output layer, 1.43, -0.4 and 0.23 . They are hard to read: they are not between 0 and 1 and do not add up to 1. Two ways to tidy them are compared in Figure 10.




Figure 10: Three raw outputs (left), after argmax (middle) and after softmax (right), while the raw output of class 1 slides from 1.43 down to -1 . Idea after StatQuest, “Neural Networks Part 5: ArgMax and SoftMax”

Softmax on those three raw outputs, one step per line:

$$\begin{aligned}
 e^{1.43} &= 4.18, & e^{-0.4} &= 0.67, & e^{0.23} &= 1.26 \\
 4.18 + 0.67 + 1.26 &= 6.11 \\
 \hat{y} &= \left(\frac{4.18}{6.11}, \frac{0.67}{6.11}, \frac{1.26}{6.11} \right) = (0.68, 0.11, 0.21)
 \end{aligned}$$

- **Argmax** sets the largest output to 1 and the rest to 0: here $(1, 0, 0)$. The result is easy to read. But in Figure 10 the argmax bars stand still while the raw output of class 1 falls from 1.43 all the way to 0.26, and then jump at once when it drops below class 3’s 0.23. A small change in a weight almost never changes the argmax output, so its slope is 0 and gradient descent gets no direction to move in.
- **Softmax** gives $(0.68, 0.11, 0.21)$: the same order, every value between 0 and 1, and a sum of 1. In Figure 10 the softmax bars move at every step, so every small change in a weight changes the loss, and gradient descent has a slope to follow.

So training uses softmax, and argmax is used only afterwards, to read off the predicted class (see section 6 of the [MNIST Note](#)).

The softmax outputs are called probabilities, but they are not a reliable measure of how sure the network should be: the MNIST network of that Note gives random noise the label 3 with a softmax output of 0.997 (its section 9.2).

The labels are [one-hot encoded](#): yes = $(1, 0, 0)$, no = $(0, 1, 0)$, maybe = $(0, 0, 1)$. A handwritten-digit network is the same idea with 10 classes (0 to 9) and 10 softmax output nodes.

9.2 Two students

Key point: The one-hot label switches off every term except the true class.

1. Student 1 (CGPA 8, IQ 80) is a *yes*: $y = (1, 0, 0)$. The network outputs $\hat{y} = (0.2, 0.3, 0.5)$. Only y_1 is 1:

$$L = -1 \times \log 0.2 = 1.609$$

2. Student 2 (6, 60) is a *no*: $y = (0, 1, 0)$, and $\hat{y} = (0.3, 0.6, 0.1)$:

$$L = -1 \times \log 0.6 = 0.511$$

Student 1 gets a large loss because the network gave its true class only 0.2. After each loss, gradient descent updates the weights.




Figure 11: Softmax outputs of the two students; only the true class (green) enters the loss

In Figure 11, the grey bars play no part in the loss: only the green bar's height matters.

10. Sparse categorical cross-entropy

Key point: The same loss as categorical cross-entropy, but the labels stay integers. No one-hot encoding is needed, and with many classes it saves memory and time.

Sparse categorical cross-entropy (G-1839) uses exactly the same network (one softmax node per class) and the same loss. Only the labels differ: *yes*, *no*, *maybe* become the integers 0, 1, 2 instead of one-hot vectors.


1. **In words:** read the true class from the integer label, pick that class's predicted probability, and take minus its log.
2. **Formula:** for an observation with label c ,

$$L = -\log \hat{y}_c$$

3. **Example:** student 1 has label 0 and $\hat{y} = (0.2, 0.3, 0.5)$, so $L = -\log 0.2 = 1.609$; student 2 has label 1, so $L = -\log 0.6 = 0.511$. These are exactly the numbers of Section 9.2.

The one-hot version multiplies the logs of all k classes by y_j , most of them by 0. The sparse version goes straight to one log (Figure 12).


Categorical: one-hot label



$$L = -(1 \cdot \log 0.2 + 0 \cdot \log 0.3 + 0 \cdot \log 0.5) = 1.609$$

k products, $k - 1$ of them multiplied by 0

Sparse: integer label



$$L = -\log \hat{y}_c = -\log 0.2 = 1.609$$

one lookup, no one-hot vector stored

Figure 12: Student 1 under both label formats: three products, two of them zero, against one lookup

With 10 classes the difference is small; with thousands of classes (words in a vocabulary, for instance) it saves both the one-hot matrix and the wasted work.

Python: The two multi-class set-ups.

```
model = keras.Sequential([
    keras.Input(shape=(2,)),
    keras.layers.Dense(3, activation="relu"),
    keras.layers.Dense(3, activation="softmax"), # one per class
])
# labels one-hot: [[1, 0, 0], [0, 1, 0], ...]
model.compile(optimizer="sgd", loss="categorical_crossentropy")
# labels as integers: [0, 1, ...]
model.compile(optimizer="sgd",
               loss="sparse_categorical_crossentropy")
```

Keras counts classes from 0. On the two students, both losses give the same average, 1.060. The Notebook checks every number of this Note with `keras.losses`.

11. Choosing a loss

Key point: Regression: MSE, or MAE with outliers, or Huber with many. Two classes: binary cross-entropy. More classes: categorical cross-entropy, or its sparse version with many classes.

Loss functions are not limited to these. Other problems have their own losses (Figure 1):

- variational autoencoders: the KL divergence (Kingma and Welling 2014);
- GANs: a discriminator (min-max) loss (Goodfellow et al. 2014);
- object detection: the focal loss (Lin et al. 2017);
- embeddings: the triplet loss (Schroff et al. 2015);
- classification, as an alternative: the SVM's [hinge loss](#). Researchers design new losses for new problems, and Keras accepts our own loss functions too. Each loss has its strengths, and a poorly chosen one keeps the network from its best solution.

12. Summary

Problem	Output layer	Loss
Regression	1 node, linear	MSE
Regression with outliers	1 node, linear	MAE
Regression, many outliers	1 node, linear	Huber
Two classes	1 node, sigmoid	binary cross-entropy
k classes, one-hot labels	k nodes, softmax	categorical cross-entropy
k classes, integer labels	k nodes, softmax	sparse categorical cross-entropy

- The loss scores one observation; the cost averages the losses over a batch or the training set.
- MSE's slope grows with the error, so far observations drive the updates; MAE's slope is always ± 1 .
- Huber loss is MSE inside $\pm\delta$ and MAE outside, and smooth everywhere.
- Cross-entropy's slope $-1/p$ grows without limit for a badly wrong probability; the squared error's slope stays within -2 .
- Argmax has slope 0, so training uses softmax; argmax only reads off the final class.
- The loss decides the output layer: linear for MSE, MAE and Huber, sigmoid for binary cross-entropy, softmax for the categorical ones.
- Sparse categorical cross-entropy is the same loss with integer labels.
- Keras names:
 - regression: `"mse"`, `"mae"`, `keras.losses.Huber(delta)`;
 - two classes: `"binary_crossentropy"`;
 - more classes: `"categorical_crossentropy"`;
 - more classes, integer labels: `"sparse_categorical_crossentropy"`.

13. Sources

Built from

- CampusX, “Loss Functions in Deep Learning | Deep Learning | CampusX”, YouTube, https://www.youtube.com/watch?v=gb5nm_3jBIo
- StatQuest with Josh Starmer, “Neural Networks Part 5: ArgMax and SoftMax”, YouTube, <https://www.youtube.com/watch?v=KpKog-L9veg> (section 9.1: why softmax and not argmax)
- StatQuest with Josh Starmer, “Neural Networks Part 6: Cross Entropy”, YouTube, <https://www.youtube.com/watch?v=6ArSys5qHAU> (section 8.1: cross-entropy against the squared error)

Other references

- Keras documentation, Losses (default reduction `"sum_over_batch_size"`: the average over the batch); `keras.losses.Huber`.

- Huber, P. J., “Robust Estimation of a Location Parameter”, *Annals of Mathematical Statistics*, 1964.
- Kingma and Welling, “Auto-Encoding Variational Bayes”, ICLR 2014.
- Goodfellow et al., “Generative Adversarial Nets”, NeurIPS 2014.
- Lin, Goyal, Girshick, He and Dollár, “Focal Loss for Dense Object Detection”, ICCV 2017.
- Schroff, Kalenichenko and Philbin, “FaceNet: A Unified Embedding for Face Recognition and Clustering”, CVPR 2015.

14. Key terms

Term	Meaning
Loss function (error function)	The error of the model on one training observation
Cost function	The average loss over a batch or the whole training set
Squared loss (L2 loss)	Another name for the mean squared error used as a training loss
L1 loss	Another name for the mean absolute error used as a training loss
Huber loss	Half the squared error for errors up to δ , a straight line beyond; MSE for small errors, MAE for large ones
δ (Huber)	The error size where Huber loss switches from squared to absolute; a hyperparameter
Categorical cross-entropy (G-349)	$-\sum_j y_j \log \hat{y}_j$ with one-hot labels; the loss for more than two classes, with a softmax output
Sparse categorical cross-entropy	Categorical cross-entropy with integer labels: $-\log \hat{y}_c$ for the true class c
KL divergence, focal loss, triplet loss	Losses for variational autoencoders, object detection and embeddings, taught with those networks