


ReLU Variants: Leaky ReLU, PReLU, ELU and SELU

Contents

1. Overview	2
2. Prerequisites	3
3. The dying ReLU problem	3
3.1 Why a dead node stops learning	3
3.2 What makes z negative	5
3.3 Why a dead node stays dead	5
3.4 Dead nodes in practice	5
4. Three ways to prevent dead nodes	7
5. Linear variants	8
5.1 Leaky ReLU	8
5.2 Parametric ReLU (PReLU)	9
6. Non-linear variants	10
6.1 ELU	10
6.2 SELU	11
6.3 GELU and SiLU: smooth versions of ReLU	13
7. Summary	15
8. Sources	15
9. Key terms	16

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- Builds on: [Vanishing gradient](#); [Activation functions](#).

1. Overview

Key point: A ReLU node whose weighted sum stays negative outputs 0 with slope 0, so it never learns again: it is dead. Four variants give negative inputs a non-zero output and slope, which keeps every node alive.

ReLU (G-1668) is the default **activation function** (G-165) for **hidden layers** (G-890; see the [activation functions](#)). ReLU's biggest weakness is the **dying ReLU problem** (G-650). This Note covers:

- what a dead node is, why it happens and why it is permanent;
- three fixes, the third being a change of activation;
- two **linear variants** (G-1098; Leaky ReLU, Parametric ReLU), which change ReLU's negative side by a straight line;
- two **non-linear variants** (G-1336; ELU, SELU), which use an exponential curve there.
- two smooth versions of ReLU used in transformers (GELU, SiLU).




Figure 1: ReLU and its four variants (left) and their derivatives (right). On the negative side ReLU's slope is 0; every variant keeps a slope above 0. The Leaky ReLU slope is drawn as 0.1 instead of 0.01 to make it visible

Figure 1 shows all five functions. The left panel plots each function (output against the input z); the right panel plots each function's slope (its derivative, how steeply the left curve rises

or falls at that z). The five agree for positive z (SELU is scaled up slightly) and differ only for negative z .

2. Prerequisites

- The [activation functions](#): ReLU, saturation, zero-centred outputs.
- The [backpropagation how](#): gradients as chain-rule products.

3. The dying ReLU problem

Key point: A dead node outputs 0 for every input. With more than half the nodes dead, the network cannot capture the patterns in the data; with all of them dead, there is no network left.

A **dead neuron** (G-551) is a node whose output is 0 for every input. The dead node's output no longer depends on the input, so it carries no information and learns nothing. Worse, it stays dead for the rest of training: in effect it has been removed from the network.

How much this matters depends on how many nodes die:


- **A few dead nodes:** the network has less **capacity** (G-344; the range of functions it can fit), but it still works.
- **More than half:** the network runs at less than half its size and cannot represent the patterns in the data well.
- **All of them:** the network outputs a constant. There is no point in training it.

3.1 Why a dead node stops learning

Key point: If $z_1 < 0$, then $a_1 = 0$ and the slope $\partial a_1 / \partial z_1$ is 0. The zero slope appears in the gradient of every weight into the node, so none of them is updated.

Take a small network for regression: two inputs, one hidden ReLU node, one output node (Figure 2). The hidden node computes

$$\begin{aligned}z_1 &= w_1 x_1 + w_2 x_2 + b_1 \\a_1 &= \max(0, z_1)\end{aligned}$$



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial a_1} \cdot \frac{\partial a_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial w_1}$$

= 0 when $z_1 < 0$, so $\partial L / \partial w_1 = \partial L / \partial w_2 = 0$:
 w_1 and w_2 never change, and z_1 stays negative

Figure 2: A ReLU node with $z_1 < 0$. The factor $\partial a_1 / \partial z_1$ is 0, so the gradients of w_1 and w_2 are 0

Here L is the loss and $\hat{y}$ is the network's prediction. **Backpropagation** (G-247) reaches w_1 through the chain $L \rightarrow \hat{y} \rightarrow a_1 \rightarrow z_1 \rightarrow w_1$:

1. **In words:** the gradient of w_1 is a product of four factors, one of which is ReLU's slope.
2. **Formula:**

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial a_1} \cdot \frac{\partial a_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial w_1}$$

The gradient of w_2 has the same first three factors.

3. **Example:** with $w_1 = -0.8$, $w_2 = -0.5$, $b_1 = 0.1$ and inputs $x_1 = 0.6$, $x_2 = 0.4$:

$$w_1 x_1 = -0.8 \times 0.6 = -0.48$$

$$w_2 x_2 = -0.5 \times 0.4 = -0.20$$

$$z_1 = -0.48 - 0.20 + 0.1 = -0.58 < 0$$

$$\frac{\partial a_1}{\partial z_1} = 0$$

The last factor is the input $x_1 = 0.6$:

$$\frac{\partial z_1}{\partial w_1} = x_1 = 0.6$$

The product for w_1 , factor by factor:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial a_1} \cdot 0 \cdot 0.6 = 0$$

So both gradients are 0:

$$\frac{\partial L}{\partial w_1} = 0$$
$$\frac{\partial L}{\partial w_2} = 0$$

The update changes nothing:

$$w_{\text{new}} = w_{\text{old}} - \eta \cdot 0$$
$$w_{\text{new}} = w_{\text{old}}$$

Neither weight changes. If z_1 is negative for every **observation** (G-1374; every record of the data), the node never gets an update again: it is dead.

3.2 What makes z negative

Key point: Two causes: a learning rate so high that one update throws the weights negative, and a large negative bias.

1. **A high learning rate** (G-1068). Suppose that for the first observation z_1 is positive, so the gradient is not 0. With a very large learning rate, $\eta \partial L / \partial w$ is a big number. Subtracting it from a small weight makes w_1 and w_2 strongly negative, and in the next round z_1 is negative for every observation.

2. **A large negative bias.** If b_1 is very negative, $z_1 = w_1 x_1 + w_2 x_2 + b_1$ is negative however the inputs vary. The bias can start out negative, or it can be pushed there by updates, again usually because the learning rate is high.

3.3 Why a dead node stays dead

Key point: The weights and bias no longer change, and the scaled inputs are too small to bring z back above 0.

Once z_1 is negative for every observation, nothing can make z_1 positive again:

- the weights and the bias get no updates, so they stay as they are;
- only the **features** (G-772) x_1, x_2 (the input variables) change from observation to observation, and they are scaled into a small range, so they cannot outweigh a large negative bias or negative weights.

For these two reasons a dead node is called permanently dead.

3.4 Dead nodes in practice

Key point: On the moons data, a learning rate of 10 kills 41% and 69% of the nodes in two hidden layers, and a bias of -1 kills 72% and 100%; accuracy drops to guessing. Leaky ReLU and ELU with the same bias reach 89% and 96%.

The Notebook trains a network with two hidden layers of 32 nodes on 500 standardised (rescaled to mean 0 and standard deviation 1) observations of `make_moons` (G-111; two features, two classes), with plain **stochastic gradient descent (SGD)** (G-1892; here one update per mini-batch of 32 observations, so 16 updates per epoch) for 200 **epochs** (G-696; full passes over the data). The Notebook then counts the nodes whose z is negative for every training observation.




Figure 3: Share of nodes whose z is negative on every training observation, after 200 epochs, with the training accuracy. For ReLU these nodes are dead; for Leaky ReLU and ELU they still pass a gradient

Figure 3 shows the results:

- **ReLU, learning rate 0.1:** almost no dead nodes (0% and 9%), accuracy 97%.
- **ReLU, learning rate 10:** 41% and 69% dead, accuracy 51%. The nodes were already dead after the first epoch; the shares did not change in the 199 epochs after it.
- **ReLU, bias -1 :** 72% and 100% dead from the very start, and still the same after 200 epochs. Accuracy 50%.

The deaths are permanent. Training the learning-rate-10 network for 20 more epochs changes the weights into its 13 dead first-layer nodes by exactly 0, as section 3.1 predicts.

Figure 4 follows the same runs through training, on a log scale (each step along the time axis multiplies the epoch by 10) so that the first epoch is visible.




Figure 4: The share of nodes negative on every training observation during training, for the runs of Figure 3. The time axis is logarithmic; the learning-rate-10 run is also shown after each batch of its first epoch.

1. **Learning rate 10 (red).** The nodes die during the very first epoch. After 8 of its 16 batches, 41% of the first-layer nodes are negative everywhere; after that the share never moves again.
2. **Bias -1 with ReLU (grey).** 72% and 100% of the nodes start negative everywhere and stay so for all 200 epochs: a flat line.
3. **Bias -1 with Leaky ReLU (blue) and ELU (green).** These are two of the variants in Figure 1, defined in sections 5.1 and 6.1. The same nodes start negative, but their slope is not 0, so their weights keep changing. The ELU shares start falling within 3 epochs; the Leaky ReLU shares from epoch 8 in the first layer and epoch 11 in the second, where they drop from 100% to 38% by epoch 27.
4. **Learning rate 0.1 (dotted).** Few nodes are ever negative everywhere.

4. Three ways to prevent dead nodes

Key point: Use a lower learning rate, start the biases at a small positive value such as 0.01, or replace ReLU with one of its variants.

1. **A lower learning rate.** In Figure 3, changing only the learning rate, from 10 to 0.1, lowers the dead shares from 41% and 69% to 0% and 9%.
2. **A positive starting bias.** Starting every bias at a small positive value, such as 0.01 or 0.1, makes it very likely that every ReLU node starts active for most inputs (Goodfellow et al. 2016, §6.3.1, suggests 0.1).
3. **A ReLU variant.** The root cause is that ReLU’s slope is exactly 0 for $z < 0$. The variants keep everything good about ReLU but give the negative side a slope.

Extra: A positive bias guards the start, not against a learning rate that is too high. Averaged over 5 seeds in the Notebook, a bias of +0.01 with learning rate 10 still ends with 49% and 80% dead nodes (bias 0: 49% and 89%), and accuracy stays near 50%. At learning rate 0.1, a starting bias of +0.1 lowers the share of second-layer nodes that start dead from 5.6% to 3.1%, and the share dead after training from 7.5% to 5.6%.

Python: Setting the starting bias in Keras.

```
keras.layers.Dense(  
    32, activation="relu",  
    bias_initializer=keras.initializers.Constant(0.01))
```

5. Linear variants

Key point: Leaky ReLU and Parametric ReLU replace the 0 on the negative side by a straight line az with a small slope a .

The **linear variants** of ReLU change only the negative side, and change it to a straight line.

5.1 Leaky ReLU

Key point: $f(z) = z$ for $z \geq 0$ and $0.01z$ for $z < 0$. Its slope on the negative side is 0.01 instead of 0, so a small gradient always flows.

1. **In words:** the **Leaky ReLU** (G-1064) keeps positive values and multiplies negative ones by 0.01 (Maas et al. 2013).
2. **Formula:**

$$f(z) = \begin{cases} z & z \geq 0 \\ 0.01z & z < 0 \end{cases}$$
$$f'(z) = \begin{cases} 1 & z \geq 0 \\ 0.01 & z < 0 \end{cases}$$

3. **Example:** $f(5) = 5$ and $f(-5) = -0.05$. At $z = -5$ ReLU's slope would be 0; Leaky ReLU's is 0.01.

Because $\partial a / \partial z$ is never 0, the gradient in section 3.1 is never exactly 0. The weights keep changing a little, and a node can climb back out of the negative region: with a starting bias of -1 , the share of first-layer nodes negative on every observation falls from 72% to 44% during training. In Figure 3, Leaky ReLU (Keras' "leaky_relu", slope 0.2) with a starting bias of -1 reaches 89% accuracy, where ReLU stayed at 50%.

Advantages:

1. **Non-saturating** (G-1339) on both sides: the output is unbounded in both directions.
2. **Easy to compute:** no exponentials.
3. **No dying ReLU problem.**

4. **Close to zero-centred** (G-2148): outputs can be negative as well as positive, though not symmetric.

The only questionable point of Leaky ReLU is the constant: why 0.01 and not some other value? The value 0.01 was chosen by experiment. Parametric ReLU lets the data choose it instead.

Extra: Keras does not use 0.01 by default (Keras documentation). The string `activation="leaky_relu"` uses slope 0.2, and the layer `keras.layers.LeakyReLU()` uses 0.3. To get 0.01, write `keras.layers.LeakyReLU(negative_slope=0.01)`.

5.2 Parametric ReLU (PReLU)

Key point: Like Leaky ReLU, but the negative slope a is a trainable parameter, learned per node along with the weights.

1. **In words:** the **Parametric ReLU (PReLU)** (G-1452) multiplies negative inputs by a slope a that is learned during training.
2. **Formula:**

$$f(z) = \begin{cases} z & z \geq 0 \\ a z & z < 0 \end{cases}$$

3. **Example:** if training sets $a = 0.25$, then $f(-2) = -0.5$; with $a = 0.01$ it would be Leaky ReLU, with $a = 0$ plain ReLU.

The slope a is a parameter like a weight, found by gradient descent from the data. The slope a is not a **hyperparameter** (G-910) that we set. Everything else, advantages included, is as for Leaky ReLU. The extra flexibility can help: the paper that introduced PReLU reports better **ImageNet** (G-920; a large benchmark of labelled photos) accuracy than with ReLU (He et al. 2015).

Python: PReLU is a separate layer after a **Dense** layer without activation.

```
keras.layers.Dense(8),
```

```
keras.layers.PReLU(),    # one learnable a per node
```

Here the PReLU layer adds 8 trainable parameters, one slope per node. In Keras each starts at 0, so the network begins as a ReLU network (He et al. 2015 started them at 0.25).

Extra: Nothing keeps the learned slope small. On the moons data the 8 slopes came out between -1.5 and 3.4 . A negative a makes the function V-shaped, which is still a valid non-linear activation.

6. Non-linear variants

Key point: ELU and SELU use an exponential curve for negative z that flattens out at a negative value. Outputs are close to zero-centred and there are no dead nodes.

The **non-linear variants** of ReLU use a curve, not a straight line, on the negative side.

6.1 ELU

Key point: $f(z) = z$ for $z \geq 0$ and $\alpha(e^z - 1)$ for $z < 0$. Smooth, close to zero-centred, no dead nodes, often better test results than ReLU; slower because of the exponential.

1. **In words:** the **ELU** (G-674; exponential linear unit; Clevert et al. 2016) is ReLU for positive z and an exponential curve for negative z that levels off at $-\alpha$.
2. **Formula:**

$$f(z) = \begin{cases} z & z \geq 0 \\ \alpha(e^z - 1) & z < 0 \end{cases}$$
$$f'(z) = \begin{cases} 1 & z \geq 0 \\ f(z) + \alpha & z < 0 \end{cases}$$

3. **Example:** with $\alpha = 1$ and $z = -1$:

$$f(-1) = e^{-1} - 1$$
$$f(-1) = 0.368 - 1 = -0.632$$
$$f'(-1) = f(-1) + 1 = -0.632 + 1 = 0.368$$

The negative-side slope comes from differentiating:

$$\frac{d}{dz} \alpha(e^z - 1) = \alpha e^z$$
$$\alpha e^z = f(z) + \alpha$$

A larger α pulls the negative side further down.

Advantages:

1. **Close to zero-centred**, so it converges faster.
2. **Better generalisation** (G-838): in experiments it often gives better results on test data than ReLU (Clevert et al. 2016).
3. **No dying ReLU problem:** the slope is above 0 for every z ($\alpha e^z > 0$ on the negative side). In Figure 3, ELU with a starting bias of -1 reaches 96%.
4. **Continuous and differentiable everywhere** (with $\alpha = 1$, which Figure 1 uses): there is no corner at 0.

Disadvantage: it needs an exponential, so it is slower to compute than ReLU. Faster convergence partly makes up for it, because fewer epochs are needed.

6.2 SELU

Key point: SELU is ELU multiplied by a fixed $\lambda \approx 1.0507$, with fixed $\alpha \approx 1.6733$. Its outputs keep mean 0 and standard deviation 1 from layer to layer: it is self-normalising.

1. **In words:** the **SELU** (G-1766; scaled exponential linear unit; Klambauer et al. 2017) is ELU with a specific α , multiplied by a scale λ .

2. **Formula:**

$$f(z) = \lambda \begin{cases} z & z \geq 0 \\ \alpha(e^z - 1) & z < 0 \end{cases}$$

$$\lambda \approx 1.0507$$

$$\alpha \approx 1.6733$$

3. **Example:**

$$f(1) = 1.0507$$

$$\lambda\alpha = 1.0507 \times 1.6733 = 1.758$$

$$f(-1) = 1.758 \times (0.368 - 1) = -1.111$$

λ and α are fixed constants, not trainable parameters. The two constants were derived so that the function has one special property.

The special property is being **self-normalising** (G-1765): the outputs of a SELU layer have mean about 0 and standard deviation about 1, and the next layer keeps them there. Normalised values between layers make the network converge fast, and SELU also generalises well in experiments.




Figure 5: Standard deviation of the activations through 30 layers of 256 nodes, with no training. SELU stays at 1; ELU and ReLU shrink layer after layer

Figure 5 shows this in the Notebook: the horizontal axis is the layer number, and the vertical axis is the standard deviation of that layer's outputs (how spread out they are). Standard-normal inputs pass through 30 layers with the same random weights for each activation. With SELU the standard deviation is 1.00 at layer 1 and still 1.00 at layer 30. With ReLU it is 0.59 at layer 1 and 4×10^{-5} at layer 30.

The disadvantage of SELU is adoption. SELU is used in few places so far, for three reasons:

- SELU is recent (2017);
- its paper has 9 pages plus a 93-page appendix of proofs;
- less research builds on it.

Extra: Self-normalisation rests on assumptions in the paper: inputs with mean 0 and variance 1, and weights drawn with variance 1/inputs (Klambauer et al. 2017). Keras' documentation for `selu` therefore asks for `kernel_initializer="lecun_normal"` (as in Figure 5) and for `AlphaDropout` instead of ordinary dropout, which the paper shows disturbs the mean and variance. Starting weights are the subject of the [weight initialisation](#).

Python: ELU and SELU in Keras.

```
keras.layers.Dense(32, activation="elu")
keras.layers.Dense(32, activation="selu",
                    kernel_initializer="lecun_normal")
```

6.3 GELU and SiLU: smooth versions of ReLU

Key point: GELU is $z \Phi(z)$ and SiLU is $z \sigma(z)$: the input times a number between 0 and 1 that grows with the input. Both look like ReLU with a smooth bend and a small dip below 0. Transformers use them.

ReLU makes a hard choice: it keeps z when z is positive and replaces it by 0 when z is negative. A softer rule makes the same choice at random, with odds that depend on z :

1. **Keep or drop.** The node's value z is kept with probability $p(z)$ and replaced by 0 otherwise. This is [dropout](#) (G-639) with a rate that depends on the input.
2. **The odds follow the input.** A large positive z is almost always kept; a very negative z is almost always dropped.
3. **Take the average.** Over many random choices the average output is:

$$\begin{aligned} z \times p(z) + 0 \times (1 - p(z)) \\ = z p(z) \end{aligned}$$

Using this average as the activation function needs no randomness at all.

The **GELU** (G-834; Gaussian error linear unit) uses for p the [standard normal cumulative distribution function](#) (G-515) $\Phi(z)$ (G-18): the probability that a standard normal value is below z . For example, $\Phi(0) = 0.50$ because half of the values are below 0, and $\Phi(1) = 0.84$ because 84 percent are below 1.

1. **In words:** the input times the probability of keeping it.
2. **Formula:**

$$\text{GELU}(z) = z \Phi(z)$$

3. **Example:**

z	-2	-1	0	1	2
$\Phi(z)$, probability of being kept	0.02	0.16	0.50	0.84	0.98
GELU, $z \Phi(z)$	-0.05	-0.16	0	0.84	1.95




Figure 6: GELU built point by point. Left: the probability $\Phi(z)$ of keeping the input. Right: the average output $z \Phi(z)$, one dot per input, then the whole curve, then ReLU and SiLU for comparison. Idea after StatQuest, “The GELU, SiLU and SwiGLU activation functions, clearly explained!!!”

Figure 6 builds the curve from these five values. Watch the red dot on the left: its height is the number that multiplies the input on the right.

The **SiLU** (G-2266) (sigmoid linear unit, also called **Swish**) uses the **sigmoid** (G-1798; the S-shaped function that turns any number into 0 to 1) for p , a curve of almost the same shape as Φ :

$$\text{SiLU}(z) = z \sigma(z)$$

For example:

$$\text{SiLU}(-1) = -1 \times 0.269 = -0.27$$

$$\text{SiLU}(1) = 1 \times 0.731 = 0.73$$

Compared with ReLU (last frame of Figure 6):

- **For large positive z** both are almost z , like ReLU, so they do not saturate there.
- **For negative z** the output is a small negative number, not 0: GELU dips to -0.17 and returns towards 0. The slope is not exactly 0, so a node with a slightly negative z still receives a gradient.
- **At 0** the curve bends smoothly; there is no corner.

GELU is the activation inside GPT’s feed-forward layers; its use there is in section 6.4 of the [GPT](#).

7. Summary

	ReLU	Leaky ReLU	PReLU	ELU	SELU
Negative side Kind	0	$0.01z$ linear variant	az , a learned linear variant	$\alpha(e^z - 1)$ non-linear variant	$\lambda\alpha(e^z - 1)$ non-linear variant
Slope for $z < 0$	0	0.01	a	$f(z) + \alpha$	$\lambda\alpha e^z$
Dead nodes	yes	no	no	no	no
Zero-centred	no	close	close	close	yes (self-normalising)
Exponential Main drawback	no dying	no slope 0.01 is arbitrary	no	yes slower	yes not widely adopted

- A ReLU node with $z < 0$ for every observation has slope 0, gets no updates and is dead for good.
- Causes: a high learning rate and a large negative bias. Fixes: a lower learning rate, a bias starting at 0.01, or a variant.
- Linear variants (Leaky ReLU, PReLU) put a straight line on the negative side; non-linear ones (ELU, SELU) a curve.
- SELU keeps activations at mean 0 and standard deviation 1 across layers.
- GELU, $z\Phi(z)$, and SiLU, $z\sigma(z)$, are smooth versions of ReLU: the input times a keep-probability that grows with the input.

8. Sources

Built from

- CampusX, “Relu Variants Explained | Leaky Relu | Parametric Relu | Elu | Selu | Activation Functions Part 2”, YouTube, <https://www.youtube.com/watch?v=2OwWs7Hzr9g>
- StatQuest with Josh Starmer, “The GELU, SiLU and SwiGLU activation functions, clearly explained!!!”, YouTube, <https://www.youtube.com/watch?v=2FaI2Fen1mQ> (section 6.3: GELU and SiLU as the average of input-dependent dropout, Figure 6)

Other references

- Clevert, D.-A., Unterthiner, T. and Hochreiter, S. (2016). Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs). ICLR 2016. arXiv:1511.07289.
- He, K., Zhang, X., Ren, S. and Sun, J. (2015). Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification. ICCV 2015. arXiv:1502.01852.
- Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*, MIT Press, section 6.3.1 (a small positive starting bias keeps ReLU units active at the start).
- Keras API documentation: `LeakyReLU`, `PReLU`, `selu` (keras.io/api).
- Klambauer, G., Unterthiner, T., Mayr, A. and Hochreiter, S. (2017). Self-Normalizing Neural Networks. NeurIPS 2017. arXiv:1706.02515.

- Maas, A. L., Hannun, A. Y. and Ng, A. Y. (2013). Rectifier Nonlinearities Improve Neural Network Acoustic Models. ICML 2013 Workshop on Deep Learning for Audio, Speech and Language Processing.

9. Key terms

Term	Meaning
Dying ReLU problem	ReLU nodes ending up with a negative weighted sum for every input, so they output 0 and stop learning
Observation	One record of the data: one row of the data table
Feature	An input variable, such as x_1 : one column of the data table
Dead neuron	A node whose output is 0 for every input; it gets no updates and stays that way
Linear variants of ReLU	ReLU variants with a straight line on the negative side: Leaky ReLU and PReLU
Non-linear variants of ReLU	ReLU variants with a curve on the negative side: ELU and SELU
Leaky ReLU	z for $z \geq 0$, $0.01z$ for $z < 0$
Parametric ReLU (PReLU)	Leaky ReLU whose negative slope a is learned during training, one per node
ELU	Exponential linear unit: z for $z \geq 0$, $\alpha(e^z - 1)$ for $z < 0$
SELU	Scaled ELU, $\lambda \approx 1.0507$ times ELU with $\alpha \approx 1.6733$; self-normalising
GELU	Gaussian error linear unit, $z \Phi(z)$: the input times the standard normal probability of a value below it; a smooth version of ReLU
SiLU (Swish) (G-2266)	Sigmoid linear unit, $z \sigma(z)$: the input times its sigmoid; close to GELU
Self-normalising	Keeping the activations of every layer at mean 0 and standard deviation 1 without a separate normalisation step