

Backpropagation in a CNN, Part 1: The Setup and the Last Layer

Contents

1. Overview	2
2. Prerequisites	3
3. A small CNN	3
3.1 The architecture and the shapes	3
3.2 The trainable parameters	4
3.3 The loss	5
4. Forward propagation	6
5. What we need: four derivatives	9
5.1 Two parts: a CNN and an ANN	10
5.2 The chains	11
6. The ANN part: $\partial L / \partial W_2$ and $\partial L / \partial b_2$	13
6.1 The first two factors	13
6.2 The last factors	15
6.3 Checking the shapes	16
7. A batch of images	17
8. Summary	19
9. Sources	19
10. Key terms	19

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Backpropagation ([Note DL-015](#)); Convolution operation and feature maps ([Note DL-042](#)).

1. Overview

Key point: A CNN learns the same way as an ordinary network: it makes a guess, measures how wrong the guess is, and nudges every weight in the direction that makes the error smaller. We cut a small CNN into a CNN part and an ANN part. For the ANN part, each output weight's nudge is the output error times the number that weight multiplied.

Backpropagation (G-247) for an ANN was taught in the [backpropagation Notes](#). A CNN adds three new operations, and the **chain rule** (G-371) must pass through each of them:

- convolution;
- max pooling;
- flatten.

This Note sets the problem up on the smallest possible CNN and finds the gradients of its last layer. The [part 2 Note](#) goes back through flatten, max pooling and convolution.




Figure 1: The small CNN as a chain of operations, with the shape of every tensor. The trainable parameters are W_1, b_1 (the filter) and W_2, b_2 (the output node). Backpropagation walks the chain from right to left

Figure 1 shows the whole network as a row of boxes. Read it left to right for the guess: a small image goes in, a probability comes out. Read it right to left for the learning: the error at the output is passed back, box by box, to every weight. Section 3 gives every box its numbers.

In practice Keras computes all of this for us. Knowing how the gradients flow through a CNN still helps us understand and use the tools better.

2. Prerequisites

- The [backpropagation Notes](#), [how it runs](#) and [why it works](#): the chain rule along a network, and the update $w \leftarrow w - \eta \partial L / \partial w$.
- The [CNN vs ANN Note](#): a **filter** (G-777) works like a node, its values are weights.
- The [convolution operation Note](#) and the [pooling Note](#).
- The [log loss Note](#) and the [sigmoid derivative Note](#).

3. A small CNN

Key point: 6×6 image $\rightarrow$ one 3×3 filter with bias $\rightarrow$ ReLU $\rightarrow 2 \times 2$ max pooling $\rightarrow$ flatten $\rightarrow$ one sigmoid node. Shapes: $6 \times 6 \rightarrow 4 \times 4 \rightarrow 4 \times 4 \rightarrow 2 \times 2 \rightarrow 4 \rightarrow 1$.

3.1 The architecture and the shapes

Key point: Convolution gives 4×4 , ReLU keeps the shape, pooling halves it to 2×2 , flatten gives 4 numbers, the output node gives 1.

The network (Figure 1) has:

1. **Input:** a 6×6 greyscale image X .
2. **Convolution:** one 3×3 filter with its bias, giving a 4×4 **feature map** (G-766).
The side is:

$$6 - 3 + 1 = 4$$

3. **ReLU** (G-1668): negatives become 0; the shape stays 4×4 .
4. **Max pooling** (G-1182): 2×2 window, stride 2, giving 2×2 .
5. **Flatten:** 4 numbers.
6. **Output:** a single node with sigmoid, giving the prediction $\hat{y}$, a number between 0 and 1.

The network is a binary classifier: is this a picture of a cat or a dog, or, in the Notebook, is this MNIST digit a 1 ($y = 1$) or a 0 ($y = 0$)? The Notebook shrinks real MNIST digits to 6×6 so that every shape matches this Note.




Figure 2: Four of the Notebook’s MNIST digits at their full 28×28 size (top) and shrunk to 6×6 (bottom), the input size of the small CNN. The first 0 is the image used in section 6.

In Figure 2, the 6×6 versions keep only the rough shape, a ring for a 0 and a vertical stroke for a 1.

3.2 The trainable parameters

Key point: 9 filter weights + 1 filter bias + 4 output weights + 1 output bias = 15 parameters.

Only two places in the network hold numbers that training can change: the filter and the output node. Before training, the Notebook draws them at random. These starting values are the ones used in every example of this Note and of the part 2 Note.

Parameter	Where	Shape	Count	Starting value (Notebook)
W_1	the filter	3×3	9	the 3×3 grid below
b_1	the filter’s bias	1×1	1	0.1
W_2	weights of the output node	1×4	4	(−0.6327, −0.3116, 0.0207, −1.1625)
b_2	bias of the output node	1×1	1	0

$$W_1 = \begin{bmatrix} 0.0629 & -0.0661 & 0.3202 \\ 0.0525 & -0.2678 & 0.1808 \\ 0.6520 & 0.4735 & -0.3519 \end{bmatrix}$$

The count of parameters, one line per place:

$$9 + 1 = 10 \quad (\text{filter})$$

$$4 + 1 = 5 \quad (\text{output node})$$

$$10 + 5 = 15$$

These 15 numbers are the **trainable parameters** (G-1065). ReLU, max pooling and flatten have none. Training means finding the 15 values that make the loss smallest.

3.3 The loss

Key point: Binary cross-entropy (log loss), the same loss as in logistic regression.

The loss is one number that says how wrong one guess is: near 0 for a confident right guess, large for a confident wrong one. Our first image is a 0, so its target is $y = 0$. Section 4 shows that the untrained network gives it the prediction $a_2 = 0.2104$, a 21 percent chance of being a 1. The guess leans the right way, so the loss should be small:

$$\begin{aligned} L &= -0 \times \log 0.2104 - (1 - 0) \times \log(1 - 0.2104) \\ &= 0 - \log 0.7896 \\ &= 0.2362 \end{aligned}$$

Here $\log$ is the natural logarithm. The general rule is the **binary cross-entropy** (G-303) (see the [log loss Note](#)), for one image with target y (0 or 1) and prediction $a_2 = \hat{y}$:

$$L = -y \log a_2 - (1 - y) \log(1 - a_2)$$

Check: $y = 0$ and $a_2 = 0.2104$ give 0.2362, as above (Notebook: 0.2362). For a batch of m images (for example $m = 32$) the loss is the average of the m single losses. For several classes we would use softmax and categorical cross-entropy instead (see the [loss functions Note](#)). The loss of a CNN is exactly the loss of an ANN.

4. Forward propagation

Key point: The image passes through six operations, each turning one grid of numbers into the next, until a single probability comes out. This walk from image to prediction is **forward propagation** (G-797).

Figure 1 is the **logical diagram** (G-1118) of the network: each arrow is one operation, each box one tensor (a grid of numbers). Figure 3 plays the walk on our first image, a 0 shrunk to 6×6 ; the steps below give the same numbers in writing.




Figure 3: The Notebook’s first image, a 0 shrunk to 6×6 , through the small CNN with the Notebook’s starting weights. The 3×3 filter slides over X to fill Z_1 ; ReLU changes nothing here, since all 16 values are positive; each 2×2 window sends its maximum to P_1 ; flatten gives F and the sigmoid gives $\hat{y} = 0.21$. Then the error $a_2 - y = 0.21$ multiplies each value of F to give $\partial L / \partial W_2$ (section 6.2)

In Figure 3, watch the grids shrink, 6×6 to 4×4 to 2×2 to 4 numbers to 1, as each operation runs.

Step 1, convolution. The input X is the 6×6 grid of pixel brightnesses, 0 for black and 1 for white. The filter W_1 sits on the top-left 3×3 window of X ; each pixel is multiplied by the filter value on top of it, and the 9 products are added (the **convolution operation** (G-481), see the [convolution operation Note](#)). The top-left window of our image is

$$\begin{bmatrix} 0 & 0 & 0.0078 \\ 0 & 0.0078 & 0.2706 \\ 0 & 0.1765 & 0.4667 \end{bmatrix}$$

and the 9 products with W_1 (section 3.2) are:

Pixel $\times$ filter value	Product
0×0.0629	0
$0 \times (-0.0661)$	0
0.0078×0.3202	0.0025
0×0.0525	0
$0.0078 \times (-0.2678)$	-0.0021
0.2706×0.1808	0.0489
0×0.6520	0
0.1765×0.4735	0.0836
$0.4667 \times (-0.3519)$	-0.1642

$$\text{sum} = 0.0025 - 0.0021 + 0.0489$$

$$+0.0836 - 0.1642$$

$$\text{sum} = -0.0313$$

$$\text{plus } b_1 : \quad -0.0313 + 0.1 = 0.0687$$

0.0687 is the top-left cell of the feature map Z_1 . The filter then slides one pixel at a time to fill all 16 cells (Notebook):

$$Z_1 = \begin{bmatrix} 0.0687 & 0.4468 & 0.2560 & 0.4006 \\ 0.3558 & 0.4947 & 0.2482 & 0.2801 \\ 0.1859 & 0.5311 & 0.7707 & 0.2888 \\ 0.0771 & 0.2393 & 0.2995 & 0.1123 \end{bmatrix}$$

Step 2, ReLU. Each negative cell becomes 0, each positive cell stays. All 16 cells of Z_1 are positive, so $A_1 = Z_1$ for this image.

Step 3, max pooling. Cut A_1 into four 2×2 windows and keep the largest number of each:

$$\text{top-left: } (0.0687, 0.4468,$$

$$0.3558, 0.4947) \rightarrow 0.4947$$

$$\text{top-right: } (0.2560, 0.4006,$$

$$0.2482, 0.2801) \rightarrow 0.4006$$

bottom-left: $(0.1859, 0.5311,$

$$0.0771, 0.2393) \rightarrow 0.5311$$

bottom-right: $(0.7707, 0.2888,$

$$0.2995, 0.1123) \rightarrow 0.7707$$

$$P_1 = \begin{bmatrix} 0.4947 & 0.4006 \\ 0.5311 & 0.7707 \end{bmatrix}$$

Step 4, flatten. Read P_1 row by row into one column of 4 numbers, f_1 to f_4 :

$$F = (0.4947, 0.4006, 0.5311, 0.7707)^T$$

The T (**transpose**, G-2012) only says that F is stored as a column.

Step 5, the output node. Each of the 4 values is multiplied by its weight in W_2 , one product per line:

$$-0.6327 \times 0.4947 = -0.3130$$

$$-0.3116 \times 0.4006 = -0.1249$$

$$0.0207 \times 0.5311 = 0.0110$$

$$-1.1625 \times 0.7707 = -0.8959$$

$$Z_2 = -0.3130 - 0.1249 + 0.0110 - 0.8959 + b_2$$

$$Z_2 = -1.3228 + 0$$

$$Z_2 = -1.3228$$

Step 6, sigmoid. The **sigmoid function** (G-1798) squeezes Z_2 into a probability:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$e^{1.3228} = 3.7538$$

$$A_2 = \frac{1}{1 + 3.7538} = 0.2104$$

The network says “21 percent a 1”; the image is a 0. The loss of section 3.3 measures this miss as 0.2362.

The same six steps as equations, for any image:

Step	Equation	Shape
Convolution	$Z_1 = X * W_1 + b_1$	4×4 (the feature map)
ReLU	$A_1 = \text{ReLU}(Z_1)$	4×4
Max pooling	$P_1 = \text{maxpool}(A_1)$	2×2
Flatten	$F = \text{flatten}(P_1)$	4×1
Output node	$Z_2 = W_2 F + b_2$	$(1 \times 4)(4 \times 1) = 1 \times 1$
Sigmoid	$A_2 = \sigma(Z_2)$	1×1 (the prediction)

Here $*$ is the convolution of step 1 and b_1 is added to every cell of the feature map. The product is written $W_2 F$, not FW_2 , so that the shapes fit:

$$(1 \times 4) \text{ times } (4 \times 1) = 1 \times 1$$

Check: the equations reproduce step by step the numbers above, ending at $A_2 = 0.2104$ (Notebook: 0.2104). With these equations we can code the forward pass and predict for any image; the Notebook does it in NumPy.

5. What we need: four derivatives

Key point: To nudge a parameter the right way, we need its slope: how much the loss changes when that parameter changes a little. We need one slope for each of the 15 parameters, so each group of slopes has the shape of its group of parameters.

Training starts from the random values of section 3.2 and repeats **gradient descent** (G-862) until the loss is small (see the [backpropagation why Note](#)). Each step moves every parameter a little against its slope.

The slope of the loss with respect to one parameter, say b_2 , is written $\partial L / \partial b_2$ and read “how much L changes per unit change of b_2 ”. It is a **partial derivative** (G-1457). Section 6.2 finds, for our image,

$$\frac{\partial L}{\partial b_2} = 0.2104$$

A positive slope means that raising b_2 raises the loss, so gradient descent lowers b_2 . With a **learning rate** (G-1068) $\eta = 0.1$, the value the part 2 Note trains with, one step is:

$$\begin{aligned} b_2 &\leftarrow b_2 - \eta \frac{\partial L}{\partial b_2} \\ &= 0 - 0.1 \times 0.2104 \\ &= -0.0210 \end{aligned}$$

Every parameter follows the same rule, with its own slope:

$$\begin{aligned} W_1 &\leftarrow W_1 - \eta \frac{\partial L}{\partial W_1} \\ b_1 &\leftarrow b_1 - \eta \frac{\partial L}{\partial b_1} \\ W_2 &\leftarrow W_2 - \eta \frac{\partial L}{\partial W_2} \\ b_2 &\leftarrow b_2 - \eta \frac{\partial L}{\partial b_2} \end{aligned}$$

W_1 and W_2 are matrices, so their derivatives are matrices of the same shape: one partial derivative per weight, 9 numbers for W_1 and 4 for W_2 (see the [Jacobian and matrix gradients Note](#)). So the whole task is to find these four derivatives.

5.1 Two parts: a CNN and an ANN

Key point: Everything up to flatten is the CNN part; everything after is an ordinary one-node ANN. Study them separately, then join them.

It helps to see the network as two networks joined together (Figure 1): a **CNN part** (convolution, ReLU, max pooling, flatten) and an **ANN part** (the output node). This Note finds the ANN part’s derivatives; the [part 2 Note](#) finds the CNN part’s.

5.2 The chains

Key point: A weight of the output node reaches the loss through 3 links; a filter weight through 7. The slope of the whole path is the product of the slopes of its links.

Think of a row of gears. Turn the first gear a little and the last gear turns too, by an amount set by every gear in between. The first weight of W_2 , call it w_1 (starting value -0.6327), is not connected to the loss directly:

1. a change in w_1 changes Z_2 ;
2. the change in Z_2 changes A_2 ;
3. the change in A_2 changes L .

Each link has its own slope, and section 6 finds each one on our image:

Link	What it asks	Value on our image
$\partial Z_2 / \partial w_1$	how much Z_2 moves per unit of w_1	$f_1 = 0.4947$
$\partial A_2 / \partial Z_2$	how much the sigmoid output moves per unit of Z_2	0.1661
$\partial L / \partial A_2$	how much the loss moves per unit of A_2	1.2664

The slope of the whole path is the product of the three, one factor per line:

$$1.2664$$

$$\times 0.1661 = 0.2104$$

$$\times 0.4947 = 0.1041$$

So raising w_1 by a small amount raises the loss by about 0.1041 times that amount. Multiplying the slopes along a path is the **chain rule** (G-371) (Goodfellow et al. 2016, §6.5.2; see section 5.3 of the [derivatives Note](#)). For all of W_2 and for b_2 :

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial A_2} \cdot \frac{\partial A_2}{\partial Z_2} \cdot \frac{\partial Z_2}{\partial W_2}$$

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial A_2} \cdot \frac{\partial A_2}{\partial Z_2} \cdot \frac{\partial Z_2}{\partial b_2}$$

Only the last factor differs between the two. Check: the first entry of $\partial L / \partial W_2$ in the Notebook is 0.1041, the product above.

For the filter the path is much longer. A change in a filter weight goes through 7 links before it reaches the loss:

Link	Goes back through	On our image	Found in
$\partial L/\partial A_2$	the loss	1.2664	section 6.1
$\partial A_2/\partial Z_2$	the sigmoid	0.1661	section 6.1
$\partial Z_2/\partial F$	the output node	the 4 weights of W_2	part 2, section 4
$\partial F/\partial P_1$	flatten	a reshape, no arithmetic	part 2, section 5
$\partial P_1/\partial A_1$	max pooling	1 at each window's maximum, 0 elsewhere	part 2, section 6
$\partial A_1/\partial Z_1$	ReLU	1 in all 16 cells (all of Z_1 is positive)	part 2, section 7
$\partial Z_1/\partial W_1$	the convolution	the pixels of X under each filter weight	part 2, section 8

Written as one product:

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial A_2} \cdot \frac{\partial A_2}{\partial Z_2} \cdot \frac{\partial Z_2}{\partial F} \cdot \frac{\partial F}{\partial P_1} \cdot \frac{\partial P_1}{\partial A_1} \cdot \frac{\partial A_1}{\partial Z_1} \cdot \frac{\partial Z_1}{\partial W_1}$$

and $\partial L/\partial b_1$ is the same chain with $\partial Z_1/\partial b_1$ as its last factor. The part 2 Note multiplies these links out on the same image.




Figure 4: The two paths from a parameter to the loss. Each arrow is one link of the chain rule. Green links are those of an ordinary ANN; red links pass through the CNN part.

Figure 4 shows the two chains side by side: the last three links are shared, so the work of section 6 is reused for the filter; the four red links, through the CNN part, are left for the part 2 Note.

Figure 5 plays the whole plan. First the forward chain is drawn, one operation per arrow. Then the gradients appear from right to left, one per tensor, each with the same shape as its

Forward: one operation per arrow

Backward: one gradient per tensor, right to left

Backward: one gradient per tensor, right to left

Backward: one gradient per tensor, right to left

Three of these factors are new: $\partial F/\partial P_1$ goes back through flatten, $\partial P_1/\partial A_1$ through max pooling, and $\partial Z_1/\partial W_1$ through the convolution. The part 2 Note finds them.

Key point: Two slopes, of the loss and of the sigmoid, multiply to a very simple number: the prediction minus the target, the **error** at the output. Each output weight's slope is that error times the value the weight multiplied; the bias's slope is the error itself.

Key point: The slope of the loss and the slope of the sigmoid multiply to the output error, prediction minus target: $0.2104 - 0 = 0.2104$ on our image.

The slope of the loss. Differentiating the log loss of section 3.3 gives

$$\frac{\partial L}{\partial a_2} = -\frac{y}{a_2} + \frac{1-y}{1-a_2}$$

On our image:

$$-\frac{0}{0.2104} = 0$$

$$\frac{1-0}{1-0.2104} = \frac{1}{0.7896} = 1.2664$$

$$\frac{\partial L}{\partial a_2} = 0 + 1.2664 = 1.2664$$

The slope of the sigmoid (see the [sigmoid derivative Note](#)) is the output times one minus the output:

$$\frac{\partial a_2}{\partial Z_2} = a_2(1-a_2)$$

$$= 0.2104 \times 0.7896$$

$$= 0.1661$$

The product of the two:

$$1.2664 \times 0.1661 = 0.2104$$

The result is exactly $a_2 - y$:

$$a_2 - y = 0.2104 - 0$$

The cancellation holds for every image (the log of the loss undoes the exponential of the sigmoid; Goodfellow et al. 2016, §6.2.2.2). First put the two terms of $\partial L / \partial a_2$ over one denominator:

$$\begin{aligned} -\frac{y}{a_2} + \frac{1-y}{1-a_2} &= \frac{-y(1-a_2) + (1-y)a_2}{a_2(1-a_2)} \\ &= \frac{-y + ya_2 + a_2 - ya_2}{a_2(1-a_2)} \end{aligned}$$

$$= \frac{a_2 - y}{a_2(1 - a_2)}$$

Then multiply by the sigmoid's slope; $a_2(1 - a_2)$ cancels:

$$\frac{a_2 - y}{a_2(1 - a_2)} \times a_2(1 - a_2) = a_2 - y$$

$$\frac{\partial L}{\partial Z_2} = a_2 - y$$

Check against the product of the two numbers above (Notebook):

$$0.2104 - 0 = 0.2104$$

6.2 The last factors

Key point: Each output weight's slope is the output error times the value of F it multiplied; the bias's slope is the error itself.

Z_2 is the sum of four products plus the bias (section 4, step 5):

$$Z_2 = w_1 f_1 + w_2 f_2 + w_3 f_3 + w_4 f_4 + b_2$$

Raise w_1 by a small amount and only the first product changes, by that amount times $f_1 = 0.4947$. So the slope of Z_2 with respect to w_1 is f_1 ; with respect to each weight w_k it is the input f_k that weight multiplies; with respect to b_2 it is 1.

On our image the error is $a_2 - y = 0.2104$ (section 6.1). Each weight's slope is the error times its input, one product per line:

$$\frac{\partial L}{\partial w_1} = 0.2104 \times 0.4947 = 0.1041$$

$$\frac{\partial L}{\partial w_2} = 0.2104 \times 0.4006 = 0.0843$$

$$\frac{\partial L}{\partial w_3} = 0.2104 \times 0.5311 = 0.1117$$

$$\frac{\partial L}{\partial w_4} = 0.2104 \times 0.7707 = 0.1621$$

$$\frac{\partial L}{\partial b_2} = 0.2104 \times 1 = 0.2104$$

All four slopes are positive, so gradient descent lowers all four weights, which lowers Z_2 and the prediction: the right move for an image of a 0. The last frames of Figure 3 show these products: every weight of W_2 gets the same error, 0.21, times the value of F it multiplied.

The general formula collects the four lines into one row, with $F^\top$ the row (f_1, f_2, f_3, f_4) :

$$\frac{\partial L}{\partial W_2} = (a_2 - y) F^\top$$

$$\frac{\partial L}{\partial b_2} = a_2 - y$$

Check against the four lines above (Notebook):

$$0.2104 \times 0.4947 = 0.1041$$

$$0.2104 \times 0.4006 = 0.0843$$

$$0.2104 \times 0.5311 = 0.1117$$

$$0.2104 \times 0.7707 = 0.1621$$

TensorFlow's `GradientTape` (G-88), which differentiates the same network automatically, gives the same numbers (largest difference 0, Notebook).

6.3 Checking the shapes

Key point: The slopes of W_2 must form a row of 4, the shape of W_2 . Turning the column F into a row (its transpose) makes the shapes fit.

A derivative is used to update its parameter, so it must have the same shape. W_2 is 1×4 . The error $a_2 - y$ is 1×1 and F is 4×1 , so we use its transpose $F^\top$, 1×4 :

$$\underbrace{(a_2 - y)}_{1 \times 1} \underbrace{F^\top}_{1 \times 4} = \underbrace{\frac{\partial L}{\partial W_2}}_{1 \times 4}$$

7. A batch of images

Key point: For a batch, each image gives its own gradient as in section 6, and the batch gradient is their average. One matrix product computes all of them at once.

With **mini-batch gradient descent** (G-1222), a batch of, say, 32 or 64 images goes forward together and backpropagation runs once for the batch (see the [gradient descent in neural networks Note](#)).

Start with the smallest batch, $m = 2$ images: our 0 and the Notebook's second image, a 1 ($y = 1$). Section 6 gives each image its own gradient:

Image	Target y	Prediction a_2	Error $a_2 - y$	F
1	0	0.2104	0.2104	(0.4947, 0.4006, 0.5311, 0.7707)
2	1	0.2910	-0.7090	(0.3766, 0.5051, 0.5133, 0.4346)

For the first weight w_1 , one product per line, then the average:

$$\text{image 1: } 0.2104 \times 0.4947 = 0.1041$$

$$\text{image 2: } -0.7090 \times 0.3766 = -0.2670$$

$$\text{sum: } 0.1041 - 0.2670 = -0.1629$$

$$\text{average: } -0.1629/2 = -0.0815$$

The same for all four weights gives $(-0.0815, -0.1369, -0.1261, -0.0730)$, and for the bias:

$$(0.2104 - 0.7090)/2 = -0.2493$$

(Notebook.) The second image, badly wrong, outweighs the first: the batch asks to raise the weights.

The matrix form stores one image per column:

- F has shape $4 \times m$, here 4×2 ;
- the predictions A_2 and the targets Y have shape $1 \times m$, here 1×2 .

The batch loss is the average of the single losses, so a factor $1/m$ appears:

$$\frac{\partial L}{\partial W_2} = \frac{1}{m} (A_2 - Y) F^\top$$

$$\frac{\partial L}{\partial b_2} = \frac{1}{m} \sum_{i=1}^m (a_{2,i} - y_i)$$

Here $\sum_{i=1}^m$ means “add the term for image 1, image 2, $\dots$, image m ”. Check: the first entry of $(A_2 - Y)F^\top$ is the row $(0.2104, -0.7090)$ times the column $(0.4947, 0.3766)$, which is the sum -0.1629 above; times $1/2$ it gives -0.0815 .

The shapes:

$$(1 \times m)(m \times 4) = 1 \times 4$$

The m cancels in the matrix product, so the derivative is 1×4 whatever the batch size, the shape of W_2 . In the Notebook, with $m = 32$, the formula and `GradientTape` agree to within 6×10^{-17} .

The matrix product adds up the 32 single-image gradients, and the $1/m$ turns the sum into an average.




Figure 6: The 32 images of the Notebook’s batch with the starting weights. Each dot is one image’s gradient for one weight of W_2 ; the black bar is the batch gradient, $\frac{1}{m}(A_2 - Y)F^\top$, their average.

In Figure 6, the images of 0s (blue) push every weight one way and the images of 1s (red) the other way, because $a_2 - y$ is positive for a 0 and negative for a 1; the batch gradient is the average of both groups.

8. Summary

Derivative	Formula (one image)	Formula (batch of m)	Shape
$\partial L / \partial Z_2$	$a_2 - y$	$A_2 - Y$	1×1 or $1 \times m$
$\partial L / \partial W_2$	$(a_2 - y)F^\top$	$\frac{1}{m}(A_2 - Y)F^\top$	1×4
$\partial L / \partial b_2$	$a_2 - y$	mean of $A_2 - Y$	1×1
$\partial L / \partial W_1, \partial L / \partial b_1$	chains of 7 factors through flatten, max pooling, ReLU and convolution		$3 \times 3, 1 \times 1$

- A CNN is trained like an ANN: forward pass, loss, chain rule, gradient descent.
- The small CNN has 15 parameters: W_1 (9), b_1 , W_2 (4), b_2 .
- Split it into a CNN part and an ANN part; the ANN part gives $\partial L / \partial W_2 = (a_2 - y)F^\top$.
- Every derivative has the shape of its parameter; shapes guide where to put transposes.
- The filter's gradients need the backward steps through flatten, max pooling and convolution: the [part 2 Note](#).

9. Sources

Built from

- CampusX, “Backpropagation in CNN | Part 1 | Deep Learning”, YouTube, <https://www.youtube.com/watch?v=RvCCFttGFM>

Other references

- Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*. MIT Press. §6.2.2.2 (sigmoid output with the log loss), §6.5.2 (the chain rule in back-propagation).
- TensorFlow API documentation: `tf.GradientTape`, [tensorflow.org/api_docs/python/tf/GradientTape](https://www.tensorflow.org/api_docs/python/tf/GradientTape)

10. Key terms

Term	Meaning
Logical diagram	A drawing of a network as a chain of tensors linked by operations
Forward propagation	Computing the prediction from the input, operation by operation
Trainable parameter (G-1065)	A number changed by gradient descent: filter values, weights and biases
Binary cross-entropy (G-303)	The log loss $-y \log a - (1 - y) \log(1 - a)$ used for two classes
$\partial L / \partial Z_2 = a_2 - y$	The error at a sigmoid output with log loss

Term	Meaning
Partial derivative (G-1457)	The slope of the loss in one parameter, the others held fixed: how much L changes per unit change of that parameter
Chain rule (G-371)	The slope along a path of links is the product of the slopes of the links
GradientTape	TensorFlow's tool that records a computation and returns its gradients automatically