

The Keras Functional API: Non-Linear Networks

Contents

1. Overview	2
2. Prerequisites	3
3. When a straight line is not enough	3
4. Building a functional model	5
4.1 Two outputs from one input	5
4.2 Two inputs, one output	8
4.3 A skip connection	10
5. A real model: age and gender from one face photo	10
5.1 The data	10
5.2 The model	11
5.3 One loss per output	12
5.4 Results	13
6. Summary	15
7. Sources	15
8. Key terms	16

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Keras workflow ([Note DL-011](#)); Vanishing gradient ([Note DL-018](#)).

1. Overview

Key point: Keras’ **Sequential** model is a straight line of layers: one input, one output. The **functional API** connects layers like a graph: a model can have several inputs, several outputs, branches that split and join, and skip connections. Each layer is called on the output of the layer before it, and `keras.Model(inputs, outputs)` collects the graph into a model.

Every network in the deep learning Notes so far, ANN or CNN, was built with Keras’ **Sequential** class, a **Sequential model** (G-1776). A Sequential model is a list of layers: data enters the first layer, passes through each layer in turn, and leaves the last. The Keras documentation calls it “appropriate for a plain stack of layers where each layer has exactly one input tensor and one output tensor” (Keras documentation, The Sequential model); Chollet describes it as “basically a Python list” and therefore “limited to simple stacks of layers” (Chollet and Watson 2025, ch. 7).

Many useful networks are not straight lines (Figure 1). The pattern of connections between a network’s layers is its **topology** (G-1991). This Note introduces the **functional API** (G-818), the Keras way of building them: “The functional API can handle models with non-linear topology, shared layers, and even multiple inputs or outputs” (Keras documentation, The Functional API).




Figure 1: (a) A Sequential model is one straight line. (b) One input, two outputs: a face photo gives both an age and an emotion. (c) Three inputs, one output: product details, a text description and a photo together give a price

This Note covers:

- when Sequential is not enough (section 3);
- how to build a functional model: two outputs, two inputs and a skip connection (section 4);
- a real model that predicts a person's age and gender from one face photo (section 5).

2. Prerequisites

- [Customer churn ANN Note](#): building, compiling and training a Keras **Sequential** model.
- [Loss functions Note](#): mean absolute error and binary cross-entropy.
- [Transfer learning Note](#): VGG16's frozen convolutional base with a new dense head.
- [Vanishing gradients Note](#): residual blocks.

3. When a straight line is not enough

Key point: Sequential assumes one input, one output and layers in a single line. Problems with several outputs or several kinds of input break that assumption.

A Sequential model rests on three assumptions:

1. it has exactly **one input**;
2. it has exactly **one output**;


3. its layers form a **single line**: each layer takes the output of the layer before it, and only that.

Two examples break them.

Two outputs from one photo. We have 25,000 face photos, and for each we want to predict two things: the person's age (a number, so **regression**, G-1655) and the emotion shown, such as happy, sad or angry (a class, so **classification**, G-395). We could train two separate networks, one per target (Figure 2a). A better design passes the photo through one CNN, which finds the facial features once, and then splits into two branches: one ends in an age, the other in an emotion (Figure 1b). A model that predicts several targets at once, each from its own output layer, is a **multi-output model** (G-1272).

Three inputs for one output. An online shop wants to predict the price of a product from three things: its details (a table of numbers such as size and weight), its text description and its photo. A naive approach trains three networks, one per input, and averages their three prices (Figure 2b). A better design gives each input the kind of layers that suits it, dense layers for the table, an RNN for the text (the [why RNNs Note](#)), a CNN for the photo, then **concatenates** the three results into one vector and predicts a single price from it (Figure 1c). A model that takes several inputs is a **multi-input model** (G-1269).

(a) Two separate networks



(b) Three networks, averaged




Figure 2: The naive designs. (a) Two separate networks, one per target: the photo passes through two CNNs. (b) Three networks, one per input, whose prices are averaged: no network sees the three inputs together. Compare Figure 1b and 1c

In Figure 2b, each network guesses a price from one kind of input alone, and the average only mixes three separate guesses. In Figure 1c, the layers after **concatenate** see all three inputs at once, so they can learn how the inputs work together.

Neither network is a single line, so neither can be built with `Sequential`. The Keras guide

lists exactly these cases: a Sequential model is not appropriate when the model has multiple inputs or multiple outputs, or when we want a non-linear topology, “e.g. a residual connection, a multi-branch model” (Keras documentation, The Sequential model). “The functional API makes it easy to manipulate multiple inputs and outputs. This cannot be handled with the Sequential API” (Keras documentation, The Functional API).

4. Building a functional model

Key point: Create an `Input`, call each layer on the output of the layer it should receive from, and finally pass the inputs and outputs to `keras.Model`.

4.1 Two outputs from one input

Key point: Three features in, two hidden layers, then a branch: one output for age (linear) and one for city (sigmoid).

Take a small table with three **features** (G-772; input variables, one per column): yearly salary, height and marital status. From them we want to predict two **targets** (G-1949; the outputs we predict) for each **observation** (G-1374; one person, one row of the table): the person’s age, and whether the person lives in Delhi or Mumbai. The network has two hidden layers of 128 and 64 nodes, then branches into two output nodes (Figure 3).




Figure 3: A functional model with one input and two outputs. Both outputs read the second hidden layer. The numbers are each layer’s parameters

Python: The model of Figure 3.

```
from keras import Model
```

```

from keras.layers import Input, Dense

x = Input(shape=(3,))                                # 3 features
hidden1 = Dense(128, activation="relu")(x)            # hidden1
    ↪ receives x
# hidden2 receives hidden1
hidden2 = Dense(64, activation="relu")(hidden1)
output1 = Dense(1, activation="linear", name="age")(hidden2)
output2 = Dense(1, activation="sigmoid", name="place")(hidden2)
model = Model(inputs=x, outputs=[output1, output2])
model.summary()
Dense(128, activation="relu") creates a layer; the second pair of brackets, (x), calls
it on x and returns its output. That call is the connection: it says where the layer's input
comes from. Model(inputs=..., outputs=...) then collects every layer between the
inputs and the outputs. With two outputs, outputs is a list.

```

The pattern is always the same: **create a layer, then call it on its input**. Figure 4 runs the code of Figure 3 line by line. Watch each call add one node and one edge to the graph, until `Model` collects them all.




Figure 4: Building the model of Figure 3, one line of code per frame. Each call adds a node (black border) and the edge from its input; the blue number is the layer’s parameter count, and the running total reaches 8,898

The age output uses a linear activation, because age is a number; the place output uses a **sigmoid** (G-1798), because it is a probability between 0 and 1 (Delhi = 0, Mumbai = 1).

`model.summary()` shows a fourth column that a Sequential summary does not have: for each layer, the layer it is connected to. It reports 8,898 **parameters** (G-1065, learnable parameters; the weights and biases that training learns) in total (Notebook), the total that Figure 4 builds up. Each comes from the parameter count of a **dense layer** (G-583).

1. **In words:** every node has one weight per input plus one bias.
2. **Formula:** parameters = inputs $\times$ nodes + nodes.
3. **Example:** hidden1 has $3 \times 128 + 128 = 512$; hidden2 has $128 \times 64 + 64 = 8,256$; each output has $64 \times 1 + 1 = 65$. Together: $512 + 8,256 + 65 + 65 = 8,898$.

Extra: Keras can draw a model as a graph:

```
keras.utils.plot_model(model, "model.png", show_shapes=True)
```

The option `show_shapes=True` adds each layer's input and output shapes (Keras documentation, The Functional API). The function needs the `pydot` package and the `Graphviz` program installed. The figures of this Note are drawn separately.

4.2 Two inputs, one output

Key point: Two inputs each go through their own dense layers; a **Concatenate** layer joins the two results end to end; the joined vector continues to one output.

The second shape has two inputs: one of 32 numbers and one of 128. Each goes through its own small stack of dense layers, and the two results are joined (Figure 5a).





Figure 5: (a) Two inputs processed separately and joined by **Concatenate**. (b) A skip connection: the input jumps over two convolution layers and is added to their output

The joining layer is **Concatenate** (G-71).

1. **In words:** concatenation places two vectors one after the other to make a longer vector. Nothing is added or multiplied. For example, joining $(1, 2)$ and $(3, 4, 5)$ gives $(1, 2, 3, 4, 5)$, of length $2 + 3 = 5$.
2. **Formula:** with $a_1, \dots, a_m$ the m numbers of the first vector and $b_1, \dots, b_n$ the n numbers of the second, $\text{concat}((a_1, \dots, a_m), (b_1, \dots, b_n)) = (a_1, \dots, a_m, b_1, \dots, b_n)$, of length $m + n$.
3. **Example:** $(0.2, 0.0, 1.3, 0.7)$ and $(0.5, 0.9, 0.0, 0.1)$ give $(0.2, 0.0, 1.3, 0.7, 0.5, 0.9, 0.0, 0.1)$: 8 numbers (Figure 6a).

(a) **Concatenate: one after the other**



(b) **Add: element by element**




Figure 6: The same two vectors joined two ways. (a) **Concatenate** places them one after the other: 8 numbers. (b) **Add** sums them position by position: 4 numbers, so the two shapes must match

Python: The model of Figure 5a.

```
from keras.layers import Concatenate
input_a = Input(shape=(32,))
input_b = Input(shape=(128,))
a = Dense(8, activation="relu")(input_a)
a = Dense(4, activation="relu")(a)
b = Dense(64, activation="relu")(input_b)
b = Dense(32, activation="relu")(b)
b = Dense(4, activation="relu")(b)
combined = Concatenate()([a, b])
# 4 + 4 = 8
↪ numbers
z = Dense(2, activation="relu")(combined)
z = Dense(1, activation="linear")(z)
model2 = Model(inputs=[input_a, input_b], outputs=z)
```

Reusing the names `a` and `b` is fine: each line takes the previous output and returns a new one. With two inputs, `inputs` is a list, and `predict` takes a list of two arrays, one per input.

The model has 10,789 parameters (Notebook). To predict, we pass both inputs together: `model2.predict([array_a, array_b])`.

4.3 A skip connection

Key point: A skip connection adds a block’s input to the block’s output. The functional API expresses it with an Add layer that receives two tensors.

A **skip connection** (G-1681; also called a **residual connection**, G-1681) lets the input of a block jump over the block’s layers and be added to their output (Figure 5b). The ResNet networks of the [pretrained models Note](#) are built from such blocks, and the [vanishing gradients Note](#) explains why they help very deep networks train. “A common use case for this is residual connections” of models that are “not connected sequentially, which the Sequential API cannot handle” (Keras documentation, The Functional API).

Python: A block of two convolutions with a skip connection.

```
from keras.layers import Conv2D, Add
inputs = Input(shape=(32, 32, 16))
y = Conv2D(16, 3, padding="same", activation="relu")(inputs)
y = Conv2D(16, 3, padding="same")(y)
# element by element: inputs + y
```

```
outputs = Add()([inputs, y])
```

```
block = Model(inputs, outputs)
```

Add adds its inputs element by element (Figure 6b), so they must have the same shape. padding="same" and 16 filters keep y at $32 \times 32 \times 16$, the shape of inputs.

The Add layer receives two tensors, inputs and y: the input reaches it by two routes, which is not a single line. The block has $2 \times (3 \times 3 \times 16 \times 16 + 16) = 4,640$ parameters, all in the two convolutions (Notebook).

5. A real model: age and gender from one face photo

Key point: One frozen VGG16 base reads the face; two dense branches predict age (mean absolute error) and gender (binary cross-entropy). It predicts age to within 9.2 years on average and gender with 86.3% accuracy, about as well as two separate models, while each photo passes through the 14.7-million-parameter base only once.

5.1 The data

Key point: UTKFace: 23,705 usable face photos with the person’s age (1 to 116) and gender in the file name.

UTKFace is a public dataset of over 20,000 face photos, with “annotations of age, gender, and ethnicity” and ages from 0 to 116 (Zhang et al. 2017; UTKFace dataset page). We use its aligned and cropped version, 200×200 pixels per face. The labels sit in each file name, [age]_[gender]_[race]_[date&time].jpg, with gender 0 for male and 1 for female. The dataset may be used for non-commercial research only, so this Note stores no face photos.

Of the 23,708 files, three lack a field in their name and are skipped. The remaining 23,705 photos have ages from 1 to 116, with a median of 29, and 47.7% of them show women

(Notebook). We split them at random into 18,964 training photos (80%) and 4,741 test photos (20%), and resize each to 128×128 .

The task has two **targets**: age, a number (regression), and gender, one of two classes (classification). It is the face example of section 3, with gender in place of emotion.

5.2 The model

Key point: VGG16's frozen convolutional base, Flatten, then two branches of 256 nodes, one ending in a linear age output and one in a sigmoid gender output.

The model combines **transfer learning** (G-2005; the [transfer learning Note](#)) with section 4: **VGG16's** (G-2088) convolutional base, frozen, reads the photo, and the network then branches into two heads (Figure 7).




Figure 7: The age and gender model: one frozen VGG16 base, Flatten, and two branches with their own output and loss

Python: The age and gender model.

```
from keras.applications.vgg16 import VGG16
conv_base = VGG16(weights="imagenet", include_top=False,
                    input_shape=(128, 128, 3))
conv_base.trainable = False

flat = layers.Flatten()(conv_base.output)
age_hidden = Dense(256, activation="relu")(flat)
gender_hidden = Dense(256, activation="relu")(flat)
age = Dense(1, activation="linear", name="age")(age_hidden)
gender = Dense(1, activation="sigmoid", name="gender")(gender_hidden)
```

```
model = Model(inputs=conv_base.input, outputs=[age, gender])
```

`conv_base.input` and `conv_base.output` connect the pretrained base into the new graph: the model's input is VGG16's input, and Flatten is called on VGG16's output. The names "age" and "gender" are used below to give each output its own loss.

For a 128×128 photo the base outputs $4 \times 4 \times 512$, which a **flatten layer** (G-788) turns into 8,192 numbers. Each branch then has $(8,192 + 1) \times 256 + (256 + 1) = 2,097,665$ parameters. The whole model has 18,910,018 parameters, of which 4,195,330 (the two branches) are trainable and the 14,714,688 of the base are frozen (Notebook).

5.3 One loss per output

Key point: Each output gets its own loss: mean absolute error for age, binary cross-entropy for gender. Training minimises their weighted sum.

A model with two outputs needs two losses, one per output. Keras lets us pass them as a dictionary keyed by the output names, together with **loss weights** (G-1131) that “modulate their contribution to the total training loss” (Keras documentation, The Functional API).

Python: Compiling and training the two-output model.

```
model.compile(optimizer="adam",
              loss={"age": "mae", "gender": "binary_crossentropy"},
              loss_weights={"age": 0.1, "gender": 1.0},
              metrics={"age": ["mae"], "gender": ["accuracy"]})
model.fit(X_train, {"age": age_train, "gender": gender_train},
         epochs=15, batch_size=64)
```

The targets are passed as a dictionary with the same keys, so Keras knows which target belongs to which output.

1. **In words:** the total loss is each output's loss times its weight, added up (Figure 8). Age errors are measured in years, which are large numbers, while binary cross-entropy is usually below 1; a small weight on age keeps it from swamping the gender loss.
2. **Formula:** $L = 0.1 \times L_{\text{age}} + 1.0 \times L_{\text{gender}}$
3. **Example:** with an age error of 9 years and a gender loss of 0.30, $L = 0.1 \times 9 + 1.0 \times 0.30 = 0.9 + 0.3 = 1.2$. Without the weight, age would contribute 9 of the 9.3.




Figure 8: The worked example as bars. Without weights, the age part (9.0) is 30 times the gender part (0.3), so training would mostly reduce the age error. With a weight of 0.1 on age, the two parts are 0.9 and 0.3

The **mean absolute error** (G-1194; MAE) is the average distance between predicted and true age, in years (the [loss functions Note](#)). The gender output uses **binary cross-entropy** (G-303), the loss for a yes-or-no target with a sigmoid output.

Extra: The base is frozen and the photos are not augmented, so the base turns each photo into the same 8,192 numbers in every epoch. The Notebook therefore runs every photo through the base once and trains only the two branches on the stored numbers. The model is the same; only the repeated work is skipped.

5.4 Results

Key point: Both outputs learn at the same time. The two-output model matches two separate single-output models: a mean age error of about 9 years against 15 for guessing, and 86% gender accuracy against 53%.

To see whether one model with two outputs loses anything, we also trained two separate models with one branch each: one for age only, one for gender only. All three were trained 15 epochs with three seeds and scored on the 4,741 test photos (Notebook):

Model (mean of 3 seeds)	Age: mean absolute error	Gender: accuracy	Trainable parameters
One model, two outputs	9.15 years (8.60–9.78)	86.3% (86.3–86.3%)	4,195,330
Age only	8.65 years (8.33–9.06)	–	2,097,665
Gender only	–	85.9% (84.4–86.8%)	2,097,665

Model (mean of 3 seeds)	Age: mean absolute error	Gender: accuracy	Trainable parameters
Guessing: the median training age; the commoner gender	15.03 years	52.5%	—




Figure 9: The two-output model on the test photos, epoch by epoch: the age error (left) and the gender accuracy (right). Thin lines: 3 seeds; thick lines: their mean; dashed: guessing

Both heads learn from the first epoch (Figure 9): after one epoch the age error is already 9.0 years and the gender accuracy 85.6%, far better than guessing (dashed lines). Later epochs change little: the mean age error moves between 8.4 and 9.6 years from epoch to epoch, and the gender accuracy between 85.0% and 86.3%. The two-output model does as well as the two separate models: its gender accuracy is slightly higher, its age error slightly higher too, and in both cases the ranges of the three seeds overlap.

The two branches in this model share no trainable layer, only the frozen base, so each head learns its own target much as a separate model would; the results agree with that. What the single model saves is the base: one photo passes once through VGG16’s 14.7 million convolution weights and gives both answers. Two separate models would each need their own copy of the base and their own pass through it, $2 \times (14,714,688 + 2,097,665) = 33,624,706$ parameters in total, against 18,910,018 for the one model.

6. Summary

Shape	Keras building block	Example in this Note
One input, one output, a straight line	Sequential	every earlier Note
Several outputs	one layer feeding several output layers	salary, height, marital status to age and city; face to age and gender
Several inputs	several Inputs joined by Concatenate	two inputs of 32 and 128 numbers
A skip connection	Add of a block's input and output	two convolutions with the input added back

- The functional API builds a model as a graph: create a layer, call it on its input, then `Model(inputs, outputs)`.
- Multi-output models take one loss per output, combined with loss weights.
- On 23,705 face photos, one VGG16-based model predicted age (mean error 9.2 years) and gender (86.3%) as well as two separate models, with one pass through the base.

7. Sources

Built from

- CampusX, “Keras Functional Model | How to build non-linear Neural Networks?”, YouTube, <https://www.youtube.com/watch?v=OvQQP1QVru8>

Other references

- Keras documentation: The Sequential model (keras.io/guides/sequential_model): when a Sequential model is and is not appropriate.
- Keras documentation: The Functional API (keras.io/guides/functional_api): non-linear topology, shared layers, multiple inputs and outputs; a model as a directed acyclic graph of layers; `plot_model`; different losses and loss weights per output; residual connections.
- Chollet, F. and Watson, M. (2025). *Deep Learning with Python*, 3rd ed. Manning. Ch. 7 (the Sequential model “limited to simple stacks of layers”; the functional API for multiple inputs, multiple outputs or a nonlinear topology).
- Zhang, Z., Song, Y. and Qi, H. (2017). Age Progression/Regression by Conditional Adversarial Autoencoder. *CVPR 2017*. arXiv:1702.08423. (The UTKFace dataset.)
- UTKFace dataset page, susanqq.github.io/UTKFace: over 20,000 face photos, ages 0 to 116, labels in the file name (gender 0 = male, 1 = female); non-commercial research use only.
- Simonyan, K. and Zisserman, A. (2015). Very Deep Convolutional Networks for Large-Scale Image Recognition. *ICLR 2015*. arXiv:1409.1556. (VGG16.)
- He, K., Zhang, X., Ren, S. and Sun, J. (2016). Deep Residual Learning for Image Recognition. *CVPR 2016*. arXiv:1512.03385. (Residual blocks: the input added to the block's output.)

8. Key terms

Term	Meaning
Feature	An input variable: one column of the data table
Target	The output we predict
Observation	One record: one row of the data table
Sequential model	A Keras model whose layers form a single line: one input, one output
Functional API	The Keras way of building a model as a graph of layers, by calling each layer on its input
Topology	The pattern of connections between a network's layers
Multi-output model	A model that predicts several targets at once, each from its own output layer
Multi-input model	A model that takes several inputs, such as a photo and a table
Input (G-97)	The Keras object that stands for a model's input and fixes its shape
Concatenate	A layer that joins vectors end to end into one longer vector
Skip (residual) connection (G-1681)	A connection that adds a block's input to the block's output
Add	A layer that adds tensors of the same shape element by element
Loss weight	A number that scales one output's loss in the total loss of a multi-output model