

RNN Architecture and Forward Propagation

Contents

1. Overview	2
2. Prerequisites	3
3. The shape of the input	3
3.1 Words become vectors	3
3.2 Time steps	4
3.3 A batch for Keras	4
4. The architecture of an RNN	5
4.1 The smallest RNN: one node	5
4.2 Two differences from an ANN	7
4.3 The network for our reviews	7
4.4 Counting the parameters	8
5. Forward propagation in an RNN	9
5.1 Unfolding through time	9
5.2 Time step by time step	10
5.3 The formulas	10
5.4 Gotcha: padded steps still run	13
6. Three properties of an RNN	14
6.1 Why “recurrent”	14
6.2 Weight sharing	14
6.3 The hidden state carries the sequence forward	14
7. The RNN as one box	15
8. Summary	17
9. Sources	17
10. Key terms	17

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Forward propagation; Exploding gradient and gradient clipping; Tanh; Sequential data; Sequence padding; Tokenization and integer encoding of text.
- **Leads to:** Types of RNN (many-to-one, one-to-many, many-to-many); Backpropagation through time (BPTT); LSTM (long short-term memory); GRU (gated recurrent unit); Deep (stacked) RNNs; Bidirectional RNNs.
- **Compare with:** Multi-layer perceptron (MLP).

1. Overview

Key point: A recurrent neural network (RNN) reads a sequence one element at a time. Its hidden layer sends its own output back to itself, so at every time step it combines the new input with a summary of everything read so far.

A **recurrent neural network** (G-1647; RNN) is a class of neural networks with a memory: it remembers past inputs, which makes it work well on sequential data. [Why an ordinary network handles sequences badly](#) was explained earlier, with [the idea an RNN uses instead](#). This Note opens the box: what an RNN looks like, how many parameters it has, and how it turns a sequence into a prediction.




Figure 1: The same recurrent layer reads “movie was good”, one word per time step. Each step mixes the new word with the previous hidden state (red arrow) and produces a new hidden state of 3 numbers (blue for positive, red for negative). The last hidden state gives the prediction

Figure 1 shows the whole idea. The three purple boxes are not three layers: they are the same layer, with the same weights, used again at each time step.

2. Prerequisites

- [Sequential data](#), and [why an ANN struggles with it](#).
- [What one node computes](#): a layer computes a weighted sum, adds a bias and applies an activation.
- [One-hot encoding](#): a category becomes a vector (a list of numbers) with a single 1.
- [Tensors and their shapes](#): shapes such as $(3, 4, 5)$, a grid with 3 layers of 4 rows and 5 columns.
- [Tanh](#) and [sigmoid](#): activation functions that squeeze a number into -1 to 1 and 0 to 1.

3. The shape of the input

Key point: An RNN takes each **observation** (one review) as a table of shape (time steps, input features): one row per word, one number per feature. Keras takes a batch of them as a 3D tensor of shape (batch size, time steps, input features).

3.1 Words become vectors

Key point: A network needs numbers, so each word becomes a vector. The simplest way is one-hot encoding over the vocabulary.

Take a sentiment analysis task: the input is a movie review, and the **target** (G-1949; the output we predict) is its sentiment, 1 for positive and 0 for negative. Each review is one **observation** (G-1374; one record of the data). Three tiny reviews:

Review	Sentiment
movie was good	1
movie was bad	0
movie was not good	0

These reviews use 5 unique words, the **vocabulary** (G-2092): movie, was, good, bad, not. With **one-hot encoding** (G-1379), each word becomes a vector of 5 numbers with a single 1:

Word	Vector
movie	[1, 0, 0, 0, 0]
was	[0, 1, 0, 0, 0]
good	[0, 0, 1, 0, 0]
bad	[0, 0, 0, 1, 0]
not	[0, 0, 0, 0, 1]

Each of the 5 positions is one input **feature** (G-772; one input variable). Real projects use a larger vocabulary and better word vectors; [integer encoding and embeddings](#) do so on real reviews.

3.2 Time steps

Key point: The words of a review enter the RNN one by one. The first word enters at **time step** (G-1976) $t = 1$, the second at $t = 2$, and so on.

We write x_{ij} for word j of review i . So the first review “movie was good” is x_{11}, x_{12}, x_{13} , and each x_{ij} is a vector of 5 numbers.

- “movie was good” has shape (3, 5): 3 time steps, 5 input features.
- “movie was bad” has shape (3, 5).
- “movie was not good” has shape (4, 5).

In general the shape of one observation is (time steps, input features).

3.3 A batch for Keras

Key point: Keras’ SimpleRNN layer takes (batch size, time steps, input features). Shorter reviews are padded to the longest one.

Keras processes several reviews at once. The three reviews above, sent together, form a **batch** (G-263): a **tensor** (G-1957; a grid of numbers with any number of axes) of shape (3, 4, 5), with 3 reviews, 4 time steps (the longest review has 4 words) and 5 input features (Figure 2). The two 3-word reviews get one zero vector as **padding** (G-1436) (padding and its cost are covered in [zero padding wastes most of the network](#)).




Figure 2: The batch of the three reviews. Each review is a table with one row per time step and one column per input feature; each row is a word’s one-hot vector. The 3-word reviews get a grey row of zeros, so the batch has shape (3, 4, 5)

Python: One-hot vectors for each review, padded into one batch.

```
import numpy as np, keras
vocab = ["movie", "was", "good", "bad", "not"]
one_hot = {w: np.eye(5)[i] for i, w in enumerate(vocab)}
reviews = ["movie was good", "movie was bad", "movie was not good"]
X = [np.array([one_hot[w] for w in r.split()]) for r in reviews]
batch = keras.utils.pad_sequences(X, dtype="float32", padding="post")
print(batch.shape)           # (3, 4, 5)
```

Padding has a cost inside the RNN too; section 5.4 shows it once the RNN's formula is known.

4. The architecture of an RNN

Key point: An RNN looks like an ANN with one hidden layer, with two differences: the input arrives one time step at a time, and the hidden layer feeds its own output back to itself at the next time step.

4.1 The smallest RNN: one node

Key point: One node with a feedback loop is already an RNN. At each step the node adds the new input, times w_i , to its own previous output, times w_h . Unrolling the loop gives one copy of the node per input, all with the same weights.

Before the network for the reviews, take the smallest case: a sequence of single numbers. A share price is low (0), medium (0.5) or high (1), and in this toy market it follows a simple rule:

Yesterday	Today	Tomorrow
low (0)	low (0)	low (0)
low (0)	medium (0.5)	high (1)
high (1)	medium (0.5)	low (0)
high (1)	high (1)	high (1)

Today's price alone is not enough: "medium" is followed by high in one row and by low in another. The prediction needs yesterday's price too, so the network must remember it.

The network is one node with a **ReLU** (G-1668; it turns negative numbers into 0 and keeps positive ones) activation and three weights, chosen by hand: $w_i = 1.8$ on the input, $w_h = -0.5$ on the **feedback loop** (the node's output, sent back into the node at the next step) and $w_o = 1.1$ on the output. All biases are 0. Figure 3 runs the third row of the table.




Figure 3: The smallest RNN on “yesterday high, today medium”. First the node with its feedback loop; then the loop unrolled into two copies of the node. The red arrow carries h_1 , times w_h , into the second copy, where it cancels today’s input, and the prediction is 0: low (idea after StatQuest, “Recurrent Neural Networks (RNNs), Clearly Explained!!!”)

1. **In words:** the prices go in one at a time, oldest first. At each step the node multiplies the new price by w_i , adds its own previous output times w_h , and applies ReLU. After the last price, the output weight turns the node’s output into the prediction.
2. **Formula:** with $h_0 = 0$,

$$h_t = \text{ReLU}(x_t w_i + h_{t-1} w_h)$$

$$\hat{y} = h_2 w_o$$

3. **Example:** yesterday high, today medium, so $x_1 = 1$ and $x_2 = 0.5$:

$$h_1 = \text{ReLU}(1 \times 1.8) = 1.8$$

$$h_2 = \text{ReLU}(0.5 \times 1.8 + 1.8 \times (-0.5))$$

$$h_2 = \text{ReLU}(0.9 - 0.9) = 0$$

$$\hat{y} = 0 \times 1.1 = 0$$

The prediction is 0: low, as the rule says.

The same three weights give all four rows:

Yesterday, today	h_1	h_2	$\hat{y}$	Rule
0, 0	0	0	0	low
0, 0.5	0	0.9	0.99	high
1, 0.5	1.8	0	0	low

Yesterday, today	h_1	h_2	$\hat{y}$	Rule
1, 1	1.8	0.9	0.99	high

Compare rows 2 and 3. Today's price is the same, 0.5, but h_1 differs, and the feedback loop carries that difference into h_2 . The loop is the network's memory.

Three points from this toy carry over to every RNN:

- **Unrolling.** Following a loop is confusing, so we draw one copy of the node per input and turn the loop into an arrow from each copy to the next. This drawing is **unfolding** (G-2041; also called unrolling). With 3 days of prices we draw 3 copies; with 50 days, 50 copies.
- **Shared weights.** Every copy uses the same w_i and w_h . Unrolling for more days adds copies to the drawing but no new weights: this network has 3 weights for any number of days.
- **Only the last output is used.** The first copy could also give an output, a prediction for today. We already know today's price, so we ignore that output and read the prediction after the last input only.

The rest of this Note builds the same thing with vectors in place of single numbers: several inputs per step, several nodes in the loop.

4.2 Two differences from an ANN

Key point: Input by time step, and a feedback connection. The feedback connection is what makes an RNN an RNN.

An ANN has an input layer, one or more hidden layers and an output layer. An RNN has the same three parts, with two differences.

1. **The input arrives one time step at a time.** An ANN takes the whole input at once. An RNN takes x_{11} at $t = 1$, then x_{12} at $t = 2$, then x_{13} at $t = 3$.
2. **The hidden layer feeds back to itself.** An ANN is a **feed-forward network** (G-775): information only moves from input to output. In an RNN, the hidden layer's output at one time step becomes an extra input to the same layer at the next time step. This fed-back output is the network's **state**.

The hidden layer with this feedback is the **recurrent layer** (G-1646), and its output at time t is the **hidden state** (G-891) h_t .

4.3 The network for our reviews

Key point: 5 input nodes (one per feature), a recurrent layer of 3 nodes, and 1 sigmoid output node. Three weight matrices: W_i (5×3), W_h (3×3) and W_o (3×1).




Figure 4: The RNN for the reviews: 5 input nodes, a recurrent layer of 3 nodes and 1 sigmoid output node. The red arrow is the feedback: the layer's 3 outputs return to all 3 nodes at the next time step, through W_h

Figure 4 shows the network.

- **Input layer: 5 nodes**, because every word is 5 numbers.
- **Recurrent layer: 3 nodes** here; any number works. Between the inputs and these nodes the layer is fully connected, as in an ANN.
- **Output layer: 1 node** with a sigmoid activation, because the task is binary classification (positive or negative).

The feedback connection is the new part. Each of the 3 recurrent nodes sends its output to **all 3** recurrent nodes at the next time step, each through its own weight.

4.4 Counting the parameters

Key point: 27 weights plus 4 biases: 31 trainable parameters.

$$15 + 9 + 3 = 27 \text{ weights}$$

$$3 + 1 = 4 \text{ biases}$$

$$27 + 4 = 31$$

Connection	Matrix	Shape	Count
input to recurrent layer	W_i	5×3	15
recurrent layer to itself (feedback)	W_h	3×3	9
recurrent layer to output	W_o	3×1	3

Connection	Matrix	Shape	Count
biases of the recurrent layer	b_h	3	3
bias of the output node	b_o	1	1
Total			31

Keras agrees, though it groups the parameters by layer, not by kind. The `SimpleRNN(3)` layer holds W_i , W_h and b_h :

$$15 + 9 + 3 = 27$$

The `Dense(1)` layer on top holds W_o and b_o :

$$3 + 1 = 4$$

So 27 appears twice by coincidence: 27 weights in the whole network, and 27 parameters in the `SimpleRNN` layer.

Python: The same network in Keras.

```
model = keras.Sequential([
    keras.Input(shape=(4, 5)),          # (time steps, features)
    keras.layers.SimpleRNN(3),
    keras.layers.Dense(1, activation="sigmoid")])
model.summary()                        # 27 + 4 = 31 parameters
for w in model.get_weights():
    print(w.shape)                     # (5, 3) (3, 3) (3,) (3, 1) (1,)
```

`get_weights()` returns the parameters in the order W_i (Keras calls it the `kernel`), W_h (the `recurrent_kernel`), b_h , W_o , b_o .

5. Forward propagation in an RNN

Key point: At each time step the recurrent layer computes $h_t = \tanh(x_t W_i + h_{t-1} W_h + b_h)$, with the same weights every time. After the last word, the output layer turns the last hidden state into the prediction.

5.1 Unfolding through time

Key point: Drawing the recurrent layer once per time step, side by side, turns the loop into a chain. The chain is the same layer used again and again.

The feedback loop of Figure 4 is hard to follow. **Unfolding** (G-2041; also called unrolling) redraws it: one copy of the recurrent layer for each time step, with an arrow carrying the hidden state from each copy to the next (Goodfellow §10.1). Figure 1 is the unfolded network for “movie was good”. Every copy uses the same W_i , W_h and b_h .

5.2 Time step by time step

Key point: Step 1 uses the first word and h_0 ; every later step uses the next word and the previous hidden state.

We feed the first review, x_{11}, x_{12}, x_{13} , and write x_t for the word at time t . Every vector is a row: x_t is 1×5 and h_t is 1×3 .

- $t = 1$: the first word goes through W_i . The shapes multiply as

$$(1 \times 5)(5 \times 3) = 1 \times 3$$

The recurrent layer applies its activation, by default **tanh** (G-1947; a function that squeezes a number into -1 to 1), to get h_1 , shape 1×3 : one output per node.

- $t = 2$: the second word enters through the same W_i . The layer also receives h_1 , through W_h :

$$(1 \times 3)(3 \times 3) = 1 \times 3$$

Both products are 1×3 , so they can be added, and tanh of the sum is h_2 .

- $t = 3$: the same again with x_3 and h_2 , giving h_3 .
- **Output:** after the last word, $h_3 W_o$ is a single number:

$$(1 \times 3)(3 \times 1) = 1 \times 1$$

The sigmoid turns it into the prediction $\hat{y}$.

At $t = 1$ there is no previous hidden state. To keep every step the same, we give the layer h_0 , a vector of zeros. Keras uses zeros by default.

5.3 The formulas

Key point: One formula for the hidden state, used at every step, and one for the output, used after the last step.

1. **In words:** multiply the new word by the input weights and the previous hidden state by the recurrent weights, add the two and the bias, then apply tanh. After the last time step T , multiply the hidden state by the output weights, add the output bias and apply the output activation g .

2. **Formula:**

$$h_t = \tanh(x_t W_i + h_{t-1} W_h + b_h)$$

$$h_0 = 0$$

$$\hat{y} = g(h_T W_o + b_o)$$

For binary classification g is the sigmoid; for several classes it is the **softmax** (G-1830; a function that turns scores into probabilities that add to 1); for regression it is linear (no activation). The recurrent layer can use another activation, such as ReLU, instead of tanh.

3. **Example:** small hand-picked weights, all biases 0, on “movie was good”.

$$W_i =$$

$$\begin{bmatrix} 0.2 & -0.1 & 0.0 \\ 0.0 & 0.1 & 0.1 \\ 0.8 & 0.3 & -0.5 \\ -0.8 & -0.3 & 0.5 \\ -0.6 & 0.2 & 0.4 \end{bmatrix}$$

$$W_h =$$

$$\begin{bmatrix} 0.5 & 0.0 & 0.1 \\ 0.2 & 0.4 & 0.0 \\ 0.0 & -0.3 & 0.5 \end{bmatrix}$$

$$W_o = \begin{bmatrix} 1.5 \\ 0.5 \\ -1.0 \end{bmatrix}$$

A one-hot vector times W_i simply picks one row of W_i : the row of that word. Figure 5 runs the three steps below; watch, at each step, the word pick its row on the left and the four rows of the computation fill in on the right.




Figure 5: The worked example, one operation per frame. Left: W_i , with the current word’s row framed. Right: the word’s row $x_t W_i$, the feedback $h_{t-1} W_h$, their sum, and h_t after tanh. Blue cells are positive and red cells negative. The last frame gives the prediction 0.83

- $t = 1$, “movie”: the word picks the first row of W_i , and h_0 is all zeros:

$$x_1 W_i = [0.2, -0.1, 0.0]$$

$$h_0 W_h = [0, 0, 0]$$

$$h_1 = \tanh([0.2, -0.1, 0.0])$$

$$h_1 = [0.197, -0.100, 0.000]$$

- $t = 2$, “was”: the word picks the second row of W_i :

$$x_2 W_i = [0.0, 0.1, 0.1]$$

The feedback, one product per line (each row of W_h times the matching number of h_1):

$$0.197 \times [0.5, 0.0, 0.1] = [0.0985, 0, 0.0197]$$

$$-0.100 \times [0.2, 0.4, 0.0] = [-0.020, -0.040, 0]$$

$$0 \times [0.0, -0.3, 0.5] = [0, 0, 0]$$

$$h_1 W_h = [0.079, -0.040, 0.020]$$

Adding $x_2 W_i$:

$$x_2 W_i + h_1 W_h = [0.079, 0.060, 0.120]$$

$$h_2 = \tanh([0.079, 0.060, 0.120])$$

$$h_2 = [0.079, 0.060, 0.119]$$

- $t = 3$, “good”: $x_3 W_i$ is the third row of W_i , so

$$x_3 W_i = [0.8, 0.3, -0.5]$$

The feedback $h_2 W_h$ with $h_2 = [0.079, 0.060, 0.119]$, one product per line (each row of W_h times the matching number of h_2):

$$0.079 \times [0.5, 0.0, 0.1] = [0.0395, 0, 0.0079]$$

$$0.060 \times [0.2, 0.4, 0.0] = [0.0120, 0.0240, 0]$$

$$0.119 \times [0.0, -0.3, 0.5] = [0, -0.0357, 0.0595]$$

$$h_2 W_h = [0.0515, -0.0117, 0.0674]$$

$$x_3 W_i + h_2 W_h$$

$$= [0.8 + 0.0515, 0.3 - 0.0117,$$

$$-0.5 + 0.0674]$$

$$= [0.851, 0.288, -0.433]$$

so

$$h_3 = \tanh([0.851, 0.288, -0.433])$$

$$h_3 = [0.692, 0.281, -0.407]$$

- **Output:**

$$0.692 \times 1.5 = 1.038$$

$$0.281 \times 0.5 = 0.141$$

$$-0.407 \times (-1.0) = 0.407$$

$$h_3 W_o = 1.038 + 0.141 + 0.407 = 1.585$$

$$\hat{y} = \sigma(1.585) = 0.83$$

Keras' `SimpleRNN`, given the same weights, returns exactly these hidden states and this prediction (Notebook). The weights here are chosen by hand, not trained, so the 0.83 only shows the computation.

Extra: Goodfellow §10.2 writes the same RNN with column vectors:

$$a^{(t)} = b + W h^{(t-1)} + U x^{(t)}$$

$$h^{(t)} = \tanh(a^{(t)})$$

$$o^{(t)} = c + V h^{(t)}$$

Their U , W and V are our W_i , W_h and W_o , transposed. We use row vectors because Keras stores its weights that way: the `kernel` of shape (input features, units) multiplies the input from the right.

5.4 Gotcha: padded steps still run

Key point: The RNN runs a step on every padded zero row, so a padded review's last hidden state is not the one after its last word. A mask makes the RNN skip padded steps.

The batch of section 3.3 gave the 3-word reviews a grey row of zeros as padding. The RNN does not know that the grey row is padding. It runs a fourth step on the zero row, so for the 3-word reviews the last hidden state is h_4 , computed from h_3 alone:

$$h_4 = \tanh(0 \cdot W_i + h_3 W_h + b_h)$$

$$h_4 = \tanh(h_3 W_h + b_h)$$

h_4 is in general not equal to h_3 . Keras' `Masking` layer (or `mask_zero=True` in an `Embedding` layer) tells the RNN to skip padded steps (Keras guide "Understanding masking and padding"). The worked example of section 5.3 feeds the 3 words without padding, so it stops at h_3 .

6. Three properties of an RNN

Key point: The layer recurs (hence the name), the weights are shared across time steps, and the last hidden state depends on every earlier input.

6.1 Why “recurrent”

Key point: The input changes at every time step, but the same layer computes every hidden state.

The hidden layer recurs: we change the input word by word and reuse one layer to compute every output. That reuse gives the network its name, recurrent neural network.

6.2 Weight sharing

Key point: Every time step uses the same W_i , W_h and b_h , so the number of parameters does not grow with the length of the sequence.

At every time step the input is new, but the weights are the same. This reuse is **parameter sharing** (G-1447; or weight sharing) across time. The network for a 3-word review and for a 300-word review has the same 31 parameters. Goodfellow §10.1 names this as an advantage of unfolding: the model has the same input size whatever the sequence length, because it is defined as a step from one state to the next.

6.3 The hidden state carries the sequence forward

Key point: h_3 contains a piece of h_2 , which contains a piece of h_1 . So the last hidden state depends on every word, in order.

Write out h_3 one step at a time (biases left out), putting each hidden state into the next:

$$h_1 = \tanh(x_1 W_i + h_0 W_h)$$

$$h_2 = \tanh(x_2 W_i + h_1 W_h)$$

$$h_3 = \tanh(x_3 W_i + h_2 W_h)$$

Replace h_2 and then h_1 in the last line, and h_3 is one nested expression:

$$h_3 = \tanh \left(x_3 W_i + \tanh \left(x_2 W_i + \tanh (x_1 W_i + h_0 W_h) W_h \right) W_h \right)$$

Every word appears, in its place in the sequence. Information from earlier words reaches the output only through the chain of W_h multiplications.

A test shows the role of W_h . Take two reviews that differ only in the middle word, “movie was good” and “movie not good”, and run the weights of section 5.3 (Notebook):

	h_3 for “movie was good”	h_3 for “movie not good”
with W_h	[0.692, 0.281, -0.407]	[0.532, 0.240, -0.336]
$W_h = 0$ (no feedback)	[0.664, 0.291, -0.462]	[0.664, 0.291, -0.462]

With the feedback connection, the middle word changes the final hidden state. Without it, h_3 sees only the last word “good”, and the two reviews look identical. The feedback connection is the memory.




Figure 6: The test of the table, step by step. Top row “movie was good”, bottom row “movie not good”, hidden states of 3 numbers (blue for positive, red for negative). With the feedback W_h (red arrows) the two h_3 differ; with $W_h = 0$ (crossed arrows) they are identical

In Figure 6, watch the middle column: “was” and “not” give different h_2 in both runs, but only the red arrows carry that difference on into h_3 .

How far back that memory reaches is limited in practice. Gradient-based training of a simple RNN struggles to learn dependencies across long spans: the probability of successful training of a traditional RNN with stochastic gradient descent rapidly reaches 0 for sequences of only length 10 or 20 (Goodfellow §10.7). [The vanishing gradient through time](#) explains why.

7. The RNN as one box

Key point: An RNN is a box with two inputs, the current word and the previous hidden state, and one output, the new hidden state. After the last word, the hidden state goes to the output layer.




Figure 7: The simplified view of a recurrent layer. The word x_t enters through W_i , the previous hidden state through W_h ; their sum, plus the bias, goes through $\tanh$ to give h_t , which is fed back at the next time step. After the last step, h_T goes through W_o and g to give $\hat{y}$

Figure 7 sums up the whole computation.

1. At every time step, the box receives the current word x_t (through W_i) and the previous hidden state h_{t-1} (through W_h).
2. It adds the two products and the bias, and applies the activation ($\tanh$ by default) to get h_t .
3. h_t goes back into the box at the next time step.
4. After the last time step T , h_T goes through W_o and the output activation g (sigmoid here) to give $\hat{y}$.

8. Summary

	ANN (feed-forward)	RNN
Input	the whole input at once	one time step at a time
Information flow	input to output only	hidden layer also feeds back to itself
Shape of one observation	(features)	(time steps, features)
Weights	one set per layer	the same W_i , W_h , b_h at every time step
Example parameters		W_i : 15, W_h : 9, W_o : 3, biases: 4, total 31

- An RNN reads a sequence one time step at a time; Keras takes (batch size, time steps, input features).
- The recurrent layer computes $h_t = \tanh(x_t W_i + h_{t-1} W_h + b_h)$, starting from $h_0 = 0$.
- After the last step, $\hat{y} = g(h_T W_o + b_o)$, with g the sigmoid, softmax or linear.
- The layer recurs with shared weights, so the parameter count does not depend on the sequence length.
- The feedback weights W_h carry earlier inputs into later hidden states: they are the network's memory.

9. Sources

Built from

- CampusX, “Recurrent Neural Network | Forward Propagation | Architecture”, YouTube, <https://www.youtube.com/watch?v=BjWqCcbusMM>
- StatQuest with Josh Starmer, “Recurrent Neural Networks (RNNs), Clearly Explained!!!”, YouTube, <https://www.youtube.com/watch?v=AsNTP8Kwu80>. 04:30–11:00 (a one-node RNN on scaled share prices; unrolling the feedback loop; ignoring the earlier outputs; weights shared across the copies).

Other references

- Goodfellow, I., Bengio, Y. and Courville, A. (2016). *Deep Learning*. MIT Press. Chapter 10, §10.1 (unfolding computational graphs), §10.2 (recurrent neural networks), §10.7 (the challenge of long-term dependencies). deeplearningbook.org/contents/rnn.html.
- Keras API documentation: **SimpleRNN** layer, keras.io/api/layers/recurrent_layers/simple_rnn (default activation tanh, initial state zeros).
- Keras guide, “Understanding masking and padding”, keras.io/guides/understanding_masking_and_padding (the **Masking** layer and **mask_zero** let sequence layers skip padded time steps).

10. Key terms

Term	Meaning
Recurrent neural network (RNN)	A neural network that reads a sequence one time step at a time and feeds its hidden layer's output back to itself
Observation	One record of the data, here one review
Feature	An input variable; here one of the 5 positions of a word vector
Target	The output we predict, here the sentiment
Vocabulary (G-2092)	The set of unique words in the data
Time step	One position in the sequence; word j enters at $t = j$
Recurrent layer	A hidden layer whose output at one time step is an input to itself at the next
Hidden state (h_t)	The recurrent layer's output at time step t ; the network's summary of the inputs so far
W_i, W_h, W_o	Input weights, recurrent (feedback) weights and output weights of an RNN
Unfolding (unrolling)	Drawing the recurrent layer once per time step, so the loop becomes a chain
Parameter sharing	Using the same weights at every time step
SimpleRNN	The Keras layer for a basic RNN; input shape (batch size, time steps, input features)