

Dot Product and Cosine Similarity

Contents

1. Overview	2
2. Products of two vectors	3
3. Computing the dot product	3
3.1 The dot product as a matrix product	4
3.2 Two laws	6
4. What the dot product is used for	7
5. The geometric meaning	8
5.1 The second formula	8
5.2 Perpendicular vectors have dot product 0	10
5.3 The angle between two vectors	10
6. Cosine similarity	11
6.1 Two phrases as two arrows	11
6.2 Definition and range	12
6.3 Recommending movies by angle	13
7. Summary	16
8. Sources	16
9. Key terms	16

Where this fits: Pipeline map steps 0 (Foundations), 8 (Model). See the [Course map](#).



- **Builds on:** Vectors and feature vectors ([Note MA-048](#)).
- **Leads to:** Equation of a hyperplane ([Note MA-051](#)); Matrix multiplication as composition ([Note MA-054](#)); Perceptron ([Note DL-004](#)); Meaning as direction in embedding space ([Note DL-072](#)).
- **Compare with:** Vector magnitude, distance and scalar operations ([Note MA-049](#)).

1. Overview

Key point: The dot product multiplies two vectors into one number; that number equals the two lengths times the cosine of the angle between them, so it measures how much two vectors point the same way.




Figure 1: The dot product of a fixed vector a and a turning vector b : positive for an acute angle, 0 at 90 degrees, negative for an obtuse angle

Figure 1 shows the whole idea. As b turns away from a , the dot product $a \cdot b$ falls from its largest value (same direction) through 0 (right angle) to negative values (pointing apart). This Note:

- computes the dot product and states its laws (section 3);
- lists what it is used for (section 4);
- explains this geometric meaning (section 5);

- turns it into cosine similarity, the standard way to compare texts in a recommender system (section 6).

Vectors, magnitudes and distances are covered in the [vectors and feature vectors Note](#) and the [magnitude, distance and scalar operations Note](#).

2. Products of two vectors

Key point: Two vectors have two products: the dot product gives a scalar, the cross product gives a vector; ML mostly uses the dot product.

Vectors can be multiplied, but not in the ordinary sense. There are two products of two vectors:

- The **dot product** (G-634), also called the **scalar product** (G-1742), because its result is a scalar.
- The **cross product** (G-506), also called the **vector product**, because its result is a vector.

ML uses the dot product almost everywhere, and the cross product rarely. This Note is about the dot product.

Extra: The usual cross product is defined for 3D vectors. For two 3D vectors it gives a third vector perpendicular to both, whose length is the area of the parallelogram they span. ML meets it only in a few places.

3. Computing the dot product

Key point: Multiply matching components, then add the products.

In plain words. A vector is a list of numbers, such as $[1, 2, 3]$. The dot product turns two equal-length lists into one number by pairing them up. Think of a shop bill. The list $a = [1, 2, 3]$ is how many pens, notebooks and erasers you buy; the list $b = [4, 5, 6]$ is the price of one of each, in rupees. The bill is the quantity times the price for each item, added up. That bill is the dot product (Figure 2 builds it one pair at a time).

Worked example, one line per item:

Item	Quantity (from a)	Price (from b)	Product
Pen	1	4	4
Notebook	2	5	10
Eraser	3	6	18
Bill			$4 + 10 + 18 = 32$

The [PCA step by step Note](#) (section 2.1) defined the dot product as “multiply matching components and add”. The standard term is the **dot product** (G-634).

The formal version. For two lists of n numbers, a and b , write a_1 for the first number of a (here $a_1 = 1$), a_2 for the second ($a_2 = 2$), and so on up to a_n ($n = 3$ here). The dot between the two vectors is the symbol of the dot product, and $\sum$ means “add up the terms for $i = 1$ to n ”:

$$a \cdot b = a_1b_1 + a_2b_2 + \cdots + a_nb_n = \sum_{i=1}^n a_ib_i$$

Check with the bill, one product per line:

$$a_1b_1 = 1 \times 4 = 4$$

$$a_2b_2 = 2 \times 5 = 10$$

$$a_3b_3 = 3 \times 6 = 18$$

$$\text{total} = 4 + 10 + 18 = 32$$

Both vectors must have the same number of components; otherwise some component would have no partner. Figure 2 builds the sum one pair at a time: each green product is one row of the table, and the last frame is the bill, 32.




Figure 2: The dot product of $[1, 2, 3]$ and $[4, 5, 6]$, one pair of components at a time: the products 4, 10 and 18 (green) add up to 32.

3.1 The dot product as a matrix product

Key point: A row of shape $1 \times n$ times a column of shape $n \times 1$ gives a 1×1 result: the dot product. So $a \cdot b = a^\top b$.

The dot product can also be written as a matrix multiplication. Write a as a row vector (shape $1 \times n$) and b as a column vector (shape $n \times 1$), as in the [vectors and feature vectors Note](#) (section 5).




Figure 3: A row of shape 1×3 times a column of shape 3×1 gives a 1×1 result, the dot product

The school rule of matrix multiplication says a $1 \times n$ matrix can multiply an $n \times 1$ matrix because the inner sizes match, and the result has the outer sizes: 1×1 , a single number (Figure 3). Each entry of the row meets its partner in the column, the products are added, and the result is exactly the dot product.

By default a vector is a column. So to put a in the row position we **transpose** (G-2012) it, and write

$$a \cdot b = a^T b$$

For $a = [1, 2, 3]$ and $b = [4, 5, 6]$, the row times the column is, one step per line:

$$\begin{aligned}
 a^T b &= \begin{bmatrix} 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix} \\
 &= 1 \times 4 + 2 \times 5 + 3 \times 6 \\
 &= 32
 \end{aligned}$$

This form, $a^T b$, appears throughout the inner workings of ML algorithms, such as $u^T x$ in the [PCA step by step Note](#).

3.2 Two laws

Key point: The order of the two vectors does not matter, and the dot product spreads over a sum.

Two rules hold for the dot product:

- **Commutative law** (G-420): $a \cdot b = b \cdot a$. Swapping the vectors only swaps the factors in each product.
- **Distributive law** (G-627): $a \cdot (b + c) = a \cdot b + a \cdot c$.

With numbers, take $a = [1, 2, 3]$, $b = [4, 5, 6]$ and $c = [7, 8, 9]$.

Commutative law:

$$a \cdot b = 1 \times 4 + 2 \times 5 + 3 \times 6 = 32$$

$$b \cdot a = 4 \times 1 + 5 \times 2 + 6 \times 3 = 4 + 10 + 18 = 32$$

Distributive law, the left side first:

$$b + c = [4 + 7, 5 + 8, 6 + 9] = [11, 13, 15]$$

$$a \cdot (b + c) = 1 \times 11 + 2 \times 13 + 3 \times 15 = 11 + 26 + 45 = 82$$

and the right side:

$$a \cdot c = 1 \times 7 + 2 \times 8 + 3 \times 9 = 7 + 16 + 27 = 50$$

$$a \cdot b + a \cdot c = 32 + 50 = 82$$

Both sides are 82.

Python: Three ways to get the same dot product.

```
import numpy as np
```

```
a = np.array([1, 2, 3])
```

```
b = np.array([4, 5, 6])
```

```
np.sum(a * b)    # 32: multiply, then add
```

```
np.dot(a, b)     # 32
```

```
a @ b           # 32: the @ operator
```

$a * b$ multiplies component by component; `np.dot` and `@` also add the products. The Notebook for this Note (`notebook.ipynb`) checks both laws.

4. What the dot product is used for

Key point: The dot product measures similarity, computes projections, and is the building block of matrix multiplication.

In linear algebra the dot product does three main jobs:

1. **Similarity.** It tells how similar two vectors are, and so which of several vectors are closest in direction. Section 6 builds this into cosine similarity.
2. **Projection** (G-1583). It gives the projection of one vector onto another: how far a reaches along the direction of b . For $a = [3, 4]$ and $b = [7, 1]$ that distance is (Figure 4):

$$\frac{a \cdot b}{\|b\|} = \frac{25}{7.07} = 3.54$$

Projections are used in the [PCA step by step Note](#) and the [SVM maths Note](#). The [dot product and duality Note](#) builds the whole dot product from this shadow.

3. **Matrix multiplication.** Every entry of a matrix product is the dot product of a row with a column.




Figure 4: The projection of $a = [3, 4]$ onto $b = [7, 1]$: the shadow of a along b (green) has length $a \cdot b / \|b\| = 3.54$.

Two uses in ML follow from these. A recommender system turns items into vectors and recommends the most similar ones, measured with the dot product. And deep learning runs on matrix multiplications, so behind the scenes it is computing dot products all the time.

5. The geometric meaning

Key point: The dot product is how far one arrow reaches along the other (its shadow), times the length of the other; the angle between the arrows controls the shadow.

5.1 The second formula

Key point: The dot product is the shadow of one vector on the other times the other's length; the second formula writes this as the two lengths times the cosine of the angle.

So far the dot product was a recipe on components. The picture below gives it a meaning.

In plain words. Hold two arrows at the origin. Shine a light straight down onto the line of b , so that a casts a **shadow** on it. This shadow is the **projection** (G-1583) of a onto b . The dot product is the length of that shadow times the length of b . If a leans the same way as b , the shadow is long and the product is large. At a right angle the shadow shrinks to a point and the product is 0. If a leans away, the shadow points backwards and the product is negative. The projection figure of Section 4 shows one shadow.

Worked example. Take $a = [3, 4]$ and $b = [7, 1]$, as in that figure.

$$\|a\| = \sqrt{3^2 + 4^2} = \sqrt{25} = 5$$

$$\|b\| = \sqrt{7^2 + 1^2} = \sqrt{50} = 7.07$$

$$a \cdot b = 3 \times 7 + 4 \times 1 = 25$$

$$\text{length of the shadow} = \frac{a \cdot b}{\|b\|} = \frac{25}{7.07} = 3.54$$

Now the dot product from the shadow:

$$\text{shadow} \times \|b\| = 3.54 \times 7.07 = 25$$

This matches the component recipe. The shadow depends on the angle θ between the arrows: its length is $\|a\| \cos \theta$. The **cosine** of an angle is a number that is 1 at 0° , 0 at 90° and -1 at 180° . Here:

$$\cos \theta = \frac{\text{shadow}}{\|a\|} = \frac{3.54}{5} = 0.71, \quad \theta = 45^\circ$$

The formal version. Combining “shadow times $\|b\|$ ” with “shadow = $\|a\| \cos \theta$ ”:

$$a \cdot b = \|a\| \|b\| \cos \theta$$

Here $\|a\|$ and $\|b\|$ are the magnitudes (lengths from the origin to the tips) and θ is the angle between a and b . Check on the worked numbers:

$$5 \times 7.07 \times 0.71 = 25$$

A second example on other numbers: for $a = [3, 4]$ and $b = [4, 3]$, the component recipe gives:

$$a \cdot b = 3 \times 4 + 4 \times 3$$

$$a \cdot b = 12 + 12 = 24$$

Both have length 5, the angle between them is 16.26° , and its cosine is 0.96:

$$\|a\| \|b\| \cos \theta = 5 \times 5 \times 0.96$$

$$\|a\| \|b\| \cos \theta = 24$$

In Figure 1, $a = [3, 1]$ has length 3.16 and b has length 2.5:

$$\|a\| = \sqrt{10} \approx 3.16$$

At $\theta = 30^\circ$ the dot product is:

$$3.16 \times 2.5 \times 0.87 \approx 6.85$$

At $\theta = 0^\circ$ it reaches its maximum:

$$3.16 \times 2.5 = 7.91$$

Extra: Why the two formulas agree. The three vectors a , b and $a - b$ form a triangle. The law of cosines gives:

$$\|a - b\|^2 = \|a\|^2 + \|b\|^2 - 2\|a\|\|b\| \cos \theta$$

Expanding the left side component by component gives:

$$\|a - b\|^2 = \|a\|^2 + \|b\|^2 - 2a \cdot b$$

Comparing the two:

$$a \cdot b = \|a\| \|b\| \cos \theta$$

5.2 Perpendicular vectors have dot product 0

Key point: For two non-zero vectors, the dot product is 0 exactly when the angle between them is 90° .

The lengths of non-zero vectors are positive. So the sign of $a \cdot b$ is the sign of $\cos \theta$ (Figure 1):

- **Acute angle** ($0^\circ \leq \theta < 90^\circ$): $\cos \theta > 0$, so $a \cdot b > 0$.
- **Right angle** ($\theta = 90^\circ$): $\cos \theta$ is 0, so the dot product is 0.
- **Obtuse angle** ($90^\circ < \theta \leq 180^\circ$): $\cos \theta < 0$, so $a \cdot b < 0$.

So if the dot product of two non-zero vectors is 0, they are perpendicular. For example:

$$[3, 4] \cdot [-4, 3] = 3 \times (-4) + 4 \times 3$$

$$[3, 4] \cdot [-4, 3] = -12 + 12 = 0$$

Perpendicular vectors are called **orthogonal** (G-1408); as data, they share nothing, and they are as dissimilar as two vectors can be without pointing apart. Algorithms such as SVM use this property heavily, and the [equation of a hyperplane Note](#) relies on it.

5.3 The angle between two vectors

Key point: Rearranging the formula gives $\cos \theta$ from the components, and the inverse cosine gives the angle.

If we know all the components of a and b , in any number of dimensions, we can find the angle between them.

1. **In words:** compute the dot product, divide by the product of the two lengths; that is $\cos \theta$. The inverse cosine (arccos) turns it into the angle.
2. **Formula:**

$$\cos \theta = \frac{a \cdot b}{\|a\| \|b\|}, \quad \theta = \cos^{-1} \left(\frac{a \cdot b}{\|a\| \|b\|} \right)$$

3. **Example:** for $a = [3, 4]$ and $b = [4, 3]$,

$$\cos \theta = \frac{24}{5 \times 5} = 0.96, \quad \theta = \cos^{-1}(0.96) \approx 16.26^\circ$$

Figure 5 draws this pair next to the perpendicular pair of section 5.2.




Figure 5: Left: $[3, 4]$ and $[4, 3]$, dot product 24, angle 16.26 degrees. Right: $[3, 4]$ and $[-4, 3]$, dot product 0, a right angle.

6. Cosine similarity

Key point: Cosine similarity is $\cos \theta$ between two vectors: 1 for the same direction, 0 for orthogonal, -1 for opposite. Cosine similarity compares direction and ignores length.

6.1 Two phrases as two arrows

Key point: Count the words of each phrase, draw the counts as an arrow, and compare the arrows by the angle between them; repeating words makes an arrow longer but does not turn it.

Take two words, *good* and *movie*, and count how often a phrase uses each. The two counts are the components of a 2D vector, so every phrase is an arrow. Figure 6 walks through four cases.

1. **Two phrases.** *good movie* has the counts $[1, 1]$ and *good* has $[1, 0]$. The angle between the two arrows is 45° , and $\cos 45^\circ = 0.71$.
2. **A longer phrase.** *good good good* has the counts $[3, 0]$. Its arrow is three times as long as the arrow of *good*, but points the same way. The angle with *good movie* is still 45° , so the cosine is still 0.71.
3. **The same words.** *good movie good movie* has the counts $[2, 2]$, on the same line as $[1, 1]$. The angle is 0° and the cosine is 1.
4. **No shared word.** *good* is $[1, 0]$ and *movie* is $[0, 1]$. The arrows are at a right angle, and the cosine is 0.

So the cosine of the angle is a score of how alike two phrases are: 1 when they use the same words in the same proportions, 0 when they share no word. The length of a phrase does not enter the score. This score is the **cosine similarity** (G-491).




Figure 6: Phrases as arrows in the plane of two word counts. Repeating “good” stretches the orange arrow without turning it, so the angle and the cosine similarity do not change. Idea after StatQuest, “Cosine Similarity, Clearly Explained!!!”.

A real text has thousands of different words, so its vector has thousands of components and cannot be drawn. The angle still exists, and the formula of section 5.3 computes its cosine from the components.

6.2 Definition and range

Key point: Cosine similarity runs from -1 to 1, and its sign tells whether the angle is acute or obtuse.

The cosine similarity of two vectors is the cosine of the angle between them,

$$\cos \theta = \frac{a \cdot b}{\|a\| \|b\|}$$

computed from the components exactly as in Section 5.3. In ML it is a standard **similarity measure** (G-1802): a number that says how alike two vectors are.




Figure 7: What each value of cosine similarity means

Figure 7 reads its values:

- **1:** $\theta = 0^\circ$. The vectors point in exactly the same direction, so they are completely similar, even if their lengths differ.
- **Between 0 and 1:** the angle is acute (0° to 90°); the closer to 1, the more similar.
- **0:** $\theta = 90^\circ$. The vectors are orthogonal: they have nothing in common.
- **Between -1 and 0:** the angle is obtuse (90° to 180°).
- **-1:** $\theta = 180^\circ$. The vectors point in opposite directions.

With numbers, take $p = [1, 2, 3]$, $q = [2, 4, 5]$ and $r = [-1, -2, 1]$:

Pair	Dot product	Lengths	Cosine similarity	Angle
p, q	$2 + 8 + 15 = 25$	3.742×6.708	$25/25.10 = 0.996$	5.1°
p, r	$-1 - 4 + 3 = -2$	3.742×2.449	$-2/9.165 = -0.218$	102.6°

p and q point almost the same way. p and r have an obtuse angle.

6.3 Recommending movies by angle

Key point: Turn every summary into a vector, compute the cosine similarity with the movie the user liked, and recommend the movies with the smallest angles.

The [vectors and feature vectors Note](#) (section 4) turned movie summaries into **bag-of-words** (G-250) vectors and recommended by Euclidean distance. Cosine similarity does the same job by angle. When a user picks a movie, we compute $\cos \theta$ between its vector and every other movie's vector, and recommend those with the smallest angle: ideally 0° , otherwise angles of 30° or 60° before ones near 90° .

For the three toy summaries $A = \text{hi how are you}$, $B = \text{my name is riya}$ and $C = \text{this is 2023}$, the vector of a summary lists how often each word of the whole vocabulary appears in it (a count vector). A word that a summary lacks has count 0.

Worked example. B and C share exactly one word, *is*. The steps:

$$B \cdot C = (\text{count of } is \text{ in } B) \times (\text{count of } is \text{ in } C) = 1 \times 1 = 1$$

All other words have count 0 in at least one of the two, so their products are 0. The lengths:

$$\|B\| = \sqrt{1^2 + 1^2 + 1^2 + 1^2} = \sqrt{4} = 2$$

$$\|C\| = \sqrt{1^2 + 1^2 + 1^2} = \sqrt{3} \approx 1.73$$

$$\cos \theta_{BC} = \frac{1}{2 \times 1.73} \approx 0.29$$

A shares no word with B, so their dot product and their cosine are both 0: orthogonal. A user who likes B gets C.

The formal version. The dot product of two count vectors counts the shared words (weighted by how often each appears); dividing by the two lengths removes the effect of length:

$$\cos \theta_{BC} = \frac{B \cdot C}{\|B\| \|C\|} = \frac{1}{2 \times 1.73}$$

This reproduces the 0.29 above.

Word counts are never negative, so for texts the cosine similarity always lies between 0 and 1. Figure 8 shows the whole table for the three summaries, plus B written twice (the Extra below).




Figure 8: Cosine similarity of the word-count vectors of A, B, C and B written twice (B2). B and C: 0.29; A and the others: 0; B and B2: 1, although B2 is twice as long.

Extra: Why cosine beats distance for text. Write B twice in a row, *my name is riya my name is riya*. Its vector is $2B$: same direction, twice as long. Its cosine similarity with B is exactly 1, but its Euclidean distance from B is 2. A long review and a short review on the same topic should count as similar; the angle sees this and the distance does not. Manning, Raghavan and Schütze (2008, §6.3.1) make the same point: two documents with very similar content can be far apart “simply because one is much longer than the other”, and the standard way to compare two documents is the cosine similarity of their vectors.

Python: scikit-learn computes all pairwise cosine similarities at once.

```
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics.pairwise import cosine_similarity
```

```
texts = ["hi how are you", "my name is riya", "this is 2023"]
```

```
X = CountVectorizer().fit_transform(texts)
```

```
cosine_similarity(X).round(2)    # 3 x 3 table; B-C is 0.29
```

Row i , column j of the result is the cosine similarity of text i and text j ; the diagonal is all 1.

7. Summary

Idea	Formula	Example
Dot product	$\sum a_i b_i$	$[1, 2, 3] \cdot [4, 5, 6] = 32$
Matrix form	$a^\top b: (1 \times n)(n \times 1) = 1 \times 1$	same 32
Laws	$a \cdot b = b \cdot a;$ $a \cdot (b + c) = a \cdot b + a \cdot c$	$32 = 32; 82 = 82$
Geometric form	$\ a\ \ b\ \cos \theta$	$5 \times 5 \times 0.96 = 24$
Orthogonal	$a \cdot b = 0$	$[3, 4] \cdot [-4, 3] = 0$
Cosine similarity	$\frac{a \cdot b}{\ a\ \ b\ }, \text{ from } -1 \text{ to } 1$	B and C: 0.29

- The dot product is one number: positive for an acute angle, 0 for a right angle, negative for an obtuse angle.
- The dot product gives similarity, projections and matrix multiplication; deep learning runs on it.
- Cosine similarity compares direction only, which makes it the usual choice for comparing texts.

8. Sources

Built from

- CampusX, “Supercharge Your ML Journey: Mastering Vectors in Linear Algebra - Part 1”, YouTube, <https://www.youtube.com/watch?v=mQewAJb8oJ8>
- Starmer, J. (StatQuest), “Cosine Similarity, Clearly Explained!!!”, YouTube, <https://www.youtube.com/watch?v=e9U0QAFbfLI>

Other references

- Manning, C. D., Raghavan, P. and Schütze, H. (2008). *Introduction to Information Retrieval*. Cambridge University Press. §6.3.1.

9. Key terms

Term	Meaning
Scalar product	Another name for the dot product, whose result is a scalar
Cross product (vector product)	A product of two 3D vectors that gives a vector perpendicular to both
Commutative law	$a \cdot b = b \cdot a$
Distributive law	$a \cdot (b + c) = a \cdot b + a \cdot c$
Orthogonal	Perpendicular; for non-zero vectors, dot product 0

Term	Meaning
Cosine similarity	The cosine of the angle between two vectors, from -1 to 1
Similarity measure	A number that says how alike two vectors are