


Online Machine Learning

Contents

1. Overview	2
2. What online learning is	2
2.1 Learning in small steps	2
2.2 The online learning workflow	3
2.3 Examples	4
3. When to use online learning	4
4. Implementing online learning	6
4.1 scikit-learn's <code>partial_fit</code>	6
4.2 Dedicated libraries	7
5. The learning rate	7
6. Out-of-core learning	9
7. Risks of online learning	10
7.1 Hard to run reliably	10
7.2 Bad data can damage the model	10
8. Batch vs online learning	11
9. Summary	13
10. Sources	13
11. Key terms	13

Where this fits: Pipeline map steps 0 (Foundations), 8 (Model), 13 (Monitor and maintain). See the [Course map](#).



- **Builds on:** Model drift ([Note ML-004](#)); Gradient descent ([Note ML-056](#)); Mini-batch gradient descent ([Note ML-059](#)).
- **Leads to:** Framing an ML problem ([Note ML-009](#)); Perceptron loss ([Note DL-006](#)); Batch size in Keras ([Note DL-020](#)).
- **Compare with:** Batch (offline) learning ([Note ML-004](#)); Batch gradient descent ([Note ML-057](#)); Mini-batch gradient descent ([Note ML-059](#)); SGD with momentum ([Note DL-034](#)).

1. Overview

Key point: In online learning, the model keeps learning after deployment. New data arrives in small mini-batches, and the model updates itself on the server after each one.

When a company says “the more you use our product, the better it gets”, it is usually describing online learning. The model is not frozen after deployment: it trains on new data while it is live.

2. What online learning is

Key point: Training is incremental: small mini-batches arrive one after another, and the model improves after each.

2.1 Learning in small steps

Key point: Many small updates instead of one big training run.

Online learning (G-1391) trains a model **incrementally** (**incremental learning**, G-931). Instead of using the whole dataset at once (as in **batch learning**, G-265), we feed the model data **sequentially**, in small groups called **mini-batches** (G-263), one after another. After each mini-batch, the model improves a little.

Each mini-batch is small, so each training step is fast and cheap. Small, cheap steps make it possible to train the model on the production server itself, while it is online. Hence the name.



Figure 1: Batch learning vs online learning

Figure 1 compares the two:

- **Batch:** all the data goes in once, and the model is then frozen. New data waits for the next retrain.
- **Online:** mini-batches stream in one by one, and the model updates after every one.

2.2 The online learning workflow

Key point: Train on a little data, deploy, then predict and learn at the same time.




Figure 2: The online learning workflow

Figure 2 shows the steps:

1. Start with a small amount of data and train the model on it.
2. Test it, then deploy it to the server.
3. New data keeps arriving. The model makes predictions on it and learns from it at the same time.

As new data arrives, the model's performance keeps improving to match it.

2.3 Examples

Key point: Chatbots, smart keyboards and video feeds adapt to us as we use them.

- **Chatbots and voice assistants** such as Google Assistant, Alexa and Siri keep learning from new conversations after deployment.
- **Smart keyboards** such as SwiftKey get better at predicting our words the more we type.
- **YouTube:** after we watch a video and go back to the feed, the feed has already changed to show related videos. The clicks we just made became new training data.




Figure 3: Three products that learn while we use them

Figure 3 shows the pattern the three share: using the product produces new data, and the deployed model learns from that data at once.

Many companies still use batch learning, but the industry is moving towards online learning.

3. When to use online learning

Key point: Use online learning when the problem keeps changing, when the data is huge, or when results are needed fast.

1. **The problem changes over time.** Some problems keep shifting: stock prices, or an e-commerce site where trends and customer behaviour change constantly. This drifting

of the problem is **concept drift** (Section 4.1 of the [batch learning Note](#)). Here the model must keep adapting, which is exactly what online learning does.

2. **Cost.** Retraining a batch model on a very large dataset is expensive. Online learning works with small mini-batches, so each step costs little.
3. **Speed.** Each training step is tiny, so the model reflects new data almost immediately.

Figure 4 tests the first reason on a problem that keeps changing: the Electricity market data of the [batch learning Note](#) (Harries 1999), where each half hour we predict whether the price goes up or down.




Figure 4: Online vs batch on the Electricity market data. Orange: a batch model trained once on the first 4 weeks. Green: an online model that starts from the same 4 weeks, then each day predicts that day's 48 half hours, and then learns from their true answers with one small update.

1. **The same start.** Both models learn from the first 4 weeks of data.
2. **Batch.** The orange model is then frozen, as in the batch learning Note.
3. **Online.** The green model keeps going. Each day it first predicts that day's 48 half hours, and then learns from their true answers with one small update.
4. **The result.** Over 130 weeks, the online model is right 71.7 percent of the time and the frozen batch model 68.0 percent. Where the market changes, for example around week 25 and after week 120, the frozen model falls behind while the online model recovers.

For problems that do not change, batch learning is still simpler and works well.

4. Implementing online learning

Key point: Use a model with `partial_fit` in scikit-learn, or a dedicated library such as River or Vowpal Wabbit.

4.1 scikit-learn’s `partial_fit`

Key point: `fit` starts from scratch on all the data; `partial_fit` continues from where the model left off.

Most scikit-learn models are trained with `fit`, which uses all the data at once. Some models also have **`partial_fit`** (G-1458), which trains on the data given and keeps what the model already learned. Calling it again with new data continues the training. scikit-learn is designed around batch learning, so only some of its models offer `partial_fit`.

One such model is **`SGDRegressor`** (G-1783). `SGDRegressor` does the same job as **linear regression** (G-1094; covered in later Notes), but learns step by step, which is what makes `partial_fit` possible.

Python: Training one observation at a time.

```
import numpy as np
from sklearn.linear_model import SGDRegressor

model = SGDRegressor()

# first row: 3 inputs, output 10
model.partial_fit(np.array([[1.0, 2.0, 3.0]]), np.array([10.0]))

# a new row arrives: training continues
model.partial_fit(np.array([[2.0, 1.0, 0.5]]), np.array([6.0]))
np.array([[...]]) is a table with one row: one observation (G-1374; one record),
with three features (G-772; input variables). np.array([...]) holds its target (G-1949),
the output value we want to predict. Each partial_fit call takes a fraction of a
second, so the model can keep learning as each new observation arrives.
```

Before the figure, the two words it needs. Each observation has one feature x (the input) and one target y (the output). The model is a straight line, and a line has two numbers. The **slope** is how much y rises when x grows by 1; the **intercept** is the value of y at $x = 0$. For the line $y = 0.67x + 0.25$ the slope is 0.67, the intercept is 0.25, and at $x = 3$ the line answers

$$0.67 \times 3 + 0.25 = 2.26$$

The true pattern in the data, $y = 2x$, has slope 2 and intercept 0, so at $x = 3$ the right answer is 6. A line with slope 0.67 is “far too flat”.

What does each call change? Figure 5 runs `partial_fit` on a stream of mini-batches of 10 observations, from data whose true pattern is $y = 2x$ plus noise (the Notebook’s data).




Figure 5: An SGDRegressor learning from a stream with `partial_fit`. Orange: the mini-batch of this call. Grey: earlier mini-batches, which the model has already used and does not keep. Blue: the model’s line after the call.

1. **Call 1.** The model has seen 10 observations. Its line is $y = 0.67x + 0.25$, far too flat.
2. **Calls 2 to 10.** Each call nudges the slope towards the data: 0.94, 1.18, and 1.66 after 10 calls.
3. **Call 50.** After 500 observations the line is $y = 1.81x + 0.51$, close to the true $y = 2x$.

The model never goes back to old observations. Each call moves the **parameters** (G-1450), the slope and the intercept, a little, starting from where the last call left them.

4.2 Dedicated libraries

Key point: River and Vowpal Wabbit are built specifically for online learning.

- **River:** a Python library for online machine learning on streaming data. River was formed by merging two earlier libraries, `creme` and `scikit-multiflow` (Montiel et al. 2021).
- **Vowpal Wabbit:** a fast library for online and interactive learning, including reinforcement learning (Vowpal Wabbit docs).

5. The learning rate

Key point: The learning rate controls how fast the model adapts. Too high and it forgets the past; too low and it is slow to learn anything new.

The **learning rate** (G-1068) sets how strongly each new mini-batch changes the model: how fast the model adapts to changing data (Géron 2019, Ch. 1).

Think of two news readers. A reader who believes every new headline changes their mind daily and forgets what they knew. A reader who ignores all headlines never learns anything new.

- **Too high:** the model changes very quickly and forgets what it learned before. It chases every bit of **noise** (G-1326): the random variation in the data that carries no pattern.

- **Too low:** the model barely changes. It remembers the past well but is slow to learn anything new.

We want a balance: the model should learn new patterns while still remembering the old ones.




Figure 6: How the learning rate changes what an online model learns (simulation)

Figure 6 draws three models as the data arrives, one point at a time. Each model follows one rule: move a fraction of the way from the current estimate towards the newest data point. That fraction is the learning rate.

Take one step by hand. The model’s current estimate is 20, and the newest data point is 60 (the true value has just jumped). The gap is

$$60 - 20 = 40$$

Each model moves its learning rate times this gap. As a rule, with the rate called r :

$$\text{change} = r \times (\text{newest point} - \text{old estimate})$$

$$\text{new estimate} = \text{old estimate} + \text{change}$$

Learning rate r	Move $r \times 40$	New estimate
0.01	0.4	20.4
0.1	4	24

Learning rate r	Move $r \times 40$	New estimate
0.7	28	48

The true value jumps from 20 to 60 at step 120. In Figure 6 the three rates behave as the table suggests:

- with a rate of 0.01, the model takes a very long time to catch up;
- with 0.7, the model reacts to every noisy point;
- with 0.1, the model follows the change quickly and stays steady.

The learning rate is one of the most important settings in online learning. If the rate is wrong, the model either forgets too fast or adapts too slowly.

6. Out-of-core learning

Key point: When data is too big to fit in memory, we split it into chunks and train on them one at a time.

Sometimes a dataset is too large to load at once. For example, a 50 GB dataset cannot be loaded on a machine with 8 GB of RAM, so it cannot be trained with batch learning.

Out-of-core learning (G-1413) solves this with the online-learning technique (scikit-learn User Guide, Strategies to scale computationally):




Figure 7: Out-of-core learning

1. Split the dataset into small chunks (Figure 7).
2. Feed the chunks to the model one at a time, training incrementally.
3. Deploy the trained model.

All of this happens offline, on our own machine. So out-of-core learning is not online learning, even though it uses the same incremental technique.

7. Risks of online learning

Key point: Online learning is hard to run reliably, and bad incoming data can damage the model. Monitor it and be ready to roll back.

7.1 Hard to run reliably

Key point: Training a model is easy; keeping it learning correctly on a live server is not.

Running online learning in production means doing three things at once:

- handling a constant stream of data;
- choosing the right learning rate;
- keeping everything working.

The job is hardest when data arrives in real time.

The tools are also young. Most are open-source libraries built by small groups, without the enterprise-grade reliability of established batch tools.

7.2 Bad data can damage the model

Key point: The model learns from whatever arrives, including bad data.

An online model changes according to the data it receives. If that data goes wrong, for example because the server is hacked and fake requests flood in, the model learns from it and becomes biased towards wrong answers.

The defences (Figure 8):




Figure 8: Protecting an online model

- **Monitoring:** watch the system constantly. An **anomaly detection** (G-201) algorithm can flag unusual incoming data.
- **Reject or go offline:** when data looks suspicious, refuse it or take the model offline.
- **Roll back** (rollback, G-1703): if damage is already done, restore the model to its last good version.

8. Batch vs online learning

Key point: Batch is simpler and proven; online adapts continuously but is harder to run.

	Batch (offline) learning	Online learning
Complexity	Low: train on our machine, deploy, predict	Higher: the model keeps changing and must be monitored
Computing power	Large training runs, occasionally	Small updates, continuously
Use in production	Easier to implement	Harder to implement
Best for	Problems that do not change (e.g. classifying dog breeds: a dog is still a dog in 10 years)	Problems that keep changing (e.g. weather, stock prices, trends)
Tools	Mature, industry-proven	Newer, still an active research area

Figure 9 puts numbers on the trade-off with the three models of Figure 4 and the [batch learning Note](#).




Figure 9: Accuracy against training work on the Electricity market data, weeks 4 to 133. Training work counts each record once each time it is fed to training.

- **Batch, never retrained:** the cheapest (1,344 records, once), and the least accurate, 68.0 percent.
- **Batch, retrained every 4 weeks:** the most accurate, 73.3 percent, but every retrain starts from scratch on all the data so far: 754 thousand records in total.

- **Online:** 71.7 percent for 72 thousand records, about a tenth of the retrained batch model's work, because after the first 4 weeks each record is used once, as it arrives.

Building a model with high accuracy is only part of the job. In industry, we also have to think about what happens after deployment: how much the server costs, and how the model reacts when the data changes.

The Notebook for this Note (`notebook.ipynb`) trains a model one row at a time with `partial_fit`, then compares a batch model and an online model when the data suddenly changes.

9. Summary

	Online learning
Training data	Small mini-batches, arriving one after another
Where training happens	On the production server, while the model is live
After deployment	The model keeps learning from new data
Key setting	The learning rate
Strong at	Changing problems, huge data, fast updates
Weak spots	Hard to run, young tools, vulnerable to bad data
Examples	Chatbots, smart keyboards, video feeds

- Online learning = incremental training on mini-batches, while the model is live.
- In scikit-learn, models with `partial_fit` can learn this way.
- The learning rate balances learning new patterns against remembering old ones.
- Out-of-core learning uses the same technique offline, for data too big for memory.
- Protect an online model with monitoring, anomaly detection and rollback.

10. Sources

Built from

- CampusX, “Online Machine Learning | Online Learning | Online Vs Offline Machine Learning”, YouTube, <https://www.youtube.com/watch?v=3oOipgCbLIk>

Other references

- Géron, A. (2019). *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*, 2nd ed. O’Reilly. Ch. 1, Online learning.
- Harries, M. (1999). *Splice-2 Comparative Evaluation: Electricity Pricing*. Technical report, University of New South Wales. The data as OpenML dataset 151, “electricity”, <https://www.openml.org/d/151>
- Montiel, J. et al. (2021). River: Machine Learning for Streaming Data in Python. *Journal of Machine Learning Research* 22(110).
- scikit-learn developers. *User Guide*, Strategies to scale computationally: bigger data. scikit-learn.org.
- Vowpal Wabbit project. *Vowpal Wabbit documentation*. vowpalwabbit.org.

11. Key terms

Term	Meaning
Online learning	Training incrementally on mini-batches while the model is live in production
Incremental training (G-931)	Training in small steps, keeping what was learned before
Sequential data	Data fed one piece after another, in order
Feature	An input variable; one column of the data table
Target	The output we want to predict
Observation	One record; one row of the data table
Mini-batch (G-263)	A small group of observations used for one training step
<code>partial_fit</code> (G-1458)	A scikit-learn method that continues training from where the model left off
<code>SGDRegressor</code> (G-1783)	A scikit-learn model that does linear regression step by step
River	A Python library for online machine learning
Vowpal Wabbit	A fast learning library that supports online learning
Learning rate (G-1068)	How strongly each new piece of data changes the model
Out-of-core learning (G-1413)	Training on data too big for memory by feeding it in chunks, offline
Biased model	A model pushed towards wrong answers, e.g. by bad data
Parameter (G-1450)	A number inside a model that training changes, such as the slope of a line
Rollback	Restoring a model to an earlier, good version