


Instance-Based vs Model-Based Learning

Contents

1. Overview	2
2. Two ways to learn	2
3. Instance-based learning	3
3.1 How it works	3
3.2 No real training	5
3.3 The vote, and why k matters	5
4. Model-based learning	7
4.1 Learning a decision boundary	7
4.2 The training data is no longer needed	9
4.3 Examples	9
5. Comparing the two	10
6. Summary	12
7. Sources	12
8. Key terms	12

Where this fits: Pipeline map steps 0 (Foundations), 5 (Engineer features), 8 (Model). See the [Course map](#).



- **Builds on:** Classification problems ([Note ML-003](#)); Feature engineering ([Note ML-007](#)); Hyperparameter tuning ([Note ML-009](#)); Feature transformation ([Note ML-022](#)); Cross-validation ([Note ML-028](#)); Curse of dimensionality ([Note ML-045](#)).
- **Leads to:** Standardization ([Note ML-009](#)); Logistic regression ([Note ML-012](#)); Normalization ([Note ML-024](#)); K-means ([Note ML-031](#)); KNN imputer ([Note ML-038](#)); PCA ([Note ML-046](#)).
- **Compare with:** Bias-variance trade-off ([Note ML-061](#)); Decision trees ([Note ML-091](#)); ANN for classification ([Note DL-011](#)).

1. Overview

Key point: Some ML algorithms memorise the training data and compare new points to it (instance-based). Others learn a general rule from the data and use only that rule (model-based).

This Note groups ML algorithms in a third way: **by how a model learns**. There are two types:

- **Instance-based learning,**
- **Model-based learning.**

We have not studied any algorithms yet. The aim here is to be able to tell, for every algorithm we meet later, which of the two it is.

2. Two ways to learn

Key point: Memorising examples vs understanding the rule behind them: people learn both ways, and so do ML algorithms.

People learn in two broad ways:

- **Memorising:** learning past answers by heart, then comparing each new question with them. Many of us studied some exam subjects this way.
- **Understanding:** learning the underlying principle, then applying it to any new question.

ML algorithms work the same way (Figure 1).




Figure 1: Two ways to learn, in people and in ML

- **Instance-based learning** (G-955) memorises: it keeps the training data and compares new points with it.
- **Model-based learning** (G-1257) understands: it extracts the underlying pattern as a mathematical function and uses that.

Both are explained below with the same example: predicting whether a student will be placed, from their IQ and CGPA. Three words describe this data:

- IQ and CGPA are the **features** (G-772; input variables, one column each in the data table);
- placed or not is the **target** (G-1949; the output we predict);
- each student is one **observation** (G-1374; one record, one row of the table).

3. Instance-based learning

Key point: Training = storing the data. Prediction = find the most similar stored points and copy their answer.

3.1 How it works

Key point: Nothing is learned until a question arrives. Then the algorithm looks at the stored points nearest to it.

In **instance-based learning**, training does nothing except store the training data. All the work happens when a new point arrives:




Figure 2: Instance-based learning step by step (KNN, $k = 3$)

Figure 2 shows the steps for a new student with IQ 94.5 and CGPA 8.3:

1. **Store:** keep every training point. Green points were placed, red were not.
2. **Measure:** compute the **distance** (G-623) from the new student to every stored student. Distance measures **similarity** (G-1805): the closer two points are, the more alike the students.
3. **Pick the nearest:** keep the k closest points. Here $k = 3$.
4. **Vote:** 2 of the 3 nearest students were placed, so the prediction is *placed*.

Step 2 in numbers, for one stored student: IQ 91.1 and CGPA 8.2, placed. Both features are first put on the same scale (the Extra below explains why): the new student becomes $(-0.89, 0.75)$ and the stored student becomes $(-1.10, 0.68)$. The distance is the straight-line distance between the two points, found in three lines:

$$\text{gap in IQ} = -1.10 - (-0.89) = -0.21 \quad \text{and} \quad \text{gap in CGPA} = 0.68 - 0.75 = -0.07$$

$$(-0.21)^2 + (-0.07)^2 = 0.044 + 0.005 = 0.049$$

$$\text{distance} = \sqrt{0.049} \approx 0.22$$

This student is one of the 3 nearest. Repeating the three lines for every stored student gives the list from which step 3 keeps the 3 smallest distances.

The idea behind step 4: points that are close together tend to share the same answer, just as where a person lives says something about them. A student who looks like placed students will probably be placed too.

This procedure is the **K-nearest neighbours (KNN)** (G-998) algorithm, covered in detail in later Notes.

Extra: IQ ranges over about 60 points, while CGPA ranges over about 5. Measured raw, distances would depend almost only on IQ. Write Δ (delta) for “the gap in”: ΔIQ is the gap between two students’ IQs. The distance is

$$\sqrt{(\Delta\text{IQ})^2 + (\Delta\text{CGPA})^2}$$

and the IQ term can reach $60^2 = 3600$ while the CGPA term reaches only $5^2 = 25$. So before measuring distances, both features are put on the same scale (**feature scaling**, G-767; see Section 7 of the [toy project Note](#)). The neighbours in Figures 2, 3 and 5 were found this way.

3.2 No real training

Key point: An instance-based model does nothing until it is asked a question.

Until the new student arrived, the algorithm did nothing with the data: it simply held on to it. Only when the question came did it look at the data and work out an answer.

So in instance-based learning, there is no real training step. Putting off the work is why instance-based learning is also called **lazy learning** (G-1057): the method puts off the work until a question arrives (Mitchell 1997, §8.6).

3.3 The vote, and why k matters

Key point: The prediction is a majority vote among the k nearest points. A different k can give a different answer.

Step 4 of Figure 2 is a **majority vote** (G-1146): each of the k nearest students casts one vote for its own class, and the class with the most votes wins. The new point is called the **query point** (G-1605).

Real data has **outliers** (G-1420): a few students with a good IQ and CGPA who were still not placed, or the reverse. Where the two classes meet, placed and not-placed students also sit side by side. In such a region the answer depends on how many neighbours vote. Figure 3 shows the vote for a query point with IQ 97.5 and CGPA 8.0.




Figure 3: The KNN vote for one new student (the star) as k grows. Outlined dots: the k nearest students. Right: the votes. Idea after StatQuest, “K-nearest neighbors, Clearly Explained”.

k	Votes for placed	Votes for not placed	Prediction
1	0	1	not placed
3	2	1	placed
5	3	2	placed
11	6	5	placed

1. **$k = 1$.** The single nearest student was not placed, so the prediction copies that one student. One neighbour decides everything.
2. **$k = 3$ and more.** The circle grows, more neighbours vote, and that one student is outvoted: the prediction becomes *placed*.

Two rules of thumb follow:

- **Use an odd k** for two classes, so the vote cannot end in a tie.
- **Small k is noisy:** a single outlier can flip the answer. **Large k is smoother,** but if k is too large, a class with few points is always outvoted.

The best k is found by trying several values on data kept aside from training. The [KNN Note](#) covers that search.

4. Model-based learning

Key point: The algorithm learns a mathematical function from the data, such as a boundary between classes. After that, it needs only the function.

4.1 Learning a decision boundary

Key point: A model learns where one class ends and the other begins, then classifies new points by which side they fall on.

In **model-based learning**, the algorithm studies the training data and builds a mathematical function that connects the features to the target.

For a **classification** (G-395) problem, the learned function is a **decision boundary** (G-555): a line (or curve) that separates the classes. Every point on one side is predicted *placed*; every point on the other side, *not placed*.

How does the algorithm find that line? Figure 4 shows **logistic regression** (G-1120), a model-based classifier, learning it on our 60 students.




Figure 4: Logistic regression learning its decision boundary on the placement data, then dropping the training data. Green: placed; red: not placed; purple: the current decision boundary.

1. **Start.** The line begins in a bad place: only 38 percent of the training students are on the correct side.
2. **One training step.** The algorithm measures how wrong the line is with a **loss function** (G-706), a single number that is large when many students are on the wrong side, then turns and shifts the line a little to make that number smaller. Repeating this step is **gradient descent** (G-862).
3. **After a few steps.** The loss falls from 1.094 to 0.259 after 2 steps, and 97 percent of the students are on the correct side. Later steps barely move the line (loss 0.142).
4. **Training data dropped.** The model keeps three numbers, $w_1 = 1.36$, $w_2 = 2.72$ and $b = -0.53$ (for IQ and CGPA on the same scale), and nothing else. The new student,

IQ 94.5 and CGPA 8.3, falls just above the line: *placed*, with probability 0.57.

The next lines show where the 0.57 comes from. On the same scale the new student is $(-0.89, 0.75)$, as in Section 3.1. The model first combines the two features and the third number into one score z , with each weight multiplying its own feature:

$$w_1 \times \text{IQ} = 1.36 \times (-0.89) = -1.21$$

$$w_2 \times \text{CGPA} = 2.72 \times 0.75 = 2.04$$

$$z = -1.21 + 2.04 + (-0.53) = 0.30$$

A score above 0 means the student is on the *placed* side of the boundary, and $z = 0$ is the boundary itself. The score is turned into a probability by the **sigmoid** (G-1798) function, which maps any number to a value between 0 and 1 (it is the subject of [Note ML-071](#)):

$$\text{probability} = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-0.30}} = \frac{1}{1 + 0.74} = 0.57$$

The **loss** of step 2 is built from the same probabilities. For one student, the loss is $-\ln$ of the probability the model gave to the student's true answer. A placed student given probability 0.8 costs $-\ln 0.8 = 0.22$; the same student given 0.2 costs $-\ln 0.2 = 1.61$. The loss of Figure 4 is the average of this cost over the 60 students, so it is large when many students get a low probability for their true answer (1.094 at the start, 0.259 after 2 steps).




Figure 5: Same data, two approaches

Figure 5 shows both approaches on the same data, classifying the same new student:

- **Left (instance-based):** the answer comes from the 3 nearest stored students.
- **Right (model-based):** the answer comes from the side of the learned boundary the student falls on. Both models were trained with scikit-learn.

4.2 The training data is no longer needed

Key point: Once the function is learned, the training data can be thrown away.

After training, a model-based algorithm keeps only the function. To classify a new student, we check which side of the decision boundary they fall on; the training points play no part.

The function is described by a few numbers called **parameters** (G-1450). For example:

- in a straight-line model, the parameters are the line's **slope** (G-1823) and **intercept** (G-960);
- in a neural network, the parameters are its **weights** (G-2106);
- in Naive Bayes ([Note ML-081](#)), the parameters are probabilities.

Instance-based keeps

IQ	CGPA	Placed?
112	8.4	Yes
91	6.1	No
127	7.9	Yes
84	8.8	No
...	...	...

every training row (1 GB of data = 1 GB stored)

Model-based keeps

$$w_1 \times \text{IQ} + w_2 \times \text{CGPA} + b$$

just 3 numbers: w_1, w_2, b

the training data can be thrown away

Figure 6: What each approach keeps after training

Figure 6 shows the difference in what is kept: the whole table for instance-based learning, just a few numbers for model-based learning.

4.3 Examples

Key point: Most algorithms are model-based; KNN is the classic instance-based one.

The instance-based examples in the table come from Mitchell (1997, Ch. 8).

Instance-based	Model-based
K-nearest neighbours (KNN)	Linear regression

Instance-based	Model-based
Locally weighted regression	Logistic regression
Radial basis function (RBF) networks	Decision trees, neural networks and most other algorithms

Extra: Model-based learning is sometimes called **eager learning** (G-655), the opposite of lazy learning: all the work is done up front, before any question arrives (Mitchell 1997, §8.6).

5. Comparing the two

Key point: Model-based learning generalises once and stores little; instance-based learning generalises at every question and stores everything.

	Model-based	Instance-based
Data preparation	Needed (e.g. handling outliers, turning categories into numbers)	Needed, the same
What training produces	A model with parameters	Nothing: the data is just stored
When it generalises	Before any question, by learning the rule	At each question, from the nearby points
How it predicts	Applies the model	Compares the new point with the training data
Training data after training	Can be thrown away	Must always be kept
Needs a similarity measure	No	Yes, e.g. distance
Storage	Small: just the parameters	Large: the whole dataset (1 GB of data = 1 GB stored)
Other name	Eager learning	Lazy learning

Figure 7 measures the storage and prediction rows of the table as the training data grows, with KNN and logistic regression on more and more students from the same data generator.




Figure 7: What each approach costs as the training data grows: numbers kept after training (left) and time to answer 10,000 new students (right), both on log scales. Timings are from one machine.

- **Storage.** KNN keeps every value of every training student: 3,000 numbers for 1,000 students, 3 million for 1 million. Logistic regression keeps 3 numbers, whatever the size of the data.
- **Prediction time.** KNN must search the stored students for each new one, so its answers slow down as the data grows (about 8 ms for 10,000 new students with 1,000 stored, about 29 ms with 1 million stored). Logistic regression computes one formula per student and takes about 0.1 ms at every size.

The Notebook for this Note ([notebook.ipynb](#)) is a small app: move a new student with sliders and change k , and see what each approach predicts.

6. Summary

- **Instance-based learning** memorises: it stores the data and, for each new point, copies the answer of the most similar stored points (KNN).
- **Model-based learning** generalises: it learns a function, such as a decision boundary, and predicts with that alone.
- Instance-based learning is lazy: no real training, all the work happens at prediction time, and all the data must be kept.
- Model-based learning keeps only a few parameters, so it is small and fast at prediction time.
- For any new algorithm, ask: does it keep the data, or does it learn a rule?

7. Sources

Built from

- CampusX, “Instance-Based Vs Model-Based Learning | Types of Machine Learning”, YouTube, <https://www.youtube.com/watch?v=ntAOq1ioTKo>
- StatQuest with Josh Starmer, “StatQuest: K-nearest neighbors, Clearly Explained”, YouTube, <https://www.youtube.com/watch?v=HVXime0nQeI> (the vote for different k , Section 3.3)

Other references

- Mitchell, T. (1997). *Machine Learning*. McGraw-Hill.

8. Key terms

Term	Meaning
Instance-based learning	Learning by storing the training data and comparing new points with it
Model-based learning	Learning a mathematical function from the data and predicting with it
Distance	A number measuring how far apart two points are; small distance = similar
Similarity	How alike two data points are
K-nearest neighbours (KNN)	Predicting from the answers of the k closest stored points
Majority vote	Predicting the class that most of the k nearest points belong to
Query point	The new point whose class we want to predict
Outlier	A value far from the rest of the data; here, a student whose result does not match similar students

Term	Meaning
Lazy learning	Another name for instance-based learning: no work until a question arrives
Eager learning	Another name for model-based learning: all the work done up front
Decision boundary	A line or curve that separates the classes in classification
Parameters (G-1450)	The numbers that describe a learned model, e.g. slope and intercept
Logistic regression	A model-based classifier that learns a straight decision boundary
Loss function (G-706)	A single number measuring how wrong a model is on the training data; training makes it smaller
Sigmoid function (G-1798)	A curve that maps any score to a value between 0 and 1, read as a probability
Gradient descent	Training by repeatedly nudging the parameters to make the loss smaller
Feature	An input variable; one column of the data table
Target	The output we predict
Observation	One record; one row of the data table
Feature scaling (G-767)	Putting features on the same scale, so no feature dominates distances