

Pandas Profiling: A Full EDA Report in One Line

Contents

1. Overview	2
2. Building the report	3
2.1 Installing the library	3
2.2 Three lines of code	3
3. Overview: the first questions, answered	5
3.1 Dataset statistics	5
3.2 Alerts	6
4. Variables: univariate analysis of every column	8
4.1 A categorical column: Sex	8
4.2 A numerical column: Age	9
5. Interactions: a scatter plot for any pair	12
6. Correlations: every pair in one heatmap	13
7. Missing values: how many, and where	15
8. Sample: a look at the rows	17
9. Using the report	17
10. Summary	19
11. Sources	19
12. Key terms	20

Where this fits: Pipeline map step 3 (Understand data). See the [Course map](#).



- **Builds on:** Setup: conda, Jupyter and Colab ([Note ML-011](#)).
- **Leads to:** Q-Q plot ([Note ML-029](#)).
- **Compare with:** Univariate analysis ([Note ML-019](#)); Skewness ([Note ML-019](#)); Bivariate and multivariate analysis ([Note ML-020](#)).

1. Overview

Key point: A profiling library runs most of the hand-made EDA (first questions, univariate, bivariate and multivariate analysis) by itself and writes the results into one web page.

So far we built EDA by hand, in three steps: the [first questions for a new dataset](#), [univariate analysis](#), and [bivariate and multivariate analysis](#). Each step took a few lines of pandas or one plot at a time. A **profiling report** does most of this work in one go: we hand it a DataFrame, and it returns a web page that describes every column and every pair of columns.

The library that builds it is known as **Pandas Profiling** (G-1440). Figure 1 shows what one call of `ProfileReport` produces: a report with six sections, each answering questions of the hand-made EDA.




Figure 1: What a profiling report contains

We use the same Titanic data as before, 891 passengers. The Notebook ([ML-021-pandas-profiling.ipynb](#)) builds the report and saves it as `titanic_report.html`; Figures 3 to 11 are screenshots of

that page.

2. Building the report

Key point: Install the library once; then three lines build the report and save it as an HTML file.

2.1 Installing the library

Key point: One `pip install` command adds the library to our Python setup.

Pandas Profiling is not part of pandas; it is a separate library. We install it once, from a terminal or from a notebook cell. The command downloads the library from **PyPI** (G-1593), the public store of Python libraries, and installs it on our machine.

Python: Installing the library.

```
pip install fg-data-profiling
```

In a Jupyter cell, put `!` in front: `!pip install fg-data-profiling`. The `!` sends the line to the terminal instead of to Python.

Extra: Version clash (October 2026). The library still requires pandas older than 3.0. In a setup that already has pandas 3, a plain `pip install fg-data-profiling` quietly downgrades pandas (and NumPy, SciPy) to older versions. To keep pandas 3, install it with `pip install --no-deps fg-data-profiling` and then install the few small helper libraries it asks for when imported. The report then works with the one-line fix in Section 2.2. In a fresh setup made just for profiling, the plain command is fine.

Extra: The library has changed its name twice. The library began as `pandas-profiling` (`import pandas_profiling`). In 2023 it became `ydata-profiling` (`import ydata_profiling`), and in 2026 `fg-data-profiling` (`import data_profiling`). The old names still appear in many tutorials; `pip install pandas-profiling` now installs only an empty package that says to use the new name, and `import ydata_profiling` warns that it gets no more updates (PyPI pages of the three packages). The report itself and the `ProfileReport` call are the same under all three names.

2.2 Three lines of code

Key point: `ProfileReport(df)` builds the report; `to_file` saves it as a web page we open in any browser.

Once installed, building a report takes three steps (Figure 2):

1. Import `ProfileReport` from the library.
2. Call `ProfileReport` on our `DataFrame`. The call creates a report object.
3. Save the object as an **HTML** (G-904) file with `to_file`.




Figure 2: The three lines of code as a flow: the CSV file becomes a DataFrame, ProfileReport builds the report object in 10 to 15 seconds, and to_file saves it as a web page

Python: Building and saving the report.

```
import pandas as pd
from data_profiling import ProfileReport
```

```
df = pd.read_csv("data/titanic_train.csv")
profile = ProfileReport(df, title="Titanic passengers")
profile.to_file("titanic_report.html")
```

title is optional; it becomes the heading of the page. Double-clicking the HTML file opens the report in a browser.

The report takes a few seconds to build. The time grows with the size of the data: behind the scenes, the library computes statistics and draws charts for every column and every pair of columns. For the 891 Titanic rows it takes about ten to fifteen seconds.

Extra: Two more options are useful:

- `ProfileReport(df, minimal=True)` skips the slow parts (correlations, interactions, missing-value charts). On data with many thousands of rows or many columns, the full report can take minutes; the minimal one stays fast.
- `profile.to_notebook_iframe()` shows the report inside Jupyter instead of saving a file, in an **iframe** (a frame that shows one web page inside another). The iframe stores the whole page in the notebook, which makes the notebook file several megabytes big.

Extra: The library officially supports pandas versions below 3. With pandas 3, building the report fails with the error `'ArrowExtensionArray' object has no attribute 'sum'`. pandas 3 stores text columns in a new format (Arrow arrays) that the library cannot add up. One line before building the report switches text back to plain Python strings:

```
pd.set_option("mode.string_storage", "python")
```

The page has six sections, listed in Figure 1. A menu at the top jumps to each one. The sections follow the order of the hand-made EDA: first the [basic questions](#), then [one column at a time](#), then [pairs of columns](#).

3. Overview: the first questions, answered

Key point: The Overview section answers the first questions about a new dataset (size, types, missing values, duplicates) in one table.

3.1 Dataset statistics

Key point: Rows, columns, missing cells, duplicate rows and memory use, all in one table.

The Overview section opens with the dataset as a whole. Figure 3 shows it for the Titanic data. These are the same first questions we asked with `df.shape`, `df.info()`, `df.isnull().sum()` and `df.duplicated().sum()` in [how big is the data](#).

Dataset statistics		Variable types	
Number of variables	12	Numeric	5
Number of observations	891	Categorical	4
Missing cells	866	Text	3
Missing cells (%)	8.1%		
Duplicate rows	0		
Duplicate rows (%)	0.0%		
Total size in memory	118.9 KiB		
Average record size in memory	136.6 B		

Figure 3: The Overview section: dataset statistics and variable types

Reading the left table from the top:

- **Number of variables:** 12. The report calls columns **variables**.
- **Number of observations:** 891. The report calls rows **observations**: one record each, here one passenger.
- **Missing cells:** 866, which is 8.1% of all cells.
- **Duplicate rows:** 0, so no passenger appears twice.
- **Total size in memory:** 118.9 KiB for the whole table; the **average record size** is 136.6 bytes, the memory one row takes.

The missing-cells percentage is worked out from the whole grid of cells.

Extra: The share of missing cells, step by step.

1. **In words:** count all cells (rows times columns), then divide the missing cells by that total.

2. Formula:

all cells = rows $\times$ columns

$$\text{missing percent} = \frac{\text{missing cells}}{\text{all cells}} \times 100$$

3. With the Titanic numbers: the grid has this many cells:

$$891 \times 12 = 10,692$$

The share of missing cells is then:

$$\frac{866}{10,692} \times 100 = 8.1 \text{ percent}$$

The 866 missing cells come from three columns: 177 in `Age`, 687 in `Cabin` and 2 in `Embarked`.

The right table counts the column types. The report sorts the 12 columns into three kinds:

Type	Count	Columns
Numeric	5	<code>PassengerId</code> , <code>Age</code> , <code>SibSp</code> , <code>Parch</code> , <code>Fare</code>
Categorical	4	<code>Survived</code> , <code>Pclass</code> , <code>Sex</code> , <code>Embarked</code>
Text	3	<code>Name</code> , <code>Ticket</code> , <code>Cabin</code>

The report decides the type itself, from the values. `Survived` and `Pclass` hold numbers, but only two or three different ones, so the report treats them as categories. The report's split matches the split into numerical and categorical columns from [numerical and categorical data](#).

Extra: “Text” is the report's name for a column of free text, where nearly every value is different (every passenger has their own name and ticket number). Older versions of the library put such columns under “Categorical” and flagged them with a **high cardinality** warning: a categorical column with very many different categories.

3.2 Alerts

Key point: Alerts are a ready-made list of possible problems, one line per column; each one is a question to check, not a verdict.

The second tab of the Overview section lists **alerts**: warnings about columns that may need attention. Figure 4 shows the twelve alerts for the Titanic data.

Alerts

Sex is highly overall correlated with Survived	High correlation
Survived is highly overall correlated with Sex	High correlation
Pclass is highly overall correlated with Fare	High correlation
Fare is highly overall correlated with Pclass	High correlation
Age has 177 (19.9%) missing values	Missing
Cabin has 687 (77.1%) missing values	Missing
PassengerId is uniformly distributed	Uniform
PassengerId has unique values	Unique
Name has unique values	Unique
SibSp has 608 (68.2%) zeros	Zeros
Parch has 678 (76.1%) zeros	Zeros
Fare has 15 (1.7%) zeros	Zeros

Figure 4: The Alerts tab of the Overview section

The alerts fall into five kinds:

- **High correlation:** Sex is strongly linked to Survived, and Pclass to Fare. These are the strongest relationships in the data, worth studying first.
- **Missing:** Age is missing for 177 passengers (19.9%) and Cabin for 687 (77.1%). Nearly a fifth of the ages must be dealt with before training a model. With more than three quarters missing, Cabin may be unusable.
- **Uniform:** every value of PassengerId occurs equally often (once each). A running number like this carries no information about the passenger.
- **Unique:** every value of PassengerId and of Name is different. A column that is different for every row cannot reveal a pattern shared by several rows.
- **Zeros:** most passengers have 0 in SibSp (68.2%) and Parch (76.1%): they travelled without siblings, spouses, parents or children. 15 fares (1.7%) are 0, which is odd for a paid ticket.

In seconds, the alerts point to the columns that need cleaning and the columns that look useless. Each alert still needs a judgement: many zeros in SibSp is simply a fact about the passengers, while 15 free tickets deserve a closer look.

Extra: The third tab, Reproduction, records when the report was made, how long it took, and the library version and settings. The Reproduction tab lets someone else rebuild exactly the same report later.

4. Variables: univariate analysis of every column

Key point: The Variables section gives each column its own block of statistics and a chart, chosen by the column's type.

The second section is the [univariate analysis](#) (one column at a time), done for every column at once. A drop-down menu picks one column, or shows all of them one below the other. What each block contains depends on the column's type.

Some columns are not worth studying. `PassengerId` is just a running number, as its alerts already said, so we skip it.

4.1 A categorical column: Sex

Key point: For a categorical column, the report counts each category and draws the counts as bars, like a count plot.

Figure 5 shows the block for `Sex`. On the left, a table gives the basic facts; on the right, a bar for each category with its count.




Figure 5: The Variables block for Sex

- **Distinct:** 2 different values, male and female. As a share of the 891 rows, that is 0.2%.
- **Missing:** 0, so every passenger has a recorded sex.
- **The bars:** 577 male and 314 female passengers. The bar chart is the [count plot](#) (one bar per category, as tall as its count), drawn sideways.

The “More details” button opens four further tabs:

- **Overview:** the length of the values in characters (here 4 for “male”, 6 for “female”) and the first five values as a sample.
- **Categories:** every category with its count and percentage (64.8% male, 35.2% female).

- **Words** and **Characters**: the words and characters that make up the values, which matter more for text columns.

4.2 A numerical column: Age

Key point: For a numerical column, the report gives the main summary numbers and a histogram; “More details” adds percentiles, spread and the extreme values.

Figure 6 shows the block for **Age**. The missing values are printed in red, because they triggered an alert.




Figure 6: The Variables block for Age

The two tables give, among others:

- **Distinct:** 88 different ages.
- **Missing:** 177 (19.9%).
- **Infinite:** 0. An infinite value can appear after a division by zero, and would break most calculations.
- **Mean:** 29.70 years. **Minimum:** 0.42 (a baby of five months). **Maximum:** 80.
- **Zeros** and **Negative:** 0 each. A negative age would be an error in the data.

On the right is the histogram of **Age**, the same shape as in [the histogram](#).

“More details” opens four tabs: Statistics, Histogram, Common values and Extreme values. The Statistics tab (Figure 7) has two tables.

Quantile statistics		Descriptive statistics	
Minimum	0.42	Standard deviation	14.526497
5-th percentile	4	Coefficient of variation (CV)	0.48912219
Q1	20.125	Kurtosis	0.17827415
median	28	Mean	29.699118
Q3	38	Median Absolute Deviation (MAD)	9
95-th percentile	56	Skewness	0.38910778
Maximum	80	Sum	21205.17
Range	79.58	Variance	211.01912
Interquartile range (IQR)	17.875	Monotonicity	Not monotonic

Figure 7: The Statistics tab of Age, under More details

Quantile statistics describe positions in the sorted data:

- **5th and 95th percentile:** 4 and 56. Only 5% of the passengers were younger than 4, and only 5% older than 56.
- **Q1, median, Q3:** 20.125, 28 and 38. The **first quartile** Q1 is the age a quarter of the passengers are below; the median splits them in half; the **third quartile** Q3 is the age three quarters are below. These cut points are **quantiles**: values that a given share of the data lies below (the 5th percentile above is one too). Together with the minimum and maximum, Q1, median and Q3 give the five-number summary drawn by a [box plot](#).
- **Range:** maximum minus minimum:

$$80 - 0.42 = 79.58$$

- **Interquartile range (IQR):** the width of the middle half of the ages, Q3 minus Q1:

$$38 - 20.125 = 17.875$$

Descriptive statistics describe the centre and the spread:

- **Standard deviation:** 14.53, and **variance:** 211.02, its square.
- **Skewness:** 0.39, slightly skewed to the right (see [skewness](#)).
- **Sum:** 21,205.17, all known ages added up.
- **Coefficient of variation, kurtosis, median absolute deviation and monotonicity** are new; the Extra boxes below explain them.

Extra: The **coefficient of variation (CV)** compares the spread with the average.

1. **In words:** divide the standard deviation by the mean. The result has no units, so

it can compare the spread of columns measured in different units (years, rupees).

2. **Formula:**

$$CV = \frac{\text{standard deviation}}{\text{mean}}$$

3. **With the Age numbers:**

$$CV = \frac{14.53}{29.70} = 0.49$$

So ages typically lie about half the mean (around 15 years) away from the mean.

Extra: The **median absolute deviation (MAD)** is a measure of spread built from medians instead of means.

1. **In words:** find the median; measure how far each value is from it; the MAD is the median of those distances.

2. **Formula:**

$$MAD = \text{median}(|x_i - \text{median}(x)|)$$

3. **With five values** 1, 2, 3, 4, 10:

- the median of the values is 3;
- the distances from 3 are 2, 1, 0, 1, 7;
- sorted, the distances are 0, 1, 1, 2, 7;
- their median, the middle one, is the MAD:

$$MAD = 1$$

The far value 10 hardly changes the MAD, while it would raise the standard deviation a lot. For **Age**, the MAD is 9 years.

Extra: Two more entries in Figure 7:

- **Kurtosis** measures how heavy the tails of a distribution are compared with a normal (bell-shaped) curve. The report gives 0 for a normal curve and a positive value when extreme values are more common than in one. **Age** has 0.18, close to normal. In its simple form, with m_2 the average of the squared distances from the mean (the variance) and m_4 the average of their fourth powers:

$$\text{kurtosis} = \frac{m_4}{m_2^2} - 3$$

Two lists of ten values show the effect. Nine values of 5 and one value of 15 have a single extreme value. Step by step:

$$\text{mean} = (9 \times 5 + 15)/10 = 6$$

The distances from 6 are -1 (nine times) and 9 (once).

$$m_2 = (9 \times 1 + 81)/10 = 9$$

$$m_4 = (9 \times 1 + 6561)/10 = 657$$

$$\text{kurtosis} = 657/81 - 3 = 5.11$$

The values 1 to 10 are spread evenly and have no extreme value; the same steps (mean 5.5, $m_2 = 8.25$, $m_4 = 120.86$) give

$$120.86/8.25^2 - 3 = -1.22$$

The extreme value is what makes the number large and positive.

- **Monotonicity** says whether the values only ever go up (or only down) from one row to the next. Ages in a passenger list are “not monotonic”; a running number like `PassengerId` is increasing. For example, 1, 2, 5, 9 only goes up (increasing), while 3, 1, 2 goes down and then up (not monotonic).

The last two tabs deal with single values:

- **Common values:** the ages that occur most often. The top three are 24 (30 passengers), 22 (27) and 18 (26).
- **Extreme values:** the five smallest and five largest values. The smallest ages are 0.42, 0.67, 0.75, 0.75 and 0.83, all babies under a year; the largest are 80, 74, 71, 71 and 70.5.

The extreme values are where outliers show up first. Here they look genuine: babies and elderly passengers really were on board.

5. Interactions: a scatter plot for any pair

Key point: The Interactions section draws a scatter plot for any two numerical columns we pick, which is bivariate analysis on demand.

The Interactions section holds a scatter plot for every pair of numerical columns. Two rows of tabs pick the pair: the top row sets the x-axis column, the second row the y-axis column. Figure 8 shows `Age` on the x-axis against `Fare` on the y-axis.

Interactions




Figure 8: The Interactions section: Age against Fare

Most fares are below 100, at every age; the dots do not rise or fall with age, so the two columns are hardly related. A few fares stand far above the rest, the highest at 512. These are the extreme fares that the box plot of **Fare** flagged as outliers in [the fare box plot](#).

Picking the same column for both axes gives no useful plot: every dot would lie on one diagonal line. Only numerical columns appear here; categorical pairs are covered by the Correlations section.

6. Correlations: every pair in one heatmap

Key point: The Correlations section shows the correlation of every pair of columns as a heatmap; dark blue is a strong positive link, dark red a strong negative one.

Correlation measures how strongly two columns move together, from -1 to $+1$, as defined in [correlation](#). The Correlations section computes it for every pair and draws the results as a heatmap. Figure 9 shows **Pearson's r**, the coefficient for straight-line relationships between numbers.




Figure 9: Pearson's r for every pair of numerical columns

Pearson's r needs numbers. **Survived** and **Pclass** are categorical, but their values are stored as numbers (0 and 1; 1, 2 and 3), so the heatmap includes them. **Sex** and **Embarked** are stored as text, so they are left out.

Each cell is the correlation of its row and its column. The colour bar on the right reads the colours: dark blue near $+1$, white near 0, dark red near -1 . The diagonal is always dark blue, since each column matches itself perfectly.

Away from the diagonal, a few cells stand out:

- **Pclass and Fare:** -0.55 . The strongest link. A higher class number (cheaper class) goes with a lower fare.
- **SibSp and Parch:** $+0.41$. Passengers with siblings or a spouse aboard often had parents or children aboard too: they travelled as families.
- **Survived and Pclass:** -0.34 , and **Survived and Fare:** $+0.26$. Passengers in better classes, who paid more, survived more often.

So at a glance we learn which **features** (input variables, one column each) relate to the **target** (**Survived**, the output we predict), and which features relate to each other. Both matter later: features linked to the target are useful, and features strongly linked to each other repeat the same information.

Python: Choosing the coefficients.

```
profile = ProfileReport(df, correlations={
```

```
"auto": {"calculate": True},
"pearson": {"calculate": True},
})
```

By default only the “Auto” tab is computed. The Notebook switches on Pearson’s r as well, as above. The same way, “spearman”, “kendall”, “phi_k” and “cramers” add more tabs.

Extra: Pearson’s r only works for two numerical columns, and only sees straight-line links. Other coefficients cover the other cases:

- **Spearman’s** and **Kendall’s** coefficients measure whether two columns rise together at all, even along a curve, by comparing the ranks of the values instead of the values (Spearman 1904; Kendall 1938). Example: $x = 1, 2, 3, 4$ and $y = 1, 4, 9, 16$ lie on a curve. Their ranks are 1, 2, 3, 4 in both lists, identical, so Spearman’s coefficient is exactly 1, while Pearson’s r is 0.98 because the points are not on a straight line.
- **Cramér’s V** measures the link between two categorical columns, from 0 (none) to 1 (complete) (Cramér 1946). For **Sex** and **Survived** the counts are 81 females who died, 233 females who survived, 468 males who died and 109 males who survived; the computer gives $V = 0.54$, a strong link.
- **Phik** (ϕ_k) works for any mix of numerical and categorical columns (Baak et al. 2020).


The default “Auto” heatmap picks a suitable coefficient for each pair: Spearman’s for two numerical columns, and Cramér’s V when a categorical column is involved. The Auto heatmap is how the report found the strong link between **Sex** and **Survived** in the alerts: **Sex** is text, so Pearson’s r cannot include it at all. Every pair above 0.5 on the Auto heatmap gets a “High correlation” alert (0.5 is the library’s default **threshold**).

Extra: Next to each Heatmap tab is a Table tab with the exact numbers behind the colours.

7. Missing values: how many, and where

Key point: The Missing values section shows, per column, how many values are present (Count) and in which rows the gaps are (Matrix).

The fifth section is all about missing values. Its first tab, Count (Figure 10), draws one bar per column; the bar’s height is the share of rows that have a value.




A simple visualization of nullity by column.

Figure 10: The Count tab of the Missing values section

The numbers above the bars are the counts of present values. Most columns are complete, with all 891. Three are not:

- **Age:** 714 present, so 177 missing.
- **Cabin:** only 204 present, so 687 missing.
- **Embarked:** 889 present, so 2 missing.

The second tab, Matrix (Figure 11), draws the whole table as a picture, a **nullity matrix** (“nullity” means missing). Each column of the data is a vertical strip, and each row of the data a thin horizontal line through all strips, from row 1 at the top to row 891 at the bottom. A filled line means the value is there; a white line means it is missing.



Nullity matrix is a data-dense display which lets you quickly visually pick out patterns in data completion.

Figure 11: The Matrix tab of the Missing values section

The more white, the more missing values. **Cabin** is mostly white, **Age** has white lines scattered all the way down, and **Embarked** has just two: the one near the top (row 62) is visible, while the second, at row 830, is too thin to show at this size. The matrix also shows whether gaps in different columns fall in the same rows; here they do not line up in any clear pattern.

Extra: The third tab, Heatmap, shows whether the gaps in one column tend to occur together with gaps in another. Missing values and how to handle them get their own Notes later, starting with [why missing values must be handled](#).

8. Sample: a look at the rows

Key point: The Sample section shows the first and last rows of the data, like `df.head()` and `df.tail()`.


The last section shows real rows, in two tabs: First rows and Last rows. They give a feel for what the values look like, as `df.head()` did in [what the data looks like](#). The last rows also show whether the end of the file differs from the start, for example a total row or a different format.

9. Using the report

Key point: We run a profiling report first on every new dataset, read it section by section, and write down what we notice.

The report covers in seconds what took the [first questions](#), [univariate analysis](#) and [bivariate and multivariate analysis](#) by hand. The report does not replace understanding: it lists facts, and we decide what they mean. A good way to use it (Figure 12):

1. Build the report as soon as we get a new dataset.
2. Read it in order: Overview and Alerts, then each variable, then Interactions, Correlations and Missing values.
3. Write down every observation, for example “**Cabin** is 77% missing” or “**Pclass** and **Fare** are strongly linked”.
4. Turn the observations into tasks: columns to drop, missing values to fill, outliers to check.



the report lists facts; we decide what they mean

Figure 12: Using the report: read the sections in order, write down what we notice (the two examples above), and turn each observation into a task

Reading reports becomes faster with practice. Running the library on three or four different datasets, and writing down observations each time, builds the habit of knowing where to look.

Extra: The report knows nothing about the meaning of the data. The report cannot tell that **Survived** is the target we want to predict, that a 0 fare might be a crew member or a free ticket, or that **PassengerId** is a label rather than a measurement. Those judgements still come from us, and from hand-made EDA when a question needs a closer look.

10. Summary

Section	What it shows	Hand-made equivalent	Titanic example
Overview	Size, types, missing cells, duplicates, memory	<code>shape</code> , <code>info</code> , <code>isnull</code> , <code>duplicated</code>	12 columns, 891 rows, 8.1% missing
Alerts (in Overview)	Possible problems, one per line	Reading every column ourselves	<code>Cabin</code> 77.1% missing
Variables	Each column: statistics and a chart	Univariate analysis	Age: mean 29.70, IQR 17.875
Interactions	Scatter plot of any two numerical columns	Scatter plots	A few fares above 500
Correlations	Correlation of every pair, as a heatmap	<code>df.corr()</code> and a heatmap	<code>Pclass</code> and <code>Fare</code> : -0.55
Missing values	Present values per column, and where the gaps are	<code>isnull().sum()</code>	Age 714, <code>Cabin</code> 204 of 891
Sample	First and last rows	<code>head</code> , <code>tail</code>	

- `ProfileReport(df).to_file("report.html")` builds a full EDA report as one web page.
- The library has had three names: `pandas-profiling`, `ydata-profiling`, and now `fg-data-profiling` (`import data_profiling`).
- Alerts point to columns to check: high correlation, missing, unique, uniform, zeros.
- “Auto” correlations cover categorical columns too; Pearson’s r only covers numbers.
- `minimal=True` keeps the report fast on big data.
- The report lists facts; we still read it, write down observations and decide what to do.

11. Sources

Built from

- CampusX, “Pandas Profiling | Day 22 | 100 Days of Machine Learning”, YouTube, <https://www.youtube.com/watch?v=E69Lg2ZgOxg>

Other references

- Baak, M., Koopman, R., Snoek, H. and Klous, S. (2020). A new correlation coefficient between categorical, ordinal and interval variables with Pearson characteristics. *Computational Statistics & Data Analysis* 152.
- Cramér, H. (1946). *Mathematical Methods of Statistics*. Princeton University Press.
- Kendall, M. (1938). A New Measure of Rank Correlation. *Biometrika* 30(1/2).
- PyPI. `pandas-profiling`, `ydata-profiling` and `fg-data-profiling` project pages. pypi.org.

- Spearman, C. (1904). The Proof and Measurement of Association between Two Things. *American Journal of Psychology* 15(1).

12. Key terms

Term	Meaning
Profiling report	An automatic EDA report describing every column and pair of columns of a dataset
Pandas Profiling	The library that builds a profiling report from a DataFrame, now named fg-data-profiling
Variable	The report's word for a column
Observation	One record of the data; the report's word for a row
Feature	An input variable, one column of the data table
Target	The output we want to predict
Average record size	The memory one row takes, on average
Alert	A warning in the report about a column that may need attention
High cardinality	A categorical column with very many different categories
Coefficient of variation (CV)	Standard deviation divided by mean: spread relative to the average
Median absolute deviation (MAD)	The median distance of the values from their median
Kurtosis	How heavy the tails of a distribution are compared with a normal curve
Monotonicity	Whether a column's values only go up, or only go down, from row to row
Pearson's r (G-1474)	The correlation coefficient for straight-line relationships between two numerical columns
Cramér's V	A measure of the link between two categorical columns, from 0 to 1
Nullity matrix	A picture of the whole table with missing values drawn as white lines