

Feature Scaling: Standardization

Contents

1. Overview	2
2. Feature scaling, in brief	3
3. Types of feature scaling	4
4. The standardization formula	4
4.1 The idea: how many standard deviations from the mean	5
4.2 The formula	6
5. What standardization does to the data	7
6. Standardization in scikit-learn	8
6.1 The data	8
6.2 Split before scaling	8
6.3 Fit on the training set, transform both	9
6.4 Checking the result with describe	10
7. Effect of scaling on the data	11
7.1 The scatter plot keeps its shape	11
7.2 The two columns become comparable	12
7.3 Each column keeps its shape	12
8. Why scaling matters: an experiment	13
9. Standardization and outliers	15
10. When to use standardization	16
11. Summary	18
12. Sources	18
13. Key terms	19

Where this fits: Pipeline map step 5 (Engineer features). See the [Course map](#).




- **Builds on:** [Feature engineering](#); [Train-test split](#); [Descriptive statistics](#); [Feature transformation](#).
- **Leads to:** [Normalization](#); [K-means](#); [Z-score outlier method](#); [PCA](#); [Gradient descent](#); [Ridge regression](#).
- **Compare with:** [Normalization](#); [Decision trees](#).

1. Overview

Key point: Standardization rescales every feature so that its mean is 0 and its standard deviation is 1: subtract the mean, then divide by the standard deviation.

Feature scaling brings **features** (G-772) (input variables, one column each of the data table) that live on very different ranges, such as age and salary, onto a similar range. Feature scaling has two main techniques: standardization and normalization. This Note covers standardization; normalization comes in the next Note.

Figure 1 shows the whole topic. Standardization is two moves done one after the other: shift the column so its mean is 0, then shrink or stretch it so its standard deviation is 1.



$$z = \frac{x - \text{mean}}{\text{std}}$$

Figure 1: Feature scaling and its two techniques; standardization is mean centring followed by scaling

2. Feature scaling, in brief

Key point: Feature scaling brings the features to a similar range, so a feature with big numbers does not drown out one with small numbers.

Feature scaling (G-767) brings the features of a dataset into a similar range. [Scaling the inputs](#) introduced the reason; here we work it through on this Note's data (Section 6.1), where each user has an age and a salary.

The plain idea. Algorithms such as KNN decide which users are alike by measuring the distance between them. A salary gap is in the thousands and an age gap is in the tens, so the salary gap decides the distance alone.

The standard term. The distance used is the **Euclidean distance** (G-715): the straight-line distance between two points. For two users it is

$$d = \sqrt{(\text{age gap})^2 + (\text{salary gap})^2}.$$

Worked example. Take a user aged 26 with a salary of 15,000 rupees, and two others:

Other user	Age gap	Salary gap	Distance
age 42, salary 53,000	16	38,000	$\sqrt{256 + 1,444,000,000} \approx 38,000$
age 35, salary 59,000	9	44,000	$\sqrt{81 + 1,936,000,000} \approx 44,000$

The age gaps add 256 and 81 to numbers above a billion: they change nothing. The 42-year-old counts as the nearer user only because the salaries are closer.

Figure 2 draws the squared distance to five users as bars, split into the age part (blue) and the salary part (orange). Watch the blue part: on raw data it is too small to see. After standardization, both parts show, and the nearest user changes from the 42-year-old to the 35-year-old, who is closer in age.




Figure 2: The squared distance from one user to five others, split into its age part and its salary part, on raw and on standardized data

Two more points about feature scaling:

- Only the features are scaled, never the **target** (G-1949) (the output we predict).
- Scaling is usually the last step of feature engineering: we first handle missing values, transform columns and deal with categories, then scale just before giving the data to the model.

3. Types of feature scaling

Key point: The two types of feature scaling are standardization and normalization.

There are two main types:

1. **Standardization** (G-1874), the subject of this Note.
2. **Normalization** (G-1349), which has several techniques of its own. The main one is **min-max scaling** (G-1217); another is the **robust scaler** (G-1699), which copes well with outliers. Both are covered in the next Note.

Standardization is sometimes called **z-score normalization** (G-2140) in ML writing, because each scaled value is a **z-score** (G-2141). The two names mean the same technique.

4. The standardization formula

Key point: Each value is replaced by its distance from the column's mean, measured in standard deviations.

4.1 The idea: how many standard deviations from the mean

Key point: A z-score answers one question about a value: how many standard deviations above or below the mean is it?

The **mean** is the average of a column. The **standard deviation** is the typical distance of the values from that mean (both are taught in [count, mean, standard deviation, minimum and maximum](#)). The **z-score** (G-2141) of a value is the number of standard deviations the value lies from the mean.

Figure 3 builds the z-scores of eight made-up ages, chosen so that the arithmetic is easy: the mean is 30 and the standard deviation is 4. Watch the numbers under the dots: the dots stay where they are, and only their numbers change, in two steps.



Figure 3: Eight ages turned into z-scores in two steps: subtract the mean, 30, then divide by the standard deviation, 4

Age	Age - 30	Divided by 4 (z-score)	In words
24	-6	-1.5	1.5 standard deviations below the mean
28	-2	-0.5	half a standard deviation below
30	0	0	exactly at the mean
34	4	1	1 standard deviation above
38	8	2	2 standard deviations above

Step 1 is **mean centring** (G-1195): subtracting the mean, so that the mean itself becomes 0. Step 2 changes the unit from years to standard deviations. Doing both steps to every value of a column is standardization.

4.2 The formula

Key point: $x' = (x - \bar{x})/\sigma$, applied to every value of the column.

Take a column **age** holding the ages of 500 customers (and the data has a salary column too). To standardize **age**, we compute a new value for every **observation** (G-1374) (one record, one row of the table), one by one.

Standardization, step by step:

1. **In words:** from each value, subtract the mean of the whole column, then divide by the standard deviation of the whole column.
2. **Formula:** for the i -th value x_i of a column with mean $\bar{x}$ and standard deviation σ ,

$$x'_i = \frac{x_i - \bar{x}}{\sigma}$$

3. **Example:** suppose the 500 ages have mean 32 and standard deviation 10. A customer aged 27 becomes

$$x' = \frac{27 - 32}{10} = -0.5,$$

and a customer aged 15 becomes

$$x' = \frac{15 - 32}{10} = -1.7.$$

Figure 4 draws the two worked ages on two number lines, one in years and one in standard deviations. Watch the mean, 32, land on 0, and each step of 10 years become a step of 1.




Figure 4: The ages 27 and 15 on the age line (mean 32, std 10) and the same two customers on the z line

Doing this for all 500 rows gives a new column of 500 numbers, such as 2.3, -1.2 and -0.5. Each one says how many standard deviations the original value lay above (positive) or below (negative) the mean.

The new column always has two fixed properties, whatever the original numbers were:

- its **mean is 0**;
- its **standard deviation is 1**.

Extra: Why the new mean is always 0 and the new standard deviation always 1. Subtracting $\bar{x}$ from every value moves the mean down by exactly $\bar{x}$, to 0, without changing the spread. Dividing every value by σ divides the spread by σ too, so the new standard deviation is $\sigma/\sigma = 1$.

5. What standardization does to the data

Key point: Standardization first moves the cloud of points so its centre sits at the origin (mean centring), then squeezes or stretches each axis until its standard deviation is 1.

Plot every row as a point, with one column on each axis. The 500 customers form a cloud somewhere away from the origin, with some spread along each axis. Standardization moves this cloud in two steps (Figure 5):

1. **Mean centring:** subtracting the mean slides the whole cloud, unchanged, until its mean sits at the origin $(0,0)$.
2. **Scaling by the standard deviation:** dividing by the standard deviation changes the spread along each axis to exactly 1.




Figure 5: Standardization in two steps: shift the mean to the origin, then make the spread 1 on each axis

In Figure 5, the red cross marks the mean and the dashed box reaches one standard deviation either side of it. Feature 1 starts with standard deviation 2 and feature 2 with 0.6.

The second step works differently for the two axes:

- **Standard deviation above 1** (feature 1): dividing by it makes the values smaller, so the points **squeeze in**.
- **Standard deviation below 1** (feature 2, here 0.6): dividing by it makes the values bigger, so the points **spread out**.

At the end, both axes have mean 0 and standard deviation 1, and the dashed box is a square from -1 to 1.

6. Standardization in scikit-learn

Key point: Split the data first, then let `StandardScaler` learn the mean and standard deviation from the training set and apply them to both sets.

6.1 The data

Key point: 400 users of a social network, with age, estimated salary and whether they bought a product.

The data is the Social Network Ads file: 400 users of a social network, and whether each one bought a product shown in an advertisement. We keep three columns:

Age	EstimatedSalary	Purchased
19	19,000	0
35	20,000	0
26	43,000	0

Age and EstimatedSalary are the features; Purchased (1 = bought, 0 = did not) is the target. So this is a classification problem.

Python: Loading the data and keeping three columns.

```
import pandas as pd
```

```
df = pd.read_csv("data/Social_Network_Ads.csv")
```

```
# drop User ID and Gender
```

```
df = df.iloc[:, 2:]
```

6.2 Split before scaling

Key point: Always do the train-test split before any feature scaling.

Whether we standardize or normalize, the **train-test split** (G-1998) comes first. The scaler should learn only from the training set, as **training and test sets** explained under **data leakage** (G-535; information from the test set leaking into training). With 30% of the rows held back for testing, the training set has 280 rows and the test set 120.

Python: Splitting the data.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    df.drop(columns="Purchased"),  # inputs
    df["Purchased"],               # target
    test_size=0.3, random_state=0)
# X_train: (280, 2), X_test: (120, 2)
```

6.3 Fit on the training set, transform both

Key point: **fit** learns the mean and standard deviation of each column; **transform** applies the formula to every value.

scikit-learn's **StandardScaler** (G-1876) class does exactly what Section 4 described: it takes every value and applies $(x_i - \bar{x})/\sigma$. Using it has two steps:

- **fit** (G-783): learn each column's mean and standard deviation from the training set, and store them.
- **transform**: apply the formula to every value, using the stored numbers.

We fit on the training set only, but transform both the training set and the test set. The test set is scaled with the training set's mean and standard deviation, never its own.

Figure 6 shows the flow with the real numbers. Watch the arrows: only the training set feeds **fit**, while both sets pass through **transform**.




Figure 6: fit learns the mean and standard deviation from the 280 training rows; transform applies them to both sets, here to the first training row

Python: Standardizing with StandardScaler.

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
# learn mean and std from the training set only
scaler.fit(X_train)
# apply them to both sets
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
scaler.mean_    # [37.86, 69807.14]
```

```
scaler.scale_   # [10.20, 34579.29]
```

After fit, the learned numbers are stored in `mean_` (the means) and `scale_` (the standard deviations), one per column.

`transform` gives back a NumPy array, which has no column names and is harder to read. We turn it back into a DataFrame.

Python: Putting the column names back.

```
X_train_scaled = pd.DataFrame(
    X_train_scaled, columns=X_train.columns)
X_test_scaled = pd.DataFrame(
    X_test_scaled, columns=X_test.columns)
```

The same formula, worked on a real row:

1. **In words:** the first training row is a 26-year-old with a salary of 15,000 rupees. We subtract each column's training mean and divide by its training standard deviation.

2. **Formula:**

$$x' = \frac{x - \text{mean}}{\text{std}}$$

3. **Example:** for the age and the salary,

$$\text{age: } \frac{26 - 37.86}{10.20} \approx -1.16$$

$$\text{salary: } \frac{15000 - 69807}{34579} \approx -1.58$$

These are exactly the numbers in the first row of `X_train_scaled`.

6.4 Checking the result with describe

Key point: After scaling, both columns have mean 0 and standard deviation 1.

`describe()` before and after scaling shows the change (rounded to one decimal place):

	Age before	Age after	Salary before	Salary after
mean	37.9	0.0	69,807.1	0.0
std	10.2	1.0	34,641.2	1.0
min	18.0	-1.9	15,000	-1.6
max	60.0	2.2	150,000	2.3

Before scaling, the two columns live in completely different ranges. After scaling, both have mean 0 and standard deviation 1, and both run from about -2 to 2.

Extra: `scaler.scale_` (34,579) and `describe()` (34,641) give slightly different salary standard deviations because pandas divides by $n - 1$ and `StandardScaler` by n ; the scaled column's std shows as 1.0 either way.

7. Effect of scaling on the data

Key point: Scaling changes the numbers on the axes, not the shape of the data.

7.1 The scatter plot keeps its shape

Key point: The cloud of points looks the same; only its centre and its units change.

Figure 7 plots every training user by age and salary, before and after scaling. The two clouds look the same: every point keeps its place relative to the others.




Figure 7: The training data before and after standardization

The difference is on the axes. Before scaling, age runs from about 20 to 60 and salary from 15,000 to 150,000. After scaling, both axes are centred on 0 and run from about -2 to 2: the data has been mean centred and given standard deviation 1.

7.2 The two columns become comparable

Key point: Before scaling, the two density curves cannot be compared; after scaling, they sit on the same range.

Figure 8 draws the **kernel density estimate (KDE)** (G-1005), a smooth curve of where the values lie, of both columns on one axis. Before scaling, age lives in a very small range (18 to 60), so its curve is a tall, thin spike near 0. Salary is spread over a huge range, so its curve is almost a flat line: its density is around 0.000003.




Figure 8: Density curves of age (blue) and salary (orange), before and after standardization

The two curves cannot be compared because their scales differ so much. After standardization, both curves sit on the same range, around 0. Any algorithm that compares the columns now treats them fairly, which helps the algorithms that measure distances or train by gradient descent (Section 8).

7.3 Each column keeps its shape

Key point: Standardization does not change the shape of a column's distribution.

Figure 9 shows each column on its own. Age before scaling and age after scaling have exactly the same curve; so do salary before and after.



Figure 9: Each column’s distribution before and after: the shape stays, the numbers change

Only the numbers under the curve change: the mean becomes 0 and the standard deviation 1. If a column was skewed (lopsided, with a long tail on one side) before, it is just as skewed after standardization.

8. Why scaling matters: an experiment

Key point: On this data, scaling lifts KNN from 82.5% to 91.7% accuracy and logistic regression from 65.8% to 86.7%, while a decision tree gets 87.5% either way.

Think of judging two runners, one timed in seconds and one in milliseconds: until the units match, the bigger numbers win by default. We train the same algorithm twice: once on the raw data, once on the standardized data. Then we compare the accuracy on the test set.

Algorithm	Raw data	Standardized data
KNN (5 neighbours)	82.5%	91.7%
Logistic regression	65.8%	86.7%
Decision tree	87.5%	87.5%

For **KNN**, which measures distances, scaling raised the accuracy from 82.5% to 91.7%. On raw data, salary differences in thousands decide the distance alone; after scaling, age counts too, as Figure 2 showed (scikit-learn examples, “Importance of Feature Scaling”).

Figure 10 shows the two KNN models on the same axes, in the original units. Each coloured area is a **decision region** (G-557): the set of points that KNN would label with that class. The line where the two colours meet is the **decision boundary** (G-555). KNN labels a

point by its 5 nearest training users. On raw data, the distance between two users is almost entirely their salary gap, because a salary gap of thousands dwarfs an age gap of tens. So the 5 nearest users are simply those with the closest salaries, whatever their age, and the raw model’s decision boundary runs flat across the plot. After standardization, one standard deviation of age counts as much as one standard deviation of salary, so the boundary bends with age as well.




Figure 10: Where KNN (5 neighbours) predicts “bought” (orange) and “did not buy” (blue), trained on raw and on standardized data; dots are the 120 test users

For **logistic regression** (G-1120), scaling raised the accuracy from about 66% to about 87%. The Extra below shows the cause: training by small downhill steps barely moves on raw data.

For a **decision tree** (G-561), scaling made no difference at all. A decision tree does not depend on the scale of the columns, for reasons that become clear in the Notes on decision trees.

Python: The experiment for logistic regression.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

lr = LogisticRegression(solver="sag")
lr_scaled = LogisticRegression(solver="sag")
lr.fit(X_train, y_train)
lr_scaled.fit(X_train_scaled, y_train)

accuracy_score(y_test, lr.predict(X_test))
# 0.658
accuracy_score(y_test, lr_scaled.predict(X_test_scaled))
```

0.867

For the decision tree, replace `LogisticRegression(solver="sag")` with `DecisionTreeClassifier()` from `sklearn.tree`.

Extra: Why `solver="sag"`? A **solver** (G-1836) is the method a model uses to find its best settings during training. "`sag`" takes many small downhill steps, like gradient descent (Section 10), and scikit-learn warns that it is only fast when the features have about the same scale (scikit-learn docs, `LogisticRegression`). Here salaries reach 150,000 while ages stay below 60, so each step is tiny and the weights hardly move from 0. The model then predicts “not purchased” for every user, and 65.8% is simply the share of non-buyers in the test set. The Notebook checks this: allowed up to 100,000 steps, `sag` stops after about 9,800 with the weights still close to 0 and the accuracy still 65.8%.

The default solver, "`lbfgs`", does find good weights on the raw data too (87.5% on raw data, 86.7% on scaled data), but needs 65 steps on raw data against 7 on scaled data. The 0.8-point gap between those two accuracies is one test user out of 120, so it is noise, not a sign that scaling hurt. Either way, unscaled data makes the training harder.

In this experiment, scaling never hurt KNN, logistic regression with `sag` or the decision tree. In general, standardizing does no harm, but for some algorithms it helps a lot.

9. Standardization and outliers

Key point: Standardization does not remove outliers or reduce their effect; they stay just as far from the rest of the data.

The original data has ages from 18 to 60 and salaries from 15,000 to 150,000 rupees. To see what happens to outliers, we add three made-up users far outside these ranges:

Age	EstimatedSalary	Purchased
5	1,000	0
90	250,000	1
95	350,000	1

After adding them, we split, standardize and plot the data again. Two of the three land in the training set (the third went to the test set), and Figure 11 shows them in red.




Figure 11: Two added outliers (red) before and after standardization

After scaling, the outliers are still outliers. They sit just as far from the main cloud as before; only the numbers on the axes changed.

So whenever we standardize a column that has **outliers** (G-1420), we must deal with the outliers separately. Later Notes cover how to detect and handle them.

Python: Adding rows to a DataFrame.

```
extra = pd.DataFrame({
    "Age": [5, 90, 95],
    "EstimatedSalary": [1000, 250000, 350000],
    "Purchased": [0, 1, 1]})
df = pd.concat([df, extra], ignore_index=True)
pd.concat joins tables one below the other; ignore_index=True rennumbers the rows 0,
1, 2, ... Older code uses df.append(...), which was removed in pandas 2.0.
```

10. When to use standardization

Key point: Standardize for algorithms that measure distances or train by gradient descent; tree-based algorithms do not need it.

Standardizing rarely does harm, but for the algorithms below we should always do it.

Algorithms that need scaling:

- **K-means** (G-996) and **KNN** (k-nearest neighbours, G-998): both compute the Euclidean distance between points. If one column's numbers are much bigger, it dominates the distance and the results are poor.
- **PCA** (G-1469) (principal component analysis): PCA looks for the directions in which the data spreads the most (the most variance). A column with big numbers looks like it has the most spread just because of its units, so the columns must be put on the same scale, and PCA also needs the data mean centred. For example, an age column

with a spread of 10 years has a variance of 100, while a salary column with a spread of 20,000 rupees has a variance of 400,000,000. PCA would find that salary “has all the spread” only because rupees are small units.

- **Gradient descent** (G-862), and every algorithm trained with it: linear regression, logistic regression and neural networks (deep learning).

Gradient descent (taught in [the idea of gradient descent](#)) improves the weights by small downhill steps. When the columns are on very different scales, some weights take big jumps while others crawl, so it struggles to settle at the minimum; with scaled columns it converges much more easily. A reason in numbers: the step for a column’s weight is proportional to the column’s value (error times value). For a user aged 25 with a salary of 50,000, one error of 1 pushes the age weight by 25 units and the salary weight by 50,000, a ratio of 2,000 to 1. Scaled, both values are near 1 and the pushes are similar.

Algorithms that do not need scaling:

- decision trees;
- random forest;
- gradient boosting;
- XGBoost.

Tree-based algorithms only compare values within one feature, asking questions like “is age greater than 40?”. They never compute a distance or add up different columns, so the scale of the columns does not matter. Scaling them does no harm, but makes no difference either.

Extra: Why a tree does not care. Suppose a tree splits on “age > 40”. After standardization, the same split becomes “scaled age > 0.21”, which is $(40 - 37.86)/10.20$. Exactly the same users fall on each side, so the tree makes exactly the same predictions.

11. Summary

	Before standardization	After standardization
Mean of each column	anything (age 37.9, salary 69,807)	0
Standard deviation	anything (age 10.2, salary 34,641)	1
Typical range	depends on the units	about -2 to 2
Shape of the distribution	some shape	the same shape
Outliers	far from the rest	still far from the rest

Algorithm	Needs scaling?	Why
KNN, K-means	yes	Euclidean distance (KNN: 82.5% to 91.7% here)
PCA	yes	looks for the largest variance; needs mean centring
Linear regression, logistic regression, neural networks	yes	trained with gradient descent
Decision tree, random forest, gradient boosting, XGBoost	no	only compare values within one column

- Feature scaling is usually the last step before the model, and only the features are scaled.
- Standardization: $x' = (x - \bar{x})/\sigma$, giving mean 0 and standard deviation 1.
- A z-score says how many standard deviations a value lies above or below the mean.
- Geometrically: mean centring, then squeezing or stretching each axis to standard deviation 1.
- Split first; fit the scaler on the training set only; transform both sets.
- Standardization keeps the shape of the data and does not remove outliers.
- On the ads data, KNN went from 82.5% to 91.7% accuracy and logistic regression from 65.8% to 86.7%; a decision tree stayed at 87.5%.

12. Sources

Built from

- CampusX, “Feature Scaling - Standardization | Day 24 | 100 Days of Machine Learning”, YouTube, <https://www.youtube.com/watch?v=1Yw9sC0PNwY>
- Khan Academy, “Z-score introduction | Modeling data distributions”, YouTube, <https://www.youtube.com/watch?v=5S-Zfa-vOXs>

Other references

- scikit-learn examples. Importance of Feature Scaling. [scikit-learn.org \(auto_examples/preprocessing\)](https://scikit-learn.org/stable/auto_examples/preprocessing/).
- scikit-learn documentation. `sklearn.linear_model.LogisticRegression` (note on the `sag` and `saga` solvers). [scikit-learn.org](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html).

- [scikit-learn documentation](#). `StandardScaler`, in the module `sklearn.preprocessing`. [scikit-learn.org](#).

13. Key terms

Term	Meaning
Feature	An input variable, one column of the data table
Target	The output we predict
Observation	One record, one row of the data table
Feature scaling (G-767)	Putting features on the same scale, so no feature dominates distances
Euclidean distance	The straight-line distance between two points
Standardization	Scaling a column by subtracting its mean and dividing by its standard deviation, so it gets mean 0 and std 1
Z-score	A value after standardization: how many standard deviations it lies from the mean
Z-score normalization	Another name for standardization
Normalization (G-1349)	The other type of feature scaling, which squeezes values into a fixed range (next Note)
Min-max scaling	The main normalization technique
Robust scaler (G-1699)	A normalization technique that copes well with outliers
Mean centring	Subtracting the mean from every value, so the column's mean becomes 0
StandardScaler	scikit-learn's class that standardizes columns with <code>fit</code> and <code>transform</code>
fit / transform	Learn the scaler's numbers from the training set / apply them to any data
Kernel density estimate (KDE)	A smooth curve that shows where a column's values lie
Data leakage	Information from the test set leaking into training
Solver	The method a model uses to find its best settings during training
Gradient descent	Finding the lowest point of a function by repeated small steps downhill