

Handling Mixed Variables

Contents

1. Overview	2
2. What a mixed variable is	3
2.1 Type 1: number and category in the same cell	3
2.2 Type 2: number in some rows, category in others	4
3. The data	4
4. Splitting a Type 2 column	5
4.1 The numerical part	6
4.2 The categorical part	6
5. Splitting a Type 1 column: Cabin	8
5.1 The number part	8
5.2 The category part	8
6. Splitting a Type 1 column: Ticket	10
7. Filling the gaps the split leaves	11
8. Summary	14
9. Sources	14
10. Key terms	14

Where this fits: Pipeline map step 5 (Engineer features). See the [Course map](#).



- **Builds on:** Feature construction and splitting ([Note ML-003](#)); Feature engineering ([Note ML-007](#)).

1. Overview


Key point: A mixed variable holds numbers and categories in the same column; we split it into one numerical feature and one categorical feature.

A **feature** (G-772) is an input variable (one column of the data table), the **target** (G-1949) is the output we predict, and an **observation** (G-1374) is one record (one row). Every feature so far has been either numerical or categorical. Some real features are both at once. This Note shows the two ways a feature can be mixed and how to split each one.

Think of a postal address such as “Flat 12B”: a number and a letter packed into one label. To sort flats by floor or by block, we first pull the two parts apart.

Figure 1 shows both types and what splitting produces. The work is all done with pandas string tools.

Type 1: number and category in the same cell



Type 2: number in some rows, category in others




Figure 1: The two types of mixed variable, and the two columns each one is split into

2. What a mixed variable is

Key point: In Type 1, one cell holds a category and a number together; in Type 2, some cells hold a number and others a category.

A **mixed variable** (G-1237) is a feature whose values contain both numerical and categorical data. A model cannot use such a feature as it is: the values are neither clean numbers nor a short list of categories.

Mixed variables come in two types.

2.1 Type 1: number and category in the same cell

Key point: Each value carries two pieces of information, a letter part and a number part.

The Titanic **Cabin** column holds values such as C85 and C123. Each value packs two facts together:

- the letter is the **deck**, a category (C);
- the digits are the **cabin number** on that deck (85, 123).

Indian Railways coach labels work the same way: in S5 or B2, the letter gives the class and the number gives the coach.




Figure 2: A Type 1 value packs a category and a number into one cell

Figure 2 pulls the two parts of each value apart.

Treating C85 as one category gives far too many categories, since almost every cabin is different. The fix is to split each value into two new features: a categorical feature for the letter and a numerical feature for the number.

2.2 Type 2: number in some rows, category in others

Key point: The column holds whole numbers in most rows and a text label in a few rows.

In the second type, each cell holds only one thing, but not always the same kind of thing. A column might read 5, 3, A, 2, 1, A: numbers in most rows, the letter A in others.

The fix is again two columns:

- where the value is a number, put it in the numerical column, and leave the categorical column empty (**NaN**, “not a number”: the marker pandas uses for a missing value, see [handling missing values](#));
- where the value is a category, do the opposite.

These two types cover most mixed data met in practice. Anything else must be handled by studying the data and writing a fix for it.

3. The data

Key point: 891 Titanic passengers with three mixed columns: **Cabin** and **Ticket** (Type 1) and **number** (Type 2).

The data has four columns for the 891 passengers of the Titanic training file:

Column	Meaning	Type
Cabin	deck letter and cabin number, e.g. C85	mixed, Type 1
Ticket	ticket code, e.g. A/5 21171	mixed, Type 1
number	how many people the passenger travelled with; A = alone	mixed, Type 2
Survived	1 if the passenger survived, 0 if not	target

Cabin and **Ticket** are the real Titanic columns. The **number** column was added by hand to build a Type 2 example, so its values are made up and say nothing about survival.

Cabin	Ticket	number	Survived
C85	PC 17599	3	1
Type 1	Type 1	Type 2	target

passenger 1 of 891; number was added by hand

Figure 3: The four columns, shown on passenger 1, with the type of each

Figure 3 tags each column with its type.

The first five rows:

	Cabin	Ticket	number	Survived
0	NaN	A/5 21171	5	0
1	C85	PC 17599	3	1
2	NaN	STON/O2. 3101282	6	1
3	C123	113803	3	1
4	NaN	373450	A	0

Python: Loading the data.

```
import pandas as pd
```

```
df = pd.read_csv("data/titanic.csv")
```

```
df.dtypes
```

```
# Cabin, Ticket, number: str    Survived: int64
```

Since pandas 3, text columns get the dtype **str** (pandas' string type) instead of the old **object** (pandas release notes, What's new in 3.0.0). All three mixed columns are text, because every one of them contains at least one letter.

4. Splitting a Type 2 column

Key point: Convert the column to numbers, letting failures become NaN; the rows that failed are the categorical part.

The **number** column has seven values: 1 to 6 and A. A (alone) is the most common value, with 139 passengers; the numbers each have 117 to 131 (left of Figure 4).




Figure 4: Left: passengers per value of the mixed column `number`. Right: passengers per deck after splitting `Cabin`

4.1 The numerical part

Key point: `pd.to_numeric(..., errors="coerce")` turns every value it can read into a number and every other value into NaN.

`pd.to_numeric` (G-1473) converts values to numbers. By default it stops with an error at the first value it cannot read, such as `A`. With `errors="coerce"` (G-707), it writes NaN for that value instead and carries on.

So `5` becomes `5`, and `A` becomes NaN. The result is exactly the numerical column we want.

4.2 The categorical part

Key point: Keep the original value only in the rows where the number conversion failed.

The categorical column is the mirror image. Wherever the numerical part is NaN, the original value was not a number, so it goes into the categorical column. Everywhere else, the categorical column is NaN.

Figure 5 splits the first six rows one at a time. Watch row 4: `A` cannot be read as a number, so the numerical column gets NaN there and the categorical column gets `A`. In every other row it is the other way round.

row 3: **3** is a number → number_numerical

row	number	number_numerical	number_categorical
0	5	5	NaN
1	3	3	NaN
2	6	6	NaN
3	3	3	NaN
4	A		
5	2		

row 4: **A** is not a number → NaN, so it goes to number_categorical

row	number	number_numerical	number_categorical
0	5	5	NaN
1	3	3	NaN
2	6	6	NaN
3	3	3	NaN
4	A	NaN	A
5	2		

Figure 5: The column number split row by row: each value that `pd.to_numeric` can read goes to `number_numerical` (blue), the value A goes to `number_categorical` (orange), and the other cell of the row stays NaN.

Python: Splitting `number` into two columns.

```
num = pd.to_numeric(df["number"], errors="coerce")
df["number_numerical"] = num.astype("Int64")
# keep the text only where num is NaN
df["number_categorical"] = df["number"].where(num.isna())
s.where(cond) keeps each value of s where cond is True and puts NaN everywhere else.
```

The result on the first six rows:

	number	number_numerical	number_categorical
0	5	5	NaN
1	3	3	NaN
2	6	6	NaN
3	3	3	NaN
4	A	<NA>	A
5	2	2	NaN

Extra: Why `.astype("Int64")`. A NumPy integer column cannot hold NaN, so `pd.to_numeric` returns floats: 5.0, 3.0, NaN. The option `downcast="integer"`, often seen in older code, cannot change this while a NaN is present, and the column stays float. **Int64** (capital I) is pandas' **nullable integer** (G-1363) type: whole numbers plus a missing marker, shown as `<NA>` (pandas user guide, "Nullable integer data type").

5. Splitting a Type 1 column: Cabin

Key point: A regular expression pulls out the digits; the first character gives the deck letter.

Cabin has 147 different values, and 687 of the 891 passengers have no cabin at all (NaN). With 147 categories for only 204 known cabins, the feature is useless as it stands.

5.1 The number part

Key point: `.str.extract(r"(\d+)")` finds the first run of digits in each value.

A **regular expression** (G-1656) (regex) is a short pattern that describes text. The pattern `\d+` means “one or more digits”, and the brackets mark the part to keep. So `.str.extract(r"(\d+)")` returns 85 from C85 and 123 from C123. Two more values, to see the rule at work: from B57 B59 it returns 57 (only the first run of digits is taken), and from F G73 it returns 73 (letters and spaces before the digits are skipped).

The `.str` (G-47) accessor applies a text method to every value of a column at once. The accessor is how pandas reaches string tools such as `extract`, `split` and `isdigit`.

5.2 The category part

Key point: `.str[0]` takes the first character of each value: the deck letter.

`.str[0]` picks character number 0 of every value: C from C85, E from E46. A missing cabin stays NaN in both new columns.

Python: Splitting Cabin.

```
num = df["Cabin"].str.extract(r"(\d+)")[0]
df["cabin_num"] = pd.to_numeric(num).astype("Int64")
df["cabin_cat"] = df["Cabin"].str[0]
str.extract returns a table with one column per bracket, so [0] takes its first column.
The r before the quotes makes it a raw string (G-1638), where \d reaches the regex
unchanged.
```

The result, on some first rows and on every unusual kind of value:

	Cabin	cabin_num	cabin_cat
0	NaN	<NA>	NaN
1	C85	85	C
3	C123	123	C
27	C23 C25 C27	23	C
75	F G73	73	F
292	D	<NA>	D




Figure 6: Splitting **Cabin** on the rows of the table above: the first character becomes the deck, the first run of digits the number

Figure 6 shows the same rows as the table: watch the unusual values, where only the first deck letter and the first number survive.

Figure 7 runs the split one row at a time. Each row's first character lands in **cabin_cat** and its first run of digits in **cabin_num**; the heading says what happens to each unusual value.

C85: letter C, digits 85

row	Cabin	cabin_cat	cabin_num
0	NaN	NaN	NaN
1	C85	C	85
3	C123		
27	C23 C25 C27		
75	F G73		
292	D		

D: a letter but no digits: cabin_num stays empty

row	Cabin	cabin_cat	cabin_num
0	NaN	NaN	NaN
1	C85	C	85
3	C123	C	123
27	C23 C25 C27	C	23
75	F G73	F	73
292	D	D	NaN

Figure 7: Cabin split row by row for the six rows of the table: the first character goes to **cabin_cat** (orange), the first run of digits to **cabin_num** (blue); a missing part stays NaN.

After the split, `cabin_cat` has only 8 categories, the decks A to G and T (right of Figure 4). C is the most common deck, then B.

Extra: Some values break the simple rule. A passenger booked into several cabins, such as C23 C25 C27, keeps only the first number, 23. F G73 gives deck F, although the cabin is G73. D and T have no number, so `cabin_num` is missing. These are 14 of the 147 values; check them before relying on the split.

Extra: `str.extract` returns text, so on its own `cabin_num` would hold the strings "85" and "123". `pd.to_numeric` turns them into numbers.

6. Splitting a Type 1 column: Ticket

Key point: Split each ticket at its spaces; the last piece is the number and the first piece is the category.

`Ticket` is messier: 681 different values in 891 rows. Most tickets are a prefix and a number, such as A/5 21171 or PC 17599, but many are digits only, such as 113803.

The rule, shown in Figure 8:

1. Split the value at its spaces into pieces.
2. **Number part:** the last piece, converted with `pd.to_numeric(..., errors="coerce")`.
3. **Category part:** the first piece, unless it is all digits; then there is no category (NaN).




Figure 8: Each ticket split at its spaces: the first piece becomes the category, the last piece the number

Python: Splitting Ticket.

```
parts = df["Ticket"].str.split()
df["ticket_num"] = pd.to_numeric(
    parts.str[-1], errors="coerce").astype("Int64")
first = parts.str[0]
df["ticket_cat"] = first.where(~first.str.isdigit())
# .str.split() turns each value into a list of pieces; .str[-1] and .str[0] then take the
```

last and first piece. `str.isdigit()` is True when a value is all digits, and `~` flips True and False.

The result:

	Ticket	ticket_num	ticket_cat
0	A/5 21171	21171	A/5
1	PC 17599	17599	PC
2	STON/O2. 3101282	3101282	STON/O2.
3	113803	113803	NaN
115	STON/O 2. 3101294	3101294	STON/O
179	LINE	<NA>	LINE

The 681 tickets have become 43 ticket categories plus a number. Four tickets read just **LINE**, so they have a category but no number.

Extra: The 43 categories still repeat themselves. **A/5**, **A/5.**, **A./5.** and **A.5.** are the same prefix written four ways, and so are **SC/Paris** and **SC/PARIS**. Removing dots and slashes and upper-casing, `.str.replace(r"[./]", "", regex=True).str.upper()`, cuts the 43 down to 29. **STON/O 2.** and **STON/O2.** still differ, because the space split the first one into two pieces.

7. Filling the gaps the split leaves

Key point: Splitting creates NaNs on purpose; fill the numerical ones with a value that makes sense and the categorical ones with a label such as “missing”.

Splitting always leaves gaps. In Type 2, every row is missing one of its two new columns. In Type 1, a value with no number or no letter leaves a gap too.

Many scikit-learn models do not accept missing values (scikit-learn User Guide, “Imputation of missing values”), so we fill them:

- **Numerical feature:** a value that makes sense for the data. In `number_numerical`, NaN came from **A**, alone, so it means 0 companions: fill with 0.
- **Categorical feature:** a new category of its own, such as “missing”.




Figure 9: The gaps of the Type 2 split (rows 0, 1 and 4), before and after filling

Figure 9 shows both fills on three rows of the `number` split.

Python: Filling the gaps.

```
df["number_numerical"] = df["number_numerical"].fillna(0)
for col in ["number_categorical", "cabin_cat",
            "ticket_cat"]:
    df[col] = df[col].fillna("missing")
```

The first five rows, with all the new columns filled:

	number_numerical	number_categorical	cabin_cat
0	5	missing	missing
1	3	missing	C
2	6	missing	missing
3	3	missing	C
4	0	A	missing

	cabin_num	ticket_num	ticket_cat
0	<NA>	21171	A/5
1	85	17599	PC
2	<NA>	3101282	STON/O2.
3	123	113803	missing
4	<NA>	373450	missing

`cabin_num` is left unfilled here: no single number stands for “no cabin”. Filling it is a missing-values question: see [imputing, filling in the gaps](#).

The new features are then ready for the usual tools: [one-hot encoding](#) for the categorical ones, [standardization](#) for the numerical ones.

Extra: The split features carry information the raw ones hid. Survival rate by deck:

Deck	B	C	D	E	no cabin
Passengers	47	59	33	32	687
Survived	74%	59%	76%	75%	30%

Passengers with a known cabin survived far more often. Even “missing” is a useful category here.

8. Summary

	Type 1	Type 2
What the cells look like	category and number together: C85, A/5 21171	a number in some rows, a category in others: 5, A
Examples in the data	Cabin, Ticket	number
Number column	the digits of every value	the value where it is a number, else NaN
Category column	the letters of every value	the value where it is not a number, else NaN
pandas tools	<code>.str.extract</code> , <code>.str[0]</code> , <code>.str.split()</code>	<code>pd.to_numeric(errors="coerce")</code> , <code>.where</code>

- A mixed variable holds numerical and categorical data in one feature; we split it into one numerical and one categorical feature.
- `pd.to_numeric(errors="coerce")` keeps what is a number and turns the rest into NaN.
- The `.str` accessor applies text methods (extract, split, first character) to a whole column at once.
- Use the nullable `Int64` type for whole numbers with gaps.
- Splitting cut `Cabin` from 147 values to 8 decks, and `Ticket` from 681 values to 43 prefixes.
- Splitting leaves NaNs: fill numbers with a sensible value and categories with "missing".
- Check the unusual values: no simple rule fits every one.

9. Sources

Built from

- CampusX, “Handling Mixed Variables | Feature Engineering”, YouTube, https://www.youtube.com/watch?v=9xiX-I5_LQY

Other references

- pandas release notes. What’s new in 3.0.0 (dedicated string data type by default). pandas.pydata.org/docs/whatsnew.
- pandas user guide. Nullable integer data type. pandas.pydata.org/docs/user_guide/integer_na.html.
- scikit-learn User Guide. Imputation of missing values. scikit-learn.org.

10. Key terms

Term	Meaning
Feature	An input variable, one column of the data table

Term	Meaning
Target	The output we predict
Observation	One record, one row of the data table
Mixed variable	A feature holding both numerical and categorical data
Type 1 mixed variable	A column whose cells each contain a category and a number together, such as C85
Type 2 mixed variable	A column with a number in some rows and a category in others
pd.to_numeric	The pandas function that converts values to numbers
errors="coerce"	The <code>pd.to_numeric</code> option that turns values it cannot convert into NaN instead of stopping
.str accessor	The pandas tool that applies a text method to every value of a column
Regular expression	A short pattern that describes text, such as <code>\d+</code> for "one or more digits"
str.extract	The pandas method that returns the part of each value matching a regular expression
Raw string	A Python string written <code>r"..."</code> , in which a backslash is kept as it is
Nullable integer (Int64)	The pandas integer type that can also hold a missing value, <code><NA></code>
str dtype	The pandas 3 type for text columns, replacing <code>object</code>