

Handling Missing Data: Random Sample Imputation, Missing Indicator and Automatic Imputer Selection

Contents

1. Overview	2
2. Random sample imputation	3
2.1 The idea	3
2.2 Why the distribution stays the same	3
2.3 Advantages	5
2.4 Disadvantages	5
3. Random sample imputation on Titanic ages	6
3.1 The data and the fill	6
3.2 Distribution, variance and outliers	6
3.3 Covariance with other columns	7
4. Same input, same fill	8
4.1 The problem	8
4.2 The fix: one seed per row	8
5. Random sample imputation on categorical columns	9
5.1 The data	9
5.2 Check 1: category shares	10
5.3 Check 2: the target per category	10
6. Missing indicator	12
6.1 The idea	12
6.2 Why it can help	12
6.3 On the Titanic data	13
6.4 The shortcut: <code>add_indicator=True</code>	15
7. Choosing the imputer automatically with grid search	16
7.1 The idea	16
7.2 The pipeline	16
7.3 Naming the settings	18
7.4 The result where the fill matters: house prices	18
7.5 When every imputer ties: the Titanic pipeline	19
8. Summary	22
9. Sources	22
10. Key terms	23

Where this fits: Pipeline map steps 4 (Clean), 5 (Engineer features), 10 (Tune). See the [Course map](#).



- **Builds on:** Poor-quality data ([Note ML-007](#)); Saving models with pickle ([Note ML-009](#)); Hyperparameter tuning ([Note ML-009](#)); Data leakage ([Note ML-012](#)); Simple imputation (mean, median, mode, constant) ([Note ML-022](#)); Column transformer ([Note ML-027](#)).
- **Leads to:** KNN imputer ([Note ML-038](#)); Iterative imputation (MICE) ([Note ML-039](#)); Bagging ([Note ML-099](#)); Random forest ([Note ML-102](#)); XGBoost ([Note ML-117](#)).
- **Compare with:** Outliers ([Note ML-022](#)); Simple imputation (mean, median, mode, constant) ([Note ML-036](#)); Optuna ([Note ML-128](#)); Bayesian optimisation ([Note ML-128](#)); Keras Tuner ([Note DL-039](#)).

1. Overview

Key point: Random sample imputation fills each gap with a real value drawn at random from the same column; a missing indicator adds a True/False column that marks where the gaps were.

Note ML-035 filled gaps in numerical features with one fixed number, and Note ML-036 filled gaps in categorical features with the mode or the word “Missing”. A **feature** (G-772) is an input variable (one column of the data table), the **target** (G-1949) is the output we predict, and an **observation** (G-1374) is one record (one row). This Note adds:

- random sample imputation (Sections 2 to 5);
- the missing indicator (Section 6);
- a way to let the computer pick the best technique for us, grid search (Section 7).

Figure 1 shows the two techniques on a small **Age** column with two gaps:

- **Random sample imputation** (G-1616) puts the known values in a bag and draws one at random for each gap. The technique works for numerical and categorical features alike.
- **Missing indicator** (G-1233) keeps the gaps for another imputer to fill, and adds a new column, **Age_NA**, that is True where the value was missing.




Figure 1: Random sample imputation draws each fill value from the known values; the missing indicator adds a column marking the gaps

Both techniques look only at the feature with the gap, never at the other features. Such a method is called **univariate imputation** (G-2051). The last section uses grid search to choose the imputation strategy automatically, together with the model’s own settings.

2. Random sample imputation

Key point: Each missing value is replaced by a value picked at random from the values the column already has.

2.1 The idea

Key point: We do not invent numbers; we reuse the column’s own values, chosen at random.

Random sample imputation (G-1616) replaces every missing value of a column with a value drawn at random from the known values of the same column. If the known ages are 22, 38, 26 and 35, a gap becomes one of those four, never some other number such as 71 or 99.

The same works for a categorical column. If a **gender** column holds “male” and “female” plus some gaps, each gap becomes “male” or “female”, drawn at random from the known rows.

2.2 Why the distribution stays the same

Key point: Common values are drawn often and rare values rarely, so the filled values follow the same shape as the known ones.

Think of every known value written on a slip of paper and put into one bag. Values that occur often have many slips, so a random draw lands on them often. Rare values have few slips and are drawn rarely.

The chance of drawing a value, step by step:

1. **In words:** count how many known values fall in a range, and divide by the number of known values.
2. **Formula:**

$$P(\text{draw lands in the range}) = \frac{\text{known values in the range}}{\text{all known values}}$$

3. **Example:** in the Titanic training set of Section 3, 174 of the 564 known ages are in their twenties:

$$\frac{174}{564} = 0.309.$$

So about 31% of the 148 gaps, around 46 of them, get an age in the twenties, the same share as in the known data. Only 20 known ages are 60 or over, so very few gaps get such an age.

Figure 2 runs the draws for the 148 gaps of Section 3. Watch the orange bars: after a few draws they jump around, but as the draws add up they settle into the grey outline of the known ages; 49 of the 148 land in the twenties.




Figure 2: Random sample imputation of the 148 missing training ages: the share of drawn ages per decade (orange) settles onto the share among the 564 known ages (grey outline)

Because every range gets its fair share of the filled values, two things stay almost the same: the **distribution** (G-626), which is how the values spread over their range, and the **variance** (G-2074), which is the average squared distance of the values from their mean. **Mean imputation** (G-1197, Note ML-035) instead piles every filled value at one point.

2.3 Advantages

Key point: Random sample imputation is easy to apply and keeps the distribution and variance intact, which suits linear models.

1. **Easy to apply.** One line of pandas per column.
2. **Keeps the distribution and the variance**, as Section 2.2 explains.
3. **Good for linear models**, such as linear and logistic regression, which are sensitive to the shape of the data.

Tree-based algorithms (G-2013), such as decision trees, gain less from it: the random values add noise to the data. The intuition: a tree splits on the order of the values, so keeping the exact shape of the distribution matters less to it.

Python: scikit-learn has no random sample imputer, so we write it in pandas.

```
pool = X_train["Age"].dropna()           # the bag
gaps = X_train["Age"].isnull()
X_train.loc[gaps, "Age"] = pool.sample(
    gaps.sum(), random_state=42).values
```

sample(n) (G-140) draws *n* values at random from a Series. **random_state** (G-128) fixes the draw, so the same code gives the same values every time. **.values** drops the row labels of the drawn values; without it, pandas would try to match them to the gaps' labels.

2.4 Disadvantages

Key point: Random sample imputation weakens the feature's relationships with other features, adds randomness, and needs the training data at prediction time.

1. **The covariance with other features changes.** The **covariance** (G-496) measures how two features move together: positive when they rise together. Each filled value is drawn without looking at the rest of the observation, so the feature's link to other features weakens.
2. **The fill is random.** Run the code twice without a fixed **random seed** (G-1617), the number that fixes a random draw, and the filled values differ.
3. **Production needs memory.** A gap in new data must be filled from the training values, so the deployed model has to keep the whole training column on the server. For a large training set, the stored column takes a lot of memory.

Random sample imputation keeps the shape however many values are missing. Section 5 shows a check that can still fail when many values are missing.

3. Random sample imputation on Titanic ages

Key point: On the Titanic data, random sample imputation of **Age** keeps the density curve and the variance almost unchanged, but the covariance with **Fare** drops.

3.1 The data and the fill

Key point: We split first, and draw the fill values for both sets from the known training ages.

The data is Kaggle's Titanic training file, 891 passengers. We use **Age** and **Fare** as inputs and **Survived** as the target. **Age** has 177 missing values (19.9%); **Fare** has none.

After an 80/20 split, the training set has 712 rows, 148 of them with no age (20.8%). The test set has 179 rows, 29 of them with no age.

Both sets are filled from the same bag: the 564 known training ages. Drawing test values from the test set itself would let test data shape the inputs, which is **data leakage** (G-535, Note ML-035).

Python: One helper for both sets.

```
def random_fill(col, pool, seed):
    out = col.copy()
    gaps = out.isnull()
    out[gaps] = pool.sample(gaps.sum(),
                           random_state=seed).values

    return out
```

```
pool = X_train["Age"].dropna()
X_train["Age_imputed"] = random_fill(X_train["Age"], pool, 42)
X_test["Age_imputed"] = random_fill(X_test["Age"], pool, 42)
```

We keep the original column next to the filled one, to compare them. The seed 42 makes the numbers below repeatable; another seed gives slightly different numbers.

3.2 Distribution, variance and outliers

Key point: The density curve and the box barely move, and the variance changes from 204.35 to 200.03.

Figure 3 compares the original **Age** (blue) with the random sample imputed one (orange). For contrast, it also shows mean imputation (green dashed, Note ML-035).




Figure 3: Age before and after imputation: the random sample curve lies on the original, while mean imputation creates a tall peak and many new outliers

The orange curve lies almost on top of the blue one. The box plots agree: the box runs from 21 to 38.25 before and from 21 to 38 after, and the outliers go from 7 to 10. Mean imputation, by contrast, narrows the box to 23 to 35.

Age in the training set	Original	Random sample	Mean (Note ML-035)
Variance	204.35	200.03	161.81
Mean	29.79	29.58	29.79
Median	28.75	28.00	29.79
Outliers (box plot)	7	10	69

3.3 Covariance with other columns

Key point: The covariance of **Age** with **Fare** falls from 71.51 to 53.26, because the drawn ages ignore the fare.

Covariance with Fare	Value
Age original	71.51
Age random sample imputed	53.26

A passenger with a high fare tended to be a little older. The drawn ages pay no attention to the fare, so for the 148 filled rows this pattern is lost, and the covariance drops by about a quarter. The drop is disadvantage 1 of Section 2.4.

4. Same input, same fill

Key point: In production, seed the draw with a value from the row itself, so the same input always gets the same fill value and the same prediction.

4.1 The problem

Key point: Without a fixed seed, the same user can get a different prediction each time they ask.

Suppose a website uses the Titanic model. A user enters a fare of 56.50 but leaves the age empty. The site draws an age at random, say 30, and shows a prediction.

The user presses the button again with the same input. This time the draw gives 4, and the prediction may change. The same question now has two answers, and the user loses trust in the model.

4.2 The fix: one seed per row

Key point: Use a value from the row, such as its fare, as the random seed.

A random draw with a fixed seed always gives the same result. So we take the seed from the row: here, the fare as a whole number. A seed taken from the row's own values is a **per-row seed** (G-1479). The same fare always gives the same seed, so the same draw, so the same age.

Python: Seeding each row's draw with its fare.

```
def fill_one(row, pool):
    if pd.isna(row["Age"]):
        seed = int(row["Fare"])
        return pool.sample(1, random_state=seed).iloc[0]
    return row["Age"]
```

```
new = pd.Series({"Age": np.nan, "Fare": 56.4958})
[fill_one(new, pool) for _ in range(3)] # [34.0, 34.0, 34.0]
```

Without the seed, `pool.sample(1)` three times gave 30, 4 and 27. Figure 4 asks eight times. Watch the two lines: the draw seeded by the fare stays at 34, while fresh draws (here 8 different seeds) jump between 17 and 52. `int()` turns a number into a whole number by dropping the decimals; `random_state` needs a whole number.




Figure 4: The same input filled eight times: a fresh draw each time (red) against a draw seeded by the row's fare (green)

Extra: The seed should come from inputs that identify the row well. `int(row["Fare"])` drops the decimals, so fares 7.25 and 7.93 share the seed 7, and every passenger with that fare and no age gets the same age. Using several inputs together, for example a hash of the whole row, spreads the seeds better.

5. Random sample imputation on categorical columns

Key point: On two house-quality columns, the category shares stay the same after random sample imputation, but in `FireplaceQu` the sale price per category changes a lot.

5.1 The data

Key point: The same two columns as Note ML-036: `GarageQual` with 5.5% missing and `FireplaceQu` with 47.3% missing.

The data is the house price file of Note ML-036: `GarageQual`, `FireplaceQu` and the target `SalePrice` for 1,460 houses. The full Kaggle file has 81 columns; `data/house_prices.csv` keeps these three.

We split 80/20 with the same seed as before. The training set has 1,168 houses: 65 with no `GarageQual` (5.6%) and 557 with no `FireplaceQu` (47.7%). Each column is filled from its own known training values.

5.2 Check 1: category shares

Key point: Each category keeps its share, even in `FireplaceQu`, where almost half the values were drawn.

Figure 5 compares each category's share before (among the known values) and after (among all rows).




Figure 5: Category shares before and after random sample imputation: no share moves by more than 0.6 percentage points

Category	GarageQual		FireplaceQu	
	before	after	before	after
Gd	1.0%	0.9%	49.4%	50.0%
TA	95.1%	95.2%	41.2%	40.7%
Fa	3.7%	3.7%	4.1%	4.0%
Po	0.1%	0.1%	2.8%	2.8%
Ex	0.1%	0.1%	2.5%	2.5%

The equal shares are the categorical version of Section 2.2: a category with half the slips gets about half the draws. Compare mode imputation in Note ML-036, where Gd in `FireplaceQu` jumped to 72.9%.

5.3 Check 2: the target per category

Key point: In `FireplaceQu`, cheap houses were spread into every category, so Gd and Ex houses now look much cheaper; this is a red flag.

The shares only tell part of the story. A second check compares the target, `SalePrice`, for each category before and after (Figure 6).




Figure 6: Sale price density per FireplaceQu category, before and after random sample imputation: the curves of Ex, Gd and TA shift left

The 557 houses with no FireplaceQu sold for a mean of 141,182 dollars, far below the Gd houses (225,668). The draws turned about half of them into Gd houses and two fifths into TA houses. Each category’s prices moved towards the cheap ones:

FireplaceQu	Mean price before	Mean price after
Ex	337,145	240,230
Gd	225,668	184,529
TA	205,749	175,367
Fa	168,994	161,279
Po	135,112	139,591

Before imputation, a better fireplace meant a pricier house. After, Gd, TA and Fa houses look much more alike, so the link between fireplace quality and price is weaker in the imputed data. Such a change in the distribution is a bigger problem for linear models.

For GarageQual, only 65 values were drawn, and the TA mean moves only from 186,766 to 182,093 dollars. So random sample imputation is fine for GarageQual, but not for FireplaceQu: too many values are missing there.

Extra: As Note ML-036 showed, a gap in FireplaceQu means “no fireplace”: all 690 such houses have zero fireplaces. Drawing Gd or TA for them invents fireplaces, which explains the price shift. For gaps with a meaning, a “Missing” category or a missing indicator fits better.

6. Missing indicator

Key point: A missing indicator adds a True/False column for each feature with gaps, so the model can learn whether “missing” itself matters.

6.1 The idea

Key point: True where the value was missing, False where it was present.

A missing indicator (G-1233) is a new column added for each feature that has gaps. The indicator holds True where that feature’s value is missing and False where it is present. For Age, the new column `Age_NA` is True for the 148 passengers with no age (Figure 1).

The indicator does not fill the gaps. The original column is still imputed in the usual way, for example with the mean. The indicator is added next to it.

6.2 Why it can help

Key point: The model can learn that observations with gaps behave differently from observations without.

The model gets a way to tell observations with a gap apart from observations without one. A doctor’s form works the same way: a blank “smoker?” box can itself be a clue, so we note that it was left blank instead of guessing an answer. If the two groups behave differently, the model can use that difference.

On the Titanic training set, they do. Of the passengers with no recorded age, 28.4% survived; of those with an age, 39.2% survived (Figure 7). Mean imputation alone hides this: a filled age of 29.79 looks like any other age.




Figure 7: Survival in the Titanic training set for passengers with and without a recorded age: the gap the indicator lets the model see

The technique became popular through machine learning competitions, where adding indicators helped winning models. The indicator does not always help, but it is cheap to try when a model needs a few more points of accuracy.

Extra: A missing indicator is most useful when the data is not missing completely at random: **MAR** (G-1158), where the gaps depend on another recorded feature, or **MNAR**, where they depend on the missing value itself (Note ML-034). When gaps are pure chance, the indicator carries no information, and the model should give it a weight near zero.

6.3 On the Titanic data

Key point: Adding `Age_NA` to mean-imputed `Age` and `Fare` lifts the test accuracy of logistic regression from 61.5% to 63.1%.

We train **logistic regression** (G-1120), a classification algorithm (Note ML-012), on `Age` and `Fare` twice, with the same split as Section 3. The score is the **accuracy** (G-162): the fraction of test passengers classified correctly.

Inputs	Test accuracy
<code>Age</code> , <code>Fare</code> (mean imputed)	61.5% (110 of 179)
<code>Age</code> , <code>Fare</code> (mean imputed) + <code>Age_NA</code>	63.1% (113 of 179)

Three more test passengers are classified correctly.

The mechanism, step by step:

1. **In words:** the model multiplies each input by a learned number, its **coefficient** (G-407), adds the results to a score z , and turns z into a chance between 0 and 1. The indicator is one more input, 1 for True and 0 for False, so it gets its own coefficient.
2. **Formula,** with the coefficients the model learned:

$$z = -0.419 - 0.019a + 0.016f - 0.295m, \quad P(\text{survived}) = \frac{1}{1 + e^{-z}}$$

Here a is the age, f the fare, and m the indicator `Age_NA` (1 or 0).

3. **Example:** two passengers with the filled age 29.79 and a fare of 60. With a recorded age (`Age_NA` = 0), $z = -0.01$ and the chance is 49.7%. With a missing age (`Age_NA` = 1), z drops by 0.295 to -0.31 and the chance is 42.4%.

The coefficient of `Age_NA` is negative (-0.30 after rounding): a missing age lowers the predicted chance of survival, as the survival rates in Section 6.2 suggest.

Figure 8 shows that coefficient at work. Two passengers have the same filled age, 29.79, and the same fare; only `Age_NA` differs. Watch the orange curve sit below the blue one at every fare, so the passenger with a missing age needs a higher fare (about 80 instead of 60) before the model predicts “survived”.




Figure 8: Predicted chance of survival against fare, for the same filled age with `Age_NA = 0` and `Age_NA = 1`

Figure 9 shows the model learning that coefficient. Training starts with every coefficient at 0, so both curves are one flat line at 50%. Watch the first iterations: the model learns the fare first, and the two curves rise together as one. From about the sixth iteration the coefficient of `Age_NA` leaves 0, and the dashed orange curve drops below the blue one until the gap settles at -0.30 .




Figure 9: The model of Section 6.3 during training: the curves for $\text{Age_NA} = 0$ and $\text{Age_NA} = 1$ start as one and separate as the coefficient of Age_NA moves from 0 to -0.30

Extra: One test set of 179 passengers is a small sample, so the Notebook also averages over 100 different splits of all 891 passengers (5-fold cross-validation, repeated 20 times). The mean accuracy is 65.7% without the indicator and 66.3% with it. The gain is small because the model has only two features, but it holds on average.

Python: scikit-learn's **MissingIndicator** (G-119).

from sklearn.impute **import** MissingIndicator

```
mi = MissingIndicator()
mi.fit(X_train)                                # which columns have gaps?
mi.features_                                  # array([0]): only Age
X_train_missing = mi.transform(X_train)
X_test_missing = mi.transform(X_test)
features_ lists the positions of the columns that had gaps during fit. transform
returns one True/False column for each of them. We then add it to the table, for example
X_train.assign(Age_NA=X_train_missing.ravel()), and impute as before.
```

6.4 The shortcut: `add_indicator=True`

Key point: `SimpleImputer(add_indicator=True)` imputes and adds the indicator columns in one step.

The missing indicator is common enough that `SimpleImputer` has a parameter for it (Note ML-035 listed it). With `add_indicator=True` (G-58), the imputer fills the gaps and appends one indicator column for each column that had gaps.

Python: Imputation and indicator in one object.

```
si = SimpleImputer(add_indicator=True)
X_train_trf = si.fit_transform(X_train)
X_test_trf = si.transform(X_test)
si.get_feature_names_out()
# ['Age', 'Fare', 'missingindicator_Age']
```

The result is the same three columns, and logistic regression again reaches 63.1%. `add_indicator=True` is the easier route: no separate object, and no joining of columns by hand.

Use the `MissingIndicator` class with an imputer that has no such parameter, such as the random sample imputation of Section 2. Otherwise, `add_indicator=True` is simpler.

7. Choosing the imputer automatically with grid search

Key point: Put the imputers inside a pipeline, and let `GridSearchCV` try every strategy and keep the one with the best cross-validated score.

7.1 The idea

Key point: The imputation strategy is a setting like any other, so grid search can tune it.

Should **Age** get the mean or the median? Should a categorical gap get the mode or a constant? Instead of deciding by hand, we can try every combination and measure which gives the best model. A setting chosen before training, such as the imputation strategy, is a **hyperparameter** (G-910).


Grid search (G-872) does this: it trains the model once for every combination of the settings we list, scores each with **cross-validation** (G-510), which trains and tests several times on different parts of the training data, and keeps the best. scikit-learn's **GridSearchCV** (G-89) runs it (see the [pipelines Note](#)). Grid search can tune the imputer only if the imputer is part of the model, so we build one **pipeline** (G-1499), a single object that bundles the processing steps and the model, from the raw data to the prediction.

7.2 The pipeline

Key point: Numerical columns are imputed and scaled, categorical columns are imputed and one-hot encoded, then logistic regression is trained.

We use the Titanic file again, dropping `PassengerId`, `Name`, `Ticket` and `Cabin`. In the training set, `Age` has 148 gaps and `Embarked`, the port where the passenger boarded, has 2.

Figure 10 shows the pipeline. Each step has a name, shown in typewriter font.



`preprocessor__num__imputer__strategy` = strategy of the imputer in num

Figure 10: The nested pipeline; a setting's full name joins the step names with double underscores

Python: The pipeline of Figure 10.

```

num = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())])
cat = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("ohe", OneHotEncoder(handle_unknown="ignore"))])
preprocessor = ColumnTransformer(transformers=[
    ("num", num, ["Age", "Fare"]),
    ("cat", cat, ["Embarked", "Sex"])]
clf = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", LogisticRegression())])

```

`handle_unknown="ignore"` turns a category never seen in training into all zeros instead of an error (Note ML-026).

Extra: A `ColumnTransformer` drops every column it is not told about (`remainder="drop"` by default). So this model uses only Age, Fare, Embarked and Sex; `Pclass`, `SibSp` and `Parch` are silently left out. Add them to a list, or set `remainder="passthrough"`, to use them.

7.3 Naming the settings

Key point: A setting deep inside a pipeline is named by its path: step names joined by two underscores, then the parameter name.

A setting inside a pipeline is named `step__parameter` (see “Hyperparameter tuning with a pipeline”, section 9 of the [pipelines Note](#)). In a nested pipeline the name is the whole path from the outer pipeline down, joined by `__`. Figure 10 shows `preprocessor__num__imputer__strategy`: the `strategy` of the `imputer` in `num`, inside `preprocessor`.

The grid lists the values to try for each name:

Setting	Values tried
<code>preprocessor__num__imputer__strategy</code>	"mean", "median"
<code>preprocessor__cat__imputer__strategy</code>	"most_frequent", "constant"
<code>classifier__C</code>	0.1, 1, 10, 100

`C` controls how strongly logistic regression is held back from fitting the training data too closely; a later Note covers it. The grid has $2 \times 2 \times 4 = 16$ combinations.

Python: The grid search.

```
param_grid = {
    "preprocessor__num__imputer__strategy": ["mean", "median"],
    "preprocessor__cat__imputer__strategy":
        ["most_frequent", "constant"],
    "classifier__C": [0.1, 1.0, 10, 100],
}
grid_search = GridSearchCV(clf, param_grid, cv=10)
grid_search.fit(X_train, y_train)
grid_search.best_params_
clf.get_params() lists every valid name, if we are unsure of a path. With cv=10, each
of the 16 combinations is trained 10 times, 160 fits in all, each on nine tenths of the
training set and scored on the other tenth.
```

7.4 The result where the fill matters: house prices

Key point: When the feature with gaps drives the target, grid search finds a clear winner: on house prices, the “Missing” category cuts the average error from 53,400 to 46,400 dollars.

Grid search can only tell two imputers apart if the feature they fill matters for the target. Think of two ways to fill a blank on a form: the choice matters only if someone reads that line. The house data of Section 5 meets that condition. `FireplaceQu` has 47.7% gaps in the training set, and those gaps mark cheaper houses (Section 5.3). We predict `SalePrice`

from `GarageQual` and `FireplaceQu` with a three-step pipeline: impute, one-hot encode (Note ML-026), then linear regression (Note ML-049). The grid tries the two categorical strategies.

The score is the **mean absolute error** (G-1194): the average size of the gap between the predicted and the real price, in dollars. Lower is better.

imputer__strategy	Mean CV error
"most_frequent"	53,447 dollars
"constant" (a "missing" category)	46,382 dollars

`best_params_` (G-66) picks "constant", and the refitted pipeline's error on the test set is 47,272 dollars. The reason is the price shift of Section 5.3, seen from the other side. Mode imputation turns the 557 cheap houses with no fireplace into Gd houses, so the model can no longer tell them apart from real Gd houses. A category of their own keeps them as a separate group with its own, lower price. The scikit-learn example on imputing before an estimator makes the same point: the choice of imputer changes the model's score, so it is worth comparing (scikit-learn docs, Imputing missing values before building an estimator).

Python: The grid search on the house data.

```
house_pipe = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("ohe", OneHotEncoder(handle_unknown="ignore")),
    ("model", LinearRegression())])
house_search = GridSearchCV(house_pipe,
    {"imputer__strategy": ["most_frequent", "constant"]},
    cv=10, scoring="neg_mean_absolute_error")
house_search.fit(Xh_train, yh_train)
house_search.best_params_      # {'imputer__strategy': 'constant'}
scoring="neg_mean_absolute_error" scores each setting by minus its mean absolute
error, because grid search always keeps the highest score. With strings, "constant" fills
the gaps with the word "missing_value" by default.
```

Extra: Averaged over 100 splits of all 1,460 houses (10-fold cross-validation, repeated 10 times), the error is 54,015 dollars with the mode and 46,863 dollars with the constant category.

7.5 When every imputer ties: the Titanic pipeline

Key point: On the Titanic data, all four imputer combinations tie at 78.8%, because the feature with gaps, `Age`, barely affects the predictions.

The Titanic pipeline of Section 7.2 shows the other side of the condition.

`best_params_` reports `C = 1`, "most_frequent" for the categorical features and "mean" for the numerical ones, with a cross-validated accuracy of 78.8%. The full table, `cv_results_` (G-73), shows that only `C` changes the score:

C	Imputer strategies	Mean CV accuracy
1, 10 or 100	any of the 4 combinations	78.8%
0.1	any of the 4 combinations	78.6%

Figure 11 puts both searches side by side. Watch the left panel, where the fill changes the error by 7,000 dollars, against the right panel, where every row is identical and only the column **C** changes the score.




Figure 11: Grid search results: on house prices the “Missing” category clearly wins; on the Titanic pipeline all four imputer combinations tie and only **C** matters

The four imputer combinations score exactly the same, so grid search reports the first of the tied combinations. The categorical fills touch only the 2 observations with no **Embarked**, and **Age** hardly matters (next Extra).

Extra: The experiment behind the tie (with **C** = 1, the same 10 folds; the last cell of the Notebook).

Change to the pipeline	Mean CV accuracy
Age filled with the mean or the median	78.8%
Age filled with 0	78.8%
Age filled with 99	78.9%
Age removed from the inputs	78.8%

Removing **Age** entirely gives the same score, fold for fold, so no fill of **Age** can change it; only the extreme fill of 99 flips a few predictions.

Extra: Grid search fits only on the training set; the test set stays unseen. After the search, `grid_search` refits the best pipeline on the whole training set, and we score it on the test set once: `grid_search.score(X_test, y_test)` gives 76.0%. The test score is the number to report, not the 78.8% from cross-validation, because the settings were chosen to maximise that one, which makes it too optimistic (Cawley and Talbot 2010).

8. Summary

	Random sample imputation	Missing indicator	Grid search over imputers
What it does	fills each gap with a random known value	adds a True/False column per column with gaps	tries every imputer setting, keeps the best
Column types	numerical and categorical	any	any
Keeps the distribution	yes	does not fill; used with another imputer	depends on the chosen imputer
Main risk	weaker links to other columns; randomness; training data needed in production	an extra column that may carry no information	slow when the grid is large
In scikit-learn	no; done in pandas with <code>sample</code>	<code>MissingIndicator</code> , or <code>SimpleImputer(add_indicator=True)</code>	<code>GridSearchCV</code> over a <code>Pipeline</code>
On this data	Age : variance 204.35 to 200.03	accuracy 61.5% to 63.1%	house prices: “Missing” beats mode, error 46,400 vs 53,400 dollars

- Random sample imputation draws each fill value from the feature’s own known values, so the shape and variance stay almost the same. The technique suits linear models.
- Random sample imputation weakens the feature’s covariance with others (**Age** with **Fare**: 71.51 to 53.26) and needs the training values at prediction time.
- In production, seed each draw with a value from the row, so the same input always gets the same fill.
- On categorical data, check the category shares and the target per category. In **FireplaceQu** the shares held, but the prices per category shifted badly: too many values were missing.
- A missing indicator lets the model learn whether a missing value is informative. The indicator is added next to any imputation.
- Grid search can tune the imputation strategy along with the model, if the imputers sit inside the pipeline. Settings are named by their path, joined with `__`. On the house prices, grid search picked the “Missing” category, which cut the average error from 53,400 to 46,400 dollars.

9. Sources

Built from

- CampusX, “Missing Indicator | Random Sample Imputation | Handling Missing Data Part 4”, YouTube, <https://www.youtube.com/watch?v=Ratcir3p03w>

Other references

- scikit-learn examples, *Imputing missing values before building an estimator* (Impute section of the example gallery).
- Cawley, G. C. and Talbot, N. L. C. (2010). On Over-fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation. *Journal of Machine Learning Research* 11, 2079–2107.

10. Key terms

Term	Meaning
Random sample imputation	Filling each gap with a value drawn at random from the column's known values
<code>sample(n)</code>	The pandas method that draws <code>n</code> values at random from a Series or DataFrame
<code>random_state</code>	A seed that fixes a random draw, so the same code gives the same result
Per-row seed	A seed taken from a row's own values, so the same input always gets the same random fill
Feature	An input variable: one column of the data table
Target	The output we predict
Observation	One record: one row of the data table
Missing indicator	A True/False column marking where a feature's value was missing
Mean absolute error (G-1194)	The average size of the gap between predicted and real values, in the target's units
<code>MissingIndicator</code>	The scikit-learn class that builds missing indicator columns; <code>features_</code> lists the columns with gaps
<code>add_indicator=True</code>	The <code>SimpleImputer</code> setting that imputes and appends missing indicators in one step
Univariate imputation	Imputation that uses only the feature with the gap
Coefficient (G-407)	The learned number that multiplies one input of a linear or logistic model
Grid search	Training a model for every combination of listed settings and keeping the best by cross-validation
<code>best_params_</code>	The best combination of settings found by <code>GridSearchCV</code>
<code>cv_results_</code>	The scores of every combination tried by <code>GridSearchCV</code>