

Multiple Linear Regression: The Normal Equation

Contents

1. Overview	2
2. The model in matrix form	3
2.1 One equation per student	3
2.2 Stacking the equations	3
3. The error in matrix form	6
4. Expanding the error	7
4.1 Two rules about transposes	8
4.2 Multiplying out	10
5. Setting the derivative to zero	11
5.1 Three rules of matrix calculus, each checked on numbers	12
5.2 The gradient, then zero	14
6. The normal equation	15
6.1 A worked example	17
6.2 The picture: a projection	20
7. The cost of the inverse	21
8. Summary	24
9. Sources	24
10. Key terms	24

Where this fits: Pipeline map step 8 (Model). See the [Course map](#).



- **Builds on:** Regression problems ([Note ML-003](#)); Simple linear regression ([Note ML-049](#)); Ordinary least squares (closed form) ([Note ML-050](#)); Batch gradient descent ([Note ML-057](#)); Linear combinations, span and basis ([Note MA-052](#)); Jacobian and matrix gradients ([Note MA-063](#)).
- **Leads to:** Polynomial regression ([Note ML-060](#)); Ridge regression ([Note ML-062](#)).
- **Compare with:** Gradient descent ([Note ML-056](#)); Moore-Penrose pseudo-inverse ([Note MA-060](#)); ANN for regression ([Note DL-013](#)).

1. Overview

Key point: With several features we still find the best coefficients the same way as with one: measure the total squared error, and find where its slope is zero in every direction. Matrices let us do that for every coefficient at once, and the answer is one formula, the normal equation.

A **feature** (G-772) is an input variable (one column of the data table), the **target** (G-1949) is the output we predict, and an **observation** (G-1374) is one record (one row). Our running example predicts a student's package (the target) from CGPA (a feature).

For **simple linear regression** (G-1808), with one feature, we found two formulas, one for the slope and one for the intercept (see [the two formulas](#)). **Multiple linear regression** (G-1279) has several features (see [the equation](#)). Each feature gets its own weight, and there is one intercept on top. These numbers are the **coefficients** (G-407). Writing a separate formula for each coefficient is hopeless when there are ten or a hundred features.

Matrices solve the problem. Like a spreadsheet formula dragged down a whole column instead of typed into every cell, a matrix lets us write one equation for all observations at once. We write all the data, all the predictions and all the coefficients as matrices. Then the same three steps as before give one formula for every coefficient at once:

1. write the error (Sections 2 and 3);
2. multiply it out and differentiate it (Sections 4 and 5);
3. set the derivative to zero and solve (Section 6).

The result, at the end of Section 6, is the **normal equation**. Section 6.1 runs it by hand on four students, with every sum and the inverse worked out on paper. [Our own class](#) then codes it from scratch. Each rule about matrices is checked on small numbers before it is used.

A note on letters. In this Note the letter m counts the features (for example $m = 3$ for CGPA, IQ and gender). In [the error function of simple linear regression](#), m was the slope. Here every slope is a coefficient, written β (the Greek letter beta).

2. The model in matrix form

Key point: Put the data in one table with an extra first column of 1s, and the coefficients in one column. One multiplication of the two then gives every prediction at once.

2.1 One equation per student

Key point: Every student gets the same equation with the same coefficients; only the feature values change.

Suppose we predict a student's package from three features: CGPA, IQ and gender. The model multiplies each feature by its coefficient and adds the intercept:

$$\hat{y} = \beta_0 + \beta_1 \cdot \text{cgpa} + \beta_2 \cdot \text{iq} + \beta_3 \cdot \text{gender}$$

Once the four numbers β_0 to β_3 are known, we can predict the package of any new student. Training means finding them.

With 100 students there are 100 such equations. To tell the values apart we give each value two labels, the student and the feature:

x_{ij} = the value of feature j for student i

x_{21} = the CGPA (feature 1) of student 2

$$\hat{y}_1 = \beta_0 + \beta_1 x_{11} + \beta_2 x_{12} + \beta_3 x_{13}$$

$$\hat{y}_2 = \beta_0 + \beta_1 x_{21} + \beta_2 x_{22} + \beta_3 x_{23}$$

and so on, down to $\hat{y}_{100}$. The coefficients are the same in every line.

In general, with n observations and m features (for example $n = 100$ students and $m = 3$ features), observation i has the prediction

$$\hat{y}_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_m x_{im}$$

and there are n such equations.

2.2 Stacking the equations

Key point: The n equations split into a table of numbers, X , times a column of coefficients, β .

A **matrix** (G-1180) is a table of numbers, and a **vector** (G-2081) is a single column of numbers. They let us write the n equations as one.

Figure 1 does it for four students with one feature, CGPA (the students of Section 6.1). Watch three things move:

1. the four left sides stack into one vector, $\hat{y}$;
2. the numbers stack into one matrix, X : a column of 1s and a column of CGPAs;
3. the coefficients, which were repeated in every line, are written once, as the vector β .

1. one prediction equation per student

$$\begin{aligned}\hat{y}_1 &= \beta_0 \cdot 1 + \beta_1 \cdot 6.89 \\ \hat{y}_2 &= \beta_0 \cdot 1 + \beta_1 \cdot 5.12 \\ \hat{y}_3 &= \beta_0 \cdot 1 + \beta_1 \cdot 7.82 \\ \hat{y}_4 &= \beta_0 \cdot 1 + \beta_1 \cdot 7.42\end{aligned}$$

3. the numbers stack into X; the coefficients are written once

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix} = \begin{bmatrix} 1 & 6.89 \\ 1 & 5.12 \\ 1 & 7.82 \\ 1 & 7.42 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

4. the first column of X is all 1s

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix} = \begin{bmatrix} 1 & 6.89 \\ 1 & 5.12 \\ 1 & 7.82 \\ 1 & 7.42 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

the column of 1s multiplies the intercept

5. all four predictions in one line

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix} = \begin{bmatrix} 1 & 6.89 \\ 1 & 5.12 \\ 1 & 7.82 \\ 1 & 7.42 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

$$\hat{y} = X \beta$$

Figure 1: Four prediction equations become one matrix equation: the predictions stack into a vector, the numbers into the matrix X with a first column of 1s, and the coefficients are written once

Figure 2 shows the same split for n observations and m features.

$$\underbrace{\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_n \end{bmatrix}}_{\hat{y}: n \times 1} = \underbrace{\begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1m} \\ 1 & x_{21} & x_{22} & \cdots & x_{2m} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{nm} \end{bmatrix}}_{X: n \times (m+1)} \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_m \end{bmatrix}}_{\beta: (m+1) \times 1}$$

an extra column of 1s carries the intercept β_0

rows: data points (n) columns: a 1, then the m inputs β : the $m + 1$ numbers to find

Figure 2: All predictions as one matrix product

- X , the **design matrix** (G-597), holds the data: one row per observation, one column per feature, plus a first column of 1s. The 1s multiply β_0 , so the intercept is treated like any other coefficient.
- β , the **coefficient vector** (G-410), holds the $m + 1$ coefficients, β_0 to β_m .
- $\hat{y}$ holds the n predictions.

Multiplying one row of X by β means: multiply each entry of the row by the matching coefficient, then add. Row i gives

$$\beta_0 \cdot 1 + \beta_1 x_{i1} + \cdots + \beta_m x_{im}$$

which is exactly the prediction for observation i . So all predictions together are

$$\hat{y} = X\beta$$

With numbers: the first student of Figure 1 has the row $[1, 6.89]$. Section 6.1 finds the coefficients

$$\beta_0 = -0.811$$

$$\beta_1 = 0.565$$

The row times β , one product per line:

$$1 \times (-0.811) = -0.811$$

$$6.89 \times 0.565 = 3.893$$

$$-0.811 + 3.893 = 3.08$$

So 3.08 is the predicted package of that student.

Extra: The shapes must fit: X is $n \times (m + 1)$ and β is $(m + 1) \times 1$. The inner sizes match, and the product is $n \times 1$, one prediction per row.

3. The error in matrix form

Key point: Subtract each prediction from the real value to get one error per student. Square the errors and add them up: that one number measures how bad the coefficients are.

For each student, the error is the real package minus the predicted one. This error on one observation is called a **residual** (G-705). The residuals of all observations form a vector:

$$e = y - \hat{y} = y - X\beta$$

The **sum of squared errors** (G-1684) E adds up the squares of the residuals:

$$E = e_1^2 + e_2^2 + \cdots + e_n^2$$

That sum is a row times a column. The **transpose** (G-2012) of e , written e^T , is the same numbers written as a row (Figure 4 in Section 4 draws it). A row times a column multiplies matching entries and adds them, which is exactly the sum of squares:

$$E = e^T e = (y - X\beta)^T (y - X\beta)$$

With numbers, take three errors:

$$e = (0.2, -0.1, 0.3)$$

$$0.2 \times 0.2 = 0.04$$

$$(-0.1) \times (-0.1) = 0.01$$

$$0.3 \times 0.3 = 0.09$$

$$e^T e = 0.04 + 0.01 + 0.09 = 0.14$$

Figure 3 draws the same idea on the four students of Section 6.1 and their best line. Each red segment is one entry of e , and the orange square on it has area e_i^2 . Watch the biggest square: the student with CGPA 7.82 and error -0.36 supplies half of $E = 0.253$, because squaring makes large errors count much more than small ones.




Figure 3: The four students of Section 6.1 and their best line: the residuals e (red) and their squares (orange); $E = e^T e = 0.253$ is the total orange area

4. Expanding the error

Key point: Multiplying out the brackets turns the error into three parts, exactly like $(a - b)^2 = a^2 - 2ab + b^2$ in ordinary algebra. Only two rules about flipping rows and columns are needed.

In ordinary algebra a square multiplies out into three parts:

$$(a - b)^2 = a^2 - 2ab + b^2$$

The error E is the matrix version of a square, and it multiplies out to the same three-part shape. To see every step on numbers, we use a tiny made-up example with two observations and one feature:

$$X = \begin{bmatrix} 1 & 2 \\ 1 & 3 \end{bmatrix}$$

$$y = \begin{bmatrix} 4 \\ 5 \end{bmatrix}$$

$$\beta = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

The first column of X is the column of 1s; the feature values are 2 and 3; the targets are 4 and 5; and we try the coefficients 1 and 1. The predictions and errors are

$$X\beta = \begin{bmatrix} 1 \times 1 + 2 \times 1 \\ 1 \times 1 + 3 \times 1 \end{bmatrix}$$

$$X\beta = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$$

$$e = y - X\beta = \begin{bmatrix} 4 - 3 \\ 5 - 4 \end{bmatrix}$$

$$e = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$E = e^\top e = 1 \times 1 + 1 \times 1 = 2$$

4.1 Two rules about transposes

Key point: Transposing a product reverses its order, and a single number is its own transpose.

Figure 4 shows both rules on the tiny example.

$$X\beta = \begin{bmatrix} 3 \\ 4 \end{bmatrix} \xrightarrow[\text{column} \rightarrow \text{row}]{\text{transpose}} (X\beta)^\top = \begin{bmatrix} 3 & 4 \end{bmatrix}$$

$$\beta^\top X^\top = \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 2 & 3 \end{bmatrix} = \begin{bmatrix} 1 + 2 & 1 + 3 \end{bmatrix} = \begin{bmatrix} 3 & 4 \end{bmatrix}$$

same row: the transpose of a product is the product of the transposes, in reverse order

Figure 4: Top: the transpose turns the column $X\beta = (3, 4)$ into the row $[3, 4]$. Bottom: multiplying the transposes in reverse order, $\beta^\top X^\top$, gives the same row

Rule 1: the transpose of a product reverses the order.

$$(AB)^\top = B^\top A^\top$$

Here A and B stand for any two matrices that can be multiplied. For our product, A is X and B is β , so the rule says

$$(X\beta)^\top = \beta^\top X^\top$$

Check on the tiny example. The left side flips the column (3, 4) into a row:

$$(X\beta)^\top = \begin{bmatrix} 3 & 4 \end{bmatrix}$$

The right side, with the transpose of X (its rows become columns):

$$X^\top = \begin{bmatrix} 1 & 1 \\ 2 & 3 \end{bmatrix}$$

$$\beta^\top X^\top = \begin{bmatrix} 1 \times 1 + 1 \times 2 & 1 \times 1 + 1 \times 3 \end{bmatrix}$$

$$\beta^\top X^\top = \begin{bmatrix} 3 & 4 \end{bmatrix}$$

The two sides match.

Rule 2: a single number is its own transpose. A 1×1 matrix has nothing to flip. The two terms below are each a single number, and each is the transpose of the other (by Rule 1), so they are equal:

$$y^\top X\beta = \beta^\top X^\top y$$

Check on the tiny example. The left side is the row of targets times the column of predictions:

$$y^\top X\beta = 4 \times 3 + 5 \times 4 = 12 + 20 = 32$$

For the right side, first multiply the transpose of X by y:

$$X^\top y = \begin{bmatrix} 1 \times 4 + 1 \times 5 \\ 2 \times 4 + 3 \times 5 \end{bmatrix}$$

$$X^\top y = \begin{bmatrix} 9 \\ 23 \end{bmatrix}$$

$$\beta^\top X^\top y = 1 \times 9 + 1 \times 23 = 32$$

Both are 32.

4.2 Multiplying out

Key point: Four products come out, and the two middle ones are equal, so the error has three parts.

Multiply the two brackets of E term by term, as with $(a - b)(a - b)$. Rule 1 turns the transpose of $y - X\beta$ into the transpose of y minus the row from Figure 4:

$$(y - X\beta)^\top = y^\top - \beta^\top X^\top$$

The four products, one per line:

$$y^\top \cdot y = y^\top y$$

$$y^\top \cdot (-X\beta) = -y^\top X\beta$$

$$(-\beta^\top X^\top) \cdot y = -\beta^\top X^\top y$$

$$(-\beta^\top X^\top) \cdot (-X\beta) = +\beta^\top X^\top X\beta$$

Adding them:

$$E = y^\top y - y^\top X\beta$$

$$- \beta^\top X^\top y + \beta^\top X^\top X\beta$$

By Rule 2 the two middle terms are equal, so they combine:

$$E = y^\top y - 2y^\top X\beta + \beta^\top X^\top X\beta$$

Check on the tiny example, one part per line. The first part:

$$y^\top y = 4 \times 4 + 5 \times 5 = 41$$

The middle part, using the 32 of Rule 2:

$$2y^\top X\beta = 2 \times 32 = 64$$

The last part is the column $X\beta = (3, 4)$ times its own transpose:

$$\beta^T X^T X \beta = (X\beta)^T (X\beta)$$

$$= 3 \times 3 + 4 \times 4 = 25$$

Together:

$$E = 41 - 64 + 25 = 2$$

the same 2 as the direct sum of squared errors above.

5. Setting the derivative to zero

Key point: At the best coefficients the error bowl is flat in every direction: its slope along each coefficient is zero. Two small rules turn that into one matrix equation.

As in simple linear regression, the best coefficients sit at the bottom of the error bowl, where the slope is zero. Now there is one slope per coefficient. Figure 5 shows the two slopes on the four students of Section 6.1. Each panel cuts the error bowl along one coefficient while the other is held fixed, and the orange tangent shows the slope there. Watch both slopes shrink as β walks to $(-0.81, 0.57)$: there both tangents lie flat at the same moment.




Figure 5: The gradient $-2X^T y + 2X^T X \beta$ on the four students: each panel is the error E along one coefficient, the orange tangent's slope is one entry of the gradient, and both reach 0 together at $\beta = (-0.81, 0.57)$

The slope along one coefficient, with the others held fixed, is a **partial derivative** (G-1457). Stacking all of them into one column gives the **gradient** (G-863). With two coefficients:

$$\frac{\partial E}{\partial \beta} = \begin{bmatrix} \partial E / \partial \beta_0 \\ \partial E / \partial \beta_1 \end{bmatrix}$$

“Both tangents flat” in Figure 5 means every entry of this column is 0.

5.1 Three rules of matrix calculus, each checked on numbers

Key point: The matrix rules are the old rules “the derivative of a constant is 0”, “of $a\beta$ is a ” and “of $a\beta^2$ is $2a\beta$ ”, applied to every coefficient at once.

Rules for differentiating with vectors and matrices are called **matrix calculus** (G-1176). E has three parts, and each needs one rule. We check each rule by writing the part out for two coefficients, β_0 and β_1 , taking the two ordinary partial derivatives, and comparing.

Rule A: a part without β has derivative 0. The first part, $y^\top y$, is 41 in the tiny example whatever β is. A constant does not change when β changes, so

$$\frac{\partial}{\partial \beta} (y^\top y) = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Rule B: a weighted sum of the coefficients has the weights as its derivative. The matrix version of “the derivative of $a\beta$ is a ”. Section 4.1 showed that the middle part is $\beta^\top X^\top y$, and that $X^\top y = (9, 23)$ in the tiny example. So, written out,

$$\beta^\top X^\top y = 9\beta_0 + 23\beta_1$$

The two partial derivatives, one per line:

$$\frac{\partial}{\partial \beta_0} (9\beta_0 + 23\beta_1) = 9$$

$$\frac{\partial}{\partial \beta_1} (9\beta_0 + 23\beta_1) = 23$$

Stacked, they are $(9, 23)$, which is $X^\top y$ itself. In general:

$$\frac{\partial}{\partial \beta} (\beta^\top X^\top y) = X^\top y$$

and so, since the middle part is twice this (Rule 2):

$$\frac{\partial}{\partial \beta} (2y^\top X\beta) = 2X^\top y$$

Rule C: a “squared” part gives two times the matrix times β . The matrix version of “the derivative of $a\beta^2$ is $2a\beta$ ”. The last part has the square matrix $X^\top X$ in the middle. In the tiny example:

$$X^\top X = \begin{bmatrix} 1 \times 1 + 1 \times 1 & 1 \times 2 + 1 \times 3 \\ 2 \times 1 + 3 \times 1 & 2 \times 2 + 3 \times 3 \end{bmatrix}$$

$$X^\top X = \begin{bmatrix} 2 & 5 \\ 5 & 13 \end{bmatrix}$$

Written out, $\beta^\top X^\top X \beta$ has one term per entry of the matrix:

$$\beta^\top X^\top X \beta = 2\beta_0\beta_0 + 5\beta_0\beta_1$$

$$+ 5\beta_1\beta_0 + 13\beta_1\beta_1$$

Collecting the two middle terms, and calling this part P :

$$P = 2\beta_0^2 + 10\beta_0\beta_1 + 13\beta_1^2$$

Its two partial derivatives:

$$\frac{\partial P}{\partial \beta_0} = 4\beta_0 + 10\beta_1$$

$$\frac{\partial P}{\partial \beta_1} = 10\beta_0 + 26\beta_1$$

Now the matrix rule’s answer, two times the matrix times β :

$$2X^\top X \beta = 2 \begin{bmatrix} 2\beta_0 + 5\beta_1 \\ 5\beta_0 + 13\beta_1 \end{bmatrix}$$

$$2X^\top X \beta = \begin{bmatrix} 4\beta_0 + 10\beta_1 \\ 10\beta_0 + 26\beta_1 \end{bmatrix}$$

The same two lines. In general:

$$\frac{\partial}{\partial \beta} (\beta^\top X^\top X \beta) = 2X^\top X \beta$$

The rule works because $X^\top X$ is a **symmetric matrix** (G-1932): it equals its own transpose, so the two 5s above are equal and add up to the 10 (MML §5.5).

Extra: For any square matrix B the general rule is (MML eq. 5.107, which writes the gradient as a row):

$$\frac{\partial(x^\top Bx)}{\partial x} = x^\top (B + B^\top)$$

With $B = X^\top X$, which equals its own transpose:

$$B + B^\top = 2X^\top X$$

5.2 The gradient, then zero

Key point: Adding the three derivatives gives the gradient; setting it to zero gives the normal equations.

Apply Rules A, B and C to the three parts of E, one per line:

$$\frac{\partial}{\partial \beta} (y^\top y) = 0$$

$$\frac{\partial}{\partial \beta} (-2y^\top X\beta) = -2X^\top y$$

$$\frac{\partial}{\partial \beta} (\beta^\top X^\top X\beta) = 2X^\top X\beta$$

Adding them:

$$\frac{\partial E}{\partial \beta} = -2X^\top y + 2X^\top X\beta$$

Check on the tiny example at $\beta = (1, 1)$, where both errors were 1:

$$-2X^\top y = \begin{bmatrix} -18 \\ -46 \end{bmatrix}$$

$$2X^\top X\beta = \begin{bmatrix} 4 + 10 \\ 10 + 26 \end{bmatrix}$$

$$2X^\top X\beta = \begin{bmatrix} 14 \\ 36 \end{bmatrix}$$

$$\frac{\partial E}{\partial \beta} = \begin{bmatrix} -18 + 14 \\ -46 + 36 \end{bmatrix}$$

$$\frac{\partial E}{\partial \beta} = \begin{bmatrix} -4 \\ -10 \end{bmatrix}$$

The same slopes come from the simple-regression derivatives of [the error function](#), minus two times the sum of the errors, and minus two times the sum of error times feature:

$$-2(1 + 1) = -4$$

$$-2(1 \times 2 + 1 \times 3) = -10$$

At the best coefficients the gradient is zero:

$$-2X^T y + 2X^T X \beta = 0$$

Dividing by 2 and moving one term across:

$$X^T X \beta = X^T y$$

These are the **normal equations** (G-1345): $m + 1$ equations, one for each coefficient.

6. The normal equation

Key point: To get the coefficients alone, we “divide” both sides by the matrix $X^T X$. For matrices, dividing means multiplying by the inverse.

With ordinary numbers, “2 times a number is 6” is solved by dividing by 2, which is the same as multiplying by one half, the number that undoes 2. Matrices have no division; instead we multiply by the **inverse** (G-968), the matrix that undoes $X^T X$. A matrix times its inverse gives the **identity matrix** (G-915), the matrix with 1s on the diagonal and 0s elsewhere, which changes nothing it multiplies.

For a 2×2 matrix the inverse has a three-step recipe (MML §2.2.2), shown in Figure 6:

1. swap the two numbers on the diagonal;
2. flip the sign of the other two;
3. divide everything by the diagonal product minus the off-diagonal product.

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \xrightarrow{\text{inverse}} \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

1. swap the two diagonal numbers
2. flip the sign of the other two
3. divide everything by $ad - bc$

$$\begin{bmatrix} 2 & 5 \\ 5 & 13 \end{bmatrix}^{-1} = \frac{1}{2 \times 13 - 5 \times 5} \begin{bmatrix} 13 & -5 \\ -5 & 2 \end{bmatrix} = \begin{bmatrix} 13 & -5 \\ -5 & 2 \end{bmatrix}$$

$$\text{check: } \begin{bmatrix} 2 & 5 \\ 5 & 13 \end{bmatrix} \begin{bmatrix} 13 & -5 \\ -5 & 2 \end{bmatrix} = \begin{bmatrix} 26 - 25 & -10 + 10 \\ 65 - 65 & -25 + 26 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \text{ the identity}$$

Figure 6: The 2×2 inverse recipe, worked on the tiny example's $X^T X$ and checked: the matrix times its inverse is the identity

The number in step 3 is the **determinant** (G-598). For bigger matrices the computer does the work; the idea is the same.

On the tiny example (bottom of Figure 6), the determinant of the transpose of X times X is

$$2 \times 13 - 5 \times 5 = 26 - 25 = 1$$

so the inverse is the swapped matrix itself:

$$\begin{aligned} & \begin{bmatrix} 2 & 5 \\ 5 & 13 \end{bmatrix}^{-1} \\ &= \begin{bmatrix} 13 & -5 \\ -5 & 2 \end{bmatrix} \end{aligned}$$

Multiply it by the column (9, 23) from Rule B:

$$\beta_0 = 13 \times 9 + (-5) \times 23 = 117 - 115 = 2$$

$$\beta_1 = (-5) \times 9 + 2 \times 23 = -45 + 46 = 1$$

The line “prediction = 2 + 1 × feature” passes through both points:

$$2 + 1 \times 2 = 4$$

$$2 + 1 \times 3 = 5$$

In general, multiplying both sides of the normal equations by the inverse leaves β alone:

$$\beta = (X^T X)^{-1} X^T y$$

The formula is the **normal equation** (G-1344). A formula that gives the answer directly is a **closed-form solution** (G-398); this one gives all coefficients of multiple linear regression at once, and the method is called **ordinary least squares**, OLS (G-1406). scikit-learn’s **LinearRegression** reaches the same solution with a numerically safer method than a literal inverse (scikit-learn docs, LinearRegression; LAPACK, DGELSD).

6.1 A worked example

Key point: On four students with one feature, the normal equation gives the same slope and intercept as the simple formulas.

Take the first four students of the placement data, with CGPA 6.89, 5.12, 7.82, 7.42 and packages 3.26, 1.98, 3.25, 3.67.

1. **Build X:** a column of four 1s next to the four CGPAs.

$$X = \begin{bmatrix} 1 & 6.89 \\ 1 & 5.12 \\ 1 & 7.82 \\ 1 & 7.42 \end{bmatrix}$$

$$y = \begin{bmatrix} 3.26 \\ 1.98 \\ 3.25 \\ 3.67 \end{bmatrix}$$

2. **Compute the transpose of X times X.** Each entry is a column of X times a column of X. The column of 1s times itself counts the students:

$$1 + 1 + 1 + 1 = 4$$

The column of 1s times the CGPAs adds the CGPAs:

$$6.89 + 5.12 + 7.82 + 7.42 = 27.25$$

The CGPAs times themselves add their squares:

$$47.4721 + 26.2144 + 61.1524 + 55.0564$$

$$= 189.8953$$

$$X^T X = \begin{bmatrix} 4 & 27.25 \\ 27.25 & 189.8953 \end{bmatrix}$$

3. **Compute the transpose of X times y.** The column of 1s times y adds the packages:

$$3.26 + 1.98 + 3.25 + 3.67 = 12.16$$

The CGPAs times the packages, one product per student:

$$6.89 \times 3.26 = 22.4614$$

$$5.12 \times 1.98 = 10.1376$$

$$7.82 \times 3.25 = 25.4150$$

$$7.42 \times 3.67 = 27.2314$$

$$22.4614 + 10.1376 + 25.4150 + 27.2314$$

$$= 85.2454$$

$$X^T y = \begin{bmatrix} 12.16 \\ 85.2454 \end{bmatrix}$$

4. **Invert with the recipe of Figure 6.** The determinant:

$$4 \times 189.8953 - 27.25 \times 27.25$$

$$= 759.5812 - 742.5625 = 17.0187$$

Swap the diagonal, flip the signs of the other two, divide by 17.0187:

$$(X^T X)^{-1} = \frac{1}{17.0187} M$$

$$M = \begin{bmatrix} 189.8953 & -27.25 \\ -27.25 & 4 \end{bmatrix}$$

$$= \begin{bmatrix} 11.158 & -1.601 \\ -1.601 & 0.235 \end{bmatrix}$$

Check the top-left entry of the inverse times the matrix, which should be 1:

$$\frac{189.8953 \times 4 - 27.25 \times 27.25}{17.0187}$$

$$= \frac{17.0187}{17.0187} = 1$$

5. **Multiply.** To keep rounding out, we multiply by the swapped matrix first and divide by 17.0187 last:

$$189.8953 \times 12.16 - 27.25 \times 85.2454$$

$$= 2309.127 - 2322.937 = -13.810$$

$$-27.25 \times 12.16 + 4 \times 85.2454$$

$$= -331.360 + 340.982 = 9.622$$

$$\beta_0 = -13.810/17.0187 = -0.811$$

$$\beta_1 = 9.622/17.0187 = 0.565$$

So $\beta_0 = -0.81$, the **intercept** (G-960), and $\beta_1 = 0.57$, the **slope** (G-1823): exactly what the simple linear regression formulas give for these four points. The simple formulas are the normal equation with a single feature.

6.2 The picture: a projection

Extra: This section goes beyond the derivation: a second, geometric way to see the same equation.

Key point: Every possible set of predictions lies on one flat plane, and the real targets stick out of it. The best predictions are the point of the plane closest to the targets, where the error meets the plane at a right angle; that right angle gives the normal equation.

The point straight below is the **projection** (G-1583) of y onto the plane.

The normal equation also has a geometric meaning (ESL §3.2, Figure 3.2; MML §3.8). Treat the n targets as one vector y with n entries, and each column of X the same way. A prediction $X\beta$ is a weighted sum of the columns, so all possible predictions fill the flat space the columns span, called the **column space** (G-414) of X . Unless the data lie exactly on a line, y is not in it.

With three observations every vector has three entries, so we can draw it. Figure 7 uses the first three students of the worked example: X has the columns $\mathbf{1} = [1, 1, 1]$ and $\text{cgpa} = [6.89, 5.12, 7.82]$, and $y = [3.26, 1.98, 3.25]$. Watch the error length as the green point moves through the plane, and the angle at the point where it stops.




Figure 7: The column space of X for three students is a plane (the two blue columns span it). The target y (orange) sticks out of it. Moving the prediction $X\beta$ through the plane, the error length is smallest, 0.36, at the foot of the perpendicular from y , where the residual meets the plane at a right angle

- **Closest point:** the error $\|y - X\hat{\beta}\|$ is the length of the dashed line, and its square is E of Section 3. It is smallest, 0.36 (its square, 0.13, is E for these three students), at $\hat{y} = X\hat{\beta} = [2.97, 2.08, 3.44]$, with $\hat{\beta} = [-0.50, 0.50]$.
- **Right angle:** there the residual $[0.29, -0.10, -0.19]$ is perpendicular to both columns: its **dot product** (G-634) with each column is 0.

The two dot products, one per line:

With the column of 1s:

$$0.29 - 0.10 - 0.19 = 0.00$$

With the cgpa column, one product per line:

$$0.29 \times 6.89 = 1.998$$

$$(-0.10) \times 5.12 = -0.512$$

$$(-0.19) \times 7.82 = -1.486$$

$$1.998 - 0.512 - 1.486 = 0.00$$

Stacked as one equation (the transpose of X holds the two columns as rows):

$$X^T(y - X\hat{\beta}) = 0$$

Multiplying out the bracket and moving one term across gives the normal equations of Section 5:

$$X^T X \hat{\beta} = X^T y$$

So the calculus of Section 5 and the right angle of Figure 7 give the same equations. The figure draws the **1** direction four times shorter so that the small residual is visible; shrinking a direction inside the plane does not change the right angle or which point is closest.

7. The cost of the inverse

Key point: Inverting $X^T X$ takes time that grows roughly with the cube of the number of features. With very many features, gradient descent is used instead.

The normal equation has one expensive step: the inverse. The more features, the bigger the matrix to invert. $X^T X$ is a square matrix with one row and one column per coefficient, $(m+1) \times (m+1)$. Inverting an $m \times m$ matrix takes on the order of m^3 operations: doubling the features makes the work about 8 times larger.

Figure 8 measures it on this computer.




Figure 8: Time to invert a matrix as its size grows

From 1,000 to 2,000 features, the time grows about 7 times, from 0.08 to 0.56 seconds (Figure 8). With tens of thousands of features, as with text or image data, the inverse becomes very slow and memory-hungry. With numbers, for 20,000 features:

$$20,000/1,000 = 20 \text{ times the features}$$

$$20^3 = 8,000 \text{ times the work}$$

$$8,000 \times 0.08 \text{ s} = 640 \text{ s} \approx 11 \text{ minutes}$$

The matrix alone would hold one number per entry, 8 bytes each:

$$20,000 \times 20,000 \times 8 \text{ bytes} = 3.2 \text{ GB}$$

The cost of the inverse is why there is a second method, **gradient descent** (G-862): it does not compute any inverse, but approaches the best coefficients step by step. Its answer is very close to the normal equation's. In scikit-learn:

- `LinearRegression` uses the closed-form (OLS) solution;
- `SGDRegressor` (G-1783) uses gradient descent.

For most tabular data the number of features is small, and `LinearRegression` is the usual choice. [Gradient descent](#) is taught next.

Extra: $X^T X$ has no inverse when one feature can be built exactly from others (**multicollinearity**, G-1273, as in [the dummy variable trap](#)). Then the normal equation has no unique answer. Libraries handle this with a “pseudo-inverse” or with regularisation ([penalising large coefficients](#)). scikit-learn’s `LinearRegression` takes the first route (LAPACK, DGELSD).

8. Summary

Step	Formula
Predictions	$\hat{y} = X\beta$ (X has a first column of 1s)
Errors	$e = y - X\beta$
Sum of squared errors	$E = e^T e$
Expanded	$E = y^T y - 2y^T X\beta + \beta^T X^T X\beta$
Derivative set to zero	$X^T X\beta = X^T y$
Normal equation	$\beta = (X^T X)^{-1} X^T y$

- A column of 1s in X lets the intercept be treated like any other coefficient.
- One formula gives all $m + 1$ coefficients at once; with one feature it reduces to the simple formulas.
- The inverse costs about m^3 operations, so very wide data uses gradient descent instead.

9. Sources

Built from

- CampusX, “Multiple Linear Regression | Part 2 | Mathematical Formulation From Scratch”, YouTube, <https://www.youtube.com/watch?v=NU37mF5q8VE>

Other references

- Deisenroth, M. P., Faisal, A. A. and Ong, C. S. (2020). *Mathematics for Machine Learning* (MML). Cambridge University Press. §2.2.2 Inverse and Transpose (the 2×2 inverse), §3.8 Orthogonal Projections and §5.5 Useful Identities for Computing Gradients.
- Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The Elements of Statistical Learning*, 2nd ed. (ESL). Springer. §3.2 and Figure 3.2 (the geometry of least squares: y projected onto the column space of X).
- scikit-learn documentation, `sklearn.linear_model.LinearRegression`, Notes section (uses `scipy.linalg.lstsq`).
- LAPACK documentation, DGELSD: minimum-norm least-squares solution using the SVD, for a matrix that may be rank-deficient. netlib.org/lapack.

10. Key terms

Term	Meaning
Feature	An input variable: one column of the data table
Target	The output we predict
Observation	One record: one row of the data table

Term	Meaning
Design matrix (X)	The data as a matrix, one row per observation, with a first column of 1s for the intercept
Coefficient vector (β)	All the coefficients of the model, β_0 to β_m , as one column
Matrix calculus	Rules for differentiating expressions with vectors and matrices
Normal equations	$X^T X \beta = X^T y$: the conditions that the best coefficients satisfy
Normal equation	$\beta = (X^T X)^{-1} X^T y$: the closed-form solution of linear regression
Inverse matrix	The matrix that undoes another: their product is the identity matrix
Determinant	For a 2×2 matrix, the diagonal product minus the off-diagonal product; the inverse divides by it
Transpose	The same numbers with rows turned into columns
Gradient	The column of partial derivatives, one per coefficient
Residual (G-705)	The error on one observation: actual minus predicted value
Closed-form solution	An answer given directly by a formula
Ordinary least squares (OLS) (G-1406)	The closed-form method for linear regression: the coefficients with the smallest sum of squared errors
Column space	All the vectors that can be built as weighted sums of a matrix's columns
Identity matrix	The square matrix with 1s on the diagonal and 0s elsewhere; multiplying by it changes nothing